eSDK Cloud Storage Plugins V2.2.16

User Guide (Kubernetes CSI)

Issue 01

Date 2022-04-13





Copyright © Huawei Technologies Co., Ltd. 2022. All rights reserved.

No part of this document may be reproduced or transmitted in any form or by any means without prior written consent of Huawei Technologies Co., Ltd.

Trademarks and Permissions

HUAWEI and other Huawei trademarks are trademarks of Huawei Technologies Co., Ltd. All other trademarks and trade names mentioned in this document are the property of their respective holders.

Notice

The purchased products, services and features are stipulated by the contract made between Huawei and the customer. All or part of the products, services and features described in this document may not be within the purchase scope or the usage scope. Unless otherwise specified in the contract, all statements, information, and recommendations in this document are provided "AS IS" without warranties, guarantees or representations of any kind, either express or implied.

The information in this document is subject to change without notice. Every effort has been made in the preparation of this document to ensure accuracy of the contents, but all statements, information, and recommendations in this document do not constitute a warranty of any kind, express or implied.

Huawei Technologies Co., Ltd.

Address: Huawei Industrial Base

Bantian, Longgang Shenzhen 518129

People's Republic of China

Website: https://e.huawei.com

About This Document

Intended Audience

This document is intended for:

- Technical support engineers
- O&M engineers
- Engineers with basic knowledge of storage and Kubernetes

Symbol Conventions

The symbols that may be found in this document are defined as follows.

Symbol	Description
▲ DANGER	Indicates a hazard with a high level of risk which, if not avoided, will result in death or serious injury.
⚠ WARNING	Indicates a hazard with a medium level of risk which, if not avoided, could result in death or serious injury.
⚠ CAUTION	Indicates a hazard with a low level of risk which, if not avoided, could result in minor or moderate injury.
NOTICE	Indicates a potentially hazardous situation which, if not avoided, could result in equipment damage, data loss, performance deterioration, or unanticipated results. NOTICE is used to address practices not related to personal injury.
□ NOTE	Supplements the important information in the main text. NOTE is used to address information not related to personal injury, equipment damage, and environment deterioration.

Change History

Issue	Date	Description
01	2022-04-13	This issue is the first official release.

Contents

About This Document	ii
1 Overview	1
2 Environmental Requirements	2
3 Restrictions	5
4 Installation and Deployment	9
4.1 Obtaining the Software Package	g
4.2 Uploading the Components in the Software Package	g
4.3 Creating a Huawei CSI Image	10
4.4 Configuring Host Multipathing	12
4.4.1 Installing Native Multipathing Software	12
4.4.1.1 Installing the Multipathing Tool Package	12
4.4.1.2 Configuring the Multipathing Service	13
4.4.2 Installing Huawei Multipathing Software	15
4.4.2.1 Installing the Huawei Multipathing Tool	15
4.5 Connecting to Enterprise Storage	15
4.5.1 Connecting to Enterprise Storage SAN over iSCSI	15
4.5.2 Connecting to Enterprise Storage SAN over FC	19
4.5.3 Connecting to Enterprise Storage NAS over NFS	23
4.5.4 Connecting to Enterprise Storage SAN over NVMe over RoCE	26
4.5.5 Connecting to Enterprise Storage SAN over NVMe over FC	29
4.6 Connecting to Distributed Storage	33
4.6.1 Connecting to Distributed Storage SAN over SCSI	33
4.6.2 Connecting to Distributed Storage SAN over iSCSI	36
4.6.3 Connecting to Distributed Storage NAS over NFS	39
4.7 Starting huawei-csi Services	42
5 Upgrade Operations	49
5.1 Uninstalling Original CSI	49
5.1.1 Uninstalling the huawei-csi-node Service	49
5.1.2 Uninstalling the huawei-csi-controller Service	50
5.1.3 Deleting the huawei-csi-configmap Object	50
5.1.4 Deleting the huawei-csi-secret Object	50

5.1.5 Deleting the RBAC Permission	51
5.1.6 Deleting the Image of the Earlier Version	52
5.2 Installing New CSI	53
6 Instructions for Use	54
6.1 (Conditionally Mandatory) Managing a StorageClass	54
6.1.1 Creating a StorageClass	54
6.1.1.1 Creating a LUN StorageClass	54
6.1.1.2 Creating a File System StorageClass	56
6.1.2 Deleting a StorageClass	58
6.2 (Conditionally Mandatory) Managing a PV	58
6.2.1 Creating a PV	59
6.2.2 Deleting a PV	62
6.3 Managing a PVC	
6.3.1 Creating a PVC	63
6.3.2 (Optional) Expanding the Capacity of a PVC	67
6.3.3 (Optional) Cloning a PVC	69
6.3.4 (Optional) Creating a PVC Using a Snapshot	
6.3.5 Deleting a PVC	73
6.4 Managing a Pod	
6.4.1 Creating a Pod	
6.4.2 Deleting a Pod	
6.5 (Optional) Managing a Snapshot	77
6.5.1 Installing the Snapshot-Dependent Component Service	
6.5.2 Managing a VolumeSnapshotClass	
6.5.2.1 Creating a VolumeSnapshotClass	
6.5.2.2 Deleting a VolumeSnapshotClass	
6.5.3 Managing a VolumeSnapshot	
6.5.3.1 Creating a VolumeSnapshot	
6.5.3.2 Deleting a VolumeSnapshot	80
7 Advanced Features	
7.1 Configuring Multiple Backends	
7.2 Creating a PVC for a Specified Backend	
7.3 Creating a PVC for a Specified Storage Pool	
7.4 Configuring ALUA	
7.4.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3	
7.4.2 Configuring ALUA for OceanStor Dorado 6.x	
7.4.3 Configuring ALUA for Distributed Storage	90
7.5 Configuring Storage Topology Awareness	
7.6 Advanced Features of Enterprise Storage	
7.6.1 Configuring QoS	
7.6.2 Configuring a vStore	
7.6.3 Configuring NAS HyperMetro	100

7.6.3.1 Prerequisites	100
7.6.3.2 Configuration Procedure	102
7.6.4 Configuring an Application Type	105
7.7 Advanced Features of Distributed Storage	106
7.7.1 Configuring QoS	106
7.7.2 Configuring a Soft Quota	107
7.7.3 Configuring an Account	108
8 Uninstalling CSI	111
8.1 (Optional) Uninstalling the Snapshot-Dependent Component Service	111
8.2 Uninstalling the huawei-csi-node Service	112
8.3 Uninstalling the huawei-csi-controller Service	112
8.4 Deleting the huawei-csi-configmap Object	113
8.5 Deleting the huawei-csi-secret Object	113
8.6 Deleting the RBAC Permission	113
8.7 Deleting the Image of the Earlier Version	114
9 Common Operations	116
9.1 Updating the User Name or Password of a Storage Device Configured on CSI	116
9.2 Updating the configmap Object of huawei-csi	117
9.3 Adding a Backend for huawei-csi	120
9.4 Updating the huawei-csi-controller Service	120
9.5 Updating the huawei-csi-node Service	121
9.6 Modifying the Log Output Mode	121
9.6.1 Modifying the Log Output Mode of the huawei-csi-controller Service	121
9.6.2 Modifying the Log Output Mode of the huawei-csi-node Service	123
10 FAQ	125
10.1 Viewing Log Information	125
10.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernet Platform Is Set Up for the First Time	
10.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported	
10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters	
10.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName	131
10.6 When a PVC Is Created, the PVC Is in the Pending State	132
10.7 Before a PVC Is Deleted, the PVC Is in the Pending State	134
10.8 When a Pod Is Created, the Pod Is in the ContainerCreating StateState	135
10.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created	136
11 Appendix	138
11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3	138
11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x	139
11.3 Example ALUA Configuration Policy of Distributed Storage	140

1 Overview

This document describes how to deploy and use the Kubernetes CSI plug-in so that Huawei enterprise and distributed storage devices provide persistent volume storage capabilities for Kubernetes.

2 Environmental Requirements

- Kubernetes has been deployed and is running properly.
- A Huawei storage device is running properly.
- Drivers required for scanning disks and mounting files (such as iSCSI, DM Multipath, and UltraPath NVMe. For details, see multipathing software selection.) must be installed on all worker hosts of Kubernetes in advance. If containers and services cannot run properly due to lack of system tools, view logs by referring to 10.1 Viewing Log Information and install the tools on the hosts.

Table 2-1 Version mappings among Kubernetes, enterprise storage products, and host machine operating systems (OSs)

Kubernetes Version	Enterprise Storage Product Version	Host Machine OS Version
1.18/1.19/1.20/ 1.21/1.22/1.23	OceanStor Dorado 6.0.0/6.0.1/6.1.0/6.1.2/6.1.3	CentOS 7.6/7.7/7.9/8.2 (x86_64)
	OceanStor Dorado V3 V300R002 OceanStor 6.1.3	SUSE 15 SP2 (x86_64) Red Hat CoreOS 4.6/4.7/4.8/4.9 (x86_64)
	OceanStor F V5/V5 V500R007/V500R007 Kunpeng OceanStor F V3/V3 V300R006	Ubuntu 18.04/20.04 (x86_64)

Table 2-2 Version mappings among Kubernetes, distributed storage products, and host machine OSs

Kubernetes Version	Distributed Storage Product Version	Host Machine OS Version	
1.18/1.19/1.20/ 1.21/1.22/1.23		CentOS 7.6/7.7/7.9/8.2 (x86_64) SUSE 15 SP2 (x86_64)	
	OceanStor Pacific series 8.1.0/8.1.1/8.1.2	Red Hat CoreOS 4.6/4.7/4.8/4.9 (x86_64) Ubuntu 18.04/20.04 (x86_64)	

NOTICE

If the host machine OS is CoreOS 4.6/4.7/4.8/4.9, see *Kubernetes CSI for Red Hat OpenShift User Guide*

Table 2-3 Features supported by Huawei CSI (√: supported; x: not supported)

Feature	1.18	1.19	1.20	1.21	1.22	1.23
Create PVC	√	√	√	√	√	√
Delete PVC	√	√	√	√	√	√
Create Pod	√	√	√	√	√	√
Delete Pod	√	√	√	√	√	√
Offline Resize	√	√	√	√	√	√
Online Resize	√	√	√	√	√	√
Create Snapshot	√	√	√	√	√	√
Delete Snapshot	√	√	√	√	√	√
Restore	√	√	√	√	√	√
Clone	√	√	√	√	√	√

Table 2-4 Mappings between host machine OS versions and multipathing software versions

Host Machine OS Version	Native Multipathing Software Version	Huawei Multipathing Software Version (Supported Only by Enterprise Storage)
CentOS 7.6/7.7/7.9 (x86_64)	Delivered with the OS, supporting FC/iSCSI	UltraPath 31.1.0, supporting FC/iSCSI
CentOS 8.2 (x86_64)	Delivered with the OS, supporting FC/iSCSI	UltraPath 31.1.0, supporting FC/iSCSI UltraPath-NVMe 31.1.RC8, supporting NVMe over RoCE/NVMe over FC
SUSE 15 SP2 (x86_64)	Delivered with the OS, supporting FC/iSCSI	UltraPath 31.1.0, supporting FC/ iSCSI UltraPath-NVMe 31.1.RC8, supporting NVMe over RoCE
CoreOS 4.6/4.7/4.8/4.9 (x86_64)	Delivered with the OS, supporting FC/iSCSI	Not supported
Ubuntu 18.04/20.04 (x86_64)	Delivered with the OS, supporting FC/iSCSI	Not supported

3 Restrictions

This chapter describes the restrictions on connecting CSI to storage devices.

Table 3-1 Restrictions

Scenario	Restriction	Supported Storage	Remarks
PersistentVolume- Claim (PVC) access mode	 ReadWriteOnc e: SAN/NAS ReadWriteMa ny: SAN (only when volumeMode is Block)/NAS ReadOnlyMan y: SAN/NAS 	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.x, OceanStor Pacific series	N/A
Creating/Deleting PVCs	You are advised to create or delete a maximum of 100 items in a batch.	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.x, OceanStor Pacific series	The maximum number of concurrent RESTful requests is 100.

Scenario	Restriction	Supported Storage	Remarks
Creating/Deleting Pods	You are advised to create or delete a maximum of 100 items in a batch.	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.x, OceanStor Pacific series	N/A
Snapshot	 NAS HyperMetro does not support snapshots. A PVC created using a static PV does not support snapshots. 	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.x	N/A
Creating a PVC using a snapshot	 The StorageClass and volumeMode of the source PVC must be the same as those of the target PVC. NAS HyperMetro is not supported. 	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5	N/A

Scenario	Restriction	Supported Storage	Remarks
Expanding the capacity of a PVC	 Only capacity expansion is supported. Capacity reduction is not supported. A PVC whose access mode is ROX does not support capacity expansion. A PVC created using a static PV does not support capacity expansion. If the Kubernetes version is 1.18 and the value of volumeMode in the PVC is Block, the PVC does not support capacity expansion. 	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.x	N/A
Cloning a PVC	 The StorageClass and volumeMode of the source PVC must be the same as those of the target PVC. A PVC created using a static PV does not support cloning. NAS HyperMetro is not supported. 	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6.x, FusionStorage 8.0.x, OceanStor Pacific series NAS: OceanStor V3/V5	N/A

Scenario	Restriction	Supported Storage	Remarks
HyperMetro	PVCs and Pods can be created only when both HyperMetro storage systems are normal. If a single storage system is faulty, only the services that have been delivered are normal and new services cannot be delivered. If both storage systems are faulty, contact Huawei technical support engineers.	NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.1.3	N/A
Residual drive letter	Due to a node fault, containerized applications are migrated to other nodes. After the node recovers, residual drive letters exist on the node. Manually clear the residual drive letters. For details, see 10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters.	SAN: OceanStor V3/V5/6.1, OceanStor Dorado V3/6. <i>x</i> , FusionStorage 8.0. <i>x</i> , OceanStor Pacific series	Condition: iSCSI/FC + Multipath
Tenant	N/A	NAS: OceanStor V3/V5/6.1, OceanStor Dorado 6.1.3	N/A

4 Installation and Deployment

- 4.1 Obtaining the Software Package
- 4.2 Uploading the Components in the Software Package
- 4.3 Creating a Huawei CSI Image
- 4.4 Configuring Host Multipathing
- 4.5 Connecting to Enterprise Storage
- 4.6 Connecting to Distributed Storage
- 4.7 Starting huawei-csi Services

4.1 Obtaining the Software Package

You can obtain Huawei Kubernetes CSI through Huawei Kubernetes CSI warehouse.

- **Step 1** Open a browser and enter https://github.com/Huawei/eSDK_K8S_Plugin/releases in the address box.
- **Step 2** Select the desired version package and download **eSDK_Cloud_Storage_Plugin_***.*.**zip. *.*.**indicates the release version number. (The version matching this document is 2.2.16.)
- **Step 3** Decompress the package.
- **Step 4** Find the package and documents in the directory generated after the decompression.

----End

4.2 Uploading the Components in the Software Package

Step 1 Decompress *eSDK_Cloud_Storage_Plugin_*.*.***:zip to obtain the software package and sample files required for installing and using CSI. Table 4-1 shows the software package structure.

-	
Component	Description
bin/huawei-csi	Implements the CSI API.
bin/secretGenerate	Encrypts plaintext passwords and produces secret objects.
bin/secretUpdate	Encrypts plaintext passwords and updates secret objects.
deploy	.yaml sample file used during CSI deployment.
examples	.yaml sample file used during CSI use.

Table 4-1 Component description

Step 2 Use a file transfer tool (such as Xftp) to upload the files generated upon decompression to the master node.

----End

4.3 Creating a Huawei CSI Image

Huawei CSI runs as a container. Currently, Huawei CSI provides only a binary package (bin/huawei-csi) which cannot be used directly. Therefore, you need to create a CSI image based on the binary file to start the Huawei CSI service.

Prerequisites

A Linux host with Docker installed is available, and the host can access the Internet (only used to download the image package).

Procedure

- **Step 1** Log in to the Linux host.
- **Step 2** Run the **mkdir image** command to create a directory (for example, **image**) on the host.

mkdir image

Step 3 Run the **cd image** command to access the **image** directory.

cd image

- **Step 4** Copy the huawei-csi component to the **image** directory.
- **Step 5** Run the following command to create a file named **Dockerfile**.

```
# cat <<EOF > ./Dockerfile
FROM busybox:stable-glibc

ADD ["huawei-csi", "/"]
RUN ["chmod", "+x", "/huawei-csi"]

ENTRYPOINT ["/huawei-csi"]

EOF
```

NOTICE

busybox:stable-glibc indicates the basic image and its tag. It is only an example. Replace it based on site requirements.

Step 6 Run the **docker build -f Dockerfile -t huawei-csi:2.2.16** . command to create an image.

docker build -f Dockerfile -t huawei-csi:2.2.16.

□ NOTE

2.2.16 indicates the plug-in version number corresponding to the software package name. It is only an example. Replace it based on site requirements. If the same image already exists in the environment, use **docker image rm** <image-id>.

Step 7 Run the **docker image ls | grep huawei-csi** command to check whether the image is created. If the following information is displayed, it is created.

docker image ls | grep huawei-csi huawei-csi 2.2.16

c8b5726118ac About a minute ago 39 MB

Step 8 Run the **docker save huawei-csi:2.2.16 -o huawei-csi.tar** command to export the image.

docker save huawei-csi:2.2.16 -o huawei-csi.tar

□ NOTE

2.2.16 indicates the plug-in version number corresponding to the software package name. It is only an example. Replace it based on site requirements.

Step 9 Run the **scp huawei-csi.tar** *<user>@<ip>:/<path>* command to copy the **huawei-csi.tar** image file to all worker nodes in the Kubernetes cluster and enter the password as prompted.

scp huawei-csi.tar <user>@<ip>:/<path>

- <user>: user name for logging in to a worker node in the Kubernetes cluster.
- <ip>: IP address for logging in to a worker node in the Kubernetes cluster.
- <path>: name of the folder to be copied to a worker node in the Kubernetes cluster.
- **Step 10** Log in to a worker node in the Kubernetes cluster.
 - If docker is used, run the **docker load -i huawei-csi.tar** command to import the image.

docker load -i huawei-csi.tar

• If containerd is used, run the ctr -n=k8s.io images import huawei-csi.tar command to import the image.

ctr -n=k8s.io images import huawei-csi.tar

- **Step 11** After the import is complete, run the following command to check whether the import is successful. If the following information is displayed, the import is successful.
 - If docker is used, run the **docker image ls | grep huawei-csi** command. # docker image ls | grep huawei-csi huawei-csi 2.2.16 14b854dba227 10 minutes ago 80MB
 - If containerd is used, run the **crictl image ls | grep huawei-csi** command.

 # crictl image ls | grep huawei-csi
 docker.io/library/huawei-csi
 2.2.16
 14b854dba2273
 93.1MB

Step 12 Repeat **Step 9** to **Step 11** to import the image to all worker nodes in the Kubernetes cluster.

----End

4.4 Configuring Host Multipathing

If you use block storage and access storage over the FC/iSCSI/NVMe over RoCE/NVMe over FC protocol, you are advised to configure host multipathing to improve storage link reliability. Currently, the following multipathing software is supported: native multipathing software (DM-Multipath) and Huawei multipathing software (UltraPath and UltraPath-NVMe).

Precautions

- For details about the host multipathing software supported by enterprise storage and distributed storage in different OSs, see **Table 2-4**.
- Table 4-2 lists the multipathing software supported by different protocols when SAN storage is used.

Table 4-2 Multipathing software supported by different protocols when SAN storage is used

SAN Storage Networking Protocol	No Multipathing Software	DM- Multipath	UltraPath	UltraPath- NVMe
VBS	√	х	х	x
iSCSI	√	√	√	√
FC	√	√	√	√
NVMe over RoCE	√	х	х	√
NVMe over FC	√	х	х	√

4.4.1 Installing Native Multipathing Software

Native multipathing software is delivered with the host system. If you need to install it, refer to this section.

4.4.1.1 Installing the Multipathing Tool Package

This section describes how to install the native multipathing tool package.

Prerequisites

Ensure that the worker nodes in the Kubernetes cluster can access the Internet (only used to download the multipathing tool package).

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node in the Kubernetes cluster through the management IP address.
- **Step 2** Install the multipathing tool package based on the OS.
 - CentOS: yum install -y device-mapper-multipath
 - SUSE: zypper install -y multipath-tools
- **Step 3** Enable the host multipathing service.
 - CentOS:

/sbin/mpathconf --enable systemctl start multipathd.service systemctl enable multipathd.service systemctl restart multipathd.service

 SUSE: systemctl restart multipath-tools.service chkconfig multipathd on

Step 4 Repeat **Step 1** to **Step 3** to install the multipathing tool on all worker nodes.

----End

4.4.1.2 Configuring the Multipathing Service

Multipathing is configured to improve the link reliability of LUNs on SAN storage. If multipathing is incorrectly configured, I/O errors will occur when a single link is faulty. As a result, the file systems or disks in the containers managed by the Kubernetes cluster are read-only or faulty, affecting I/O delivery.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node in the Kubernetes cluster through the management IP address.
- Step 2 Run the vi /etc/multipath.conf command to modify the multipath.conf file. If the file does not exist, configure or generate it by referring to the storage host connectivity guide. Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press Esc and enter :wq! to save the modification. This document uses Red Hat as an example. For details about other OSs, see the storage host connectivity guide.

NOTICE

Load balancing mode: During service read and write, the I/O paths from a host to all controllers on a storage device are the same. For details, see **Configuring**Multipathing > Concepts in *Huawei SAN Storage Host Connectivity Guide for Red Hat.*

Local preferred mode: When a host delivers I/Os to controllers, the storage device with better performance is accessed because the service link distances from storage devices are different. For details, see **Configuring Multipathing** > **Concepts** in *Huawei SAN Storage Host Connectivity Guide for Red Hat*.

 If enterprise storage and the load balancing mode are used, you are advised to add the following content to the devices field in the multipathing configuration file (/etc/multipath.conf). For details, see OceanStor Dorado Host Connectivity Guide for Red Hat and OceanStor Dorado Host Connectivity Guide for SUSE.

```
defaults {
     user_friendly_names yes
     find_multipaths no
devices {
  device {
                               "HUAWEI"
          vendor
          product
                                "XSG1"
          path_grouping_policy
                                    multibus
          path_checker
                                 tur
          prio
                              const
          path_selector
                                 "service-time 0"
          failback
                               immediate
          no_path_retry
      }
}
```

If enterprise storage and the local preferred mode are used, you are advised to add the following content to the **devices** field in the multipathing configuration file (/etc/multipath.conf). For details, see *Huawei SAN Storage Host Connectivity Guide for Red Hat* and *Huawei SAN Storage Host Connectivity Guide for SUSE*.

```
defaults {
     user_friendly_names yes
     find_multipaths no
devices {
  device {
                            "HUAWEI"
         vendor
         product
                           "XSG1"
         path_grouping_policy group_by_prio
         path_checker
                             tur
         prio
                          alua
         path_selector
                             "round-robin 0"
         failback
                           immediate
         no_path_retry
}
```

 If distributed storage is used, you are advised to add the following content to the devices field in the multipathing configuration file (/etc/multipath.conf).
 The configuration varies according to the OS. For details, see Configuring Multipathing for an Application Server (Red Hat or CentOS) in FusionStorage 8.0.1 Block Storage Basic Service Configuration Guide 08.

```
defaults {
     user_friendly_names yes
     find_multipaths no
devices {
  device {
                                "Huawei"
          vendor
                                "VBS fileIO"
          product
          path_grouping_policy
                                    multibus
          path_checker
                                  tur
          prio
                              const
          path_selector
                                 "service-time 0"
          failback
                               immediate
          no_path_retry
                                  "10"
```

Step 3 After the configuration is complete, run the following command to restart the multipathd service.

systemctl restart multipathd.service

Step 4 Repeat **Step 1** to **Step 3** to configure the multipathing service for all worker nodes.

----End

4.4.2 Installing Huawei Multipathing Software

Huawei multipathing software is provided by Huawei. If you need to install it, refer to this section.

4.4.2.1 Installing the Huawei Multipathing Tool

This section describes how to install the Huawei multipathing tool package.

Prerequisites

Select Huawei multipathing software according to the host OS, networking, and version mappings. For enterprise users, log in to https://support.huawei.com/enterprise. For carrier users, log in to https://support.huawei.com. Then search for UltraPath to obtain the software package and user guide.

Installation Procedure

Install Huawei multipathing software according to the obtained Huawei multipathing software user quide.

Configuration Procedure

Configure Huawei multipathing software according to the obtained Huawei multipathing software user guide.

○ NOTE

If multiple multipathing software products coexist, see the Huawei multipathing software user guide for compatibility and configuration methods.

4.5 Connecting to Enterprise Storage

This section describes how to connect the huawei-csi plug-in to Huawei enterprise storage.

When the same SAN storage is connected to Kubernetes, you cannot configure multiple data protocols (iSCSI, FC, NVMe over RoCE, and NVMe over FC) on one worker node.

4.5.1 Connecting to Enterprise Storage SAN over iSCSI

Perform this operation when you want to connect to enterprise storage SAN over iSCSI.

Prerequisites

- An iSCSI client has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes. For details, see 4.4 Configuring Host Multipathing.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-iscsi.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-3.

Table 4-3 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-4.

Table 4-4 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the block storage service.

Configuration Item	Format	Description	Remarks
parameters	Dictionary	Variable parameters in scenarios where iSCSI is used. This parameter is mandatory.	In scenarios where iSCSI is used, set the protocol parameter to a fixed value: iscsi . Set the portals parameter to the iSCSI service IP addresses of the storage device. Use commas (,) to separate multiple iSCSI service IP addresses. You can log in to DeviceManager to obtain the iSCSI service IP addresses. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is iSCSI. (For other series, see the corresponding operation description.)

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.5.2 Connecting to Enterprise Storage SAN over FC

Perform this operation when you want to connect to enterprise storage SAN over FC.

Restrictions

To connect to enterprise storage SAN over FC, ensure that no residual drive letter exists on the host. If any residual drive letter exists, clear the drive letter by referring to 10.4 After a Worker Node in the Cluster Breaks Down and

Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters.

Prerequisites

- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes can communicate with the storage device to be connected over FC.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes of Kubernetes. For details, see 4.4 Configuring Host Multipathing.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-fc.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-5.

Table 4-5 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-6.

Table 4-6 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-).
			NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the block storage service.
parameters	Dictionary	Variable parameters in scenarios where FC is used. This parameter is mandatory.	In scenarios where FC is used, set the protocol parameter to a fixed value: fc .

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.5.3 Connecting to Enterprise Storage NAS over NFS

Perform this operation when you want to connect to enterprise storage NAS over NFS.

Prerequisites

- An NFS client tool has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes communicate properly with the NFS logical port of the storage device to be connected.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-nfs.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-7.

Table 4-7 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-8.

Table 4-8 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage NAS is connected, the value is fixed to oceanstor-nas.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the file storage service.

Configuration Item	Format	Description	Remarks
parameters	Dictionary	Variable parameters in scenarios where NFS is used. This parameter is mandatory.	The protocol parameter is fixed to nfs. portals: logical port IP address or DNS zone of the storage device. Only one IP address or DNS zone can be
			configured. You can log in to DeviceManager to obtain the logical port IP address. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is NFS. (For other series, see the corresponding operation description.)

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.5.4 Connecting to Enterprise Storage SAN over NVMe over RoCE

Perform this operation when you want to connect to enterprise storage SAN over NVMe over RoCE.

Prerequisites

- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected.

- The nyme-cli tool has been installed on all worker nodes of Kubernetes, and the tool version is 1.9 or later.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes of Kubernetes. For details, see 4.4 Configuring Host Multipathing.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-roce.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 4-9**.

Table 4-9 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single backend storage device, see Table 4-10.

Table 4-10 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.

Configuration Item	Format	Description	Remarks
parameters	Dictionary	Variable parameters in scenarios where the NVMe over RoCE protocol is used. This parameter is mandatory.	In scenarios where the NVMe over RoCE protocol is used, set the protocol parameter to a fixed value: roce .
			Set portals to the IP addresses of the logical ports when data protocol type of the storage device is NVMe over RoCE. Use commas (,) to separate the IP addresses.
			You can log in to DeviceManager to obtain the logical port IP address. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical
			Ports and obtain the IP address whose data protocol is NVMe over RoCE. (For other series, see the corresponding operation description.)

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.5.5 Connecting to Enterprise Storage SAN over NVMe over FC

Perform this operation when you want to connect to enterprise storage SAN over NVMe over FC.

Restrictions

To connect to enterprise storage SAN over NVMe over FC, ensure that no residual drive letter exists on the host. If any residual drive letter exists, clear the drive letter by referring to 10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters.

Prerequisites

- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes can communicate with the storage device to be connected over NVMe over FC.
- The nyme-cli tool has been installed on all worker nodes of Kubernetes, and the tool version is 1.9 or later.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes of Kubernetes. For details, see 4.4 Configuring Host Multipathing.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-fc-nvme.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-11.

Table 4-11 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-12.

Table 4-12 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the block storage service.
parameters	Dictionary	Variable parameters in scenarios where the NVMe over FC protocol is used. This parameter is mandatory.	In scenarios where the NVMe over FC protocol is used, set the protocol parameter to a fixed value: fc-nvme .

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n kube-system | grep huawei-csi-configmap
huawei-csi-configmap 1 5s
```

----End

4.6 Connecting to Distributed Storage

This section describes how to connect the huawei-csi plug-in to Huawei distributed storage.

When the same SAN storage is connected to Kubernetes, you cannot configure multiple data protocols (SCSI and iSCSI) on one worker node.

4.6.1 Connecting to Distributed Storage SAN over SCSI

Perform this operation when you want to connect to distributed storage SAN over SCSI.

Prerequisites

- The distributed storage VBS client has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-scsi.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-13.

} }

Table 4-13 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-14.

Table 4-14 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage SAN is connected, the value is fixed to fusionstorage-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.

Configuration Item	Format	Description	Remarks
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.
parameters	Dictionary	Variable parameters in scenarios where SCSI is used. This parameter is mandatory.	The protocol parameter is fixed to scsi. Set portals to a pair list of host names and VBS node IP addresses. The format is [{"hostname":"****"}], where hostname indicates the host name of a worker node and **** indicates the management IP address of a distributed storage block client (only IPv4 addresses are supported currently). If there are multiple worker nodes, configure them in dictionary format and separate them with commas (,). In the preceding example, hostname01 is the host name of a worker node in Kubernetes, and 192.168.125.21 is the management IP address of a VBS node after VBS is created for the worker node.

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.6.2 Connecting to Distributed Storage SAN over iSCSI

Perform this operation when you want to connect to distributed storage SAN over iSCSI.

Prerequisites

- An iSCSI client has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes.

Precautions

- The host name of a Kubernetes worker node consists of digits, letters, underscores (_), hyphens (-), periods (.), and colons (:), and must start with a digit, letter, or underscore (_). The name length cannot exceed 31 characters.
- Only FusionStorage 8.0.0 and later versions support iSCSI networking configuration.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-iscsi.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-15.

kind: ConfigMap apiVersion: v1 metadata: name: huawei-csi-configmap namespace: kube-system

Table 4-15 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-16.

Table 4-16 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage SAN is connected, the value is fixed to fusionstorage-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-).
			NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.
parameters	Dictionary	Variable parameters in scenarios where iSCSI is used. This parameter is mandatory.	In scenarios where iSCSI is used, set the protocol parameter to a fixed value: iscsi . Set the portals parameter to the iSCSI service IP addresses of the storage device. Use commas (,) to separate multiple them. You can log in to DeviceManager to obtain them. You can log in to DeviceManager to obtain the iSCSI service IP addresses. Take OceanStor Pacific series as an example. On DeviceManager, choose Resources > Access > Service Network . (For other series, see the corresponding operation description.)

Step 4 Run the **kubectl create** -**f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

----End

4.6.3 Connecting to Distributed Storage NAS over NFS

Perform this operation when you want to connect to distributed storage NAS over NFS.

Prerequisites

- An NFS client tool has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected.
- All worker nodes of Kubernetes communicate properly with the IP address of the NFS logical port on the storage device to be connected.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-nfs.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-17.

Table 4-17 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-18.

Table 4-18 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage NAS is connected, the value is fixed to fusionstorage-nas.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.

Format	Description	Remarks
List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.
Dictionary	Variable parameters in scenarios where NFS is used. This parameter is mandatory.	portals: logical port IP address the specified storage device. You can log in to DeviceManager to obtain it. Only one IP address can be configured. You can log in to DeviceManager to
		obtain the logical port IP address. Take OceanStor Pacific series as an example. On DeviceManager, choose Resources > Access > Service Network and click the name of a zone. On the page that is displayed, click the IP Address Pool tab. The IP Address/Mask column indicates the logical port IP address. For other series, see the corresponding
	List	List Name of a storage pool used on the storage device to be connected. This parameter is mandatory. Dictionary Variable parameters in scenarios where NFS is used. This parameter is

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.7 Starting huawei-csi Services

This section describes how to start huawei-csi services.

Precautions

An image may need to be downloaded during the procedure. Therefore, worker nodes in the Kubernetes cluster must be able to access external networks. In an intranet environment, obtain the image package in other ways and manually import it into all worker nodes. For details about the image package list, see **Table 4-21**.

Prerequisites

- You have obtained the user name and password of the storage device from the administrator.
- For details about the supported user types and requirements when different storage devices are connected, see **Table 4-19**.

■ NOTE

When vStore users are used, only NAS storage can be connected. For details about the supported storage models, see **3 Restrictions**.

Table 4-19 Details about users supported when different storage devices are connected to CSI

Storage Type	User Type	Role	Level	Туре
OceanStor	System user	Administrator	Administrator	Local user
V3/V5	vStore user	vStore administrator	Administrator	Local user
OceanStor Dorado V3	System user	Administrator	Administrator	Local user
OceanStor 6.1	System user	Administrator	N/A	Local user
OceanStor	System user	Administrator	N/A	Local user
Dorado 6.1.3	vStore user	vStore administrator	N/A	Local user
OceanStor Pacific series	System user	Administrator	N/A	Local user

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Copy the **secretGenerate** tool in the Kubernetes CSI component package to any directory on the master node. For details about the tool path, see **4.2 Uploading the Components in the Software Package**.
- **Step 3** Use an encryption tool to enter the user name and password of the storage device.
 - 1. Run the **chmod +x secretGenerate** command to grant the execute permission on the secretGenerate tool.

chmod +x secretGenerate

 Run the ./secretGenerate command to run the secretGenerate tool and enter the ID of the backend to be configured as prompted. If Configured is false, the backend is not configured. If Configured is true, the backend is configured.

```
# ./secretGenerate
Getting backend configuration information.....
Number Configured BackendName
                                       Urls
     false
              strage-backend [https://192.168.125.25:8088]
              strage-backend-02 [https://192.168.125.26:8088]
     false
2
3
     false
              strage-backend-03 [https://192.168.125.27:8088]
4
     false
              strage-backend-04 [https://192.168.125.28:8088]
              strage-backend-05
     false
                                 [https://192.168.125.29:28443]
              strage-backend-06 [https://192.168.125.30:28443]
     false
Please enter the backend number to configure (Enter 'exit' to exit):3
```

3. Enter the user name and password as prompted to create a **secret** object.

4. After the configuration is complete, enter **exit** to exit and save the configuration.

```
Please enter the backend number to configure (Enter 'exit' to exit): exit
Saving configuration. Please wait......
The configuration is saved successfully.
```

5. Run the **kubectl get secret -n kube-system | grep huawei-csi-secret** command to check whether the **secret** object is successfully created.

```
# kubectl -n kube-system get secret huawei-csi-secret
NAME TYPE DATA AGE
huawei-csi-secret Opaque 1 8d
```

Step 4 Check the Kubernetes version. Go to the corresponding directory according to the Kubernetes version. (The example directory is **deploy/v1.20-v1.23**.)

```
# kubectl version --short=true
Client Version: v1.22.1
Server Version: v1.22.1
```

Step 5 Run the following command to create the RBAC permission.

kubectl apply -f huawei-csi-rbac.yaml

- **Step 6** Start the controller service.
 - 1. Run the following command to deploy the snapshot-crd service. # kubectl apply -f huawei-csi-snapshot-crd.yaml
 - 2. Run the **vi** huawei-csi-controller.yaml command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify the following parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

◯ NOTE

In the image configuration item under huawei-csi-driver in the sample .yaml file, huawei-csi:*.** must be replaced with <Name>:<Version> of the created Huawei CSI image. For details, see view the image. The docker is used as an example. containers:

 name: huawei-csi-driver image: huawei-csi:2.2.16

3. Run the following command to start the controller service.

kubectl apply -f huawei-csi-controller.yaml

Step 7 Start the node service.

1. Run the **vi** *huawei-csi-node.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

□ NOTE

(Mandatory) In the image configuration item under huawei-csi-driver in the sample .yaml file, *huawei-csi:*.*** must be replaced with *<Name>:<Version>* of the created Huawei CSI image. For details, see view the image. The docker is used as an example.

containers:

- name: huawei-csi-driver image: huawei-csi:2.2.16
- (Optional) In the args section of huawei-csi-driver in the .yaml file, --volumeuse-multipath indicates that multipathing is enabled by default. The following shows how to change the value.

args:

- "--endpoint=/csi/csi.sock" "--containerized"
- "--driver-name=csi.huawei.com"
- "--volume-use-multipath=false"
- (Optional) In the args section of huawei-csi-driver in the .yaml file, --connectorthreads indicates the number of concurrent operations on the drive letter on the host. The value is an integer ranging from 1 to 10, and the default value is 4. To change the value, refer to the following.

args:

- "--endpoint=/csi/csi.sock"
- "--containerized"
- "--driver-name=csi.huawei.com"
- "--volume-use-multipath=true"
- "--connector-threads=5"
- (Optional) In the args section of huawei-csi-driver in the .yaml file, --scanvolume-timeout indicates the timeout of waiting for multipathing aggregation when DM-multipath is used on the host. The value is an integer ranging from 1 to 600, and the default value is 3. To change the value, refer to the following. args:
 - "--endpoint=/csi/csi.sock"
 - "--containerized"
 - "--driver-name=csi.huawei.com"
 - "--volume-use-multipath=true"
 - "--connector-threads=4"
 - "--scan-volume-timeout=3"
- (Optional) In the args section of huawei-csi-driver in the .yaml file, for enterprise storage, if --volume-use-multipath is set to true, you can configure the multipathing type according to the networking mode. For details, see Table 4-20. args:
 - "--endpoint=/csi/csi.sock"
 - "--containerized"
 - "--driver-name=csi.huawei.com"
 - "--connector-threads=4"
 - "--volume-use-multipath=true"
 - "--scsi-multipath-type=DM-multipath"
 - "--nvme-multipath-type=HW-UltraPath-NVMe"

Storage Protocol	Parameter	Description	Remarks
iSCSI/FC	scsi-multipath- type	The value can be: - DM-multipath - HW-UltraPath - HW- UltraPath- NVMe The default value is DM- multipath.	 DM- multipath: native multipathing software of the OS HW- UltraPath: Huawei UltraPath multipathing
NVMe over RoCE/NVMe over FC	nvme- multipath-type	The default value is HW-UltraPath-NVMe and only HW-UltraPath-NVMe can be configured.	software - HW- UltraPath- NVMe: Huawei UltraPath- NVMe multipathing software

Table 4-20 Parameters for configuring enterprise storage multipathing

2. Run the following command to start the node service.
kubectl apply -f huawei-csi-node.yaml

Step 8 After the huawei-csi services are deployed, run the **kubectl get pod -A | grep huawei** command to check whether the services are started.

kubectl get pod -A | grep huawei kube-system huawei-csi-controller-695b84b4d8-tg64l 7/7 **Running** 0 14s kube-system huawei-csi-node-g6f7z 3/3 **Running** 0 14s

Ⅲ NOTE

The Pod of *huawei-csi-controller-695b84b4d8-tg64l* has seven containers, including liveness-probe, csi-provisioner, csi-attacher, csi-resizer, csi-snapshotter, shapshot-controller, and huawei-csi-driver. Each container has its own image repository and functions. For details about the containers, see **Table 4-21**.

The Pod of *huawei-csi-node-g6f7z* has three containers, including liveness-probe, csi-node-driver-registrar, and huawei-csi-driver. Each container has its own image repository and functions. For details about the containers, see **Table 4-21**.

Table 4-21 Container description

Containe r Name	Container Image (v1.18 to v1.19)	Container Image (v1.20 to v1.23)	Feature Description	Rema rks
liveness- probe	k8s.gcr.io/sig- storage/ livenessprobe:v 2.5.0	k8s.gcr.io/sig- storage/ livenessprobe:v 2.5.0	Monitors the health status of CSI and reports it to Kubernetes so that Kubernetes can automatically detect CSI program problems and restart the Pod to rectify the problems.	View detail s
csi- provision er	k8s.gcr.io/sig- storage/csi- provisioner:v2.2 .2	k8s.gcr.io/sig- storage/csi- provisioner:v3.0 .0	 Calls the CSI Controller service to create a LUN or file system on the storage system as a PV and bind the PV to a PVC when creating a PVC. Calls the CSI Controller service to unbind a PV from a PVC and delete the LUN or file system corresponding to the PV when deleting a PVC. 	View detail s
csi- attacher	k8s.gcr.io/sig- storage/csi- attacher:v3.3.0	k8s.gcr.io/sig- storage/csi- attacher:v3.3.0	Calls the CSI Controller service to perform the "Publish/Unpublish Volume" operation when creating or deleting a Pod.	View detail s
csi-resizer	k8s.gcr.io/sig- storage/csi- resizer:v1.3.0	k8s.gcr.io/sig- storage/csi- resizer:v1.3.0	Calls CSI to provide more storage space for a PVC when expanding the capacity of the PVC.	View detail s
csi- snapshott er	k8s.gcr.io/sig- storage/csi- snapshotter:v3. 0.3	k8s.gcr.io/sig- storage/csi- snapshotter:v4. 2.1	Calls CSI to create or delete a snapshot on the storage system when creating or deleting a VolumeSnapshot.	View detail s

Containe r Name	Container Image (v1.18 to v1.19)	Container Image (v1.20 to v1.23)	Feature Description	Rema rks
shapshot- controller	k8s.gcr.io/sig- storage/ snapshot- controller:v3.0.	k8s.gcr.io/sig- storage/csi- snapshotter:v4. 2.1	Listens to the VolumeSnapshot and VolumeSnapshotContent objects in the Kubernetes API and triggers csi-snapshotter to create a snapshot on the storage system when creating or deleting a VolumeSnapshot.	View detail s
csi-node- driver- registrar	k8s.gcr.io/sig- storage/csi- node-driver- registrar:v2.3.0	k8s.gcr.io/sig- storage/csi- node-driver- registrar:v2.3.0	Obtains CSI information and registers a node with kubelet using the plugin registration mechanism of kubelet so that Kubernetes can detect the connection between the node and Huawei storage.	View detail s
huawei- csi-driver	The name and tag of huaweicsi-driver are specified in 4.3 Creating a Huawei CSI Image.	The name and tag of huaweicsi-driver are specified in 4.3 Creating a Huawei CSI Image.	Connects to the Kubernetes platform to provide Huawei storage (centralized or distributed storage) resources for containers.	Versio n mapp ings Restri ctions

----End

5 Upgrade Operations

Prerequisites

NOTICE

- During the upgrade, CSI cannot be used to deliver new resources.
- During the upgrade, do not uninstall the snapshot-dependent component service.
- The CSI upgrade does not affect delivered resources such as PVCs, snapshots, and Pods.
- 5.1 Uninstalling Original CSI
- 5.2 Installing New CSI

5.1 Uninstalling Original CSI

Perform this operation when you want to uninstall CSI.

Preparations

Before uninstalling CSI, run the **kubectl get configmap huawei-csi-configmap -n kube-system -o yaml >> huawei-csi-configmap.yaml.bak** command to back up the content of the **huawei-csi-configmap** file. (During the CSI upgrade, the **backends** parameter in **huawei-csi-configmap.yaml** must be the same as the existing value of **configmap**.)

5.1.1 Uninstalling the huawei-csi-node Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command:

kubectl delete daemonset huawei-csi-node -n kube-system

Step 3 Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get daemonset huawei-csi-node -n kube-system

----End

5.1.2 Uninstalling the huawei-csi-controller Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following commands in sequence.

 # kubectl delete deployment huawei-csi-controller -n kube-system
- **Step 3** Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get deployment huawei-csi-controller -n kube-system

----End

5.1.3 Deleting the huawei-csi-configmap Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command.

kubectl delete configmap huawei-csi-configmap -n kube-system

Step 3 Run the following command to check whether the object is successfully deleted. If **NotFound** is displayed, the object is successfully deleted.

kubectl get configmap huawei-csi-configmap -n kube-system

----End

5.1.4 Deleting the huawei-csi-secret Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command.

kubectl delete secret huawei-csi-secret -n kube-system

Step 3 Run the following command to check whether the object is successfully deleted. If **NotFound** is displayed, the object is successfully deleted.

kubectl get secret huawei-csi-secret -n kube-system

----End

5.1.5 Deleting the RBAC Permission

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Delete the RBAC permission.
 - If the huawei-csi version is later than 2.2.15, run the following command to delete the permission.
 - # kubectl -n kube-system -l provisioner=csi.huawei.com delete ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding
 - If the huawei-csi version is 2.2.15 or earlier, perform the following operations to delete the permission.
 - a. Run the following command to create a file named **delete-huawei-csi-rbac.sh**.

```
# cat <<EOF > delete-huawei-csi-rbac.sh
kubectl delete ServiceAccount huawei-csi-controller -n kube-system
kubectl delete ServiceAccount huawei-csi-node -n kube-system
kubectl delete ClusterRole huawei-csi-attacher-runner -n kube-system
kubectl delete ClusterRole huawei-csi-driver-registrar-runner -n kube-system
kubectl delete ClusterRole huawei-csi-provisioner-runner -n kube-system
kubectl delete ClusterRole huawei-csi-resizer-runner -n kube-system
kubectl delete ClusterRole huawei-csi-snapshotter-runner -n kube-system
kubectl delete ClusterRole snapshot-controller-runner -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-attacher-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-driver-registrar-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-provisioner-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-resizer-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-snapshotter-role -n kube-system
kubectl delete ClusterRoleBinding snapshot-controller-role -n kube-system
kubectl delete Role huawei-csi-resizer-cfg -n kube-system
kubectl delete Role huawei-csi-snapshotter-leaderelection -n kube-system
kubectl delete Role snapshot-controller-leaderelection -n kube-system
kubectl delete RoleBinding huawei-csi-resizer-role-cfg -n kube-system
kubectl delete RoleBinding huawei-csi-snapshotter-leaderelection -n kube-system
kubectl delete RoleBinding snapshot-controller-leaderelection -n kube-system
```

Run the following command to delete the RBAC permission. If the **NotFound** error is reported, ignore it.
 # sh delete-huawei-csi-rbac.sh

Step 3 Check whether the RBAC permission has been deleted.

- If the huawei-csi version is later than 2.2.15, run the following command. If
 No resources found is displayed, the permission is successfully deleted.
 # kubectl -n kube-system -l provisioner=csi.huawei.com get
 ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding
- If the huawei-csi version is 2.2.15 or earlier, perform the following operations to check whether the RBAC permission is successfully deleted.
 - Run the following command to create a file named check-huawei-csirbac.sh.

```
# cat <<EOF > check-huawei-csi-rbac.sh
kubectl get ServiceAccount -n kube-system | grep huawei-csi
kubectl get ClusterRole -n kube-system | grep huawei-csi
kubectl get ClusterRoleBinding -n kube-system | grep huawei-csi
kubectl get Role -n kube-system | grep huawei-csi
kubectl get RoleBinding -n kube-system | grep huawei-csi
kubectl get RoleBinding -n kube-system | grep huawei-csi
kubectl get ClusterRole snapshot-controller-runner -n kube-system --ignore-not-found=true
```

kubectl get ClusterRoleBinding snapshot-controller-role -n kube-system --ignore-not-found=true kubectl get Role snapshot-controller-leaderelection -n kube-system --ignore-not-found=true kubectl get RoleBinding snapshot-controller-leaderelection -n kube-system --ignore-not-found=true

FOI

b. Run the following command. If no command output is displayed, the RBAC permission has been successfully deleted.

sh check-huawei-csi-rbac.sh

----End

5.1.6 Deleting the Image of the Earlier Version

To delete the **huawei-csi** image from the cluster, you need to perform the deletion operation on all worker nodes.

To delete the image from a single node, perform the following steps.

Prerequisites

The container service that depends on the image has been stopped. Otherwise, the image cannot be deleted.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node through the management IP address.
- **Step 2** Run the following command to view all existing versions.
 - If docker is used, run the **docker image ls | grep huawei-csi** command.

docker image ls | grep huawei-csi
REPOSITORY TAG IMAGE ID CREATED SIZE
huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
huawei-csi 2.2.16 14b854dba227 2 weeks ago 79.6MB

• If containerd is used, run the **crictl image ls | grep huawei-csi** command.

crictl image ls | grep huawei-csi
REPOSITORY TAG IMAGE ID CREATED SIZE
docker.io/library/huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
docker.io/library/huawei-csi 2.2.16 14b854dba227 2 weeks ago 79.6MB

- **Step 3** Run the following command to delete the image of the earlier version:
 - If docker is used, run the **docker rmi** *<REPOSITORY>:<TAG>* command. # docker rmi huawei-csi:2.2.15
 - If containerd is used, run the crictl rmi <REPOSITORY>:<TAG> command.
 # crictl rmi huawei-csi:2.2.15
- **Step 4** Run the following command again to check whether the image is successfully deleted. If the target version is not displayed, the image of the version is successfully deleted.
 - If docker is used, run the docker image ls | grep huawei-csi command.
 # docker image ls | grep huawei-csi
 huawei-csi
 2.2.16
 14b854dba227
 10 minutes ago
 80MB
 - If containerd is used, run the crictl image ls | grep huawei-csi command.
 # crictl image ls | grep huawei-csi docker.io/library/huawei-csi
 2.2.16
 14b854dba2273
 93.1MB

----End

5.2 Installing New CSI

After the uninstallation is complete, you need to reinstall the CSI.

Prerequisites

The **huawei-csi-configmap.yaml** file of the original CSI has been backed up.

Precautions

If the template of **huawei-csi-configmap.yaml** has changed, ensure that the following parameter settings are the same as those before the upgrade. Otherwise, huawei-csi services cannot be started and created resources cannot be managed.

- The values of storage, name, and pools must be the same as those in the huawei-csi-configmap.yaml.bak file backed up in Prerequisites in 5.1 Uninstalling Original CSI.
- For details about urls and parameters, see the huawei-csi-configmap.yaml.bak file backed up in Prerequisites in 5.1 Uninstalling Original CSI and set them based on the huawei-csi-configmap.yaml template of the current version. For details about the template, see 4.5 Connecting to Enterprise Storage and 4.6 Connecting to Distributed Storage. The following command output is only an example.

Procedure

- **Step 1** Obtain the CSI software package of the new version. For details, see **4.1 Obtaining the Software Package**.
- **Step 2** Create a CSI image of the new version. For details, see **4.3 Creating a Huawei CSI Image**.
- **Step 3** Create **huawei-csi-configmap**. For details, see **4.5 Connecting to Enterprise Storage** or **4.6 Connecting to Distributed Storage**.
- **Step 4** Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.

----End

6 Instructions for Use

This chapter describes how to manage StorageClasses, PVCs, Pods, and snapshots after connecting Kubernetes to Huawei storage.

- 6.1 (Conditionally Mandatory) Managing a StorageClass
- 6.2 (Conditionally Mandatory) Managing a PV
- 6.3 Managing a PVC
- 6.4 Managing a Pod
- 6.5 (Optional) Managing a Snapshot

6.1 (Conditionally Mandatory) Managing a StorageClass

A PV can be provisioned in either of the following modes: static provisioning and dynamic provisioning. In static provisioning mode, storage resources (LUNs/shares) are created on storage devices in advance and then statically provisioned to Kubernetes. In dynamic provisioning mode, CSI automatically creates storage resources (LUNs/shares).

Perform this operation when you want to use the dynamic PV function.

6.1.1 Creating a StorageClass

A StorageClass is a set of capabilities that can be selected when you apply for block storage resources. Kubernetes cluster users can create PVCs based on a StorageClass.

6.1.1.1 Creating a LUN StorageClass

This section describes how to create a LUN StorageClass.

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Run the **vi** *StorageClass.yaml* command to create a file named *StorageClass.yaml*. # vi StorageClass.yaml
- **Step 3** Press I or Insert to enter the editing mode and enter the following information in the *StorageClass.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *StorageClass.yaml* file. You can also refer to the **examples/lun-sc-for-csi-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-1**.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: "mysc" provisioner: "csi.huawei.com" parameters: volumeType: "lun" allocType: "thin"

Table 6-1 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a StorageClass object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
provisioner	provisioner identifier.	The value is fixed to csi.huawei.com.
parameters.volumeType	Type of the volume to be created.	The value is fixed to lun .
parameters.allocType	Allocation type of the volume to be created.	This parameter is optional. The value can be thin or thick , and the default value is thin .
parameters.cloneSpeed	Clone speed.	This parameter is optional. The value ranges from 1 to 4 and the default value is 3. 4 indicates the highest speed. This parameter is available when you clone a PVC or create a PVC using a snapshot. For details, see 6.3.3 (Optional) Cloning a PVC or 6.3.4 (Optional) Creating a PVC Using a Snapshot.

Parameter	Description	Remarks
parameters.fsType	File system type.	This parameter is optional. The value can be ext2, ext3, ext4, or xfs, and the default value is ext4.
		NOTICE CSI does not verify whether the value of the fsType parameter is valid. Ensure that the value is correct.

- **Step 4** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorageClass.yaml
- **Step 5** Run the following command to view the information about the created StorageClass.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION AGE
mysc csi.huawei.com Delete Immediate false 87s
```

----End

6.1.1.2 Creating a File System StorageClass

This section describes how to create a file system StorageClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to create a file named *StorageClass.yaml*. # vi StorageClass.yaml
- **Step 3** Press I or Insert to enter the editing mode and enter the following information in the *StorageClass.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *StorageClass.yaml* file. You can also refer to the **examples/fs-sc-for-csi-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-2**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "mysc"
provisioner: "csi.huawei.com"
parameters:
volumeType: "fs"
allocType: "thin"
authClient: "*"
```

Table 6-2 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a StorageClass object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
provisioner	provisioner identifier.	The value is fixed to csi.huawei.com.
parameters.volu meType	Type of the volume to be created.	The value is fixed to fs .
parameters.auth Client	Client that can access the volume.	This parameter is mandatory. OceanStor Dorado 6.x is used as an example. You can enter the client host name (a fully qualified domain name (FQDN) is recommended), client IP address, or client IP address segment, or use an asterisk (*) to represent all client IP addresses. The IP addresses can be IPv4
		addresses, IPv6 addresses, or a combination of IPv4 and IPv6 addresses.
		You can enter multiple host names, IP addresses, or IP address segments and separate them with semicolons (;) or spaces or by pressing Enter .
		Example: 192.168.0.10;192.168.0.0/24;*
parameters.alloc Type	Allocation type of the volume to be created.	This parameter is optional. The value can be thin or thick , and the default value is thin .
parameters.clone Speed	Clone speed.	This parameter is optional. The value ranges from 1 to 4 and the default value is 3 . 4 indicates the highest speed. This parameter is available when you clone a PVC or create a PVC using a snapshot. For details, see 6.3.3 (Optional) Cloning a PVC or 6.3.4 (Optional) Creating a PVC Using a Snapshot.

Parameter	Description	Remarks
parameters.fsTyp e	File system type.	This parameter is optional. The value can be ext2, ext3, ext4, or xfs, and the default value is ext4.
		NOTICE CSI does not verify whether the value of the fsType parameter is valid. Ensure that the value is correct.

- **Step 4** Run the following command to create a StorageClass based on the .yaml file.
 - # kubectl create -f StorageClass.yaml
- **Step 5** Run the following command to view the information about the created StorageClass.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION AGE
mysc csi.huawei.com Delete Immediate false 34s
```

----End

6.1.2 Deleting a StorageClass

This section describes how to delete a StorageClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query StorageClasses in the cluster.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION
AGE
huawei-nas csi.huawei.com Delete Immediate false 3s
mysc csi.huawei.com Delete Immediate false 16s
```

Step 3 Run the following command to delete a StorageClass. For example, delete the StorageClass named *mysc*.

```
# kubectl delete sc mysc
storageclass.storage.k8s.io "mysc" deleted
```

Step 4 Run the following command to query StorageClasses in the cluster. If the command output does not contain the name of the StorageClass you want to delete, it is successfully deleted.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION
AGE
huawei-nas csi.huawei.com Delete Immediate false 3s
```

----End

6.2 (Conditionally Mandatory) Managing a PV

A PV can be provisioned in either of the following modes: static provisioning and dynamic provisioning. In static provisioning mode, storage resources (LUNs/shares)

are created on storage devices in advance and then statically provisioned to Kubernetes. In dynamic provisioning mode, CSI automatically creates storage resources (LUNs/shares).

Perform this operation when you want to use the static PV function.

6.2.1 Creating a PV

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *PersistentVolume.yaml* command to create a file named *PersistentVolume.yaml*.

vi PersistentVolume.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *PersistentVolume.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *PersistentVolume.yaml* file. You can also refer to the **examples/static-pv-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-3**.

```
kind: PersistentVolume
apiVersion: v1
metadata:
name: mypv
spec:
volumeMode: Block
storageClassName: ""
accessModes:
- ReadWriteOnce
csi:
driver: csi.huawei.com
volumeHandle: <backendName>.<volume-name>
fsType: <string>
capacity:
storage: 100Gi
```

Table 6-3 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a PV object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.

Parameter	Description	Remarks
spec.volumeMode	Volume mode.	This parameter is optional. The value can be Filesystem or Block. The default value is Filesystem. This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.
spec.storageClassNa me	Name of the StorageClass object.	This parameter is mandatory. Set it to an empty string, that is, enter "".
spec.persistentVolu meReclaimPolicy	Volume reclamation policy.	This parameter is optional. The value can be Retain (manual reclamation) or Delete (deleting associated storage resources). The default value is Retain . For details, see Table 6-5.
spec.accessModes	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce and ReadOnlyMany.
		If the volume mode is Block, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany. When ReadWriteMany is set, the Pod service must ensure data consistency. File system volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany.
spec.csi.driver	CSI driver name.	The value is fixed to csi.huawei.com.

Parameter	Description	Remarks
spec.csi.volumeHan dle	Unique identifier of a storage resource.	It consists of two parts: • <backendname>: indicates the name of each backend in configmap. You can run the following command to obtain configmap: kubectl get configmap huawei-csi-configmap -n kube-system -o yaml • <volume-name>: indicates the name of a resource (LUN/file system) on the storage. You can obtain the value from DeviceManager.</volume-name></backendname>
spec.csi.fsType	File system type.	This parameter is optional. The value can be ext2, ext3, ext4, or xfs, and the default value is ext4. This parameter is valid only when volumeMode is set to Filesystem. NOTICE CSI does not verify whether the value of the fsType parameter is valid. Ensure that the value is correct.
spec.capacity.storag e	Volume size.	When creating a PV, ensure that its capacity is the same as that of the corresponding resource on the storage. Kubernetes will not invoke CSI to check whether the value of this parameter is correct. Therefore, the PV can be successfully created even if its capacity is inconsistent with that of the corresponding resource on the storage.

Step 4 Run the following command to create a PV based on the .yaml file.

kubectl create -f PersistentVolume.yaml

Step 5 After a period of time, run the following command to view the information about the created PV.

kubectl get pv
NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS
REASON AGE
mypv 100Gi RWO Retain Available 4s

Table 6-4 PV status description

Status	Description
Available	Not bound to a PVC.
Bound	Bound to a PVC.
Released	A PVC has been deleted, but its resource has not been reclaimed by the cluster.
Failed	A volume fails to be automatically reclaimed.
Terminating	The PV object has been marked as "deleted" by the Kubernetes controller, but the PV cannot be deleted because it is being used.

Table 6-5 Reclamation policy description

Reclamation Policy	Description
Delete	After the PVC object is deleted, the PV still exists and the corresponding data volume is regarded as "Failed".
	After the PV is deleted, the associated storage assets are also removed from the back-end storage device.
Retain	You can manually reclaim resources. After the PVC object is deleted, the PV still exists and the corresponding data volume is regarded as "Released". However, the PV cannot be used for another PVC because the data of the previous PVC remains on the PV.
	After the PV is deleted, the associated storage assets on the back-end storage device still exist.

----End

6.2.2 Deleting a PV

This section describes how to delete a PV.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query PVs in the cluster.

kubectl get pv
NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS
REASON AGE
mypv 100Gi RWO Retain Available 11m

Step 3 Run the following command to delete a PV. For example, delete the PV named *mvpv*.

kubectl delete pv *mypv* persistentvolume "mypv" deleted

Step 4 Run the following command to query PVs in the cluster. If the command output does not contain the name of the PV you want to delete, it is successfully deleted.

kubectl get pv No resources found in default namespace.

----End

6.3 Managing a PVC

6.3.1 Creating a PVC

Perform this operation when you want to create a PVC.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *PersistentVolumeClaim.yaml* command to create a file named *PersistentVolumeClaim.yaml*.

vi PersistentVolumeClaim.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *PersistentVolumeClaim.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *PersistentVolumeClaim.yaml* file. You can also refer to the **examples/pvc-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-6**.

kind: PersistentVolumeClaim apiVersion: v1 metadata: name: "mypvc" spec: accessModes: - ReadWriteOnce volumeMode: Filesystem volumeName: mypv storageClassName: "mysc" resources: requests: storage: 100Gi

Table 6-6 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a PVC object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.volumeMode	Volume mode.	This parameter is optional. The value can be Filesystem or Block . The default value is Filesystem . This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.
spec.volumeName	Name of the PV object.	This parameter is mandatory when a PVC is created statically.
spec.storageClass Name	Name of the StorageClass object.	 When creating a PVC dynamically, enter the name of the StorageClass object created in 6.1 (Conditionally Mandatory) Managing a StorageClass. When creating a PVC
		statically, set it to an empty string, that is, enter "".

Parameter	Description	Remarks
spec.resources.req uests.storage	Size of the volume to be created.	The value format is ***Gi. The unit is GiB. • The PVC capacity depends on storage specifications and host specifications. The following uses the connection between OceanStor Dorado 6.1.2/OceanStor Pacific series 8.1.0 and CentOS 7 as an example. See Table 6-7 and Table 6-8.
		 For other storage devices and hosts, check the specifications according to the value of VolumeType in StorageClass.
		- If the value of volumeType is lun, refer to the storage specifications. For details, see https:// info.support.huawei.com/ storage/spec/#/home. In addition, refer to the host connectivity guide at https:// support.huawei.com/ enterprise/en/doc/ EDOC1100113070/ e067543b.
		 If the value of volumeType is fs, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home.
		If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications.
		When a PVC is created using a static PV and the PVC capacity is smaller than the capacity of the bound PV, the PVC capacity is set to the capacity of the bound PV. If the PVC capacity is greater than the capacity of the

Parameter	Description	Remarks
		bound PV, the PVC cannot be created.
spec.accessModes	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce and ReadOnlyMany.
		If the volume mode is Block, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany. When ReadWriteMany is set, the Pod service must ensure data consistency.
		File system volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany.

Table 6-7 PVC capacity specifications (ext4)

volumeTyp e	Storage Type	Storage Specificati ons	ext4 Specification s	CSI Specification s
lun	OceanStor Dorado 6.1.2	512 Ki to 256 Ti	50 Ti	512 Ki to 50 Ti
	OceanStor Pacific series 8.1.0	64 Mi to 512 Ti	50 Ti	64 Mi to 50 Ti
fs	OceanStor Dorado 6.1.2	1 Gi to 32 Pi	N/A	1 Gi to 32 Pi
	OceanStor Pacific series 8.1.0	1 Ki to 256 Pi	N/A	1 Ki to 256 Pi

Table 6-8 PVC capacity specifications (xfs)

volumeTyp e	Storage Type	Storage Specificati ons	xfs Specification s	CSI Specificatio ns
lun	OceanStor Dorado 6.1.2	512 Ki to 256 Ti	500 Ti	512 Ki to 500 Ti

volumeTyp e	Storage Type	Storage Specificati ons	xfs Specification s	CSI Specificatio ns
	OceanStor Pacific series 8.1.0	64 Mi to 512 Ti	500 Ti	64 Mi to 500 Ti
fs	OceanStor Dorado 6.1.2	1 Gi to 32 Pi	N/A	1 Gi to 32 Pi
	OceanStor Pacific series 8.1.0	1 Ki to 256 Pi	N/A	1 Ki to 256 Pi

Step 4 Run the following command to create a PVC based on the .yaml file.

kubectl create -f PersistentVolumeClaim.yaml

Step 5 After a period of time, run the following command to view the information about the created PVC.

```
# kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWX mysc 12s
```

□ NOTE

After the PVC is created, if the PVC is in the **Pending** state, see **10.6 When a PVC Is Created**, **the PVC Is in the Pending State**.

----End

6.3.2 (Optional) Expanding the Capacity of a PVC

This section describes how to expand the capacity of a PVC.

NOTICE

If the Kubernetes version is 1.18 and the value of **volumeMode** in the PVC is **Block**, the PVC does not support capacity expansion.

Prerequisites

- A PVC has been created, and the backend where the PVC is located supports capacity expansion. For details about the storage devices that support capacity expansion, see 3 Restrictions.
- The Kubernetes version is later than 1.16. You can run the **kubectl get node** command to view the version.
- The huawei-csi services are running properly.

 "Indicate and Advantage are running properly."

```
# kubectl get pod -A | grep huawei
kube-system huawei-csi-controller-fd5f97768-qlldc 7/7 Running 0 16s
kube-system huawei-csi-node-25txd 3/3 Running 0 15s
```

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

Step 2 Check whether the csi-resizer service is enabled for huawei-csi-controller.

kubectl describe deploy huawei-csi-controller -n kube-system | grep csi-resizer csi-resizer:
Image: k8s.gcr.io/sig-storage/csi-resizer:v1.3.0

- If the preceding information is displayed, the csi-resizer service is enabled. In this case, go to **Step 3**.
- Otherwise, the csi-resizer service is not enabled. In this case, you need to upgrade huawei-csi to 2.2.15 or later.
- **Step 3** Run the **kubectl get pvc** *mypvc* command to query the StorageClass name of the PVC. In the preceding command, *mypvc* indicates the name of the PVC to be expanded.

```
# kubectl get pvc mypvc
NAME STATUS VOLUME CAPACITY ACCESS MODES

STORAGECLASS AGE

mypvc Bound pvc-3383be36-537c-4cb1-8f32-a415fa6ba384 2Gi RWX

mysc 145m
```

Step 4 Run the **kubectl get sc** *mysc* command to check the StorageClass supports capacity expansion. In the preceding command, *mysc* indicates the name of the StorageClass to be queried.

```
# kubectl get sc mysc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE

ALLOWVOLUMEEXPANSION AGE

mysc csi.huawei.com Delete Immediate true 172m
```

- If the value of **ALLOWVOLUMEEXPANSION** is **false**, the current StorageClass does not support capacity expansion. In this case, go to **Step 5**.
- If the value of **ALLOWVOLUMEEXPANSION** is **true**, the current StorageClass supports capacity expansion. In this case, go to **Step 6**.
- **Step 5** (Optional) Run the following command to change the value of **allowVolumeExpansion** to **true**. In the preceding command, *mysc* indicates the name of the StorageClass to be modified.

kubectl patch sc mysc --patch '{"allowVolumeExpansion":true}'

Step 6 Run the following command to expand the capacity.

```
# kubectl patch pvc mypvc -p '{"spec":{"resources":{"requests":{"storage":"120Gi"}}}}'
```

In the preceding command, *mypvc* indicates the name of the PVC to be expanded, and *120Gi* indicates the capacity after expansion. Change the values based on the site requirements.

MOTE

- The PVC capacity depends on storage specifications and host specifications. The following uses the connection between OceanStor Dorado 6.1.2 or OceanStor Pacific series 8.1.0 and CentOS 7 as an example. For details, see **Table 6-9**.
- For other storage devices and hosts, check the specifications according to the value of VolumeType in 6.1.1 Creating a StorageClass.
 - If the value of volumeType is lun, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home. In addition, refer to the host connectivity guide at https://support.huawei.com/enterprise/en/doc/EDOC1100113070/e067543b.
 - If the value of volumeType is fs, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home.
- If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications.

1 Ki to 256 Pi

volumeTyp **Storage Type** Storage ext4 Specificati Specification **Specification** e ons OceanStor Dorado 512 Ki to 50 512 Ki to 50 Ti lun 6.1.2 256 Ti OceanStor Pacific 50 Ti 64 Mi to 50 64 Mi to series 8.1.0 512 Ti Ti OceanStor Dorado 1 Gi to 32 1 Gi to 32 Pi fs N/A 6.1.2 Ρi

1 Ki to 256

Ρi

N/A

Table 6-9 PVC capacity specifications (ext4)

Table 6-10 PVC capacity specifications (xfs)

series 8.1.0

OceanStor Pacific

volumeTyp e	Storage Type	Storage Specificati ons	xfs Specification s	CSI Specificatio ns
lun	OceanStor Dorado 6.1.2	512 Ki to 256 Ti	500 Ti	512 Ki to 500 Ti
	OceanStor Pacific series 8.1.0	64 Mi to 512 Ti	500 Ti	64 Mi to 500 Ti
fs	OceanStor Dorado 6.1.2	1 Gi to 32 Pi	N/A	1 Gi to 32 Pi
	OceanStor Pacific series 8.1.0	1 Ki to 256 Pi	N/A	1 Ki to 256 Pi

Step 7 Run the following command to check whether the capacity changes.

kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 120Gi RWX mysc 24s

----End

6.3.3 (Optional) Cloning a PVC

Perform this operation when you want to clone an existing PVC on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *clone.yaml* command to create a file named *clone.yaml*.

 # vi clone.yaml

Step 3 Configure the *clone.yaml* file. The **examples**/*<Kubernetes version*>/clone.yaml template file is as follows. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-11**.

kind: PersistentVolumeClaim apiVersion: v1 metadata: name: myclone spec: storageClassName: mysc dataSource: name: mypvc kind: PersistentVolumeClaim volumeMode: Filesystem accessModes: - ReadWriteOnce resources: requests: storage: 2Gi

Table 6-11 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a new PVC object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.storageClass Name	Name of the StorageClass object.	Enter the name of the StorageClass object created in 6.1 (Conditionally Mandatory) Managing a StorageClass. The value must be the same as the name of the StorageClass in dataSource.
spec.dataSource. name	Name of the source PVC object.	-
spec.volumeMod e	Volume mode.	This parameter is optional. The value can be Filesystem or Block. The default value is Filesystem. The value must be the same as the volumeMode value of the source PVC object. This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.

Parameter	Description	Remarks
spec.accessMode s	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce and ReadOnlyMany.
		If the volume mode is Block , LUN volumes support ReadWriteOnce , ReadOnlyMany , and ReadWriteMany . When ReadWriteMany is set, the Pod service must ensure data consistency.
		File system volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany.
spec.resources.re quests.storage	Size of the volume to be created.	The value must be greater than or equal to the size of the source PVC. The value format is ***Gi. The unit is GiB.

Step 4 Run the following command to create a PVC based on the .yaml file.

kubectl create -f clone.yaml

----End

6.3.4 (Optional) Creating a PVC Using a Snapshot

Perform this operation when you want to create a PVC for an existing snapshot on Kubernetes

Prerequisites

A snapshot has been created. For details, see **6.5 (Optional) Managing a Snapshot**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** restore.yaml command to create a file named restore.yaml. # vi restore.yaml
- **Step 3** Configure the *restore.yaml* file. The **examples**/*<Kubernetes version*>/restore.yaml template file is as follows. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 6-12.

kind: PersistentVolumeClaim apiVersion: v1 metadata:

```
name: ***
spec:
storageClassName: ***
dataSource:
name: ***
kind: VolumeSnapshot
apiGroup: snapshot.storage.k8s.io
volumeMode: Filesystem
accessModes:
- ReadWriteOnce
resources:
requests:
storage: ***Gi
```

Table 6-12 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a new PVC object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.storageClass Name	Name of the StorageClass object.	Enter the name of the StorageClass object created in 6.1 (Conditionally Mandatory) Managing a StorageClass. The value must be the same as the name of the StorageClass of the original PVC in dataSource.
spec.dataSource.n ame	Name of the source VolumeSnapshot object.	-
spec.volumeMode	Volume mode.	This parameter is optional. The value can be Filesystem or Block. The default value is Filesystem. The value must be the same as the volumeMode value of the source PVC object. This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.

Parameter	Description	Remarks
spec.accessModes	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce and ReadOnlyMany.
		If the volume mode is Block, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany. When ReadWriteMany is set, the Pod service must ensure data consistency.
		File system volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteMany.
spec.resources.req uests.storage	Size of the volume to be created.	The value must be greater than or equal to the size of the source VolumeSnapshot. The value format is *** Gi . The unit is GiB.

Step 4 Run the following command to create a PVC based on the .yaml file.

kubectl create -f restore.yaml

----End

6.3.5 Deleting a PVC

This section describes how to delete a PVC.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query PVCs in the cluster.

kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWX mysc 12s

□ NOTE

Before deleting a PVC, if the PVC is in the **Pending** state, you are not advised to directly delete the PVC. To delete the PVC, see **10.7 Before a PVC Is Deleted, the PVC Is in the Pending State**.

Step 3 Run the following command to delete a PVC. For example, delete the PVC named *mypvc*.

kubectl delete pvc *mypvc* persistentvolumeclaim "*mypvc*" deleted

Step 4 Run the following command to query PVCs in the cluster. If the command output does not contain the name of the PVC you want to delete, it is successfully deleted.

```
# kubectl get pvc
No resources found in default namespace.
```

----End

6.4 Managing a Pod

6.4.1 Creating a Pod

A Pod is an original storage pool or storage function set. It function as the container of virtual volumes, which means only the storage container allocates storage space to virtual volumes. This operation enables you to quickly obtain specified storage resources.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *pod.yaml* command to create a file named *pod.yaml*.

 # vi pod.yaml
- **Step 3** Press **I** or **Insert** to enter the editing mode and enter the following information in the *pod.yaml* file. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - If **volumeMode** is set to **Filesystem** in **Step 3**, the following shows a template of the *pod.yaml* file. You can also refer to the **examples/pod-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-13**.

```
kind: Pod
apiVersion: v1
metadata:
name: "mypod"
spec:
containers:
- name: "mycontainer"
image: "***"
volumeMounts:
- name: mypv
mountPath: "/mnt/path/in/container"
volumes:
- name: mypv
persistentVolumeClaim:
claimName: "mypvc"
```

Table 6-13 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a Pod object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.nam e	User-defined container name.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.imag e	Container image.	Set this parameter based on the site requirements.
spec.containers.imag e.volumeMounts.mo untPath	Volume mount path in the container.	-
spec.volumes.persist entVolumeClaim.clai mName	Name of the PVC object.	Enter the name of the PVC object created in 6.3.1 Creating a PVC.

• If **volumeMode** is set to **Block** in **Step 3**, the following shows a template of the *pod.yaml* file. You can also refer to the **examples/pod-rbd-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-14**. kind: Pod

```
apiVersion: v1
metadata:
name: "mypod"
spec:
containers:
    - name: "mycontainer"
    image: "***"
    volumeDevices:
    - name: mypv
    devicePath: "/dev/xvda"
    volumes:
    - name: mypv
    persistentVolumeClaim:
    claimName: "mypvc"
```

Parameter	Description	Remarks
metadata.name	User-defined name of a Pod object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.nam e	User-defined container name.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.imag e	Container image.	Set this parameter based on the site requirements.
spec.containers.imag e.volumeDevices.devi cePath	Volume device path in the container.	-
spec.volumes.persiste ntVolumeClaim.claim Name	Name of the PVC object.	Enter the name of the PVC object created in 6.3.1 Creating a PVC .

Table 6-14 Parameter description

Step 4 Run the following command to create a Pod based on the .yaml file. # kubectl create -f pod.yaml

Step 5 Run the following command to view the information about the created Pod.

kubectl get pod NAME READY STATUS RESTARTS AGE mypod 1/1 Running 0 37s

----End

6.4.2 Deleting a Pod

This section describes how to delete a Pod.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query Pods in the cluster.

kubectl get pod NAME READY STATUS RESTARTS AGE mypod 1/1 Running 0 14h

Step 3 Run the following command to delete a Pod. For example, delete the Pod named *mypod*.

kubectl delete pod *mypod* pod "*mypod*" deleted

Step 4 Run the following command to query Pods in the cluster. If the command output does not contain the name of the Pod you want to delete, it is successfully deleted.

kubectl get pod No resources found in default namespace.

----End

6.5 (Optional) Managing a Snapshot

CSI supports snapshot v1beta1 for Kubernetes v1.17 to v1.19, and completes GA in Kubernetes v1.20. Therefore, Kubernetes v1.20 is recommended. For details, see https://kubernetes-csi.github.io/docs/external-snapshotter.html.

6.5.1 Installing the Snapshot-Dependent Component Service

Perform the following steps only on any master node.

Prerequisites

Before the installation, run the **kubectl api-resources** | **grep snapshot** | **awk** '{**print \$1**}' command on the master node to check whether the snapshot-related resource service has been installed. If the following information is displayed, you do not need to install it again.

kubectl api-resources | grep snapshot | awk '{print \$1}' volumesnapshotclasses volumesnapshotcontents volumesnapshots

If the command output does not contain any of the preceding services, perform the following operations to install the service.

Step 1 Check the Kubernetes version.

kubectl version --short=true Client Version: v1.22.1 Server Version: v1.22.1

Step 2 Go to the corresponding directory according to the Kubernetes version, and run the following command to install the snapshot-dependent component service.

kubectl apply -f huawei-csi-snapshot-crd.yaml

----End

6.5.2 Managing a VolumeSnapshotClass

6.5.2.1 Creating a VolumeSnapshotClass

This section describes how to create a VolumeSnapshotClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *volume-snapshot-class.yaml* command to create a file named *volume-snapshot-class.yaml*.

vi volume-snapshot-class.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *volume-snapshot-class.yaml* file. After the modification is complete, press Esc and enter :wq! to save the modification.

The following shows a template of the *volume-snapshot-class.yaml* file. You can also refer to the **examples**/*<Kubernetes version>*/**volume-snapshot-class.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-15**.

apiVersion: snapshot.storage.k8s.io/v1beta1 kind: VolumeSnapshotClass metadata: name: mysnapclass driver: csi.huawei.com deletionPolicy: Delete

Table 6-15 Parameter description

Parameter	Description	Remarks
metadata.nam e	User-defined name of a VolumeSnapshotClass object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
driver	driver identifier.	The value is fixed to csi.huawei.com.
deletionPolicy	Handles the VolumeSnapshotContent policy when a VolumeSnapshot is deleted.	This parameter is mandatory. The value can be Delete or Retain .

Step 4 Run the following command to create a VolumeSnapshotClass based on the .yaml file.

kubectl create -f volume-snapshot-class.yaml

Step 5 Run the following command to view the information about the created VolumeSnapshotClass.

kubectl get volumesnapshotclass NAME DRIVER DELETIONPOLICY AGE mysnapclass csi.huawei.com Delete 25s

----End

6.5.2.2 Deleting a VolumeSnapshotClass

This section describes how to delete a VolumeSnapshotClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query VolumeSnapshotClasses in the cluster.

```
# kubectl get volumesnapshotclass
NAME DRIVER DELETIONPOLICY AGE
mysnapclass csi.huawei.com Delete 52s
```

Step 3 Run the following command to delete a VolumeSnapshotClass. For example, delete the VolumeSnapshotClass named *mysnapclass*.

```
# kubectl delete volumesnapshotclass mysnapclass volumesnapshotclass.snapshot.storage.k8s.io "mysnapclass" deleted
```

Step 4 Run the following command to query VolumeSnapshotClasses in the cluster. If the command output does not contain the name of the VolumeSnapshotClass you want to delete, it is successfully deleted.

```
# kubectl get volumesnapshotclass
No resources found in default namespace.
```

----End

6.5.3 Managing a VolumeSnapshot

6.5.3.1 Creating a VolumeSnapshot

This section describes how to create a VolumeSnapshot.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *volume-snapshot.yaml* command to create a file named *volume-snapshot.yaml*.

vi volume-snapshot.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *volume-snapshot.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *volume-snapshot.yaml* file. You can also refer to the **examples**/*<Kubernetes version>*/**volume-snapshot.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-16**.

```
apiVersion: snapshot.storage.k8s.io/v1beta1 kind: VolumeSnapshot metadata: name: mysnapshot spec: volumeSnapshotClassName: mysnapclass
```

source: persistentVolumeClaimName: mypvc

Table 6-16 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a VolumeSnapshot object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.volumeSnapsho tClassName	Name of the VolumeSnapshotClass object.	Enter the name of the VolumeSnapshotClass object created in 6.5.2.1 Creating a VolumeSnapshotClass.
spec.source.persisten tVolumeClaimName	Name of the source PVC object.	Enter the name of the PVC object created in 6.3.1 Creating a PVC.

- **Step 4** Run the following command to create a VolumeSnapshot based on the .yaml file. # kubectl create -f volume-snapshot.yaml
- **Step 5** Run the following command to view the information about the created VolumeSnapshot.

```
# kubectl get volumesnapshot
NAME READYTOUSE SOURCEPVC SOURCESNAPSHOTCONTENT RESTORESIZE
SNAPSHOTCLASS SNAPSHOTCONTENT CREATIONTIME AGE
mysnapshot true mypvc 100Gi mysnapclass
snapcontent-1009af0a-24c2-4435-861c-516224503f2d <invalid> 78s
```

----End

6.5.3.2 Deleting a VolumeSnapshot

This section describes how to delete a VolumeSnapshot.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query VolumeSnapshots in the cluster.

```
# kubectl get volumesnapshot
NAME READYTOUSE SOURCEPVC SOURCESNAPSHOTCONTENT RESTORESIZE
SNAPSHOTCLASS SNAPSHOTCONTENT CREATIONTIME AGE
mysnapshot true mypvc 100Gi mysnapclass
snapcontent-1009af0a-24c2-4435-861c-516224503f2d <invalid> 78s
```

Step 3 Run the following command to delete a VolumeSnapshot. For example, delete the VolumeSnapshot named *mysnapshot*.

kubectl delete volumesnapshot *mysnapshot* volumesnapshot.snapshot.storage.k8s.io "*mysnapshot*" deleted

Step 4 Run the following command to query VolumeSnapshots in the cluster. If the command output does not contain the name of the VolumeSnapshot you want to delete, it is successfully deleted.

kubectl get volumesnapshot No resources found in default namespace.

----End

Advanced Features

This chapter describes how to configure advanced features of Huawei storage.

- 7.1 Configuring Multiple Backends
- 7.2 Creating a PVC for a Specified Backend
- 7.3 Creating a PVC for a Specified Storage Pool
- 7.4 Configuring ALUA
- 7.5 Configuring Storage Topology Awareness
- 7.6 Advanced Features of Enterprise Storage
- 7.7 Advanced Features of Distributed Storage

7.1 Configuring Multiple Backends

Huawei CSI supports multiple backends. Perform this operation when you want to configure multiple backends.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Configure the **huawei-csi-configmap.yaml** file. The following shows a template of the **huawei-csi-configmap.yaml** file. Set related parameters based on the site requirements and save the file in yaml format.

Multiple backends are separated by commas (,). For details about each backend, see 4.5 Connecting to Enterprise Storage or 4.6 Connecting to Distributed Storage.

```
kind: ConfigMap
apiVersion: v1
metadata:
name: huawei-csi-configmap
namespace: kube-system
data:
csi.json: |
{
```

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

Step 5 Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.

----End

7.2 Creating a PVC for a Specified Backend

When multiple backends are configured, you can perform the following operations to create a PVC for a specified backend.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify parameters in the following fields. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - Add the **backend** configuration item under **parameters**.
 - The value of **metadata.name** is the user-defined name of a StorageClass object.

Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.

• The value of **parameters.backend** is the name of a backend in **huawei-csi-configmap.yaml**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
backend: "***"
```

Step 3 Run the following command to create a StorageClass based on the .yaml file.

kubectl create -f StorageClass.yaml

----End

7.3 Creating a PVC for a Specified Storage Pool

When multiple storage pools are configured, you can perform the following operations to create a PVC for a specified storage pool.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify parameters in the following fields. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - The value of **metadata.name** is the user-defined name of a StorageClass object.

Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.

- Add the **pool** configuration item under **parameters**.
- The value of **pool** is the name of a storage pool in **huawei-csi-configmap.yaml**.

Ⅲ NOTE

The volume to be created using the StorageClass will be created in the specified storage pool. The existing PVC will not change the storage pool information.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
pool: "***"
```

Step 3 Run the following command to create a StorageClass based on the .yaml file.

kubectl create -f StorageClass.yaml

----End

7.4 Configuring ALUA

7.4.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3

This section describes how to configure ALUA if multipathing is used during the connection to block storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.5 Connecting to Enterprise Storage or 4.6 Connecting to Distributed Storage.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-1**.

```
{
  "backends": [
  {
    "storage": "oceanstor-san",
    ...
    "parameters": {..., "ALUA": {"<HostName>": {"MULTIPATHTYPE": "*", "FAILOVERMODE": "*",
    "SPECIALMODETYPE": "*", "PATHTYPE": "*"}, "<HostName>": {...}}}
  }
}
```

Table 7-1 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node.	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
MULTIPATHTYPE	Multipathing type. The value can be: • 0: default • 1: third-party multipathing	-
FAILOVERMODE	Initiator switchover mode. The value can be: • 0: early-version ALUA • 1: common ALUA • 2: ALUA not used • 3: special ALUA	This parameter needs to be specified only when third-party multipathing is used. All OceanStor V5 models do not support early-version ALUA.

Parameter	Description	Remarks
SPECIALMODETYPE	Special mode type of the initiator. The value can be: • 0: special mode 0 • 1: special mode 1 • 2: special mode 2 • 3: special mode 3	This parameter needs to be specified only when the initiator switchover mode is special ALUA.
PATHTYPE	Initiator path type. The value can be: • 0: preferred path • 1: non-preferred path	This parameter needs to be specified only when third-party multipathing is used.

□ NOTE

- The ALUA configuration may vary according to the OS. Visit https://support.huawei.com/enterprise/en/index.html, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS and configure ALUA based on the recommended configurations in the guide.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- The value of **HostName** is a regular expression. For details about how to configure it, see **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used. When you configure **HostName**, the number of host connections is limited. For details about the limitation, see **Specifications Query** and search for **Maximum number of iSCSI connections per controller enclosure**. If the number of host connections is less than or equal to the specifications, you are advised to use the general configuration. If the number of host connections is greater than the specifications, you are advised to use the common configuration.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- In the general configuration, if you need to exactly match a host, refer to example 3 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- OceanStor V3/V5 and OceanStor Dorado V3 use this configuration mode. For details about related parameters, see **Table 7-2**.

Scenario	Host Type	Whether the Storage Has the Preferred Path	Recommended ALUA Configuration
HyperMetro storage	CentOS/RHEL host	Yes	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="0"
		No	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="1"
	SUSE/Debian host	Yes	ALUA="1" FAILOVERMODE="1" PATHTYPE="0"
		No	ALUA="1" FAILOVERMODE="1" PATHTYPE="1"
Non- HyperMetro storage	CentOS/RHEL host	N/A	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="0"
	SUSE/Debian host	N/A	ALUA="1" FAILOVERMODE="1" PATHTYPE="0"

Table 7-2 Recommended ALUA parameter configurations for OceanStor V3/V5 and OceanStor Dorado V3

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

Step 5 Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.

----End

7.4.2 Configuring ALUA for OceanStor Dorado 6.x

This section describes how to configure ALUA if multipathing is used during the connection to block storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.5 Connecting to Enterprise Storage or 4.6 Connecting to Distributed Storage.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-3**.

```
{
  "backends": [
  {
    "storage": "oceanstor-san",
    ...
    "parameters": {..., "ALUA": {"<HostName>": {"accessMode": "*", "hyperMetroPathOptimized": "*"},
  "<HostName>": {...}}}
  }
}
```

Table 7-3 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node.	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
accessMode	Host access mode. The value can be:	
	• 0 : balanced mode	
	• 1: asymmetric mode	
hyperMetroPathOp- timized	Whether the path of the host on the current storage array is preferred in HyperMetro scenarios. The value can be:	This parameter needs to be specified only when the host access mode is set to asymmetric.
	• 1 : yes	
	• 0 : no	

□ NOTE

- The ALUA configuration may vary according to the OS. Visit https://support.huawei.com/enterprise/en/index.html, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS and configure ALUA based on the recommended configurations in the guide.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- The value of **HostName** is a regular expression. For details about how to configure it, see **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used. When you configure **HostName**, the number of host connections is limited. For details about the limitation, see **Specifications Query** and search for **Maximum number of iSCSI connections per controller enclosure**. If the number of host connections is less than or equal to the specifications, you are advised to use the general configuration. If the number of host connections is greater than the specifications, you are advised to use the common configuration.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- In the general configuration, if you need to exactly match a host, refer to example
 3 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- If a host uses only OceanStor Dorado 6.x all-flash storage, see **Table 7-4** for detailed parameters.
- If you add OceanStor Dorado 6.x all-flash storage to a host that uses OceanStor converged storage, see **Table 7-5** for detailed parameters.

Table 7-4 Recommended ALUA parameter configurations for OceanStor Dorado 6.*x* all-flash storage

Scenario	Host Type	Host Access Mode	Recommended ALUA Configuration
HyperMetro storage	CentOS/ RHEL/SUSE/	Load balancing mode	ALUA not required
	Debian host	Asymmetric mode + Storage with the preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the non-preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="0"
Non- HyperMetro storage	CentOS/ RHEL/SUSE/ Debian host	N/A	ALUA not required

Scenario	Host Type	Host Access Mode	Recommended ALUA Configuration
HyperMet ro storage	CentOS/RHEL/SUSE/ Debian host	Load balancing mode	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the non- preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="0"
Non- HyperMet ro storage	CentOS/RHEL/SUSE/ Debian host	N/A	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"

Table 7-5 Recommended ALUA parameter configurations for hybrid OceanStor V3/V5, OceanStor Dorado V3, and OceanStor Dorado 6.*x*

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

Step 5 Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.

----End

7.4.3 Configuring ALUA for Distributed Storage

This section describes how to configure ALUA if multipathing is used during the connection to block storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.5 Connecting to Enterprise Storage or 4.6 Connecting to Distributed Storage.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-6**.

```
{
  "backends": [
    {
        "storage": "fusionstorage-san",
        ...
        "parameters": {..., "ALUA": {"<HostName>": {"switchoverMode": "*", "pathType": "*"},
        "<HostName>": {...}}}
    }
}
```

Table 7-6 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node.	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
switchoverMode	Switchover mode. The value can be: • Disable_alua: disables	
	ALUA.	
	• Enable_alua: enables ALUA.	
pathType	Path type. The value can be:	
	optimal_path: preferred path	
	 non_optimal_path: non- preferred path 	

□ NOTE

- Only the iSCSI scenario of distributed storage is supported.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- When the value of **HostName** is a regular expression, configure it by referring to **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- In the general configuration, if you need to exactly match a host, refer to example
 3 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- **Step 3** Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

Step 5 Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.

----End

7.5 Configuring Storage Topology Awareness

In the Kubernetes cluster, resources can be scheduled and provisioned based on the topology labels of nodes and the topology capabilities supported by storage backends.

Prerequisites

- Kubernetes v1.17 and later versions support the topology awareness feature.
- You need to configure topology labels on worker nodes in the cluster. The method is as follows:
 - a. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
 - b. Run the **kubectl get node** command to view information about worker nodes in the current cluster.

```
# kubectl get node
NAME STATUS ROLES AGE VERSION
node01 Ready controlplane,etcd,worker 42d v1.19.3
node02 Ready worker 42d v1.19.3
node03 Ready worker 42d v1.19.3
```

c. Run the kubectl label node <nodename> topology.kubernetes.io/ <key>= <value> command to configure a topology label for a worker node. In the preceding command, <nodename> indicates the name of a worker node. For details about the key and value parameters, see Table 7-7.

kubectl label node node01 topology.kubernetes.io/zone=ChengDu node/node01 labeled

Table 7-7 Parameter description

Parameter	Description	Remarks
<key></key>	Unique identifier of a topology label.	The value can be zone , region , or protocol . < <i>protocol</i> >.
		<pre><pre><pre><pre><pre><pre>can be set to iscsi, nfs, fc, or roce.</pre></pre></pre></pre></pre></pre>
<value></value>	Value of a topology label.	If key is set to zone or region , value is a user-defined parameter.
		If key is set to protocol . <pre><pre>/protocol.</pre><pre>/protocol></pre>, value is fixed to csi.huawei.com.</pre>

- A topology label must start with topology.kubernetes.io. Topology label examples:
 - Example 1: topology.kubernetes.io/region=China-west
 - Example 2: topology.kubernetes.io/zone=ChengDu
 - Example 3: topology.kubernetes.io/protocol.iscsi=csi.huawei.com
 - Example 4: topology.kubernetes.io/protocol.fc=csi.huawei.com
- A key in a topology label on a node can have only one value.
- If multiple protocols are configured in a topology label on a node, when you select a backend, the backend needs to meet only one of the protocols.
- If both the region and the zone are configured in a topology label on a node, when you select a backend, the backend must meet both of them.
- d. Run the kubectl get nodes -o=json path='{range .items[*]} [{.metadata.name}, {.metadata.labels}]{"\n"}{end}' | grep --color "topology.kubernetes.io" command to view the label information about all worker nodes in the current cluster.

kubectl get nodes -o=jsonpath='{range .items[*]}{{.metadata.name}, {.metadata.labels}]{"\n"} {end}' | grep --color "topology.kubernetes.io" [node01, {"beta.kubernetes.io/arch":"amd64","beta.kubernetes.io/os":"linux","kubernetes.io/arch":"amd64","kubernetes.io/hostname":"node01","kubernetes.io/os":"linux","node-role.kubernetes.io/controlplane":"true","node-role.kubernetes.io/etcd":"true","node-role.kubernetes.io/worker":"true","topology.kubernetes.io/zone":"ChengDu"}]

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.5 Connecting to Enterprise Storage or 4.6 Connecting to Distributed Storage.

Add the **supportedTopologies** field under the **backends** section in the *huawei-csi-configmap.yaml* file to configure the topology information supported by each backend. The following is a backend example.

```
{
    "backends":[
    {
        "storage": "oceanstor-san",
        ...
        "parameters": {"protocol": "iscsi", "portals": ["192.168.125.22", "192.168.125.23"]},
        "supportedTopologies": [
            {"topology.kubernetes.io/region": "China-west", "topology.kubernetes.io/zone": "ChengDu"},
            {"topology.kubernetes.io/region": "China-south", "topology.kubernetes.io/zone": "ShenZhen"}]
    }
}
```

□ NOTE

- **supportedTopologies** is a list. Each element in the list is a dictionary.
- Only topology.kubernetes.io/region or topology.kubernetes.io/zone can be configured for each element in the list. The parameter value must be the same as the topology label set in the prerequisites. (topology.kubernetes.io/protocol.protocol>
 does not need to be configured.)
- **Step 3** Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

```
# kubectl create -f huawei-csi-configmap.yaml
```

Step 4 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

- **Step 5** Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.
- **Step 6** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and add related parameters in the .yaml file. For details about the parameters, see **Table 7-8**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add the following configuration items to the StorageClass.yaml file.

• Example 1: Configure zone and region information in the StorageClass. kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: example-storageclass

provisioner: csi.huawei.com

parameters: volumeType: lun allocType: thin

volume Binding Mode: Wait For First Consumer

allowedTopologies:

- matchLabelExpressions:
- key: topology.kubernetes.io/zone values:
- ChengDu
- key: topology.kubernetes.io/region values:
- China-west
- Example 2: Configure protocol information in the StorageClass.

kind: StorageClass

apiVersion: storage.k8s.io/v1

metadata:

name: protocol-example-storageclass

provisioner: csi.huawei.com

parameters: volumeType: lun allocType: thin

volumeBindingMode: WaitForFirstConsumer allowedTopologies:

- matchLabelExpressions:
- key: topology.kubernetes.io/protocol.iscsi values:
- csi.huawei.com

Table 7-8 Parameter description

Parameter	Description	Remarks
volumeBindi ngMode	PersistentVolume binding mode, used to control the time when PersistentVolume resources are dynamically allocated and bound.	You can set this parameter to WaitForFirstConsumer or Immediate. WaitForFirstConsumer: indicates that the binding and allocation of the PersistentVolume are delayed until a Pod that uses the PVC is created.
		Immediate: The PersistentVolume is bound and allocated immediately after a PVC is created.
allowedTopol ogies.matchL abelExpressio ns	Topology information label, which is used to filter CSI backends and Kubernetes nodes. If the matching fails, PVCs or Pods cannot be created. Both key and value must be configured in a fixed format.	key: This parameter can be set to topology.kubernetes.io/zone or topology.kubernetes.io/region. topology.kubernetes.io/ protocol. <protocol>: <protocol> indicates the protocol type and can be iscsi, fc, or nfs.</protocol></protocol>

Parameter	Description	Remarks
	•	value: If key is topology.kubernetes.io/ zone or topology.kubernetes.io/ region, value must be the same as the topology label set in the prerequisites. If key is topology.kubernetes.io/
		protocol. <pre>col. <pre>c</pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre></pre>

- **Step 7** Run the following command to create a StorageClass based on the .yaml file.

 # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC with the topology capability. For details, see **6.3.1 Creating a PVC**.
- **Step 9** Use the PVC to create a Pod. For details, see **6.4.1 Creating a Pod**.

----End

7.6 Advanced Features of Enterprise Storage

7.6.1 Configuring QoS

This section describes how to create a LUN/file system volume that supports QoS.

Precautions

- The QoS feature is not a standard feature of Kubernetes and is customized by storage vendors.
- A QoS policy can be specified only when a StorageClass is created. Once the QoS policy is created, it cannot be modified because the StorageClass cannot be modified on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - The value of **metadata.name** is the user-defined name of a StorageClass object. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-).
 - Add the qos configuration item under parameters. For other information about parameters, see 6.1.1.1 Creating a LUN StorageClass or 6.1.1.2 Creating a File System StorageClass.

- The value of the **qos** section is JSON character strings in dictionary format. A character string is enclosed by single quotation marks and the dictionary key by double quotation marks.
 - For details about the parameters of OceanStor V3/OceanStor V5 series storage devices, see Table 7-9.
 - For details about the parameters of OceanStor Dorado V3 series storage devices, see Table 7-10.
 - For details about the parameters of OceanStor Dorado 6.x series storage devices, see Table 7-11.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: "***" provisioner: "csi.huawei.com" parameters:

qos: '{"IOTYPE": 2, "MINIOPS": 1000}'

Table 7-9 QoS parameters supported by OceanStor V3/OceanStor V5

Parameter	Description	Remarks
IOTYPE	Read/write type.	This parameter is optional. If it is not specified, the default value of the storage backend is used. For details, see related storage documents.
		The value can be:
		• 0 : read I/O
		• 1: write I/O
		• 2: read and write I/Os
MAXBANDWIDTH	Maximum bandwidth. This is a restriction policy parameter.	The value is an integer greater than 0, expressed in MB/s.
MINBANDWIDTH	Minimum bandwidth. This is a protection policy parameter.	The value is an integer greater than 0, expressed in MB/s.
MAXIOPS	Maximum IOPS. This is a restriction policy parameter.	The value is an integer greater than 0
MINIOPS	Minimum IOPS. This is a protection policy parameter.	The value is an integer greater than 0
LATENCY	Maximum latency. This is a protection policy parameter.	The value is an integer greater than 0, expressed in ms.

MAXIOPS

The value is an integer ranging

from 100 to 999999999.

ParameterDescriptionRemarksIOTYPERead/write type.The value can be:
• 2: read and write I/OsMAXBANDWIDTHMaximum bandwidth.
This is a restriction
policy parameter.The value is an integer ranging
from 1 to 999999999, expressed in
MB/s.

Table 7-10 QoS parameters supported by OceanStor Dorado V3

Table 7-11 QoS parameters supported by OceanStor Dorado 6.x

Maximum IOPS. This

is a restriction policy

parameter.

Parameter	Description	Remarks
IOTYPE	Read/write type.	The value can be:
		• 2: read and write I/Os
MAXBANDWIDTH	Maximum bandwidth.	The value is an integer ranging from 1 to 999999999, expressed in MB/s.
MINBANDWIDTH	Minimum bandwidth.	The value is an integer ranging from 1 to 999999999, expressed in MB/s.
MAXIOPS	Maximum IOPS.	The value is an integer ranging from 100 to 999999999.
MINIOPS	Minimum IOPS.	The value is an integer ranging from 100 to 999999999.
LATENCY	Maximum latency.	The value can be 0.5 or 1.5 , expressed in ms.

□ NOTE

- Different protection policy parameters or restriction policy parameters can be specified at the same time. However, protection policy parameters cannot coexist with restriction policy parameters.
- vStore users do not support QoS policies.
- The QoS configuration takes effect only on the newly created PVC. QoS cannot be added automatically for PVCs with the same StorageClass name that have been provisioned.

Step 3 Run the following command to create a StorageClass based on the .yaml file.

kubectl create -f StorgeClass.yaml

Step 4 Use the StorageClass to create a PVC with the QoS capability. For details, see **6.3.1** Creating a PVC.

----End

7.6.2 Configuring a vStore

This section describes how to configure a vStore.

Precautions

This feature does not support SAN storage. For details about the supported models, see **3 Restrictions**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.5 Connecting to Enterprise Storage.

• If the vStore uses OceanStor V3/V5 storage, add the **vstoreName** parameter to the backend configuration and set the parameter to the vStore name of the storage device.

```
{
    "backends": [
    {
        ...
        "vstoreName": "***"
    }
    ]
}
```

• If the vStore uses OceanStor 6.1 or OceanStor Dorado 6.x storage, configure the backend by referring to 4.5 Connecting to Enterprise Storage. Note that urls indicates the logical management ports of the vStore, and pools and portals must be available storage pools and logical data ports of the current vStore respectively.

□ NOTE

After configuring **huawei-csi-configmap.yaml**, restart huawei-csi-controller and huawei-csi-node. Otherwise, the configuration does not take effect.

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 Run the **kubectl get configmap -n kube-system | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n kube-system | grep huawei-csi-configmap
huawei-csi-configmap 1 5s
```

Step 5 Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services**.

□ NOTE

- When starting the huawei-csi services, enter the user name and password of the storage device vStore entered in **Step 3** in **4.7 Starting huawei-csi Services**.
- For details about the storage models supported by this feature, 3 Restrictions.

----End

7.6.3 Configuring NAS HyperMetro

Perform this operation when you want to configure NAS HyperMetro.

7.6.3.1 Prerequisites

Configuring the Storage HyperMetro Relationship

Before configuring NAS HyperMetro, you need to configure the HyperMetro relationship between two storage devices, including the remote device and HyperMetro domain. For details, see the product documentation of the corresponding storage model.

Configuring HyperMetro vStore Information

Before configuring NAS HyperMetro, you need to create vStores, vStore pairs, vStore users, and corresponding logical ports for the two storage devices. For details, see the product documentation of the desired storage model. This section uses OceanStor Dorado 6.1.3 as an example.

- **Step 1** Create a vStore on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > vStore Service > vStores.
 - 3. Click **Create**, specify **Name** (for example, **vStoreA**), and click **OK**.
- **Step 2** Create a vStore on the secondary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the secondary storage device in the HyperMetro domain.
 - 2. Choose Services > vStore Service > vStores.
 - 3. Click **Create**, specify **Name** (for example, **vStoreB**), and click **OK**.
- **Step 3** Create a HyperMetro vStore pair on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - Choose Data Protection > Protection Entities > vStores > vStores.
 - Select the vStore for which you want to create a HyperMetro vStore pair (for example, vStoreA created in Step 1) and click Create HyperMetro vStore Pair.

- 4. Set **Pair Creation** to **Manual**, set **Remote vStore** to the vStore of the secondary storage device (for example, **vStoreB** created in **Step 2**), set other parameters based on site requirements, and click **OK**.
- **Step 4** Create a logical management port for the local storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > Network > Logical Ports.
 - Click Create, specify Name (for example, manageA), set Role to Management, Owning vStore to vStoreA, and Owning Site to Local device, set other parameters based on site requirements, and click OK.
- **Step 5** Create a logical management port for the remote storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > Network > Logical Ports.
 - Click Create, specify Name (for example, manageB), set Role to Management, Owning vStore to vStoreA, and Owning Site to Remote device, set other parameters based on site requirements, and click OK.
 - 4. On DeviceManager of the secondary storage device in the HyperMetro domain, check whether the logical port is correct.
 - For example, the **Home Port** parameter value of the logical port and that of the logical port corresponding to the primary storage device in the HyperMetro domain are on the same network plane.
- **Step 6** Create a logical service port for the local storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > Network > Logical Ports.
 - Click Create, specify Name (for example, dataA), set Role to Data, Data Protocol to NFS, Owning vStore to vStoreA, and Owning Site to Local device, set other parameters based on site requirements, and click OK.
- **Step 7** Create a logical service port for the remote storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > Network > Logical Ports.
 - 3. Click **Create**, specify **Name** (for example, **dataB**), set **Role** to **Data**, **Data Protocol** to **NFS**, **Owning vStore** to **vStoreA**, and **Owning Site** to **Remote device**, set other parameters based on site requirements, and click **OK**.
 - 4. On DeviceManager of the secondary storage device in the HyperMetro domain, check whether the logical port is correct.
 - For example, the **Home Port** parameter value of the logical port and that of the logical port corresponding to the primary storage device in the HyperMetro domain are on the same network plane.

Step 8 Create a vStore user on the primary storage device in the HyperMetro domain.

- 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
- 2. Choose Services > vStore Service > vStores.
- 3. Click the name of **vStoreA**, click **User Management**, and click **Create**. Set **Type** to **Local user** and **Role** to **vStore administrator**, set other parameters based on site requirements, and click **OK**.

Step 9 Create a vStore user on the secondary storage device in the HyperMetro domain.

- 1. Log in to DeviceManager of the secondary storage device in the HyperMetro domain.
- 2. Choose Services > vStore Service > vStores.
- 3. Click the name of vStoreB, click User Management, and click Create. Set Type to Local user and Role to vStore administrator, set other parameters based on site requirements, and click OK.

----End

7.6.3.2 Configuration Procedure

Precautions

For details about the resource objects that support NAS HyperMetro and the feature description, see **Table 7-12**.

Table 7-12 Feature description

Resource Object	Operation	Supporte d	Remarks
PVC	Creation	Yes	This feature can be used
	Deletion	Yes	together with other features (except remote replication), such as QoS.
	Capacity expansion	Yes	-
	Synchronizing a HyperMetro pair	No	Storage supports these operations. Because Kubernetes cannot detect
	Pausing a HyperMetro pair	No	HyperMetro pairs, Kubernetes does not support these operations.
	Preferred site No switchover		
Pod	Creation	Yes	Primary and secondary file
	Deletion	Yes	systems can be mounted to a Pod at the same time.

Resource Object	Operation	Supporte d	Remarks
VolumeSnapshot	Creation	Yes	Snapshots can be operated
	Deletion	Yes	only for primary storage file systems.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

In the **backends** section of the *huawei-csi-configmap.yaml* file, add two backends with a HyperMetro relationship. For details about the configuration items for each backend, see **4.5 Connecting to Enterprise Storage**. You need to add some additional configuration parameters in the HyperMetro scenario. For details, see **Table 7-13**.

```
kind: ConfigMap
apiVersion: v1
metadata:
 name: huawei-csi-configmap
 namespace: kube-system
data:
 csi.json: |
      "backends": [
             "storage": "oceanstor-nas",
             "name": "hyperMetro1",
            "urls": ["https://192.168.125.20:8088", "https://192.168.125.21:8088"],
            "pools": ["storagepool01", "storagepool02"],
"parameters": {"protocol": "nfs", "portals": ["192.168.125.22"]},
"vstoreName": "vStore1",
            "metrovStorePairID": "f09838237b93c000",
             "metroBackend": "hyperMetro2"
            "storage": "oceanstor-nas",
             "name": "hyperMetro2",
            "urls": ["https://192.168.125.24:8088", "https://192.168.125.25:8088"],
            "pools": ["storagepool01", "storagepool02"], "parameters": {"protocol": "nfs", "portals": ["192.168.125.26"]}, "vstoreName": "vStore1",
            "metrovStorePairID": "f09838237b93c000",
             "metroBackend": "hyperMetro1"
      ]
```

Configuration Item	Format	Description	Remarks
vstoreName	String	vStore name. This parameter is conditionally mandatory (mandatory for OceanStor V3/V5).	Only vStore users support NAS HyperMetro. For details about how to configure vStore users, see 7.6.3.1 Prerequisites.
metrovStorePa irID	String	ID of the HyperMetro vStore pair to which a vStore belongs. This parameter is mandatory.	For example, the parameter of OceanStor Dorado 6.x/ OceanStor 6.1 is displayed as ID on DeviceManager.
metroBackend	String	Name of a peer end in HyperMetro. The two backends form a HyperMetro relationship. This parameter is mandatory.	The peer end of hyperMetro1 is hyperMetro2, and the peer end of hyperMetro2 is hyperMetro1.

Table 7-13 HyperMetro configuration items of a back-end storage device

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n kube-system** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n kube-system | grep huawei-csi-configmap huawei-csi-configmap 1 5s

- **Step 5** Start huawei-csi services. For details, see **4.7 Starting huawei-csi Services** (use the vStore user created in **7.6.3.1 Prerequisites** for configuration).
- **Step 6** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and add the **hyperMetro** parameter under **parameters** in the .yaml file. For details about the parameters, see **Table 7-14**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: "***" provisioner: "csi.huawei.com" parameters: ... volumeType: fs hyperMetro: "true"

Parameter	Description	Remarks
parameters.hyperMetro	Whether a HyperMetro volume is to be created.	If this parameter is set to true , a HyperMetro volume is to be created. If this parameter is not set or set to false , no HyperMetro volume is to be created.

Table 7-14 Parameter description

- **Step 7** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC with the NAS HyperMetro capability. For details, see **6.3.1 Creating a PVC**.

----End

7.6.4 Configuring an Application Type

This section describes how to create a LUN/file system volume that supports different application types.

Precautions

- The application type feature is not a standard feature of Kubernetes and is customized by storage vendors.
- An application type can be specified only when a PVC is created.
- A created PVC cannot be modified on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add the applicationType configuration item under parameters. The value of applicationType is a character string. For details, see Table 7-15. For other information about parameters, see 6.1.1.1 Creating a LUN StorageClass or 6.1.1.2 Creating a File System StorageClass.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
volumeType: "***"
applicationType: "***"
```

A STATE OF THE PROPERTY OF THE				
Parameter	Description	Remarks		
parameters.applic ationType	Application type name on the storage device. The value is a character string.	If the value of volumeType is lun, log in to DeviceManager and choose Services > Block Service > LUN Groups > LUNs > Create to obtain the application type name. If the value of volumeType is fs, log in to DeviceManager and choose Services > File Service > File Systems > Create to obtain the application type name.		

Table 7-15 Parameter description of applicationType

□ NOTE

This feature applies only to OceanStor Dorado 6.x series storage systems.

- **Step 3** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorgeClass.yaml
- **Step 4** Use the StorageClass to create a PVC with the application type capability. For details, see **6.3.1 Creating a PVC**.

----End

7.7 Advanced Features of Distributed Storage

7.7.1 Configuring QoS

This section describes how to create a LUN volume that supports QoS.

Precautions

- The QoS feature is not a standard feature of Kubernetes and is customized by storage vendors.
- A QoS policy can be specified only when a StorageClass is created. Once the QoS policy is created, it cannot be modified because the StorageClass cannot be modified on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - Add the qos configuration item under parameters. For other information about parameters, see 6.1.1.1 Creating a LUN StorageClass.

The value of the qos section is JSON character strings in dictionary format. A
character string is enclosed by single quotation marks and the dictionary key
by double quotation marks. For details about the parameters, see Table 7-16.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
qos: '{"maxMBPS": 999, "maxIOPS": 999}'
```

Table 7-16 Parameters in qos

Parameter	Description	Remarks
maxMBPS	Maximum bandwidth.	This parameter is mandatory. The value is an integer greater than 0, expressed in MB/s.
maxIOPS	Maximum IOPS.	This parameter is mandatory. The value is an integer greater than 0

Step 3 Run the following command to create a PVC based on the .yaml file.

kubectl create -f StorageClass.yaml

----End

7.7.2 Configuring a Soft Quota

This section describes how to create a PVC that supports soft quotas.

Precautions

- This feature is supported only by OceanStor Pacific series 8.1.0 and later versions
- This feature can be configured only when a storage pool of the file system type is connected.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - Add the storageQuota configuration item under parameters. For other information about parameters, see 6.1.1.2 Creating a File System StorageClass.
 - The value of the storageQuota section is JSON character strings in dictionary format. A character string is enclosed by single quotation marks and the dictionary key by double quotation marks. For details about the parameters, see Table 7-17.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
volumeType: "fs"
...
storageQuota: '{"spaceQuota": "softQuota", "gracePeriod": 100}'
```

Table 7-17 Parameters in storageQuota

Parameter	Description	Remarks
spaceQuota	File quota type.	This parameter is mandatory. Only softQuota or hardQuota can be configured.
gracePeriod	Grace period allowed when the soft quota is configured.	This parameter is conditionally optional only when spaceQuota is set to softQuota .
		The value is an integer ranging from 0 to 4294967294.

- **Step 3** Run the following command to create a PVC based on the .yaml file.
 - # kubectl create -f StorageClass.yaml
- **Step 4** Configure a StorageClass in the PVC according to **6.3.1 Creating a PVC** to finish creating a PVC.

----End

7.7.3 Configuring an Account

This section describes how to configure a backend for a specified account and create a PVC for the backend.

Precautions

This feature supports only distributed NAS series storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

For details about the backend configuration, see 4.6.3 Connecting to Distributed Storage NAS over NFS. Add the accountName configuration item to the backends section in the huawei-csi-configmap.yaml file. accountName indicates the account name to be specified and its value is a character string. For details about the parameter, see Table 7-18. portals must be set to the logical ports owned by the account.

```
{
  "backends":[
  {
    "storage": "fusionstorage-nas",
    "name": "***",
    ...
  "accountName": "***",
    "parameters": {"protocol": "nfs", "portals": ["*.*.*.*"]}
  }
}
```

■ NOTE

portals must be the IP addresses under the configured account. To view its value, log in to DeviceManager and choose **Resources** > **Access** > **Service Network**.

Table 7-18 Parameter description of accountName

Parameter	Description	Remarks
backends.accou ntName	Account name to be specified. The value is a character string.	• If accountName is set to system or this parameter is not specified, the PVC created using the backend will be created under the system account (user name: system).
		If accountName is set to another user, the PVC created using the backend will be created only under the user.

- **Step 3** Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.
 - # kubectl create -f huawei-csi-configmap.yaml
- **Step 4** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 5** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 6** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add the **backend** configuration item under **parameters**. The value of **backend** is a character string and is the backend name of the account specified in **Step 2**. For other information about **parameters**, see **6.1.1.2 Creating a File System StorageClass**. For details about how to create a PVC for a specified backend, see **7.2 Creating a PVC for a Specified Backend**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
volumeType: "***"
backend: "***"
```

- **Step 7** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC for a specified user. For details, see **6.3.1** Creating a PVC.

----End

8 Uninstalling CSI

Perform this operation when you want to uninstall all CSIs.

- 8.1 (Optional) Uninstalling the Snapshot-Dependent Component Service
- 8.2 Uninstalling the huawei-csi-node Service
- 8.3 Uninstalling the huawei-csi-controller Service
- 8.4 Deleting the huawei-csi-configmap Object
- 8.5 Deleting the huawei-csi-secret Object
- 8.6 Deleting the RBAC Permission
- 8.7 Deleting the Image of the Earlier Version

8.1 (Optional) Uninstalling the Snapshot-Dependent Component Service

NOTICE

- Do not uninstall the snapshot-dependent component service when snapshots exist. Otherwise, Kubernetes will automatically delete all user snapshots and they cannot be restored. Exercise caution when performing this operation. For details, see **Delete a CustomResourceDefinition**.
- Do not uninstall the snapshot-dependent component service during the CSI upgrade.

Scenario Description

- 1. Currently, only Huawei CSI is available in the Kubernetes cluster, and Huawei CSI is no longer used.
- 2. Before the uninstallation, use Huawei CSI to clear VolumeSnapshots in the Kubernetes cluster. For details, see **6.5.3.2 Deleting a VolumeSnapshot**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to uninstall the snapshot-dependent component service.
 - # kubectl delete crd volumesnapshotclasses.snapshot.storage.k8s.io volumesnapshotcontents.snapshot.storage.k8s.io volumesnapshots.snapshot.storage.k8s.io
- **Step 3** Run the following command to check whether the service is successfully uninstalled.

If the command output is empty, the uninstallation is successful. # kubectl get crd | grep snapshot.storage.k8s.io

----End

8.2 Uninstalling the huawei-csi-node Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command:

kubectl delete daemonset huawei-csi-node -n kube-system

Step 3 Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get daemonset huawei-csi-node -n kube-system

----End

8.3 Uninstalling the huawei-csi-controller Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following commands in sequence.

kubectl delete deployment huawei-csi-controller -n kube-system

Step 3 Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get deployment huawei-csi-controller -n kube-system

----End

8.4 Deleting the huawei-csi-configmap Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command.

kubectl delete configmap huawei-csi-configmap -n kube-system

Step 3 Run the following command to check whether the object is successfully deleted. If **NotFound** is displayed, the object is successfully deleted.

kubectl get configmap huawei-csi-configmap -n kube-system

----End

8.5 Deleting the huawei-csi-secret Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command.

kubectl delete secret huawei-csi-secret -n kube-system

Step 3 Run the following command to check whether the object is successfully deleted. If **NotFound** is displayed, the object is successfully deleted.

kubectl get secret huawei-csi-secret -n kube-system

----End

8.6 Deleting the RBAC Permission

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Delete the RBAC permission.
 - If the huawei-csi version is later than 2.2.15, run the following command to delete the permission.

kubectl -n kube-system -l provisioner=csi.huawei.com delete ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding

- If the huawei-csi version is 2.2.15 or earlier, perform the following operations to delete the permission.
 - a. Run the following command to create a file named **delete-huawei-csi-rhac.sh**

cat <<EOF > delete-huawei-csi-rbac.sh kubectl delete ServiceAccount huawei-csi-controller -n kube-system

```
kubectl delete ServiceAccount huawei-csi-node -n kube-system
kubectl delete ClusterRole huawei-csi-attacher-runner -n kube-system
kubectl delete ClusterRole huawei-csi-driver-registrar-runner -n kube-system
kubectl delete ClusterRole huawei-csi-provisioner-runner -n kube-system
kubectl delete ClusterRole huawei-csi-resizer-runner -n kube-system
kubectl delete ClusterRole huawei-csi-snapshotter-runner -n kube-system
kubectl delete ClusterRole snapshot-controller-runner -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-attacher-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-driver-registrar-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-provisioner-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-resizer-role -n kube-system
kubectl delete ClusterRoleBinding huawei-csi-snapshotter-role -n kube-system
kubectl delete ClusterRoleBinding snapshot-controller-role -n kube-system
kubectl delete Role huawei-csi-resizer-cfg -n kube-system
kubectl delete Role huawei-csi-snapshotter-leaderelection -n kube-system
kubectl delete Role snapshot-controller-leaderelection -n kube-system
kubectl delete RoleBinding huawei-csi-resizer-role-cfg -n kube-system
kubectl delete RoleBinding huawei-csi-snapshotter-leaderelection -n kube-system
kubectl delete RoleBinding snapshot-controller-leaderelection -n kube-system
```

b. Run the following command to delete the RBAC permission. If the **NotFound** error is reported, ignore it.

sh delete-huawei-csi-rbac.sh

Step 3 Check whether the RBAC permission has been deleted.

- If the huawei-csi version is later than 2.2.15, run the following command. If
 No resources found is displayed, the permission is successfully deleted.
 # kubectl -n kube-system -l provisioner=csi.huawei.com get
 ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding
- If the huawei-csi version is 2.2.15 or earlier, perform the following operations to check whether the RBAC permission is successfully deleted.
 - a. Run the following command to create a file named **check-huawei-csi-rbac.sh**.

```
# cat <<EOF > check-huawei-csi-rbac.sh kubectl get ServiceAccount -n kube-system | grep huawei-csi kubectl get ClusterRole -n kube-system | grep huawei-csi kubectl get ClusterRoleBinding -n kube-system | grep huawei-csi kubectl get Role -n kube-system | grep huawei-csi kubectl get RoleBinding -n kube-system | grep huawei-csi kubectl get RoleBinding -n kube-system | grep huawei-csi kubectl get ClusterRole snapshot-controller-runner -n kube-system --ignore-not-found=true kubectl get ClusterRoleBinding snapshot-controller-role -n kube-system --ignore-not-found=true kubectl get RoleBinding snapshot-controller-leaderelection -n kube-system --ignore-not-found=true kubectl get RoleBinding snapshot-controller-leaderelection -n kube-system --ignore-not-found=true
```

b. Run the following command. If no command output is displayed, the RBAC permission has been successfully deleted.

sh check-huawei-csi-rbac.sh

----End

8.7 Deleting the Image of the Earlier Version

To delete the **huawei-csi** image from the cluster, you need to perform the deletion operation on all worker nodes.

To delete the image from a single node, perform the following steps.

Prerequisites

The container service that depends on the image has been stopped. Otherwise, the image cannot be deleted.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node through the management IP address.
- **Step 2** Run the following command to view all existing versions.
 - If docker is used, run the **docker image ls | grep huawei-csi** command.

```
# docker image ls | grep huawei-csi

REPOSITORY TAG IMAGE ID CREATED SIZE
huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
huawei-csi 2.2.16 14b854dba227 2 weeks ago 79.6MB
```

• If containerd is used, run the **crictl image ls | grep huawei-csi** command.

# crictl image ls grep huawei-csi					
REPOSITORY	TAG	IMAGE ID	CREATED	SIZE	
docker.io/library/huawei-csi	2.2.15	b30b3a8b5959	2 weeks ago	79.7MB	
docker.io/library/huawei-csi	2.2.16	14b854dba227	2 weeks ago	79.6MB	

- **Step 3** Run the following command to delete the image of the earlier version:
 - If docker is used, run the **docker rmi** *<REPOSITORY>*:*<TAG>* command. # docker rmi huawei-csi:2.2.15
 - If containerd is used, run the **crictl rmi** *<REPOSITORY>*: *<TAG>* command. # crictl rmi huawei-csi:2.2.15
- **Step 4** Run the following command again to check whether the image is successfully deleted. If the target version is not displayed, the image of the version is successfully deleted.
 - If docker is used, run the **docker image ls | grep huawei-csi** command.

 # docker image ls | grep huawei-csi
 huawei-csi
 2.2.16
 14b854dba227
 10 minutes ago
 80MB
 - If containerd is used, run the **crictl image ls | grep huawei-csi** command.

 # crictl image ls | grep huawei-csi
 docker.io/library/huawei-csi
 2.2.16
 14b854dba2273
 93.1MB

----End

9 Common Operations

- 9.1 Updating the User Name or Password of a Storage Device Configured on CSI
- 9.2 Updating the configmap Object of huawei-csi
- 9.3 Adding a Backend for huawei-csi
- 9.4 Updating the huawei-csi-controller Service
- 9.5 Updating the huawei-csi-node Service
- 9.6 Modifying the Log Output Mode

9.1 Updating the User Name or Password of a Storage Device Configured on CSI

When the user name or password of a storage device changes, you need to update the configuration information on CSI. Otherwise, huawei-csi services cannot work properly.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **chmod** +**x secretUpdate** command to grant the execute permission on the secretUpdate tool.

chmod +x secretUpdate

Step 3 Run the ./secretUpdate command to run the secretUpdate tool and enter the ID of the backend to be configured as prompted. If **Configured** is **false**, the backend is not configured. If **Configured** is **true**, the backend is configured.

Step 4 Enter the user name and password as prompted to update the **secret** object.

Step 5 After the configuration is complete, enter **exit** to exit and save the configuration.

```
Please enter the backend number to configure (Enter 'exit' to exit): exit
Saving configuration. Please wait.......
The configuration is saved successfully.
```

- **Step 6** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 7** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 8** Run the **kubectl get pod -A | grep huawei** command to check whether the services are restarted successfully.

```
# kubectl get pod -A | grep huawei
kube-system huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
kube-system huawei-csi-node-g6f7z 3/3 Running 0 14s
```

----End

9.2 Updating the configmap Object of huawei-csi

Perform this operation when you want to add a storage pool to an existing backend or change an existing service IP address.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the kubectl edit configmap huawei-csi-configmap -n kube-system command, press I or Insert to enter the editing mode, and modify related parameters. The iSCSI networking is used as an example. For details about the parameters, see Table 9-1. After the modification is complete, press Esc and enter :wq! to save the modification.

Table 9-1 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json". backends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of back-end storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 9-2.

Table 9-2 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage SAN is connected, the value is fixed to fusionstorage-san .
name	String	Storage backend name.	 User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). This parameter cannot be modified.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https://192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.

Configuration Item	Format	Description	Remarks
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	 One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. Currently, only storage pools can be added. You can log in to DeviceManager to obtain the storage pools that support the block storage service.
parameters	Dictionary	Variable parameters in scenarios where iSCSI is used.	In scenarios where iSCSI is used, set the protocol parameter to a fixed value: iscsi . Set the portals parameter to the iSCSI service IP addresses of the storage device. Use commas (,) to separate multiple iSCSI service IP addresses. You can log in to DeviceManager to obtain the iSCSI service IP addresses. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is iSCSI. (For other series, see the corresponding operation description.)

- Step 3 If the storage, name, or urls parameter is modified, you need to update the user name or password of the storage device. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.
- **Step 4** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 5** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 6** Run the **kubectl get pod -A | grep huawei** command to check whether the services are restarted successfully.

```
# kubectl get pod -A | grep huawei
kube-system huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
kube-system huawei-csi-node-g6f7z 3/3 Running 0 14s
```

----End

9.3 Adding a Backend for huawei-csi

Perform this operation when you want to add a storage device or a storage pool as an independent backend.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Configure multiple backends. For details, see **7.1 Configuring Multiple Backends**.
- Step 3 Configure accounts for the new backends. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.
- **Step 4** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 5** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=kube-system | kubectl replace --force -f -
- **Step 6** Run the **kubectl get pod -A | grep huawei** command to check whether the services are restarted successfully.

```
# kubectl get pod -A | grep huawei
kube-system huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
kube-system huawei-csi-node-g6f7z 3/3 Running 0 14s
```

----End

9.4 Updating the huawei-csi-controller Service

Perform this operation when you need to update the huawei-csi-controller service, for example, adding the snapshot or the capacity expansion function.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-controller service. For details, see **8.3 Uninstalling the** huawei-csi-controller Service.
- **Step 3** Delete the RBAC permission. For details, see **8.6** Deleting the RBAC Permission.
- **Step 4** Create the RBAC permission. For details, see **Step 5**.
- **Step 5** Start the controller service. For details, see **Step 6**.
- **Step 6** After the huawei-csi service is deployed, run the **kubectl get pod -A | grep huawei-csi-controller** command to check whether the service is started.

kubectl get pod -A | grep huawei-csi-controller kube-system huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s

----End

9.5 Updating the huawei-csi-node Service

Perform this operation when you need to update the huawei-csi-node service.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-node service. For details, see **8.2 Uninstalling the** huawei-csi-node Service.
- **Step 3** Install the new huawei-csi-node service. For details, see **4.7 Starting huawei-csi Services**.
- **Step 4** After the huawei-csi service is deployed, run the **kubectl get pod -A | grep huawei-csi-node** command to check whether the service is started.

kubectl get pod -A | grep huawei-csi-node kube-system huawei-csi-node-g6f7z 3/3 Running 0 14s

----End

9.6 Modifying the Log Output Mode

huawei-csi supports two log output modes: **file** and **console**. **file** indicates that logs are output to the fixed directory (**/var/log/huawei**), and **console** indicates that logs are output to the standard directory of the container. You can set the log output mode as required. The default mode is **file**.

9.6.1 Modifying the Log Output Mode of the huawei-csi-controller Service

Perform this operation when you want to set the log output mode of the huawei-csi-controller service.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-controller service. For details, see **8.3 Uninstalling the** huawei-csi-controller Service.
- **Step 3** Check the Kubernetes version. Go to the corresponding directory according to the Kubernetes version. (The example directory is **deploy/v1.20-v1.23**.)

kubectl version --short=true Client Version: v1.22.1 Server Version: v1.22.1 **Step 4** Run the **vi** *huawei-csi-controller.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify the following parameters. For details, see **Table 9-3**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

args: - "--

- "--endpoint=\$(CSI_ENDPOINT)"
- "--controller"
- "--containerized"
- "--driver-name=csi.huawei.com"
- "--loggingModule=file"
- "--logLevel=info"
- "--logFileDir=/var/log/huawei"
- "--logFileSize=20M"
- "--maxBackups=9"

Table 9-3 Description of log output parameters

Configuration Item	Description	Remarks
loggingModule	huawei-csi log output mode.	The value can be file or console . The default value is file .
logLevel	huawei-csi log output level.	Supported levels are debug , info , warning , error , and fatal . The default level is info .
logFileDir	huawei-csi log directory in file output mode.	This parameter is available only when loggingModule is set to file . The default log directory is /var/log/huawei .
logFileSize	Size of a single huawei-csi log file in file output mode.	This parameter is available only when loggingModule is set to file . The default log file size is 20 MiB.
maxBackups	Maximum number of huawei-csi log file backups in file output mode.	This parameter is available only when loggingModule is set to file . The default number of log file backups is 9.

Step 5 Run the following command to start the controller service.

kubectl create -f huawei-csi-controller-snapshot-v1.yaml

Step 6 After the huawei-csi service is deployed, run the **kubectl get pod -A -o wide** | **grep huawei** command to check whether the service is started.

```
# kubectl get pod -A -o wide | grep huawei
kube-system huawei-csi-controller-b59577886-qqzm8 7/7 Running 0 18h 10.244.1.67
node <none> <none>
```

- **Step 7** View the logs of the huawei-csi-controller service.
 - If **loggingModule** is set to **file**, log in to the node, go to the log directory specified by **logFileDir**, and run the following command to view the log of huawei-csi-controller.

tail -f huawei-csi-controller

• If **loggingModule** is set to **console**, run the following command to view the log of huawei-csi-controller.

kubectl logs *huawei-csi-controller* -c huawei-csi-driver -n kube-system

----End

9.6.2 Modifying the Log Output Mode of the huawei-csi-node Service

Perform this operation when you want to set the log output mode of the huawei-csi-node service.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-node service. For details, see **8.2 Uninstalling the** huawei-csi-node Service.
- Step 3 Run the vi huawei-csi-node.yaml command to modify the .yaml file. Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press Esc and enter :wq! to save the modification. Compile the huawei-csi-node.yaml file. For details, see sample file deploy/huawei-csi-node.yaml in the software package. For details about the parameters, see Table 9-4.

args:

- "--endpoint=/csi/csi.sock"
- "--containerized"
- "--driver-name=csi.huawei.com"
- "--volume-use-multipath=false"
- "--loggingModule=file"
- "--logLevel=info"
- "--logFileDir=/var/log/huawei"
- "--logFileSize=20M"
- "--maxBackups=9"

Table 9-4 Description of log output parameters

Configuration Item	Description	Remarks
loggingModule	huawei-csi log output mode.	The value can be file or console . The default value is file .
logLevel	huawei-csi log output level.	Supported levels are debug , info , warning , error , and fatal . The default level is info .
logFileDir	huawei-csi log directory in file output mode.	This parameter is available only when loggingModule is set to file . The default log directory is /var/log/huawei .
logFileSize	Size of a single huawei-csi log file in file output mode.	This parameter is available only when loggingModule is set to file . The default log file size is 20 MiB.

Configuration Item	Description	Remarks
maxBackups	Maximum number of huawei-csi log file backups in file output mode.	This parameter is available only when loggingModule is set to file . The default number of log file backups is 9.

Step 4 Run the following command to start the node service.

kubectl create -f huawei-csi-node.yaml

Step 5 After the huawei-csi service is deployed, run the **kubectl get pod -A -o wide | grep huawei-csi-node** command to check whether the service is started.

```
# kubectl get pod -A | grep huawei-csi-node kube-system huawei-csi-node-4sfwr 3/3 Running 0 18h 10.244.1.68 node <none> <none>
```

- **Step 6** View the logs of the huawei-csi-node service.
 - If **loggingModule** is set to **file**, log in to the node, go to the log directory specified by **logFileDir**, and run the following command to view the log of huawei-csi-node.

tail -f huawei-csi-node

 If loggingModule is set to console, run the following command to view the log of huawei-csi-node.

kubectl logs *huawei-csi-node* -c huawei-csi-driver -n kube-system

----End

10_{FAQ}

10.1 Viewing Log Information

10.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernetes Platform Is Set Up for the First Time

10.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported

10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters

10.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName

10.6 When a PVC Is Created, the PVC Is in the Pending State

10.7 Before a PVC Is Deleted, the PVC Is in the Pending State

10.8 When a Pod Is Created, the Pod Is in the ContainerCreating State

10.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created

10.1 Viewing Log Information

Viewing Logs Generated When the secret Object Is Configured

Step 1 Run the **cd /var/log/huawei** command to go to the log directory.

cd /var/log/huawei

Step 2 Run the following command to view the logs of huawei-csi-install.

vi huawei-csi-install

----Fnd

Viewing Logs of the huawei-csi-controller Service

Step 1 Run the following command to obtain the node where huawei-csi-controller is located.

```
# kubectl get pod -A -o wide | grep huawei
kube-system huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s <host1-ip>
<host1-name> <none> <none>
```

- **Step 2** Use a remote access tool, such as PuTTY, to log in to the huawei-csi-controller node in the Kubernetes cluster through the management IP address.
- **Step 3** Run the **cd /var/log/huawei** command to go to the log directory.

 # cd /var/log/huawei
- **Step 4** Run the following command to view the customized output logs of the container.

 # vi huawei-csi-controller
- **Step 5** Run the **cd /var/log/containers** command to go to the container directory. # cd /var/log/containers
- **Step 6** Run the following command to view the standard output logs of the container.

 # vi huawei-csi-controller-<name>_kube-system_huawei-csi-driver-<contrainer-id>.log

----End

Viewing Logs of the huawei-csi-node Service

- **Step 1** Run the following command to obtain the node where huawei-csi-node is located.

 # kubectl get pod -A -o wide | grep huawei kube-system huawei-csi-node-g6f7z 3/3 Running 0 14s <host2-ip> <host2-name> <none>
- **Step 2** Use a remote access tool, such as PuTTY, to log in to the huawei-csi-node node in the Kubernetes cluster through the management IP address.
- **Step 3** Run the **cd /var/log/huawei** command to go to the log directory.

 # cd /var/log/huawei
- **Step 4** Run the following command to view the customized output logs of the container.

 # vi huawei-csi-node
- **Step 5** Run the **cd /var/log/containers** command to go to the container directory.

 # cd /var/log/containers
- **Step 6** Run the following command to view the standard output logs of the container. # vi huawei-csi-node-<name>_kube-system_huawei-csi-driver-<contrainer-id>.log

----End

10.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernetes Platform Is Set Up for the First Time

Symptom

When you create a Pod, error Cannot connect ISCSI portal *.*.*: libkmod: kmod_module_insert_module: could not find module by name='iscsi_tcp' is reported in the /var/log/huawei-csi-node log.

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

The iscsi_tcp service may be stopped after the Kubernetes platform is set up and the iscsi service is installed. You can run the **lsmod | grep iscsi | grep iscsi_tcp** command to check whether the service is stopped.

```
# Ismod | grep iscsi | grep iscsi_tcp
iscsi_tcp 18333 6
libiscsi_tcp 25146 1 iscsi_tcp
libiscsi 57233 2 libiscsi_tcp,iscsi_tcp
scsi_transport_iscsi 99909 3 iscsi_tcp,libiscsi
```

Solution or Workaround

Run the following command to manually load the iscsi_tcp service.

```
# modprobe iscsi_tcp
# lsmod | grep iscsi | grep iscsi_tcp
iscsi_tcp 18333 6
libiscsi_tcp 25146 1 iscsi_tcp
```

10.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported

Symptom

The huawei-csi-node service cannot be started. When you run the **kubectl describe daemonset huawei-csi-node -n kube-system** command, error message "/var/lib/iscsi is not a directory" is reported.

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

The /var/lib/iscsi directory does not exist in the huawei-csi-node container.

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the following command to delete the huawei-csi-node service (huawei-csi-node.yaml is the configuration file in Step 7 in 4.7 Starting huawei-csi Services).

 # kubectl delete -f huawei-csi-node.yaml
- **Step 3** Run the following command to view the huawei-csi-node service. If no command output is displayed, the deletion is complete.

kubectl get pod -A | grep huawei-csi-node

- Step 4 Run the vi huawei-csi-node.yaml command to modify the .yaml file. Press I or Insert to enter the editing mode, set path in huawei-csi-node.yaml > volumes > iscsi-dir > hostPath to /var/lib/iscsi and delete the type line. After the modification is complete, press Esc and enter :wq! to save the modification. Compile the huawei-csi-node.yaml file. For details, see sample file deploy/huawei-csi-node.yaml in the software package.
- **Step 5** Run the following command to start the node service.

kubectl create -f huawei-csi-node.yaml

Step 6 After the huawei-csi service is deployed, run the **kubectl get pod -A | grep huawei-csi-node** command to check whether the service is started.

kubectl get pod -A | grep huawei-csi-node kube-system huawei-csi-node-g6f7z 3/3 Running 0 14s

----End

10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters

Symptom

A Pod is running on worker node A, and an external block device is mounted to the Pod through CSI. After worker node A is powered off abnormally, the Kubernetes platform detects that the node is faulty and switches the Pod to worker node B. After worker node A recovers, the drive letters on worker node A change from normal to faulty.

Environment Configuration

Kubernetes version: 1.18 or later

Storage type: block storage

Root Cause Analysis

After worker node A recovers, Kubernetes initiates an unmapping operation on the storage, but does not initiate a drive letter removal operation on the host. After Kubernetes completes the unmapping, residual drive letters exist on worker node A.

Solution or Workaround

Currently, you can only manually clear the residual drive letters on the host. Alternatively, restart the host again and use the disk scanning mechanism during the host restart to clear the residual drive letters. The specific method is as follows:

Step 1 Check the residual drive letters on the host.

1. Run the **multipath -ll** command to check whether a DM multipathing device with abnormal multipathing status exists.

As shown in the following figure, the path status is **failed faulty running**, the corresponding DM multipathing device is **dm-12**, and the associated SCSI disks are **sdi** and **sdj**. If multiple paths are configured, multiple SCSI disks exist. Record these SCSI disks.

- If yes, go to **Step 1.2**.
- If no, no further action is required.
- 2. Check whether the residual DM multipathing device is readable.

Run the **dd if=/dev/***dm-xx* **of=/dev/null count=1 bs=1M iflag=direct** command.

dm-xx indicates the device ID obtained in Step 1.1.

If the returned result is **Input/output error** and the read data is **0 bytes (0 B) copied**, the device is unreadable.

```
#dd if=/dev/dm-12 of=/dev/null count=1 bs=1M iflag=direct
dd: error reading '/dev/dm-12': Input/output error
0+0 records in
0+0 records out
0 bytes (0 B) copied, 0.0236862 s, 0.0 kB/s
```

- If yes, record the residual dm-xx device and associated disk IDs (for details, see Step 1.1) and perform the clearing operation.
- If the command execution is suspended, go to **Step 1.3**.
- If other cases, contact technical support engineers.
- 3. Log in to the node again in another window.
 - a. Run the following command to view the suspended process.

 # ps -ef | grep dm-12 | grep -w dd

 root 21725 9748 0 10:33 pts/10 00:00:00 dd if=/dev/dm-12 of=/dev/null count=1 bs=10M iflag=direct
 - b. Kill the pid. # kill -9 *pid*
 - c. Record the residual *dm-xx* device and associated disk IDs (for details, see **Step 1.1**) and perform the clearing operation.

Step 2 Clear the residual drive letters on the host.

1. Run the **multipath** -**f** /**dev**/*dm*-* command to delete residual multipathing aggregation device information according to the DM multipathing device obtained in **Step 1**.

```
# multipath -f /dev/dm-12
```

If an error is reported, contact technical support engineers.

2. Run the following command to clear the residual SCSI disks according to the drive letters of the residual disks obtained in the troubleshooting method. echo 1 > /sys/block/xxxx/device/delete

When multiple paths are configured, clear the residual disks based on the drive letters. The residual paths are **sdi** and **sdj**.

```
# echo 1 > /sys/block/sdi/device/delete
# echo 1 > /sys/block/sdj/device/delete
```

If an error is reported, contact technical support engineers.

3. Check whether the DM multipathing device and SCSI disk information has been cleared.

Run the **multipath** -**ll**, **ls** -**l** /**sys/block**/, and **ls** -**l** /**dev/disk/by-id**/ commands in sequence to query the path and disk information. If the residual **dm-12** device and SCSI disks **sdi** and **sdj** are cleared, the clearing is complete.

```
# multipath -ll
mpathb (3618cf24100f8f457014a764c000001f6) dm-3 HUAWEI ,XSG1
size=100G features='0' hwhandler='0' wp=rw
-+- policy='service-time 0' prio=-1 status=active
|- 39:0:0:1
               sdd 8:48 active ready running
 `- 38:0:0:1
               sde 8:64 active ready running
mpathn (3618cf24100f8f457315a764c000001f6) dm-5 HUAWEI ,XSG1
size=100G features='0' hwhandler='0' wp=rw
-+- policy='service-time 0' prio=-1 status=active
|- 39:0:0:2

`- 38:0:0:2
               sdc 8:32 active ready running
               sdb 8:16 active ready running
# ls -l /sys/block/
total 0
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-0 -> ../devices/virtual/block/dm-0
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-1 -> ../devices/virtual/block/dm-1
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-2 -> ../devices/virtual/block/dm-2
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-3 -> ../devices/virtual/block/dm-3
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdb -> ../devices/platform/host35/session2/
target35:0:0/35:0:0:1/block/sdb
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdc -> ../devices/platform/host34/
target34:65535:5692/34:65535:5692:0/block/sdc
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdd -> ../devices/platform/host39/session6/
target39:0:0/39:0:0:1/block/sdd
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sde -> ../devices/platform/host38/session5/
target38:0:0/38:0:0:1/block/sde
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdh -> ../devices/platform/host39/session6/
target39:0:0/39:0:0:3/block/sdh
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdi -> ../devices/platform/host38/session5/target38:0:0/38:0:0:3/
block/sdi
ls -l /dev/disk/by-id/
total 0
lrwxrwxrwx 1 root root 10 Aug 11 19:57 dm-name-mpathb -> ../../dm-3
lrwxrwxrwx 1 root root 10 Aug 11 19:58 dm-name-mpathn -> ../../dm-5
lrwxrwxrwx 1 root root 10 Aug 11 19:57 dm-uuid-mpath-3618cf24100f8f457014a764c000001f6 -> ../../
lrwxrwxrwx 1 root root 10 Aug 11 19:58 dm-uuid-mpath-3618cf24100f8f457315a764c000001f6 -> ../../
dm-5
lrwxrwxrwx 1 root root 9 Aug 11 19:57 scsi-3618cf24100f8f457014a764c000001f6 -> ../../sdd
lrwxrwxrwx 1 root root 9 Aug 11 19:57 scsi-3618cf24100f8f45712345678000103e8 -> ../../sdi
lrwxrwxrwx 1 root root 9 Aug 3 15:17 scsi-3648435a10058805278654321ffffffff -> ../../sdb
lrwxrwxrwx 1 root root 9 Aug 2 14:49 scsi-36888603000020aff44cc0d060c987f1 -> ../../sdc
lrwxrwxrwx 1 root root 9 Aug 11 19:57 wwn-0x618cf24100f8f457014a764c000001f6 -> ../../sdd
lrwxrwxrwx 1 root root 9 Aug 11 19:57 wwn-0x618cf24100f8f45712345678000103e8 -> ../../sdi
lrwxrwxrwx 1 root root 9 Aug 3 15:17 wwn-0x648435a10058805278654321ffffffff -> ../../sdb
lrwxrwxrwx 1 root root 9 Aug 2 14:49 wwn-0x68886030000020aff44cc0d060c987f1 -> ../../sdc
```

----End

10.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName

Symptom

The huawei-csi services (huawei-csi-controller or huawei-csi-node) cannot be started. After the **kubectl get pod -A | grep huawei** command is executed, the command output shows that the service status is **InvalidImageName**.

```
# kubectl get pod -A | grep huawei kube-system huawei-csi-controller-fd5f97768-qlldc 6/7 InvalidImageName 0 16s kube-system huawei-csi-node-25txd 2/3 InvalidImageName 0 15s
```

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

In the .yaml configuration files of the controller and node, the Huawei CSI image version number is incorrect. For example:

```
...
- name: huawei-csi-driver
image: huawei-csi:2.2.16
...
```

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to modify the configuration file of the huawei-csi-node service. Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

kubectl edit daemonset huawei-csi-node -o yaml -n=kube-system

◯ NOTE

• In the **image** configuration item under **huawei-csi-driver** in the sample .yaml file, huawei-csi:*:**must be replaced with <Name>:<Version> of the created Huawei CSI image.

```
containers:
```

- name: huawei-csi-driver image: huawei-csi:2.2.16
- **Step 3** Run the following command to modify the configuration file of the huawei-csi-controller service: Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

kubectl edit deployment huawei-csi-controller -o yaml -n=kube-system

 In the image configuration item under huawei-csi-driver in the sample .yaml file, huawei-csi: *:** must be replaced with <Name>:<Version> of the created Huawei CSI image.

containers:

 name: huawei-csi-driver image: huawei-csi:2.2.16

- **Step 4** Wait until the huawei-csi-node and huawei-csi-controller services are started.
- **Step 5** Run the following command to check whether the huawei-csi services are started.

```
# kubectl get pod -A | grep huawei
kube-system huawei-csi-controller-58799449cf-zvhmv 7/7 Running 0 2m29s
kube-system huawei-csi-node-7fxh6 3/3 Running 0 12m
```

----End

10.6 When a PVC Is Created, the PVC Is in the Pending State

Symptom

A PVC is created. After a period of time, the PVC is still in the **Pending** state.

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

Cause 1: A StorageClass with the specified name is not created in advance. As a result, Kubernetes cannot find the specified StorageClass name when a PVC is created.

Cause 2: The storage pool capability does not match the StorageClass capability. As a result, huawei-csi fails to select a storage pool.

Cause 3: An error code (for example, 50331651) is returned by a RESTful interface of the storage. As a result, huawei-csi fails to create a PVC.

Cause 4: The storage does not return a response within the timeout period set by huawei-csi. As a result, huawei-csi returns a timeout error to Kubernetes.

Cause 5: Other causes.

Solution or Workaround

When a PVC is created, if the PVC is in the **Pending** state, you need to take different measures according to the following causes.

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- Step 2 Run the following command to view details about the PVC.

 # kubectl describe pvc mypvc
- **Step 3** Perform the corresponding operation according to the **Events** information in the detailed PVC information.
 - If the PVC is in the **Pending** state due to cause 1, perform the following steps.

 Events:

 Type Reason Age From Message

Warning ProvisioningFailed 0s (x15 over 3m24s) persistentvolume-controller storageclass.storage.k8s.io "mysc" not found

- a. Delete the PVC. For details, see 6.3.5 Deleting a PVC.
- b. Create a StorageClass. For details, see **6.1.1 Creating a StorageClass**.
- c. Create a PVC. For details, see 6.3.1 Creating a PVC.
- If the PVC is in the **Pending** state due to cause 2, perform the following steps.

Events:			
Type From	Reason	Age	Message

Normal Provisioning 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = **failed to select pool**, the capability filter failed, error: failed to select pool, the final filter field: **replication**, parameters map[allocType:thin replication:True size: 1099511627776 volumeType:lun]. please check your storage class

- Delete the PVC. For details, see 6.3.5 Deleting a PVC.
- b. Delete the StorageClass. For details, see 6.1.2 Deleting a StorageClass.
- c. Modify the **StorageClass.yaml** file based on the **Events** information.
- d. Create a StorageClass. For details, see 6.1.1 Creating a StorageClass.
- e. Create a PVC. For details, see **6.3.1 Creating a PVC**.
- If the PVC is in the **Pending** state due to cause 3, contact Huawei engineers.

Events:				
Type From	Reason	Age		Message
				wessage

Normal Provisioning 63s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 62s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = Create volume map[ALLOCTYPE:1 CAPACITY:20 DESCRIPTION:Created from Kubernetes CSI NAME:pvc-63ebfda5-4cf0-458e-83bd-ecc PARENTID:0] error: 50331651

• If the PVC is in the **Pending** state due to cause 4, perform the following steps.

Type From	Reason	Age				Mes	ssage	
							3	
			, ,	\				

Normal Provisioning 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = context deadline exceeded (Client.Timeout exceeded while awaiting headers)

- a. Wait for 10 minutes and check the PVC details again by referring to this section
- b. If it is still in the **Pending** state, contact Huawei engineers.
- If the PVC is in the **Pending** state due to cause 5, contact Huawei engineers.

----End

10.7 Before a PVC Is Deleted, the PVC Is in the Pending State

Symptom

Before a PVC is deleted, the PVC is in the **Pending** state.

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

Cause 1: A StorageClass with the specified name is not created in advance. As a result, Kubernetes cannot find the specified StorageClass name when a PVC is created.

Cause 2: The storage pool capability does not match the StorageClass capability. As a result, huawei-csi fails to select a storage pool.

Cause 3: An error code (for example, 50331651) is returned by a RESTful interface of the storage. As a result, huawei-csi fails to create a PVC.

Cause 4: The storage does not return a response within the timeout period set by huawei-csi. As a result, huawei-csi returns a timeout error to Kubernetes.

Cause 5: Other causes.

Solution or Workaround

To delete a PVC in the **Pending** state, you need to take different measures according to the following causes.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to view details about the PVC. # kubectl describe pvc mypvc
- **Step 3** Perform the corresponding operation according to the **Events** information in the detailed PVC information.
 - If the PVC is in the **Pending** state due to cause 1, run the **kubectl delete pvc** *mypvc* command to delete the PVC.

Events:					
Type	Reason	Age	From	Message	

Warning ProvisioningFailed 0s (x15 over 3m24s) persistentvolume-controller storageclass.storage.k8s.io "mysc" not found

If the PVC is in the **Pending** state due to cause 2, run the **kubectl delete pvc** mypvc command to delete the PVC.

Events:
Type Reason Age
From Message

Normal Provisioning 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = **failed to select pool**, the capability filter failed, error: failed to select pool, the final filter field: *replication*, parameters map[allocType:thin replication:True size: 1099511627776 volumeType:lun]. please check your storage class

• If the PVC is in the **Pending** state due to cause 3, run the **kubectl delete pvc** *mypvc* command to delete the PVC.

Events:				
Type	Reason	Age		
From			N	Лessage

Normal Provisioning 63s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 62s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = Create volume map[ALLOCTYPE:1 CAPACITY:20 DESCRIPTION:Created from Kubernetes CSI NAME:pvc-63ebfda5-4cf0-458e-83bd-ecc PARENTID:0] error: 50331651

• If the PVC is in the **Pending** state due to cause 4, contact Huawei engineers.

Events: Type	Reason	Age	
From	Reason	Age	Message

Normal Provisioning 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = context deadline exceeded (Client.Timeout exceeded while awaiting headers)

If the PVC is in the **Pending** state due to cause 5, contact Huawei engineers.

----End

10.8 When a Pod Is Created, the Pod Is in the ContainerCreating State

Symptom

A Pod is created. After a period of time, the Pod is still in the **ContainerCreating** state. Check the log information (for details, see **10.1 Viewing Log Information**). The error message "Fibre Channel volume device not found" is displayed.

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

This problem occurs because residual disks exist on the host node. As a result, disks fail to be found when a Pod is created next time.

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query information about the node where the Pod resides.

```
# kubectl get pod -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE
READINESS GATES
mypod 0/1 ContainerCreating 0 51s 10.244.1.224 node1 <none>
```

- **Step 3** Delete the Pod. For details, see **6.4.2 Deleting a Pod**.
- **Step 4** Use a remote access tool, such as PuTTY, to log in to the *node1* node in the Kubernetes cluster through the management IP address. *node1* indicates the node queried in **Step 2**.
- **Step 5** Clear the residual drive letters. For details, see **Solution or Workaround**.

----End

10.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created

Symptom

When a Pod is being created, the Pod is in the **ContainerCreating** state for a long time. Check the huawei-csi-node log (for details, see **10.1 Viewing Log Information**). No Pod creation information is recorded in the huawei-csi-node log. After the **kubectl get volumeattachment** command is executed, the name of the PV used by the Pod is not displayed in the **PV** column. After a long period of time (more than ten minutes), the Pod is normally created and the Pod status changes to **Running**.

Environment Configuration

Kubernetes version: 1.18 or later

Root Cause Analysis

The kube-controller-manager component of Kubernetes is abnormal.

Solution or Workaround

Contact container platform engineers to rectify the fault.

11 Appendix

- 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3
- 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x
- 11.3 Example ALUA Configuration Policy of Distributed Storage

11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3

Example 1: The configuration file content is as follows:

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.1**Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3, the priority of the second configuration section (where HostName is node1) is higher than that of the first configuration section (where HostName is *).

Example 2: The configuration file content is as follows:

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

According to the configuration policy rules in **7.4.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. A matches the beginning of a character string, and \$ matches the end of a character string.

11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x

Example 1: The configuration file content is as follows:

```
...
"parameters": {..., "ALUA": {
    "*": {"accessMode": "1", "hyperMetroPathOptimized": "1"},
    "node1": {"accessMode": "1", "hyperMetroPathOptimized": "0"}}}
...
```

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2 Configuring ALUA for OceanStor Dorado 6.x**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2 Configuring ALUA for OceanStor Dorado 6.x**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

According to the configuration policy rules in **7.4.2 Configuring ALUA for OceanStor Dorado 6.x**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.

11.3 Example ALUA Configuration Policy of Distributed Storage

Example 1: The configuration file content is as follows:

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.3**Configuring ALUA for Distributed Storage, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.3 Configuring ALUA for Distributed Storage**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

According to the configuration policy rules in **7.4.3 Configuring ALUA for Distributed Storage**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.