eSDK Enterprise Storage Plugins V3.0.0

User Guide (Kubernetes CSI)

Issue 01

Date 2022-07-15





Copyright © Huawei Technologies Co., Ltd. 2022. All rights reserved.

No part of this document may be reproduced or transmitted in any form or by any means without prior written consent of Huawei Technologies Co., Ltd.

Trademarks and Permissions

HUAWEI and other Huawei trademarks are trademarks of Huawei Technologies Co., Ltd. All other trademarks and trade names mentioned in this document are the property of their respective holders.

Notice

The purchased products, services and features are stipulated by the contract made between Huawei and the customer. All or part of the products, services and features described in this document may not be within the purchase scope or the usage scope. Unless otherwise specified in the contract, all statements, information, and recommendations in this document are provided "AS IS" without warranties, guarantees or representations of any kind, either express or implied.

The information in this document is subject to change without notice. Every effort has been made in the preparation of this document to ensure accuracy of the contents, but all statements, information, and recommendations in this document do not constitute a warranty of any kind, express or implied.

Huawei Technologies Co., Ltd.

Address: Huawei Industrial Base

Bantian, Longgang Shenzhen 518129

People's Republic of China

Website: https://e.huawei.com

About This Document

Intended Audience

This document is intended for:

- Technical support engineers
- O&M engineers
- Engineers with basic knowledge of storage and Kubernetes

Symbol Conventions

The symbols that may be found in this document are defined as follows.

Symbol	Description
▲ DANGER	Indicates a hazard with a high level of risk which, if not avoided, will result in death or serious injury.
⚠ WARNING	Indicates a hazard with a medium level of risk which, if not avoided, could result in death or serious injury.
⚠ CAUTION	Indicates a hazard with a low level of risk which, if not avoided, could result in minor or moderate injury.
NOTICE	Indicates a potentially hazardous situation which, if not avoided, could result in equipment damage, data loss, performance deterioration, or unanticipated results. NOTICE is used to address practices not related to personal injury.
□ NOTE	Supplements the important information in the main text. NOTE is used to address information not related to personal injury, equipment damage, and environment deterioration.

Change History

Issue	Date	Description	
01	2022-07-15	This issue is the first official release.	

Contents

About This Document	ii
1 Overview	1
2 Features and Compatibility	2
2.1 Compatibility	
2.2 Kubernetes Features	4
2.3 Storage Features	6
3 Installation Preparations	8
3.1 Prerequisites	8
3.2 Obtaining the Software Package	8
3.3 Uploading the Components in the Software Package	9
3.4 Creating a Huawei CSI Image	9
3.5 Uploading a Huawei CSI Image	10
3.6 Configuring Host Multipathing	
3.6.1 Installing Native Multipathing Software	12
3.6.1.1 Installing the Multipathing Tool Package	12
3.6.1.2 Configuring the Multipathing Service	13
3.6.2 Installing Huawei Multipathing Software	
3.6.2.1 Installing the Huawei Multipathing Tool	15
4 Installing Huawei CSI	16
4.1 Installing Huawei CSI Using Helm	16
4.1.1 Installing Helm	16
4.1.2 Installing Huawei CSI	17
4.2 Manually Installing Huawei CSI	29
4.2.1 Connecting to Enterprise Storage	
4.2.1.1 Connecting to Enterprise Storage SAN over iSCSI	
4.2.1.2 Connecting to Enterprise Storage SAN over FC	
4.2.1.3 Connecting to Enterprise Storage NAS over NFS	
4.2.1.4 Connecting to Enterprise Storage SAN over NVMe over RoCE	
4.2.1.5 Connecting to Enterprise Storage SAN over NVMe over FC	
4.2.2 Connecting to Distributed Storage	
4.2.2.1 Connecting to Distributed Storage SAN over SCSI	
4.2.2.2 Connecting to Distributed Storage SAN over iSCSI	49

4.2.2.3 Connecting to Distributed Storage NAS over NFS	52
4.2.2.4 Connecting to Distributed Storage NAS over DPC	55
4.2.3 Starting huawei-csi Services	57
5 Upgrade Operations	64
5.1 Upgrading Huawei CSI Using Helm	64
5.1.1 Upgrading CSI	64
5.1.2 Rolling Back CSI	65
5.2 Manually Upgrading Huawei CSI	66
5.2.1 Uninstalling Original CSI	66
5.2.1.1 Uninstalling the huawei-csi-node Service	66
5.2.1.2 Uninstalling the huawei-csi-controller Service	67
5.2.1.3 Deleting the huawei-csi-configmap Object	67
5.2.1.4 Deleting the huawei-csi-secret Object	67
5.2.1.5 Deleting the RBAC Permission	68
5.2.1.6 Deleting the Image of the Earlier Version	69
5.2.2 Installing New CSI	70
6 Instructions for Use	71
6.1 (Conditionally Mandatory) Managing a StorageClass	
6.1.1 Creating a StorageClass	71
6.1.1.1 Creating a LUN StorageClass	
6.1.1.2 Creating a File System StorageClass	
6.1.2 Deleting a StorageClass	
6.2 (Conditionally Mandatory) Managing a PV	76
6.2.1 Creating a PV	76
6.2.2 Deleting a PV	80
6.3 Managing a PVC	80
6.3.1 Creating a PVC	81
6.3.2 (Optional) Expanding the Capacity of a PVC	85
6.3.3 (Optional) Cloning a PVC	87
6.3.4 (Optional) Creating a PVC Using a Snapshot	89
6.3.5 Deleting a PVC	92
6.4 Managing a Pod	92
6.4.1 Creating a Pod	92
6.4.2 Deleting a Pod	95
6.5 (Optional) Managing a Snapshot	95
6.5.1 Installing the Snapshot-Dependent Component Service	95
6.5.2 Managing a VolumeSnapshotClass	96
6.5.2.1 Creating a VolumeSnapshotClass	96
6.5.2.2 Deleting a VolumeSnapshotClass	98
6.5.3 Managing a VolumeSnapshot	98
6.5.3.1 Creating a VolumeSnapshot	98
6.5.3.2 Deleting a VolumeSnapshot	90

7 Advanced Features	101
7.1 Configuring Multiple Backends	101
7.1.1 Configuring Multiple Backends Using Helm	101
7.1.2 Manually Configuring Multiple Backends	102
7.2 Creating a PVC for a Specified Backend	103
7.3 Creating a PVC for a Specified Storage Pool	104
7.4 Configuring ALUA	
7.4.1 Configuring ALUA Using Helm	105
7.4.1.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3 Using Helm	105
7.4.1.2 Configuring ALUA for OceanStor Dorado 6.x Using Helm	108
7.4.1.3 Configuring ALUA for Distributed Storage Using Helm	111
7.4.2 Manually Configuring ALUA	113
7.4.2.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3	114
7.4.2.2 Configuring ALUA for OceanStor Dorado 6.x	117
7.4.2.3 Configuring ALUA for Distributed Storage	120
7.5 Configuring Storage Topology Awareness	122
7.5.1 Configuring Storage Topology Awareness Using Helm	123
7.5.2 Manually Configuring Storage Topology Awareness	126
7.6 Advanced Features of Enterprise Storage	128
7.6.1 Configuring QoS	128
7.6.2 Configuring a vStore	131
7.6.2.1 Configuring a vStore Using Helm	131
7.6.2.2 Manually Configuring a vStore	
7.6.3 Configuring NAS HyperMetro	
7.6.3.1 Prerequisites	134
7.6.3.2 Configuring NAS HyperMetro Using Helm	136
7.6.3.3 Manually Configuring NAS HyperMetro	138
7.6.4 Configuring an Application Type	
7.7 Advanced Features of Distributed Storage	
7.7.1 Configuring QoS	
7.7.2 Configuring a Soft Quota	
7.7.3 Configuring an Account	
7.7.3.1 Configuring an Account Using Helm	
7.7.3.2 Manually Configuring an Account	
7.7.4 Configuring Mount Parameters in the DPC Scenario	147
8 Uninstalling CSI	150
8.1 Uninstalling huawei-csi Using Helm	
8.2 Manually Uninstalling huawei-csi	151
8.2.1 Uninstalling the huawei-csi-node Service	151
8.2.2 Uninstalling the huawei-csi-controller Service	151
8.2.3 Deleting the huawei-csi-configmap Object	151
8.2.4 Deleting the huawei-csi-secret Object	152

8.2.5 Deleting the RBAC Permission	152
8.2.6 Deleting the Image of the Earlier Version	153
8.3 (Optional) Uninstalling the Snapshot-Dependent Component Service	154
9 Common Operations	. 156
9.1 Updating the User Name or Password of a Storage Device Configured on CSI	156
9.2 Updating the configmap Object of huawei-csi	157
9.2.1 Updating the configmap Object Using Helm	157
9.2.2 Manually Updating the configmap Object	158
9.3 Adding a Backend for huawei-csi	161
9.3.1 Adding a Backend Using Helm	161
9.3.2 Manually Adding a Backend	162
9.4 Updating the huawei-csi-controller Service	162
9.4.1 Updating the controller Service Using Helm	162
9.4.2 Manually Updating the controller Service	164
9.5 Updating the huawei-csi-node Service	164
9.5.1 Updating the node Service Using Helm	164
9.5.2 Manually Updating the node Service	165
9.6 Modifying the Log Output Mode	166
9.6.1 Modifying the Log Output Mode of the controller or node Service Using Helm	166
9.6.2 Manually Modifying the Log Output Mode of the huawei-csi-controller Service	167
9.6.3 Manually Modifying the Log Output Mode of the huawei-csi-node Service	168
9.7 Enabling the ReadWriteOncePod Feature Gate	170
9.8 Configuring Access to the Kubernetes Cluster as a Non-root User	171
10 FAQ	.172
10.1 Viewing Log Information	172
10.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernete Platform Is Set Up for the First Time	
10.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported	174
10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters	
10.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName	177
10.6 When a PVC Is Created, the PVC Is in the Pending State	179
10.7 Before a PVC Is Deleted, the PVC Is in the Pending State	180
10.8 When a Pod Is Created, the Pod Is in the ContainerCreating State	182
10.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created	183
11 Appendix	. 184
11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3	
11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x	
11.3 Example ALUA Configuration Policy of Distributed Storage	186

1 Overview

This document describes how to deploy and use the Kubernetes CSI plug-in so that Huawei enterprise and distributed storage devices provide persistent volume storage capabilities for Kubernetes.

2 Features and Compatibility

This chapter describes the features and compatibility of Huawei CSI.

- 2.1 Compatibility
- 2.2 Kubernetes Features
- 2.3 Storage Features

2.1 Compatibility

Table 2-1 Supported Huawei storage products

Storage Product	Version
OceanStor Dorado V6	6.0.1, 6.1.0, 6.1.2, 6.1.3
OceanStor Dorado V3	V300R002
OceanStor V6	6.1.3
OceanStor V5/F V5	V500R007, V500R007 Kunpeng
OceanStor V3/F V3	V300R006
FusionStorage	V100R006C30
FusionStorage Block	8.0.0, 8.0.1
OceanStor Pacific	8.1.0, 8.1.1, 8.1.2

Table 2-2 Supported container platforms and operating systems (OSs)

Container platform/OS	Version
Kubernetes	1.21, 1.22, 1.23, 1.24
CentOS	7.6 x86_64, 7.7 x86_64, 7.9 x86_64, 8.2 x86_64

Container platform/OS	Version
SUSE	15 SP2 x86_64
Red Hat CoreOS	4.6 x86_64, 4.7 x86_64, 4.8 x86_64, 4.9 x86_64, 4.10 x86_64
Ubuntu	18.04 x86_64, 20.04 x86_64
Kylin V10	SP1 Arm, SP2 Arm

NOTICE

If the host machine OS is CoreOS 4.6/4.7/4.8/4.9/4.10, see *Kubernetes CSI for Red Hat OpenShift User Guide*

Table 2-3 Mappings between host machine OS versions and multipathing software versions

Host Machine OS Version	Native Multipathing Software Version	Huawei Multipathing Software Version (Supported Only by Enterprise Storage)
CentOS 7.6/7.7/7.9 (x86_64)	Delivered with the OS, supporting FC/iSCSI	UltraPath 31.1.0, supporting FC/iSCSI
CentOS 8.2 (x86_64)	Delivered with the OS, supporting FC/iSCSI	UltraPath 31.1.0, supporting FC/ iSCSI UltraPath-NVMe 31.1.RC8, supporting NVMe over RoCE/NVMe over FC
SUSE 15 SP2 (x86_64)	Delivered with the OS, supporting FC/iSCSI	UltraPath 31.1.0, supporting FC/iSCSI UltraPath-NVMe 31.1.RC8, supporting NVMe over RoCE
CoreOS 4.6/4.7/4.8/4.9/4.1 0 (x86_64)	Delivered with the OS, supporting FC/iSCSI	Not supported
Ubuntu 18.04/20.04 (x86_64)	Delivered with the OS, supporting FC/ iSCSI	Not supported
Kylin V10 SP1/SP2 (Arm)	Delivered with the OS, supporting FC/ iSCSI	Not supported

2.2 Kubernetes Features

Table 2-4 Kubernetes features supported by enterprise storage

Feature	V3	tor V3/F tor V5/F	OceanS tor Dorado V3	OceanS	tor V6	OceanS Dorado	
-	SAN	NAS	-	SAN	NAS	SAN	NAS
Dynamic Provisioning	√	√	√	√	√	√	√
Static Provisioning	√	√	√	√	√	√	√
Expand Persistent Volume	√	√	√	√	√	√	√
Create VolumeSnaps hot	√	√	√	√	√	√	√
Create Volume from Snapshot	√	√	√	√	×	√	×
Delete Snapshot	√	√	√	√	√	√	√
CSI Volume Cloning	√	√	√	√	×	√	×
CSI Raw Block Volume	√	×	√	√	×	√	×
Topology	√	√	√	√	√	√	√

Table 2-5 Kubernetes features supported by distributed storage

Feature	FusionStorag e	FusionStorag e Block	OceanStor Pacific	
-	-	-	SAN	NAS
Dynamic Provisioning	√	√	√	√
Static Provisioning	√	√	√	√

Feature	FusionStorag e	FusionStorag e Block	OceanStor Pacific	
Expand Persistent Volume	√	√	√	×
Create VolumeSnapsho t	√	√	√	×
Create Volume from Snapshot	√	√	√	×
Delete Snapshot	√	√	√	×
CSI Volume Cloning	√	√	√	×
CSI Raw Block Volume	√	√	√	×
Topology	√	√	√	√

Table 2-6 Restrictions on supported Kubernetes features

Feature	Restrictions
PVC Access Modes	 RWO/ROX/RWOP: supported by all types of volumes. The RWOP access mode is supported only by Kubernetes 1.22 and later versions. RWX: supported only by Raw Block (SAN) volumes and NAS volumes.
Create VolumeSnapshot	 NAS HyperMetro does not support snapshot creation. A PVC created using a static PV does not support snapshot creation.
Create Volume from Snapshot	 The StorageClass and volumeMode of the source PVC must be the same as those of the target PVC. NAS HyperMetro does not support volume creation using snapshots.
Expand Persistent Volume	 Only capacity expansion is supported. Capacity reduction is not supported. A PVC whose access mode is ROX does not support capacity expansion. A PVC created using a static PV does not support capacity expansion.

Feature	Restrictions
CSI Volume Cloning	The StorageClass and volumeMode of the source PVC must be the same as those of the target PVC.
	A PVC created using a static PV does not support cloning.
	NAS HyperMetro does not support cloning.

2.3 Storage Features

Table 2-7 Supported enterprise storage features

Standard Feature	OceanSt V3/V3 OceanSt V5/V5		OceanS tor Dorado V3	OceanSt	or V6	OceanS Dorado	
-	SAN	NAS	-	SAN	NAS	SAN	NAS
QoS	√	√	√	√	√	√	√
Application type	×	×	×	×	×	√	√
HyperMetr o	×	×	×	×	√	×	√ (6.1.3)
Multi- tenant	×	√	×	×	√	×	√ (6.1.3)
NFS 3.0	×	√	×	×	√	×	√
NFS 4.0 ¹	×	√	×	×	√	×	√
NFS 4.1 ¹	×	√	×	×	√	×	√
Note 1: Some enterprise storage devices do not support NES 4.0 or NES 4.1							

Note 1: Some enterprise storage devices do not support NFS 4.0 or NFS 4.1.

Table 2-8 Supported distributed storage features

Standard Feature	FusionStorag e	FusionStorage Block	OceanStor Pacific	
-	-	-	Block	File
QoS	√	√	√	×
Soft quota	×	×	×	√

Standard Feature	FusionStorag e	FusionStorage Block	OceanStor Pacific	
Account	×	×	×	√
DPC	×	×	×	√ (8.1.2)
NFS 3.0	×	×	×	√
NFS 4.1	×	×	×	√ (8.1.2)

Table 2-9 Restrictions on supported storage features

Feature	Restriction
HyperMetro	 PVCs and Pods can be created only when both HyperMetro storage devices are normal. If a single storage device is faulty, only delivered services can be normal and new services cannot be delivered. If both storage devices are faulty, contact Huawei technical support engineers. The working mode of the file system HyperMetro domain must be HyperMetro.
DPC	 For details about the host machine OSs supported by DPC, see Huawei Storage Interoperability Navigator.
NFS 4.0	• When a Pod is created using a PVC which is created using a static PV, the NFS 4.0 protocol cannot be specified for mounting. However, you can modify the host configurations to use the NFS 4.0 protocol for mounting by default.
NFS 4.1	• When a Pod is created using a PVC which is created using a static PV, the NFS 4.1 protocol cannot be specified for mounting. However, you can modify the host configurations to use the NFS 4.1 protocol for mounting by default.

3 Installation Preparations

- 3.1 Prerequisites
- 3.2 Obtaining the Software Package
- 3.3 Uploading the Components in the Software Package
- 3.4 Creating a Huawei CSI Image
- 3.5 Uploading a Huawei CSI Image
- 3.6 Configuring Host Multipathing

3.1 Prerequisites

- Kubernetes has been deployed and is running properly.
- A Huawei storage device is running properly.
- Drivers required for scanning disks and mounting files (such as iSCSI, DM-Multipath, and UltraPath NVMe. For details about how to select multipathing software, see multipathing software selection.) have been installed on all worker hosts in the Kubernetes cluster. If containers and services cannot run properly due to lack of system tools, view logs by referring to 10.1 Viewing Log Information and install the tools on the hosts.

3.2 Obtaining the Software Package

Procedure

- **Step 1** Open a browser and enter https://github.com/Huawei/eSDK_K8S_Plugin/releases in the address box.
- **Step 2** Select the desired version package and download **eSDK_EnterPrise_Storage_Plugin_***.*.**.zip. *.*.** indicates the release version number. (The version matching this document is 3.0.0.)
- **Step 3** Decompress the package and obtain the required components and documents.

----End

3.3 Uploading the Components in the Software Package

Procedure

Step 1 Decompress *eSDK_EnterPrise_Storage_Plugin_*.*.***:zip to obtain the software package and sample files required for installing and using CSI. **Table 3-1** shows the software package structure.

Table 3-1 Component description

Component	Description
bin/huawei-csi	Implements the CSI API.
bin/secretGenerate	Encrypts plaintext passwords and produces secret objects.
bin/secretUpdate	Encrypts plaintext passwords and updates secret objects.
Helm	Helm component used to deploy Huawei CSI.
deploy	.yaml sample file used during CSI deployment.
examples	.yaml sample file used during CSI use.
tools	Script used to upload images.

Step 2 Use a file transfer tool (such as Xftp) to upload the files generated upon decompression to the master node.

----End

3.4 Creating a Huawei CSI Image

Huawei CSI runs as a container. Currently, Huawei CSI provides only a binary file (**bin/huawei-csi**) which cannot be used directly. Therefore, you need to create a CSI image using the binary file to start the Huawei CSI service.

Prerequisites

A Linux host with Docker installed is available, and the host can access the Internet (only used to download the image package).

Procedure

- **Step 1** Log in to the Linux host.
- **Step 2** Run the **mkdir image** command to create a directory (for example, **image**) on the host.

mkdir image

Step 3 Run the **cd image** command to access the **image** directory.

cd image

- **Step 4** Copy the huawei-csi component to the **image** directory.
- **Step 5** Run the following command to create a file named **Dockerfile**.

```
# cat <<EOF > ./Dockerfile
FROM busybox:stable-glibc

ADD ["huawei-csi", "/"]
RUN ["chmod", "+x", "/huawei-csi"]

ENTRYPOINT ["/huawei-csi"]

EOF
```

NOTICE

busybox:stable-glibc indicates the basic image and its tag. It is only an example. Replace it based on site requirements.

Step 6 Run the **docker build -f Dockerfile -t huawei-csi:3.0.0** . command to create an image.

docker build -f Dockerfile -t huawei-csi:3.0.0.

□ NOTE

3.0.0 indicates the plug-in version number corresponding to the software package name. It is only an example. Replace it based on site requirements. If the same image already exists in the environment, use **docker image rm** <*image-id*>.

Step 7 Run the **docker image ls | grep huawei-csi** command to check whether the image is created. If the following information is displayed, it is created.

```
# docker image ls | grep huawei-csi
huawei-csi 3.0.0 c8b5726118ac About a minute ago 39 MB
```

Step 8 Run the **docker save huawei-csi:3.0.0 -o huawei-csi.tar** command to export the image.

docker save huawei-csi:3.0.0 -o huawei-csi.tar

◯ NOTE

3.0.0 indicates the plug-in version number corresponding to the software package name. It is only an example. Replace it based on site requirements.

----End

3.5 Uploading a Huawei CSI Image

This section describes how to use the image upload script to upload a Huawei CSI image to all worker nodes and load the image to the container runtime. If the image repository is used, upload the image to the repository and skip this section.

Prerequisites

- The host where the script is to be executed can communicate with all hosts to which the image is to be imported using SSH.
- The **expect**, **sshpass**, and **scp** software packages have been installed on the host where the script is to be executed.

Procedure

- **Step 1** Run the **vi** *worker-list.txt* command to create the **worker-list.txt** configuration file. # vi worker-list.txt
- **Step 2** Configure the **worker-list.txt** file. The template of the **worker-list.txt** file is as follows. Press **I** or **Insert** to enter the editing mode and add node information. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
# ip
192.168.128.16
192.168.128.17
```

- **Step 3** Upload and import an image.
 - If containerd is used as the container runtime, run the ./containerd-upload.sh worker-list.txt huawei-csi.tar command, enter the user name and password as prompted, and upload and import an image. The worker-list.txt file is the node information file configured in Step 2, and huawei-csi.tar is the image package exported in Step 8.

 # ./containerd-upload.sh worker-list.txt huawei-csi.tar
 - If Docker is used as the container runtime, run the ./docker-upload.sh worker-list.txt huawei-csi.tar command, enter the user name and password as prompted, and upload and import an image. The worker-list.txt file is the node information file configured in Step 2, and huawei-csi.tar is the image package exported in Step 8.

./docker-upload.sh worker-list.txt huawei-csi.tar

- **Step 4** Check whether the image is successfully imported.
 - 1. If **All images are uploaded successfully** is displayed at the end of the script, the image has been successfully imported to all nodes.

 All images are uploaded successfully
 - 2. If the following information is displayed at the end of the script, the image fails to be imported to the nodes in the list. In this case, check the script output carefully to locate the failure cause.

```
List of nodes to which the image fails to be imported: 192.168.128.16 192.168.128.17
```

----End

3.6 Configuring Host Multipathing

If you use block storage and access storage over the FC/iSCSI/NVMe over RoCE/NVMe over FC protocol, you are advised to configure host multipathing to improve storage link reliability. Currently, the following multipathing software is supported: native multipathing software (DM-Multipath) and Huawei multipathing software (UltraPath and UltraPath-NVMe).

Precautions

- For details about the host multipathing software supported by enterprise storage and distributed storage in different OSs, see **Table 2-3**.
- Table 3-2 lists the multipathing software supported by different protocols when SAN storage is used.

Table 3-2 Multipathing software supported by different protocols when SAN storage is used

SAN Storage Networking Protocol	No Multipathing Software	DM- Multipath	UltraPath	UltraPath- NVMe
VBS	√	х	х	х
iSCSI	√	√	√	√
FC	√	√	√	√
NVMe over RoCE	√	х	х	√
NVMe over FC	√	х	х	√

3.6.1 Installing Native Multipathing Software

Native multipathing software is delivered with the host system. If you need to install it, refer to this section.

3.6.1.1 Installing the Multipathing Tool Package

This section describes how to install the native multipathing tool package.

Prerequisites

Ensure that the worker nodes in the Kubernetes cluster can access the Internet (only used to download the multipathing tool package).

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node in the Kubernetes cluster through the management IP address.
- **Step 2** Install the multipathing tool package based on the OS.
 - CentOS: yum install -y device-mapper-multipath
 - SUSE: zypper install -y multipath-tools
- **Step 3** Enable the host multipathing service.

CentOS:

/sbin/mpathconf --enable systemctl start multipathd.service systemctl enable multipathd.service systemctl restart multipathd.service

 SUSE: systemctl restart multipath-tools.service chkconfig multipathd on

Step 4 Repeat **Step 1** to **Step 3** to install the multipathing tool on all worker nodes.

----End

3.6.1.2 Configuring the Multipathing Service

Multipathing is configured to improve the link reliability of LUNs on SAN storage. If multipathing is incorrectly configured, I/O errors will occur when a single link is faulty. As a result, the file systems or disks in the containers managed by the Kubernetes cluster are read-only or faulty, affecting I/O delivery.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node in the Kubernetes cluster through the management IP address.
- Step 2 Run the vi /etc/multipath.conf command to modify the multipath.conf file. If the file does not exist, configure or generate it by referring to the storage host connectivity guide. Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press Esc and enter :wq! to save the modification. This document uses Red Hat as an example. For details about other OSs, see the storage host connectivity guide.

NOTICE

Load balancing mode: During service read and write, the I/O paths from a host to all controllers on a storage device are the same. For details, see **Configuring**Multipathing > Concepts in *Huawei SAN Storage Host Connectivity Guide for Red Hat*.

Local preferred mode: When a host delivers I/Os to controllers, the storage device with better performance is accessed because the service link distances from storage devices are different. For details, see **Configuring Multipathing** > **Concepts** in *Huawei SAN Storage Host Connectivity Guide for Red Hat*.

 If enterprise storage and the load balancing mode are used, you are advised to add the following content to the devices field in the multipathing configuration file (/etc/multipath.conf). For details, see OceanStor Dorado Host Connectivity Guide for Red Hat and OceanStor Dorado Host Connectivity Guide for SUSE.

```
product "XSG1"
path_grouping_policy multibus
path_checker tur
prio const
path_selector "service-time 0"
failback immediate
no_path_retry 15
}
```

If enterprise storage and the local preferred mode are used, you are advised to add the following content to the **devices** field in the multipathing configuration file (/etc/multipath.conf). For details, see *Huawei SAN Storage Host Connectivity Guide for Red Hat* and *Huawei SAN Storage Host Connectivity Guide for SUSE*.

```
defaults {
     user friendly names yes
     find_multipaths no
devices {
  device {
                            "HUAWEI"
         vendor
                            "XSG1"
          product
          path_grouping_policy group_by_prio
          path_checker
                             tur
         prio
          path_selector
                             "round-robin 0"
          failback
                           immediate
         no_path_retry
                              15
}
```

 If distributed storage is used, you are advised to add the following content to the devices field in the multipathing configuration file (/etc/multipath.conf).
 The configuration varies according to the OS. For details, see Configuring Multipathing for an Application Server (Red Hat or CentOS) in FusionStorage 8.0.1 Block Storage Basic Service Configuration Guide 08.

```
defaults {
     user_friendly_names yes
     find_multipaths no
devices {
  device {
          vendor
                               "Huawei"
                               "VBS fileIO"
          product
          path_grouping_policy
                                   multibus
          path_checker
                                 tur
          prio
                              const
          path_selector
                                 "service-time 0"
          failback
                               immediate
          no_path_retry
                                  "10"
```

Step 3 After the configuration is complete, run the following command to restart the multipathd service.

systemctl restart multipathd.service

Step 4 Repeat **Step 1** to **Step 3** to configure the multipathing service for all worker nodes.

----End

3.6.2 Installing Huawei Multipathing Software

Huawei multipathing software is provided by Huawei. If you need to install it, refer to this section.

3.6.2.1 Installing the Huawei Multipathing Tool

This section describes how to install the Huawei multipathing tool package.

Prerequisites

Select Huawei multipathing software according to the host OS, networking, and version mappings. For enterprise users, log in to https://support.huawei.com/ enterprise. For carrier users, log in to https://support.huawei.com. Then search for UltraPath to obtain the software package and user guide.

Installation Procedure

Install Huawei multipathing software according to the obtained Huawei multipathing software user guide.

Configuration Procedure

Configure Huawei multipathing software according to the obtained Huawei multipathing software user guide.

□ NOTE

If multiple multipathing software products coexist, see the Huawei multipathing software user guide for compatibility and configuration methods.

4 Installing Huawei CSI

This chapter describes how to install Huawei CSI. You are advised to install Huawei CSI using Helm. If you need to manually install Huawei CSI, see 4.2 Manually Installing Huawei CSI. The two installation methods are retained in the current version. The manual installation method will be deleted in V3.2.0.

Huawei CSI can be installed as the root user or a non-root user. When installing Huawei CSI as a non-root user, ensure that the current user can access the API Server of the Kubernetes cluster. For details about how to configure access to the Kubernetes cluster as a non-root user, see 9.8 Configuring Access to the Kubernetes Cluster as a Non-root User.

Huawei CSI must be run as the root user.

- 4.1 Installing Huawei CSI Using Helm
- 4.2 Manually Installing Huawei CSI

4.1 Installing Huawei CSI Using Helm

This section describes how to install Huawei CSI using Helm 3.

4.1.1 Installing Helm

Prerequisites

Ensure that the master node in the Kubernetes cluster can access the Internet.

Procedure

- **Step 1** Run the following command to download the Helm 3 installation script.

 # curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
- **Step 2** Run the following command to modify the permission on the Helm 3 installation script.
 - # chmod 700 get helm.sh
- Step 3 Determine the Helm version to be installed based on the version mapping between Helm and Kubernetes. For details about the version mapping, see Helm Version Support Policy. Then run the following command to change the

DESIRED_VERSION environment variable to the Helm version to be installed and run the installation command.

DESIRED_VERSION=v3.9.0 ./get_helm.sh

Step 4 Run the following command to check whether Helm 3 of the specified version is successfully installed.

helm version version:"v3.9.0", GitCommit:"7ceeda6c585217a19a1131663d8cd1f7d641b2a7", GitTreeState:"clean", GoVersion:"go1.17.5"}

----End

4.1.2 Installing Huawei CSI

Component Description

When installing huawei-csi-controller, Helm installs the following components in the Pod of the Deployment type in the specified namespace:

- huawei-csi-driver: Huawei CSI driver
- Kubernetes External Provisioner: used to provide volumes
- Kubernetes External Attacher: used to attach volumes
- (Optional) Kubernetes External Snapshotter: used to provide snapshot support (installed as CRD)
- (Optional) Kubernetes External Resizer: used to expand the capacity of volumes

When installing huawei-csi-node, Helm installs the following components in the Pod of the DaemonSet type in the specified namespace:

- huawei-csi-driver: Huawei CSI driver
- Kubernetes Node Registrar: used to process driver registration

Prerequisites

- The **huawei-csi image** has been created and imported to all worker nodes (or pulled online from the image repository).
- The **sidecar** image may need to be downloaded during the installation. Therefore, worker nodes in the Kubernetes cluster must be able to access external networks. In an intranet environment, obtain the image package in other ways and manually import it into all worker nodes. For details about the image package list, see **Table 4-23**.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.
- You have obtained the user name and password of the storage device from the administrator.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected. (Huawei CSI uses the **ping** command to check.)
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected. (Huawei CSI uses the **ping** command to check.)

- Software clients required by the corresponding protocol, such as iSCSI and NFS clients, have been installed on all worker nodes of Kubernetes.
- For details about the supported user types and requirements when different storage devices are connected, see **Table 4-21**.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes. For details, see 3.6 Configuring Host Multipathing.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Copy the **helm** directory in the Kubernetes CSI component package to any directory on the master node. For details about the Helm tool path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Run the **kubectl create namespace** *huawei-csi* command to create a namespace for deploying Huawei CSI. *huawei-csi* is a user-defined namespace. If Huawei CSI and Helm are deployed in the same namespace, skip this step and create a namespace by referring to **Step 7**.

kubectl create namespace huawei-csi

- **Step 4** Run the **cd** *helm/esdk* command to go to the **helm/esdk** working directory.

 # cd helm/esdk
- **Step 5** Run the **vi** *values.yaml* command to modify the **values.yaml** file. # vi values.yaml
- **Step 6** Configure the **values.yaml** file. For details about the parameters, see **Table 4-1**. For mandatory parameters without default values, you need to set them based on site requirements. For parameters with default values, you do not need to modify them.

If multiple backends need to be configured, add them in the **backends** list. For details about backend configuration examples of different storage types and protocols, see **Table 4-2**.

```
# An array of storages with the access info
backends:
  storage: "oceanstor-nas"
  # support upper&lower characters, numeric and [-_].
  name: "nfs-155"
   - "https://192.168.129.155:8088"
  pools:
    - "StoragePool001"
  parameters:
   protocol: "nfs"
    portals:
     - "192.168.128.155"
 - storage: "fusionstorage-san"
  # support upper&lower characters, numeric and [-_].
  name: "iscsi-156"
  urls:
    - "https://192.168.129.156:8088"
  pools:
    - "StoragePool001"
  parameters:
    protocol: "iscsi"
    portals:
```

```
- "192.168.128.160"
     - "192.168.128.161"
kubernetes:
 namespace: huawei-csi
images:
 # The image name and tag for the attacher, provisioner and registrar sidecars. These must match the
appropriate Kubernetes version.
 sidecar:
  csiAttacher: k8s.gcr.io/sig-storage/csi-attacher:v3.3.0
  csiProvisioner: k8s.gcr.io/sig-storage/csi-provisioner:v3.0.0
  csiResizer: k8s.gcr.io/sig-storage/csi-resizer:v1.3.0
  registrar: k8s.gcr.io/sig-storage/csi-node-driver-registrar:v2.3.0
  livenessProbe: k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
  csiSnapshotter: k8s.gcr.io/sig-storage/csi-snapshotter:v4.2.1
  snapshotController: k8s.gcr.io/sig-storage/snapshot-controller:v4.2.1
 # The image name and tag for the Huawei CSI Service container
 # Replace the appropriate tag name
 huaweiCSIService: huawei-csi:3.0.0
# The CSI driver parameter configuration
 driverName: csi.huawei.com # It is strongly recommended not to modify this parameter
 endpoint: /csi/csi.sock # It is strongly recommended not to modify this parameter
 connectorThreads: 4
 volumeUseMultipath: true # Flag to enable or disable volume multipath access
 scsiMultipathType: DM-multipath # Required, if volume-use-multipath is set to TRUE
 nvmeMultipathType: HW-UltraPath-NVMe # Required, if volume-use-multipath is set to TRUE
 scanVolumeTimeout: 3
 backendUpdateInterval: 60
 controllerLogging:
  module: file
  level: info
  fileDir: /var/log/huawei
  fileSize: 20M
  maxBackups: 9
 nodeLogging:
  module: file
  level: info
  fileDir: /var/log/huawei
  fileSize: 20M
  maxBackups: 9
huaweiCsiController:
 replicas: 1 # Currently, the value can only be set to 1.
# Default image pull policy for sidecar container images
sidecarImagePullPolicy: "IfNotPresent"
# Default image pull policy for Huawei plugin container images
huaweiImagePullPolicy: "IfNotPresent"
# Flag to enable or disable snapshot (Optional)
snapshot:
 enable: true
# Flag to enable or disable resize (Optional)
resizer:
 enable: true
```

Table 4-1 Parameters in the values.yaml file

Parameter	Description	Mandatory	Default Value
backends.storage	Storage device type. • For enterprise SAN storage, set this parameter to oceanstor-san.	Yes	-
	 For enterprise NAS storage, set this parameter to oceanstor-nas. 		
	 For distributed SAN storage, set this parameter to fusionstorage-san. 		
	 For distributed NAS storage, set this parameter to fusionstorage-nas. 		
backends.name	Storage backend name. The value can contain uppercase letters, lowercase letters, digits, and the following special characters: []	Yes	-
	If multiple storage backends need to be configured, ensure that the storage backend name is unique.		
backends.urls	Management URLs of storage device. The value format is a list. One or more management URLs (multiple controllers) of the same storage device are supported. Currently, only IPv4 addresses are supported.	Yes	-

Parameter	Description	Mandatory	Default Value
backends.pools	Storage pools of storage devices. The value format is a list. One or more storage pools on the same storage device are supported.	Yes	
backends.paramet ers.protocol	 Storage protocol. The value is a character string. For enterprise SAN storage, the value can be iscsi, fc, roce, or fc-nvme. For enterprise NAS storage, the value can be nfs. For distributed SAN storage, the value can be iscsi or scsi. For distributed NAS storage, the value can be iscsi or scsi. For distributed NAS storage, the value can be nfs or dpc. 	Yes	
backends.paramet ers.portals	Service port. The value is a character string. Multiple ports can be configured if the protocol is iscsi or roce . Only one port can be configured if the protocol is nfs . Service ports do not need to be configured if the protocol is fc , fc-nvme , or dpc . If the protocol is scsi , the port is in dictionary format where the key indicates the host name and the value indicates the IP address.	Conditionally mandatory	-

Parameter	Description	Mandatory	Default Value
kubernetes.names pace	Namespace, which can be customized. The name must start and end with a letter or digit, and contain letters, digits, or hyphens (-). For example, my-name and 123-abc.	Yes	huawei-csi
images.sidecar.live nessProbe	livenessprobe sidecar image.	Yes	k8s.gcr.io/sig- storage/ livenessprobe :v2.5.0
images.sidecar.csiP rovisioner	csi-provisioner sidecar image.	Yes	k8s.gcr.io/sig- storage/csi- provisioner:v 3.0.0
images.sidecar.csi Attacher	csi-attacher sidecar image.	Yes	k8s.gcr.io/sig- storage/csi- attacher:v3.3.
images.sidecar.csiR esizer	csi-resizer sidecar image.	Yes	k8s.gcr.io/sig- storage/csi- resizer:v1.3.0
images.sidecar.csiS napshotter	csi-snapshotter sidecar image.	Yes	k8s.gcr.io/sig- storage/csi- snapshotter:v 4.2.1
images.sidecar.sna pshotController	snapshot-controller sidecar image.	Yes	k8s.gcr.io/sig- storage/ snapshot- controller:v4. 2.1
images.sidecar.regi strar	csi-node-driver-registrar sidecar image.	Yes	k8s.gcr.io/sig- storage/csi- node-driver- registrar:v2.3.
images.huaweiCSI Service	huawei-csi image.	Yes	-
csi_driver.endpoint	Communication endpoint. You are advised not to change the value.	Yes	/csi/csi.sock

Parameter	Description	Mandatory	Default Value
csi_driver.driverNa me	Registered driver name. You are advised not to change the value.	Yes	csi.huawei.co m
csi_driver.connecto rThreads	Maximum number of disks that can be concurrently scanned/ detached. The value is an integer ranging from 1 to 10.	Yes	4
csi_driver.volumeU seMultipath	Whether to use multipathing software. The value is a Boolean value.	Yes	true
csi_driver.scsiMulti pathType	Multipathing software used by fc/iscsi. The following parameter values can be configured: DM-multipath HW-UltraPath HW-UltraPath-NVMe	Mandatory when volumeUseMu ltipath is set to TRUE.	DM- multipath
csi_driver.nvmeMu ltipathType	Multipathing software used by roce/fc-nvme. Only HW-UltraPath- NVMe is supported.	Mandatory when volumeUseMu ltipath is set to TRUE.	HW- UltraPath- NVMe
csi_driver.scanVolu meTimeout	Timeout interval for waiting for multipathing aggregation when DM-multipath is used on the host. The value ranges from 1 to 600 seconds.	Yes	3
csi_driver.backend UpdateInterval	Interval for updating backend capabilities. The value ranges from 60 to 600 seconds.	Yes	60
csi_driver.controlle rLogging.module	Record type of the controller log. The following parameter values can be configured: • file • console	Yes	file

Parameter	Description	Mandatory	Default Value
csi_driver.controlle rLogging.level	Output level of the controller log. The following parameter values can be configured: • debug • info • warning • error • fatal	Yes	info
csi_driver.controlle rLogging.fileDir	Directory of the controller log in file output mode.	Yes	/var/log/ huawei
csi_driver.controlle rLogging.fileSize	Size of a single controller log file in file output mode.	Yes	20M
csi_driver.controlle rLogging.maxBack ups	Maximum number of controller log file backups in file output mode.	Yes	9
csi_driver.nodeLog ging.module	Record type of the node log. The following parameter values can be configured: • file • console	Yes	file
csi_driver.nodeLog ging.level	Output level of the node log. The following parameter values can be configured: • debug • info • warning • error • fatal	Yes	info
csi_driver.nodeLog ging.fileDir	Directory of the node log in file output mode.	Yes	/var/log/ huawei
csi_driver.nodeLog ging.fileSize	Size of a single node log file in file output mode.	Yes	20M
csi_driver.nodeLog ging.maxBackups	Maximum number of node log file backups in file output mode.	Yes	9

Parameter	Description	Mandatory	Default Value
sideCarImagePull- Policy	Pull policy of the sidecar image.	Yes	IfNotPresent
huweiCSIImagePul lPolicy	Pull policy of the huaweicsi image.	Yes	IfNotPresent
huaweiCsiControl- ler.replicas	Quantity of huawei-csi- controller. The value can only be 1 .	Yes	1
snapshot.enable	Whether to enable the snapshot feature. Kubernetes volume snapshot CRD is stored in the helm/crd directory. If this parameter is set to true, the snapshot CRD resource will be installed when the helm install command is executed. NOTICE Helm is not responsible for the CRD lifecycle. If you need to update the CRD version, you must update it manually.	Yes	true
resizer.enable	Whether to enable the capacity expansion feature.	Yes	true

Table 4-2 Configuration reference template of backends

Storage Type	Protocol	Backend Configuration Reference Template	Remarks
Enterprise storage SAN	iSCSI	backends: - storage: "oceanstor-san" # support upper&lower characters, numeric and []. name: "iscsi-155" urls: - "https://192.168.129.155:8088" - "https://192.168.129.156:8088" pools: - "StoragePool001" parameters: protocol: "iscsi" portals: - "192.168.128.120" - "192.168.128.121"	Ensure that an iSCSI client has been installed on all worker nodes.

Storage Type	Protocol	Backend Configuration Reference Template	Remarks
Enterprise storage SAN	FC	backends: - storage: "oceanstor-san" # support upper&lower characters, numeric and []. name: "fc-155" urls: - "https://192.168.129.155:8088" - "https://192.168.129.156:8088" pools: - "StoragePool001" parameters: protocol: "fc"	-
Enterprise storage NAS	NFS	backends: - storage: "oceanstor-nas" # support upper&lower characters, numeric and []. name: "nfs-155" urls: - "https://192.168.129.155:8088" - "https://192.168.129.156:8088" pools: - "StoragePool001" parameters: protocol: "nfs" portals: - "192.168.128.155"	Ensure that an NFS client tool has been installed on all worker nodes.
Enterprise storage SAN	NVMe over RoCE	backends: - storage: "oceanstor-san" # support upper&lower characters, numeric and []. name: "roce-155" urls: - "https://192.168.129.155:8088" - "https://192.168.129.156:8088" pools: - "StoragePool001" parameters: protocol: "roce" portals: - "192.168.128.120" - "192.168.128.121"	Ensure that the nvme-cli tool has been installed on all worker nodes and its version is 1.9 or later.
Enterprise storage SAN	NVMe over FC	backends: - storage: "oceanstor-san" # support upper&lower characters, numeric and []. name: "fc-nvme-155" urls: - "https://192.168.129.155:8088" - "https://192.168.129.156:8088" pools: - "StoragePool001" parameters: protocol: "fc-nvme"	Ensure that the nvme-cli tool has been installed on all worker nodes and its version is 1.9 or later.

Storage Type	Protocol	Backend Configuration Reference Template	Remarks
Distributed storage SAN	SCSI	backends: - storage: "fusionstorage-san" # support upper&lower characters, numeric and []. name: "scsi-155" urls: - "https://192.168.129.155:28443" pools: - "StoragePool001" parameters: protocol: "scsi" portals: - {"hostname01": "192.168.125.21"} - {"hostname02": "192.168.125.22"}	Ensure that the distributed storage VBS client has been installed on all worker nodes.
Distributed storage SAN	iSCSI	backends: - storage: "fusionstorage-san" # support upper&lower characters, numeric and []. name: "iscsi-155" urls: - "https://192.168.129.155:28443" - "https://192.168.129.156:28443" pools: - "StoragePool001" parameters: protocol: "iscsi" portals: - "192.168.128.120" - "192.168.128.121"	Ensure that an iSCSI client has been installed on all worker nodes.
Distributed storage NAS	NFS	backends: - storage: "fusionstorage-nas" # support upper&lower characters, numeric and []. name: "nfs-155" urls: - "https://192.168.129.155:28443" - "https://192.168.129.156:28443" pools: - "StoragePool001" parameters: protocol: "nfs" portals: - "192.168.128.120"	Ensure that an NFS client tool has been installed on all worker nodes.
Distributed storage NAS	DPC	backends: - storage: "fusionstorage-nas" # support upper&lower characters, numeric and []. name: "dpc-155" urls: - "https://192.168.129.155:28443" - "https://192.168.129.156:28443" pools: - "StoragePool001" parameters: protocol: "dpc"	Ensure that all worker nodes have been added as DPC compute nodes on the storage device to be connected.

Step 7 Run the **helm install** *helm-release-name* ./ -n *huawei-csi* --create-namespace command to install Huawei CSI.

In the preceding command, *helm-release-name* indicates the custom chart name, ./ indicates that the Helm project in the current directory is used, and *huawei-csi* indicates the custom namespace.

helm install helm-huawei-csi ./ -n huawei-csi --create-namespace NAME: helm-huawei-csi LAST DEPLOYED: Wed Jun 8 11:50:28 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 1 TEST SUITE: None

- **Step 8** Copy the **secretGenerate** tool in the Huawei CSI component package to any directory on the master node. For details about the tool path, see **3.3 Uploading the Components in the Software Package**.
- **Step 9** Use an encryption tool to enter the user name and password of the storage device.
 - 1. Run the **chmod +x secretGenerate** command to grant the execute permission on the secretGenerate tool.

chmod +x secretGenerate

2. Run the ./secretGenerate --namespace=huawei-csi --logFileDir=/var/log/huawei command to run the secretGenerate tool. Change the value of namespace to the actual namespace. If this parameter value is not used, the huawei-csi namespace is used by default. Change the value of logFileDir to the actual log directory. If this parameter value is not used, the /var/log/huawei directory is used by default. Then enter the ID of the backend to be configured as prompted. If Configured is false, the backend is not configured. If Configured is true, the backend is configured.

3. Enter the user name and password as prompted to create a **secret** object.

4. After the configuration is complete, enter **exit** to exit and save the configuration.

Please enter the backend number to configure (Enter 'exit' to exit): exit
Saving configuration. Please wait......
The configuration is saved successfully.

5. Run the **kubectl get secret -n huawei-csi | grep huawei-csi-secret** command to check whether the **secret** object is successfully created.

kubectl -n huawei-csi get secret huawei-csi-secret NAME TYPE DATA AGE huawei-csi-secret Opaque 1 8d

Step 10 After the huawei-csi service is deployed, run the **kubectl get pod -n huawei-csi** command to check whether the service is started.

```
# kubectl get pod -n huawei-csi
NAME READY STATUS RESTARTS AGE
```

huawei-csi-controller-764bd64c97-kr2nm 7/7 Running 0 144m huawei-csi-node-7m48s 3/3 Running 0 144m

----End

4.2 Manually Installing Huawei CSI

This section describes how to manually install Huawei CSI.

4.2.1 Connecting to Enterprise Storage

This section describes how to connect the huawei-csi plug-in to Huawei enterprise storage.

Restrictions

When the same SAN storage is connected to Kubernetes, you cannot configure multiple data protocols (iSCSI, FC, NVMe over RoCE, and NVMe over FC) on one worker node.

4.2.1.1 Connecting to Enterprise Storage SAN over iSCSI

Perform this operation when you want to connect to enterprise storage SAN over iSCSI.

Prerequisites

- An iSCSI client has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address
 of the storage device to be connected. (Huawei CSI uses the ping command
 to check.)
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected. (Huawei CSI uses the **ping** command to check.)
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes. For details, see 3.6 Configuring Host Multipathing.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-iscsi.yaml** example file in the

software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 4-3**.

Table 4-3 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-4.

Table 4-4 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.

Configuration Item	Format	Description	Remarks
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the block storage service.

Configuration Item	Format	Description	Remarks
parameters	Dictionary	Variable parameters in scenarios where iSCSI is used. This parameter is mandatory.	In scenarios where iSCSI is used, set the protocol parameter to a fixed value: iscsi . Set the portals parameter to the iSCSI service IP addresses of the storage device. Use commas (,) to separate multiple iSCSI service IP addresses. You can log in to DeviceManager to obtain the iSCSI service IP addresses. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is iSCSI. (For other series, see the corresponding operation description.)

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.1.2 Connecting to Enterprise Storage SAN over FC

Perform this operation when you want to connect to enterprise storage SAN over FC.

Restrictions

To connect to enterprise storage SAN over FC, ensure that no residual drive letter exists on the host. If any residual drive letter exists, clear the drive letter by referring to 10.4 After a Worker Node in the Cluster Breaks Down and

Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters.

Prerequisites

- All Kubernetes nodes communicate properly with the management IP address
 of the storage device to be connected. (Huawei CSI uses the ping command
 to check.)
- All worker nodes of Kubernetes can communicate with the storage device to be connected over FC. (Huawei CSI uses the **ping** command to check.)
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes of Kubernetes. For details, see 3.6 Configuring Host Multipathing.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-fc.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-5.

Table 4-5 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-6.

Table 4-6 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-).
			NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the block storage service.
parameters	Dictionary	Variable parameters in scenarios where FC is used. This parameter is mandatory.	In scenarios where FC is used, set the protocol parameter to a fixed value: fc .

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.1.3 Connecting to Enterprise Storage NAS over NFS

Perform this operation when you want to connect to enterprise storage NAS over NFS.

Prerequisites

- An NFS client tool has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected. (Huawei CSI uses the ping command to check.)
- All worker nodes of Kubernetes communicate properly with the NFS logical port of the storage device to be connected. (Huawei CSI uses the ping command to check.)
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-nfs.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in vaml format. For details, see Table 4-7.

```
kind: ConfigMap
apiVersion: v1
metadata:
 name: huawei-csi-configmap
 namespace: huawei-csi
data:
 csi.json: |
      "backends": [
         {
            "storage": "oceanstor-nas",
            "name": "storage",
             "urls": ["https://192.168.125.20:8088", "https://192.168.125.21:8088"],
             "pools": ["storagepool01", "storagepool02"],
"parameters": {"protocol": "nfs", "portals": ["192.168.125.22"]}
         }
      ]
  }
```

Table 4-7 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-8.

Table 4-8 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage NAS is connected, the value is fixed to oceanstor-nas.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the file storage service.

Configuration Item	Format	Description	Remarks
parameters	Dictionary	Variable parameters in scenarios where NFS is used. This parameter is mandatory.	The protocol parameter is fixed to nfs. portals: logical port IP address or DNS zone of the storage device. Only one IP address or DNS zone can be
			configured. You can log in to DeviceManager to obtain the logical port IP address. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is NFS. (For other series, see the corresponding operation description.)

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.1.4 Connecting to Enterprise Storage SAN over NVMe over RoCE

Perform this operation when you want to connect to enterprise storage SAN over NVMe over RoCE.

Prerequisites

- All Kubernetes nodes communicate properly with the management IP address
 of the storage device to be connected. (Huawei CSI uses the ping command
 to check.)
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected. (Huawei CSI uses the **ping** command to check.)

- The nyme-cli tool has been installed on all worker nodes of Kubernetes, and the tool version is 1.9 or later.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes of Kubernetes. For details, see 3.6 Configuring Host Multipathing.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-roce.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-9.

Table 4-9 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-10.

Table 4-10 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.

Configuration Item	Format	Description	Remarks
parameters	Dictionary	Variable parameters in scenarios where the NVMe over RoCE protocol is used. This parameter is mandatory.	In scenarios where the NVMe over RoCE protocol is used, set the protocol parameter to a fixed value: roce . Set portals to the IP addresses of the logical ports when data protocol type of the storage device is NVMe over RoCE. Use commas (,) to separate the IP addresses.
			You can log in to DeviceManager to obtain the logical port IP address. Take OceanStor Dorado 6.x series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is NVMe over RoCE. (For other series, see the corresponding operation description.)

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.1.5 Connecting to Enterprise Storage SAN over NVMe over FC

Perform this operation when you want to connect to enterprise storage SAN over NVMe over FC.

Restrictions

To connect to enterprise storage SAN over NVMe over FC, ensure that no residual drive letter exists on the host. If any residual drive letter exists, clear the drive

letter by referring to 10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters.

Prerequisites

- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected. (Huawei CSI uses the ping command to check.)
- All worker nodes of Kubernetes can communicate with the storage device to be connected over NVMe over FC. (Huawei CSI uses the **ping** command to check.)
- The nyme-cli tool has been installed on all worker nodes of Kubernetes, and the tool version is 1.9 or later.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes of Kubernetes. For details, see 3.6 Configuring Host Multipathing.
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-oceanstor-fc-nvme.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-11.

Table 4-11 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-12.

Table 4-12 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the enterprise storage SAN is connected, the value is fixed to oceanstor-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device has multiple management URLs.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools that support the block storage service.
parameters	Dictionary	Variable parameters in scenarios where the NVMe over FC protocol is used. This parameter is mandatory.	In scenarios where the NVMe over FC protocol is used, set the protocol parameter to a fixed value: fc-nvme .

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n huawei-csi | grep huawei-csi-configmap
huawei-csi-configmap 1 5s
```

----End

4.2.2 Connecting to Distributed Storage

This section describes how to connect the huawei-csi plug-in to Huawei distributed storage.

Restrictions

When the same SAN storage is connected to Kubernetes, you cannot configure multiple data protocols (SCSI and iSCSI) on one worker node.

4.2.2.1 Connecting to Distributed Storage SAN over SCSI

Perform this operation when you want to connect to distributed storage SAN over SCSI.

Prerequisites

- The distributed storage VBS client has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address of the storage device to be connected. (Huawei CSI uses the ping command to check.)
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-scsi.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-13.

Table 4-13 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-14.

Table 4-14 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage SAN is connected, the value is fixed to fusionstorage-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.

Configuration Item	Format	Description	Remarks
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.
parameters	Dictionary	Variable parameters in scenarios where SCSI is used. This parameter is mandatory.	The protocol parameter is fixed to scsi. Set portals to a pair list of host names and VBS node IP addresses. The format is [{"hostname":"*.*.**"}], where hostname indicates the host name of a worker node and *.*.** indicates the management IP address of a distributed storage block client (only IPv4 addresses are supported currently). If there are multiple worker nodes, configure them in dictionary format and separate them with commas (,). In the preceding example, hostname01 is the host name of a worker node in Kubernetes, and 192.168.125.21 is the management IP address of a VBS node after VBS is created for the worker node.

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.2.2 Connecting to Distributed Storage SAN over iSCSI

Perform this operation when you want to connect to distributed storage SAN over iSCSI.

Prerequisites

- An iSCSI client has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address
 of the storage device to be connected. (Huawei CSI uses the ping command
 to check.)
- All worker nodes of Kubernetes communicate properly with the service IP address of the storage device to be connected. (Huawei CSI uses the ping command to check.)
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes.

Precautions

- The host name of a Kubernetes worker node consists of digits, letters, underscores (_), hyphens (-), periods (.), and colons (:), and must start with a digit, letter, or underscore (_). The name length cannot exceed 31 characters.
- Only FusionStorage 8.0.0 and later versions support iSCSI networking configuration.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-iscsi.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-15.

kind: ConfigMap apiVersion: v1 metadata:

Table 4-15 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".ba ckends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single backend storage device, see Table 4-16.

Table 4-16 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage SAN is connected, the value is fixed to fusionstorage-san.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-).
			NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.

Configuration Item	Format	Description	Remarks
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.
parameters	Dictionary	Variable parameters in scenarios where iSCSI is used. This parameter is mandatory.	In scenarios where iSCSI is used, set the protocol parameter to a fixed value: iscsi . Set the portals parameter to the iSCSI service IP addresses of the storage device. Use commas (,) to separate multiple them. You can log in to DeviceManager to obtain them. You can log in to DeviceManager to obtain the iSCSI service IP addresses. Take OceanStor Pacific series as an example. On DeviceManager, choose Resources > Access > Service Network . (For other series, see the corresponding operation description.)

Step 4 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n huawei-csi | grep huawei-csi-configmap
huawei-csi-configmap 1 5s
```

----End

4.2.2.3 Connecting to Distributed Storage NAS over NFS

Perform this operation when you want to connect to distributed storage NAS over NFS.

Prerequisites

- An NFS client tool has been installed on all worker nodes of Kubernetes.
- All Kubernetes nodes communicate properly with the management IP address
 of the storage device to be connected. (Huawei CSI uses the ping command
 to check.)
- All worker nodes of Kubernetes communicate properly with the IP address of the NFS logical port on the storage device to be connected. (Huawei CSI uses the ping command to check.)
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-nfs.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 4-17.

Table 4-17 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of backend storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-18.

Table 4-18 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage NAS is connected, the value is fixed to fusionstorage-nas.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.

	Description	Remarks
List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pools.
Dictionary	Variable parameters in scenarios where NFS is used. This parameter is mandatory.	portals: logical port IP address the specified storage device. You can log in to DeviceManager to obtain it. Only one IP address can be configured. You can log in to DeviceManager to
		obtain the logical port IP address. Take OceanStor Pacific series as an example. On DeviceManager, choose Resources > Access > Service Network and click the name of a zone. On the page that is displayed, click the IP Address Pool tab. The IP Address/Mask column indicates the logical port IP address. For other series, see the corresponding
		storage pool used on the storage device to be connected. This parameter is mandatory. Dictionary Variable parameters in scenarios where NFS is used. This parameter is

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.2.4 Connecting to Distributed Storage NAS over DPC

Perform this operation when you want to connect to distributed storage NAS over DPC.

Prerequisites

- All worker nodes of Kubernetes have been added as DPC compute nodes on the storage device to be connected.
- All Kubernetes nodes communicate properly with the management IP address
 of the storage device to be connected. (Huawei CSI uses the ping command
 to check.)
- You have obtained the IP address, login account, and password of any master node in the Kubernetes cluster from the administrator.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to create a file named *huawei-csi-configmap.yaml*.

vi huawei-csi-configmap.yaml

Step 3 Configure the *huawei-csi-configmap.yaml* file. The following shows a template of the *huawei-csi-configmap.yaml* file. You can also refer to the **deploy/huawei-csi-configmap/huawei-csi-configmap-fusionstorage-nfs.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 4-19**.

Table 4-19 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".backends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of back-end storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 4-20.

Table 4-20 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage NAS is connected, the value is fixed to fusionstorage-nas.
name	String	Storage backend name.	User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). NOTE If multiple storage backends need to be configured, ensure that the storage backend name is unique.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	For FusionStorage, only one management URL can be configured.

Configuration Item	Format	Description	Remarks
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. You can log in to DeviceManager to obtain the storage pool name.
parameters	Dictionary	Variable parameters in scenarios where DPC is used. This parameter is mandatory.	In scenarios where DPC is used, set the protocol parameter to a fixed value: dpc .

kubectl create -f huawei-csi-configmap.yaml

Step 5 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

----End

4.2.3 Starting huawei-csi Services

This section describes how to start huawei-csi services.

Precautions

An image may need to be downloaded during the procedure. Therefore, worker nodes in the Kubernetes cluster must be able to access external networks. In an intranet environment, obtain the image package in other ways and manually import it into all worker nodes. For details about the image package list, see **Table 4-23**.

Prerequisites

- You have obtained the user name and password of the storage device from the administrator.
- For details about the supported user types and requirements when different storage devices are connected, see **Table 4-21**.

□ NOTE

When vStore users are used, only NAS storage can be connected. For details about the supported storage models, see **Table 2-7**.

Table 4-21 Details about users supported when different storage devices are connected to CSI

Storage Type	User Type	Role	Level	Туре
OceanStor	System user	Administrator	Administrator	Local user
V3/V5	vStore user	vStore administrator	Administrator	Local user
OceanStor Dorado V3	System user	Administrator	Administrator	Local user
OceanStor 6.1	System user	Administrator	N/A	Local user
OceanStor Dorado 6.1.3	System user	Administrator	N/A	Local user
	vStore user	vStore administrator	N/A	Local user
OceanStor Pacific series	System user	Administrator	N/A	Local user

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Copy the **secretGenerate** tool in the Kubernetes CSI component package to any directory on the master node. For details about the tool path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Use an encryption tool to enter the user name and password of the storage device.
 - 1. Run the **chmod** +x **secretGenerate** command to grant the execute permission on the secretGenerate tool.
 - # chmod +x secretGenerate
 - 2. Run the ./secretGenerate --namespace=huawei-csi --logFileDir=/var/log/huawei command to run the secretGenerate tool. Change the value of namespace to the actual namespace. If this parameter value is not used, the huawei-csi namespace is used by default. Change the value of logFileDir to the actual log directory. If this parameter value is not used, the /var/log/huawei directory is used by default. Then enter the ID of the backend to be configured as prompted. If Configured is false, the backend is not configured. If Configured is true, the backend is configured.

```
# ./secretGenerate huawei-csi
Getting backend configuration information.......

Number Configured BackendName Urls
false strage-backend [https://192.168.125.25:8088]
```

```
false
              strage-backend-02 [https://192.168.125.26:8088]
3
     false
              strage-backend-03
                                   [https://192.168.125.27:8088]
4
     false
              strage-backend-04
                                  [https://192.168.125.28:8088]
              strage-backend-05
                                  [https://192.168.125.29:28443]
     false
5
6
     false
              strage-backend-06 [https://192.168.125.30:28443]
Please enter the backend number to configure (Enter 'exit' to exit):3
```

3. Enter the user name and password as prompted to create a **secret** object.

4. After the configuration is complete, enter **exit** to exit and save the configuration.

Please enter the backend number to configure (Enter 'exit' to exit): exit
Saving configuration. Please wait......
The configuration is saved successfully.

5. Run the **kubectl get secret -n huawei-csi | grep huawei-csi-secret** command to check whether the **secret** object is successfully created.

```
# kubectl -n huawei-csi get secret huawei-csi-secret

NAME TYPE DATA AGE
huawei-csi-secret Opaque 1 8d
```

Step 4 Run the following command to create the RBAC permission.

kubectl apply -f huawei-csi-rbac.yaml

Step 5 Start the controller service.

- 1. Run the following command to deploy the snapshot-crd service.

 # kubectl apply -f huawei-csi-snapshot-crd.yaml
- Run the vi huawei-csi-controller.yaml command to modify the .yaml file. Press
 I or Insert to enter the editing mode and modify the following parameters.
 After the modification is complete, press Esc and enter :wq! to save the modification.

□ NOTE

 (Mandatory) In the image configuration item under huawei-csi-driver in the sample .yaml file, change huawei-csi:*.**to <Name>:<Version> of the Huawei CSI image created and uploaded in the previous chapter. The docker is used as an example.

```
containers:
...
- name: huawei-csi-driver
image: huawei-csi:3.0.0
```

 (Optional) The namespace configuration item under metadata in the sample .yaml file indicates the namespace where the huawei-csi-controller service is installed. If you need to change the value, ensure that the namespaces in the config.yaml, rbac.yaml, and node.yaml files are the same. The change method is as follows.

```
as follows.
metadata:
...
name: huawei-csi-controller
namespace: huawei-csi
```

3. Run the following command to start the controller service.

kubectl apply -f huawei-csi-controller.yaml

Step 6 Start the node service.

Run the vi huawei-csi-node.yaml command to modify the .yaml file. Press I or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

□ NOTE

(Mandatory) In the **image** configuration item under **huawei-csi-driver** in the sample .yaml file, change huawei-csi: *.** to <Name>< <Version> of the Huawei CSI image created and uploaded in the previous chapter. The docker is used as an example.

containers:

- name: huawei-csi-driver image: huawei-csi:3.0.0

(Optional) The namespace configuration item under metadata in the sample .yaml file indicates the namespace where the huawei-csi-node service is installed. If you need to change the value, ensure that the namespaces in the config.yaml, rbac.yaml, and controller.yaml files are the same. The change method is as follows.

metadata: name: huawei-csi-controller namespace: huawei-csi

(Optional) In the args section of huawei-csi-driver in the .yaml file, --volumeuse-multipath indicates that multipathing is enabled by default. The following shows how to change the value.

args:

- "--endpoint=/csi/csi.sock"
- "--containerized"
- "--driver-name=csi.huawei.com"
- "--volume-use-multipath=false"
- (Optional) In the args section of huawei-csi-driver in the .yaml file, --connectorthreads indicates the number of concurrent operations on the drive letter on the host. The value is an integer ranging from 1 to 10, and the default value is 4. To change the value, refer to the following.

args:

- "--endpoint=/csi/csi.sock"
- "--containerized"
- "--driver-name=csi.huawei.com"
- "--volume-use-multipath=true"
- "--connector-threads=5"
- (Optional) In the args section of huawei-csi-driver in the .yaml file, --scanvolume-timeout indicates the timeout of waiting for multipathing aggregation when DM-multipath is used on the host. The value is an integer ranging from 1 to 600, and the default value is 3. To change the value, refer to the following. args:
 - "--endpoint=/csi/csi.sock"
 - "--containerized"
 - "--driver-name=csi.huawei.com"
 - "--volume-use-multipath=true"
 - "--connector-threads=4"
 - "--scan-volume-timeout=3"
- (Optional) In the args section of huawei-csi-driver in the .yaml file, for enterprise storage, if --volume-use-multipath is set to true, you can configure the multipathing type according to the networking mode. For details, see Table 4-22. args:
 - "--endpoint=/csi/csi.sock" "--containerized"

 - "--driver-name=csi.huawei.com"
 - "--connector-threads=4"
 - "--volume-use-multipath=true"
 - "--scsi-multipath-type=DM-multipath"
 - "--nvme-multipath-type=HW-UltraPath-NVMe"

Storage Protocol	Parameter	Description	Remarks
iSCSI/FC	scsi-multipath- type	The value can be: - DM-multipath - HW-UltraPath - HW- UltraPath- NVMe The default value is DM- multipath.	 DM- multipath: native multipathing software of the OS HW- UltraPath: Huawei UltraPath multipathing
NVMe over RoCE/NVMe over FC	nvme- multipath-type	The default value is HW-UltraPath-NVMe and only HW-UltraPath-NVMe can be configured.	software - HW- UltraPath- NVMe: Huawei UltraPath- NVMe multipathing software

Table 4-22 Parameters for configuring enterprise storage multipathing

2. Run the following command to start the node service.
kubectl apply -f huawei-csi-node.yaml

Step 7 After the huawei-csi services are deployed, run the **kubectl get pod -A | grep huawei** command to check whether the services are started.

kubectl get pod -A | grep huawei huawei-csi huawei-csi-controller-695b84b4d8-tg64l 7/7 **Running** 0 14s huawei-csi huawei-csi-node-g6f7z 3/3 **Running** 0 14s

MOTE

The Pod of *huawei-csi-controller-695b84b4d8-tg64l* has seven containers, including liveness-probe, csi-provisioner, csi-attacher, csi-resizer, csi-snapshotter, shapshot-controller, and huawei-csi-driver. Each container has its own image repository and functions. For details about the containers, see **Table 4-23**.

The Pod of *huawei-csi-node-g6f7z* has three containers, including liveness-probe, csi-node-driver-registrar, and huawei-csi-driver. Each container has its own image repository and functions. For details about the containers, see **Table 4-23**.

Table 4-23 Container description

Container Name	Container Image	Feature Description	Remark s
liveness- probe	k8s.gcr.io/sig- storage/ livenessprobe:v2.5.0	Monitors the health status of CSI and reports it to Kubernetes so that Kubernetes can automatically detect CSI program problems and restart the Pod to rectify the problems.	View details
csi- provisioner	k8s.gcr.io/sig- storage/csi- provisioner:v3.0.0	 Calls the CSI Controller service to create a LUN or file system on the storage system as a PV and bind the PV to a PVC when creating a PVC. Calls the CSI Controller service to unbind a PV from a PVC and delete the LUN or file system corresponding to the PV when deleting a PVC. 	View details
csi-attacher	k8s.gcr.io/sig- storage/csi- attacher:v3.3.0	Calls the CSI Controller service to perform the "Publish/ Unpublish Volume" operation when creating or deleting a Pod.	View details
csi-resizer	k8s.gcr.io/sig- storage/csi- resizer:v1.3.0	Calls CSI to provide more storage space for a PVC when expanding the capacity of the PVC.	View details
csi- snapshotter	k8s.gcr.io/sig- storage/csi- snapshotter:v4.2.1	Calls CSI to create or delete a snapshot on the storage system when creating or deleting a VolumeSnapshot.	View details
shapshot- controller	k8s.gcr.io/sig- storage/snapshot- controller:v4.2.1	Listens to the VolumeSnapshot and VolumeSnapshotContent objects in the Kubernetes API and triggers csi-snapshotter to create a snapshot on the storage system when creating or deleting a VolumeSnapshot.	View details

Container Name	Container Image	Feature Description	Remark s
csi-node- driver- registrar	k8s.gcr.io/sig- storage/csi-node- driver- registrar:v2.3.0	Obtains CSI information and registers a node with kubelet using the plug-in registration mechanism of kubelet so that Kubernetes can detect the connection between the node and Huawei storage.	View details
huawei-csi- driver	The name and tag of huawei-csi-driver are specified in 3.4 Creating a Huawei CSI Image.	Connects to the Kubernetes platform to provide Huawei storage (centralized or distributed storage) resources for containers.	Version mappin gs

----End

5 Upgrade Operations

NOTICE

- During the upgrade, CSI cannot be used to deliver new resources.
- During the upgrade, do not uninstall the snapshot-dependent component service.
- The CSI upgrade does not affect delivered resources such as PVCs, snapshots, and Pods.
- 5.1 Upgrading Huawei CSI Using Helm
- 5.2 Manually Upgrading Huawei CSI

5.1 Upgrading Huawei CSI Using Helm

NOTE

This section describes how to upgrade or roll back Huawei CSI that is installed using Helm. For details about how to upgrade (uninstall or install) Huawei CSI that is manually installed, see 5.2 Manually Upgrading Huawei CSI.

5.1.1 Upgrading CSI

Prerequisites

When upgrading CSI, ensure that the parameter configurations in the **backends** field in the **values.yaml** file are the same as those in the **huawei-csi-configmap.yaml** file configured during Huawei CSI installation. Otherwise, CSI upon the upgrade cannot manage the previously provisioned resources such as PVCs and Pods.

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.

cp values.yaml values.yaml.bak

Step 4 Run the **vi** *values.yaml* command to open the file, and change the parameter values based on the upgrade requirements. After the change is complete, press **Esc** and enter :**wq!** to save the change. You can modify all parameters provided in **Table 4-1**. Exercise caution when modifying backend parameters.

vi values.yaml

Step 5 Run the helm upgrade helm-huawei-csi./ -n huawei-csi-f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

5.1.2 Rolling Back CSI

The Helm rollback command rolls back the current version to the specified version. If no version number is specified, it will be rolled back to the previous version.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Copy the **helm** directory in the Kubernetes CSI component package to any directory on the master node. For details about the Helm tool path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Check the version of Huawei CSI installed using Helm. You can run the **helm list A** command to query all Helm versions in all namespaces, and the **helm list n** *huawei-csi* command to query the Helm version in the specified namespace.

helm list -A
NAME NAMESPACE REVISION UPDATED STATUS CHART
APP VERSION
helm-test huawei-csi 1 2022-06-08 08:48:30.038729177 +0000 UTC deployed

esdk-1.0.0 3.0.0
helm list -n huawei-csi
NAME NAMESPACE REVISION UPDATED STATUS CHART
APP VERSION
helm-test huawei-csi 1 2022-06-08 08:48:30.038729177 +0000 UTC deployed
esdk-1.0.0 3.0.0

Step 4 Run the **helm rollback** *release-name revision-number* **-n** *huawei-csi* command to roll back Huawei CSI to the specified version. In the following example, the version is rolled back to **1**.

In the preceding command, *release-name* indicates the name of the chart to be rolled back, *revision-number* indicates the target version of the rollback, and *huawei-csi* indicates the namespace where the chart is located.

helm rollback helm-test 1 -n huawei-csi Rollback was a success! Happy Helming!

----End

5.2 Manually Upgrading Huawei CSI

5.2.1 Uninstalling Original CSI

Perform this operation when you want to uninstall CSI.

Preparations

Before uninstalling CSI, run the **kubectl get configmap huawei-csi-configmap -n huawei-csi -o yaml >> huawei-csi-configmap.yaml.bak** command to back up the content of the **huawei-csi-configmap** file. (During the CSI upgrade, the **backends** parameter in **huawei-csi-configmap.yaml** must be the same as the existing **configmap** configuration.)

5.2.1.1 Uninstalling the huawei-csi-node Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete daemonset huawei-csi-node -n** *huawei-csi* command to uninstall the huawei-csi-node service. Replace *huawei-csi* with the actual namespace.

kubectl delete daemonset huawei-csi-node -n huawei-csi

Step 3 Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get daemonset huawei-csi-node -n huawei-csi

----End

5.2.1.2 Uninstalling the huawei-csi-controller Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete deployment huawei-csi-controller -n** *huawei-csi* command to uninstall the huawei-csi-controller service. Replace *huawei-csi* with the actual namespace.
 - # kubectl delete deployment huawei-csi-controller -n huawei-csi
- **Step 3** Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get deployment huawei-csi-controller -n huawei-csi

----End

5.2.1.3 Deleting the huawei-csi-configmap Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete configmap** *huawei-csi-configmap* **-n** *huawei-csi* command to delete the **configmap** object. *huawei-csi-configmap* is the name of the **configmap** object, and *huawei-csi* is the namespace where the object is located.
 - # kubectl delete configmap huawei-csi-configmap -n huawei-csi
- **Step 3** Run the following command to check whether the object is successfully deleted. If **NotFound** is displayed, the object is successfully deleted.

kubectl get configmap huawei-csi-configmap -n huawei-csi

----End

5.2.1.4 Deleting the huawei-csi-secret Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete secret** *huawei-csi-secret* **-n** *huawei-csi* command to delete the **secret** object. *huawei-csi-secret* is the name of the **secret** object, and *huawei-csi* is the namespace where the **secret** object is located.
 - # kubectl delete secret huawei-csi-secret -n huawei-csi
- **Step 3** Run the following command to check whether the **secret** object is successfully deleted. If **NotFound** is displayed in the command output, the **huawei-csi-secret** object is successfully deleted.

kubectl get secret huawei-csi-secret -n huawei-csi Error from server (NotFound): secrets "huawei-csi-secret" not found

----End

5.2.1.5 Deleting the RBAC Permission

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Delete the RBAC permission.
 - If the huawei-csi version is later than 2.2.15, run the following command to delete the permission. -n indicates the namespace. Change it based on site requirements.
 - # kubectl -n huawei-csi -l provisioner=csi.huawei.com delete ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding
 - If the huawei-csi version is 2.2.15 or earlier, perform the following operations to delete the permission.
 - Run the following command to create a file named delete-huawei-csi-rbac.sh. -n indicates the namespace. Change it based on site requirements.

```
# cat <<EOF > delete-huawei-csi-rbac.sh
kubectl delete ServiceAccount huawei-csi-controller -n huawei-csi
kubectl delete ServiceAccount huawei-csi-node -n huawei-csi
kubectl delete ClusterRole huawei-csi-attacher-runner -n huawei-csi
kubectl delete ClusterRole huawei-csi-driver-registrar-runner -n huawei-csi
kubectl delete ClusterRole huawei-csi-provisioner-runner -n huawei-csi
kubectl delete ClusterRole huawei-csi-resizer-runner -n huawei-csi
kubectl delete ClusterRole huawei-csi-snapshotter-runner -n huawei-csi
kubectl delete ClusterRole snapshot-controller-runner -n huawei-csi
kubectl delete ClusterRoleBinding huawei-csi-attacher-role -n huawei-csi
kubectl delete ClusterRoleBinding huawei-csi-driver-registrar-role -n huawei-csi
kubectl delete ClusterRoleBinding huawei-csi-provisioner-role -n huawei-csi
kubectl delete ClusterRoleBinding huawei-csi-resizer-role -n huawei-csi
kubectl delete ClusterRoleBinding huawei-csi-snapshotter-role -n huawei-csi
kubectl delete ClusterRoleBinding snapshot-controller-role -n huawei-csi
kubectl delete Role huawei-csi-resizer-cfg -n huawei-csi
kubectl delete Role huawei-csi-snapshotter-leaderelection -n huawei-csi
kubectl delete Role snapshot-controller-leaderelection -n huawei-csi
kubectl delete RoleBinding huawei-csi-resizer-role-cfg -n huawei-csi
kubectl delete RoleBinding huawei-csi-snapshotter-leaderelection -n huawei-csi
kubectl delete RoleBinding snapshot-controller-leaderelection -n huawei-csi
```

- Run the following command to delete the RBAC permission. If the NotFound error is reported, ignore it.
 # sh delete-huawei-csi-rbac.sh
- **Step 3** Check whether the RBAC permission has been deleted.
 - If the huawei-csi version is later than 2.2.15, run the following command. -n indicates the namespace. Change it based on site requirements. If No resources found is displayed, the permission is successfully deleted.
 # kubectl -n huawei-csi -l provisioner=csi.huawei.com get

kubectt -n nuawei-csi -t provisioner=csi.nuawei.com get ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding

- If the huawei-csi version is 2.2.15 or earlier, perform the following operations to check whether the RBAC permission is successfully deleted.
 - a. Run the following command to create a file named **check-huawei-csi-rbac.sh**. **-n** indicates the namespace. Change it based on site requirements.

```
# cat <<EOF > check-huawei-csi-rbac.sh
kubectl get ServiceAccount -n huawei-csi | grep huawei-csi
```

kubectl get ClusterRole -n huawei-csi | grep huawei-csi kubectl get ClusterRoleBinding -n huawei-csi | grep huawei-csi kubectl get Role -n huawei-csi | grep huawei-csi kubectl get RoleBinding -n huawei-csi | grep huawei-csi kubectl get RoleBinding -n huawei-csi | grep huawei-csi kubectl get ClusterRole snapshot-controller-runner -n huawei-csi --ignore-not-found=true kubectl get ClusterRoleBinding snapshot-controller-role -n huawei-csi --ignore-not-found=true kubectl get RoleBinding snapshot-controller-leaderelection -n huawei-csi --ignore-not-found=true kubectl get RoleBinding snapshot-controller-leaderelection -n huawei-csi --ignore-not-found=true

b. Run the following command. If no command output is displayed, the RBAC permission has been successfully deleted.

sh check-huawei-csi-rbac.sh

----End

5.2.1.6 Deleting the Image of the Earlier Version

To delete the **huawei-csi** image from the cluster, you need to perform the deletion operation on all worker nodes.

To delete the image from a single node, perform the following steps.

Prerequisites

The container service that depends on the image has been stopped. Otherwise, the image cannot be deleted.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node through the management IP address.
- **Step 2** Run the following command to view all existing versions.
 - If docker is used, run the **docker image ls | grep huawei-csi** command.

docker image ls | grep huawei-csi
REPOSITORY TAG IMAGE ID CREATED SIZE
huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
huawei-csi 3.0.0 14b854dba227 2 weeks ago 79.6MB

If containerd is used, run the crictl image ls | grep huawei-csi command.

crictl image ls | grep huawei-csi
REPOSITORY TAG IMAGE ID CREATED SIZE
docker.io/library/huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
docker.io/library/huawei-csi 3.0.0 14b854dba227 2 weeks ago 79.6MB

- **Step 3** Run the following command to delete the image of the earlier version:
 - If docker is used, run the **docker rmi** *<REPOSITORY>*: *<TAG>* command. # docker rmi huawei-csi:2.2.15
 - If containerd is used, run the crictl rmi <REPOSITORY>:<TAG> command.
 # crictl rmi huawei-csi:2.2.15
- **Step 4** Run the following command again to check whether the image is successfully deleted. If the target version is not displayed, the image of the version is successfully deleted.
 - If docker is used, run the **docker image ls | grep huawei-csi** command.
 # docker image ls | grep huawei-csi
 huawei-csi 3.0.0 14b854dba227 10 minutes ago 80MB

• If containerd is used, run the **crictl image ls | grep huawei-csi** command.

crictl image ls | grep huawei-csi
docker.io/library/huawei-csi
3.0.0 14b854dba2273 93.1MB

----End

5.2.2 Installing New CSI

After the uninstallation is complete, you need to reinstall the CSI.

Prerequisites

The **huawei-csi-configmap.yaml** file of the original CSI has been backed up.

Precautions

If the template of **huawei-csi-configmap.yaml** has changed, ensure that the following parameter settings are the same as those before the upgrade. Otherwise, huawei-csi services cannot be started and created resources cannot be managed.

- The values of storage, name, and pools must be the same as those in the huawei-csi-configmap.yaml.bak file backed up in Prerequisites in 5.2.1 Uninstalling Original CSI.
- For details about urls and parameters, see the huawei-csi-configmap.yaml.bak file backed up in Prerequisites in 5.2.1 Uninstalling Original CSI and set them based on the huawei-csi-configmap.yaml template of the current version. For details about the template, see 4.2.1 Connecting to Enterprise Storage and 4.2.2 Connecting to Distributed Storage. The following command output is only an example.

```
"backends": [

{
    "storage": "oceanstor-san",
    "name": "***",
    "urls": ["https://*.*.*:8088", "https://*.*.*:8088"],
    "pools": ["***", "***"],
    "parameters": {"protocol": "iscsi", "portals": ["*.*.*", "*.*.*"]}
}

]
```

Procedure

- **Step 1** Obtain the CSI software package of the new version. For details, see **3.2 Obtaining the Software Package**.
- **Step 2** Create a CSI image of the new version. For details, see **3.4 Creating a Huawei CSI Image**.
- **Step 3** Create **huawei-csi-configmap**. For details, see **4.2.1 Connecting to Enterprise Storage** or **4.2.2 Connecting to Distributed Storage**.
- **Step 4** Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.

----End

6 Instructions for Use

This chapter describes how to use Huawei CSI.

- 6.1 (Conditionally Mandatory) Managing a StorageClass
- 6.2 (Conditionally Mandatory) Managing a PV
- 6.3 Managing a PVC
- 6.4 Managing a Pod
- 6.5 (Optional) Managing a Snapshot

6.1 (Conditionally Mandatory) Managing a StorageClass

A PV can be provisioned in either of the following modes: static provisioning and dynamic provisioning. In static provisioning mode, storage resources (LUNs/shares) are created on storage devices in advance and then statically provisioned to Kubernetes. In dynamic provisioning mode, CSI automatically creates storage resources (LUNs/shares).

Perform this operation when you want to use the dynamic PV function.

6.1.1 Creating a StorageClass

A StorageClass is a set of capabilities that can be selected when you apply for block storage resources. Kubernetes cluster users can create PVCs based on a StorageClass.

6.1.1.1 Creating a LUN StorageClass

This section describes how to create a LUN StorageClass.

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Run the **vi** *StorageClass.yaml* command to create a file named *StorageClass.yaml.* # vi StorageClass.yaml
- **Step 3** Press I or Insert to enter the editing mode and enter the following information in the *StorageClass.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *StorageClass.yaml* file. You can also refer to the **examples/lun-sc-for-csi-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-1**.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: "mysc" provisioner: "csi.huawei.com" parameters: volumeType: "lun" allocType: "thin"

Table 6-1 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a StorageClass object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
provisioner	provisioner identifier.	The value is fixed to csi.huawei.com.
parameters.volumeType	Type of the volume to be created.	The value is fixed to lun .
parameters.allocType	Allocation type of the volume to be created.	This parameter is optional. The value can be thin or thick , and the default value is thin .
parameters.cloneSpeed	Clone speed.	This parameter is optional. The value ranges from 1 to 4 and the default value is 3. 4 indicates the highest speed. This parameter is available when you clone a PVC or create a PVC using a snapshot. For details, see 6.3.3 (Optional) Cloning a PVC or 6.3.4 (Optional) Creating a PVC Using a Snapshot.

Parameter	Description	Remarks
parameters.fsType	File system type.	This parameter is optional. The value can be ext2, ext3, ext4, or xfs, and the default value is ext4. NOTICE CSI does not verify whether the value of the fsType parameter is valid. Ensure that the value is correct.

- **Step 4** Run the following command to create a StorageClass based on the .yaml file.

 # kubectl create -f StorageClass.yaml
- **Step 5** Run the following command to view the information about the created StorageClass.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION AGE
mysc csi.huawei.com Delete Immediate false 87s
```

----End

6.1.1.2 Creating a File System StorageClass

This section describes how to create a file system StorageClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to create a file named *StorageClass.yaml*. # vi StorageClass.yaml
- **Step 3** Press I or Insert to enter the editing mode and enter the following information in the *StorageClass.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *StorageClass.yaml* file. You can also refer to the **examples/fs-sc-for-csi-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-2**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "mysc"
provisioner: "csi.huawei.com"
mountOptions:
- nfsvers=3
parameters:
volumeType: "fs"
allocType: "thin"
authClient: "*"
```

Table 6-2 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a StorageClass object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
provisioner	provisioner identifier.	The value is fixed to csi.huawei.com.
mountOptions	Mount option.	Optional parameter after the -o parameter when the mount command is executed on the host. The value is in list format.
		If the NFS version is specified for mounting, use the nfsvers=* parameter. Currently, NFS 3, 4.0, and 4.1 protocols are supported (the protocol must be supported and enabled on storage devices). If nfsvers is set to 4 , the latest protocol version NFS 4 may be used for mounting due to different OS configurations, for example, 4.1. If protocol 4.0 is required, you are advised to set nfsvers to 4.0 .
parameters.volu meType	Type of the volume to be created.	The value is fixed to fs .
parameters.auth Client	Client that can access the volume.	This parameter is mandatory. OceanStor Dorado 6.x is used as an example. You can enter the client host name (a fully qualified domain name (FQDN) is recommended), client IP address, or client IP address segment, or use an asterisk (*) to represent all client IP addresses. The IP addresses can be IPv4 addresses, IPv6 addresses, or a combination of IPv4 and IPv6 addresses. You can enter multiple host names, IP addresses, or IP address segments and separate them with semicolons (;) or spaces or by pressing Enter.
		Example: 192.168.0.10;192.168.0.0/24;*

Parameter	Description	Remarks
parameters.alloc Type	Allocation type of the volume to be created.	This parameter is optional. The value can be thin or thick , and the default value is thin .
parameters.clone Speed	Clone speed.	This parameter is optional. The value ranges from 1 to 4 and the default value is 3. 4 indicates the highest speed. This parameter is available when you clone a PVC or create a PVC using a snapshot. For details, see 6.3.3 (Optional) Cloning a PVC or 6.3.4 (Optional) Creating a PVC Using a Snapshot.
parameters.fsTyp e	File system type.	This parameter is optional. The value can be ext2, ext3, ext4, or xfs, and the default value is ext4. NOTICE CSI does not verify whether the value of the fsType parameter is valid. Ensure that the value is correct.

- **Step 4** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorageClass.yaml
- **Step 5** Run the following command to view the information about the created StorageClass.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION AGE
mysc csi.huawei.com Delete Immediate false 34s
```

----End

6.1.2 Deleting a StorageClass

This section describes how to delete a StorageClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the following command to query StorageClasses in the cluster.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION
AGE
huawei-nas csi.huawei.com Delete Immediate false 3s
mysc csi.huawei.com Delete Immediate false 16s
```

Step 3 Run the following command to delete a StorageClass. For example, delete the StorageClass named *mysc*.

kubectl delete sc *mysc* storageclass.storage.k8s.io "*mysc*" deleted **Step 4** Run the following command to query StorageClasses in the cluster. If the command output does not contain the name of the StorageClass you want to delete, it is successfully deleted.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION
AGE
huawei-nas csi.huawei.com Delete Immediate false 3s
```

----End

6.2 (Conditionally Mandatory) Managing a PV

A PV can be provisioned in either of the following modes: static provisioning and dynamic provisioning. In static provisioning mode, storage resources (LUNs/shares) are created on storage devices in advance and then statically provisioned to Kubernetes. In dynamic provisioning mode, CSI automatically creates storage resources (LUNs/shares).

Perform this operation when you want to use the static PV function.

6.2.1 Creating a PV

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *PersistentVolume.yaml* command to create a file named *PersistentVolume.yaml*.

vi PersistentVolume.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *PersistentVolume.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *PersistentVolume.yaml* file. You can also refer to the **examples/static-pv-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-3**.

```
kind: PersistentVolume
apiVersion: v1
metadata:
name: mypv
spec:
volumeMode: Block
storageClassName: ""
accessModes:
- ReadWriteOnce
csi:
driver: csi.huawei.com
volumeHandle: <backendName>.<volume-name>
fsType: <string>
capacity:
storage: 100Gi
```

Table 6-3 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a PV object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.volumeMode	Volume mode.	This parameter is optional. The value can be Filesystem or Block . The default value is Filesystem . This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.
spec.storageClassNa me	Name of the StorageClass object.	This parameter is mandatory. Set it to an empty string, that is, enter "".
spec.persistentVolu meReclaimPolicy	Volume reclamation policy.	This parameter is optional. The value can be Retain (manual reclamation) or Delete (deleting associated storage resources). The default value is Retain . For details, see Table 6-5.

Parameter	Description	Remarks
spec.accessModes	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteOncePod.
		If the volume mode is Block , LUN volumes support ReadWriteOnce , ReadOnlyMany , ReadWriteMany , and ReadWriteOncePod . When ReadWriteMany is set, the Pod service must ensure data consistency.
		File system volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteMany, and ReadWriteOncePod.
		NOTICE The ReadWriteOncePod access mode is an alpha feature in Kubernetes v1.22+. Therefore, before using it, you need to enable the ReadWriteOncePod feature for the cluster. For details, see 9.7 Enabling the ReadWriteOncePod Feature Gate.
spec.csi.driver	CSI driver name.	The value is fixed to csi.huawei.com.
spec.csi.volumeHan dle	Unique identifier of a storage resource.	It consists of two parts: • <backernel name="">: indicates the name of each backend in configmap. You can run the following command to obtain configmap: kubectl get configmap huawei-csi-configmap -n huawei-csi-o yaml • <volume-name>: indicates the name of a resource (LUN/file system) on the storage. You can obtain the value from DeviceManager.</volume-name></backernel>

Parameter	Description	Remarks
spec.csi.fsType	File system type.	This parameter is optional. The value can be ext2, ext3, ext4, or xfs, and the default value is ext4. This parameter is valid only when volumeMode is set to Filesystem.
		NOTICE CSI does not verify whether the value of the fsType parameter is valid. Ensure that the value is correct.
spec.capacity.storag e	Volume size.	When creating a PV, ensure that its capacity is the same as that of the corresponding resource on the storage. Kubernetes will not invoke CSI to check whether the value of this parameter is correct. Therefore, the PV can be successfully created even if its capacity is inconsistent with that of the corresponding resource on the storage.

Step 4 Run the following command to create a PV based on the .yaml file.

kubectl create -f PersistentVolume.yaml

Step 5 After a period of time, run the following command to view the information about the created PV.

kubectl get pv

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS
REASON AGE

mypv 100Gi RWO Retain Available 4s

Table 6-4 PV status description

Status	Description
Available	Not bound to a PVC.
Bound	Bound to a PVC.
Released	A PVC has been deleted, but its resource has not been reclaimed by the cluster.
Failed	A volume fails to be automatically reclaimed.
Terminating	The PV object has been marked as "deleted" by the Kubernetes controller, but the PV cannot be deleted because it is being used.

Reclamation Description **Policy** • After the PVC object is deleted, the PV still exists, the Delete corresponding data volume is in the **Failed** status, and the PV cannot be bound to other PVCs. After the PV object is deleted, the corresponding file system or LUN in the back-end storage device will also be deleted. Retain • You can manually reclaim resources. After the PVC object is deleted, the PV still exists and the corresponding data volume is in the **Released** status. However, the PV cannot be bound to other PVCs because the data of the previous PVC remains on the PV. • After the PV object is deleted, the corresponding file system or LUN in the back-end storage device will be

Table 6-5 Reclamation policy description

----End

6.2.2 Deleting a PV

This section describes how to delete a PV.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query PVs in the cluster.

kubectl get pv

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS
REASON AGE

mypv 100Gi RWO Retain Available 11m

Step 3 Run the following command to delete a PV. For example, delete the PV named *mypv*.

kubectl delete pv *mypv* persistentvolume "mypv" deleted

Step 4 Run the following command to query PVs in the cluster. If the command output does not contain the name of the PV you want to delete, it is successfully deleted.

kubectl get pv No resources found in default namespace.

----End

6.3 Managing a PVC

6.3.1 Creating a PVC

This section describes how to create a PVC.

Restrictions

The storage limits the maximum number of concurrent RESTful requests to 100. Therefore, you are advised to create or delete a maximum of 100 PVCs in a batch.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *PersistentVolumeClaim.yaml* command to create a file named *PersistentVolumeClaim.yaml*.

vi PersistentVolumeClaim.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *PersistentVolumeClaim.yaml* file. After the modification is complete, press **Esc** and enter :wq! to save the modification.

The following shows a template of the *PersistentVolumeClaim.yaml* file. You can also refer to the **examples/pvc-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-6**.

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
name: "mypvc"
spec:
accessModes:
- ReadWriteOnce
volumeMode: Filesystem
volumeName: mypv
storageClassName: "mysc"
resources:
requests:
storage: 100Gi
```

Table 6-6 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a PVC object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.

Parameter	Description	Remarks
spec.volumeMode	Volume mode.	This parameter is optional. The value can be Filesystem or Block . The default value is Filesystem . This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.
spec.volumeName	Name of the PV object.	This parameter is mandatory only when a PVC is created statically.
spec.storageClass Name	Name of the StorageClass object.	When creating a PVC dynamically, enter the name of the StorageClass object created in 6.1 (Conditionally Mandatory) Managing a StorageClass. When creating a PVC
		When creating a PVC statically, set it to an empty string, that is, enter "".

Parameter	Description	Remarks
spec.resources.req uests.storage	Size of the volume to be created.	The value format is ***Gi. The unit is GiB. • The PVC capacity depends on storage specifications and host specifications. The following uses the connection between OceanStor Dorado 6.1.2/OceanStor Pacific series 8.1.0 and CentOS 7 as an example. See Table 6-7 and Table 6-8.
		 For other storage devices and hosts, check the specifications according to the value of VolumeType in StorageClass.
		 If the value of volumeType is lun, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home. In addition, refer to the host connectivity guide at https://support.huawei.com/enterprise/en/doc/EDOC1100113070/e067543b.
		 If the value of volumeType is fs, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home.
		If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications.
		When a PVC is created using a static PV and the PVC capacity is smaller than the capacity of the bound PV, the PVC capacity is set to the capacity of the bound PV. If the PVC capacity is greater than the capacity of the

Parameter	Description	Remarks
		bound PV, the PVC cannot be created.
spec.accessModes	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteOncePod.
		If the volume mode is Block, LUN volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteMany, and ReadWriteOncePod. When ReadWriteMany is set, the Pod service must ensure data consistency.
		File system volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteMany, and ReadWriteOncePod.
		NOTICE The ReadWriteOncePod access mode is an alpha feature in Kubernetes v1.22+. Therefore, before using it, you need to enable the ReadWriteOncePod feature for the cluster. For details, see 9.7 Enabling the ReadWriteOncePod Feature Gate.

Table 6-7 PVC capacity specifications (ext4)

volumeTyp e	Storage Type	Storage Specificati ons	ext4 Specification s	CSI Specification s
lun	OceanStor Dorado 6.1.2	512 Ki to 256 Ti	50 Ti	512 Ki to 50 Ti
	OceanStor Pacific series 8.1.0	64 Mi to 512 Ti	50 Ti	64 Mi to 50 Ti
fs	OceanStor Dorado 6.1.2	1 Gi to 32 Pi	N/A	1 Gi to 32 Pi
	OceanStor Pacific series 8.1.0	1 Ki to 256 Pi	N/A	1 Ki to 256 Pi

volumeTyp e	Storage Type	Storage Specificati ons	xfs Specification s	CSI Specificatio ns
lun	OceanStor Dorado 6.1.2	512 Ki to 256 Ti	500 Ti	512 Ki to 500 Ti
	OceanStor Pacific series 8.1.0	64 Mi to 512 Ti	500 Ti	64 Mi to 500 Ti
fs	OceanStor Dorado 6.1.2	1 Gi to 32 Pi	N/A	1 Gi to 32 Pi
	OceanStor Pacific series 8.1.0	1 Ki to 256 Pi	N/A	1 Ki to 256 Pi

Table 6-8 PVC capacity specifications (xfs)

Step 4 Run the following command to create a PVC based on the .yaml file.

kubectl create -f PersistentVolumeClaim.yaml

Step 5 After a period of time, run the following command to view the information about the created PVC.

```
# kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWX mysc 12s
```


After the PVC is created, if the PVC is in the **Pending** state, see **10.6 When a PVC Is Created**, **the PVC Is in the Pending State**.

----End

6.3.2 (Optional) Expanding the Capacity of a PVC

This section describes how to expand the capacity of a PVC.

Prerequisites

- A PVC has been created, and the backend where the PVC is located supports capacity expansion. For details about the storage devices that support capacity expansion, see Table 2-4 and Table 2-5.
- The huawei-csi services are running properly.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-fd5f97768-qlldc 7/7 Running 0 16s
huawei-csi huawei-csi-node-25txd 3/3 Running 0 15s
```

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Check whether the csi-resizer service is enabled for huawei-csi-controller.

```
# kubectl describe deploy huawei-csi-controller -n huawei-csi | grep csi-resizer csi-resizer:

Image: k8s.gcr.io/sig-storage/csi-resizer:v1.3.0
```

- If the preceding information is displayed, the csi-resizer service is enabled. In this case, go to **Step 3**.
- Otherwise, the csi-resizer service is not enabled. In this case, you need to upgrade huawei-csi to 2.2.15 or later.
- **Step 3** Run the **kubectl get pvc** *mypvc* command to query the StorageClass name of the PVC. In the preceding command, *mypvc* indicates the name of the PVC to be expanded.

```
# kubectl get pvc mypvc
NAME STATUS VOLUME CAPACITY ACCESS MODES
STORAGECLASS AGE
mypvc Bound pvc-3383be36-537c-4cb1-8f32-a415fa6ba384 2Gi RWX
mysc 145m
```

Step 4 Run the **kubectl get sc** *mysc* command to check the StorageClass supports capacity expansion. In the preceding command, *mysc* indicates the name of the StorageClass to be gueried.

```
# kubectl get sc mysc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE

ALLOWVOLUMEEXPANSION AGE

mysc csi.huawei.com Delete Immediate true 172m
```

- If the value of **ALLOWVOLUMEEXPANSION** is **false**, the current StorageClass does not support capacity expansion. In this case, go to **Step 5**.
- If the value of **ALLOWVOLUMEEXPANSION** is **true**, the current StorageClass supports capacity expansion. In this case, go to **Step 6**.
- **Step 5** (Optional) Run the following command to change the value of **allowVolumeExpansion** to **true**. In the preceding command, *mysc* indicates the name of the StorageClass to be modified.

kubectl patch sc mysc --patch '{"allowVolumeExpansion":true}'

Step 6 Run the following command to expand the capacity.

```
# kubectl patch pvc mypvc-p '{"spec":{"resources":{"requests":{"storage":"120Gi'}}}}'
```

In the preceding command, *mypvc* indicates the name of the PVC to be expanded, and *120Gi* indicates the capacity after expansion. Change the values based on the site requirements.

- The PVC capacity depends on storage specifications and host specifications. The following uses the connection between OceanStor Dorado 6.1.2 or OceanStor Pacific series 8.1.0 and CentOS 7 as an example. For details, see Table 6-9.
- For other storage devices and hosts, check the specifications according to the value of VolumeType in 6.1.1 Creating a StorageClass.
 - If the value of volumeType is lun, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home. In addition, refer to the host connectivity guide at https://support.huawei.com/enterprise/en/doc/EDOC1100113070/e067543b.
 - If the value of volumeType is fs, refer to the storage specifications. For details, see https://info.support.huawei.com/storage/spec/#/home.
- If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications.

volumeTyp **Storage Type** Storage ext4 Specificati Specification **Specification** e ons OceanStor Dorado 512 Ki to 50 512 Ki to 50 Ti lun 6.1.2 256 Ti OceanStor Pacific 50 Ti 64 Mi to 50 64 Mi to series 8.1.0 512 Ti Τi OceanStor Dorado 1 Gi to 32 1 Gi to 32 Pi fs N/A 6.1.2 Ρi OceanStor Pacific 1 Ki to 256 N/A 1 Ki to 256 Pi series 8.1.0 Ρi

Table 6-9 PVC capacity specifications (ext4)

Table 6-10 PVC capacity specifications (xfs)

volumeTyp e	Storage Type	Storage Specificati ons	xfs Specification s	CSI Specificatio ns
lun	OceanStor Dorado 6.1.2	512 Ki to 256 Ti	500 Ti	512 Ki to 500 Ti
	OceanStor Pacific series 8.1.0	64 Mi to 512 Ti	500 Ti	64 Mi to 500 Ti
fs	OceanStor Dorado 6.1.2	1 Gi to 32 Pi	N/A	1 Gi to 32 Pi
	OceanStor Pacific series 8.1.0	1 Ki to 256 Pi	N/A	1 Ki to 256 Pi

Step 7 Run the following command to check whether the capacity changes.

kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 120Gi RWX mysc 24s

----End

6.3.3 (Optional) Cloning a PVC

Perform this operation when you want to clone an existing PVC on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *clone.yaml* command to create a file named *clone.yaml*.

 # vi clone.yaml

Step 3 Configure the *clone.yaml* file. The **examples/clone.yaml** template file is as follows. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-11**.

kind: PersistentVolumeClaim apiVersion: v1 metadata: name: myclone spec: storageClassName: mysc dataSource: name: mypvc kind: PersistentVolumeClaim volumeMode: Filesystem accessModes: - ReadWriteOnce resources: requests: storage: 2Gi

Table 6-11 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a new PVC object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.storageClass Name	Name of the StorageClass object.	Enter the name of the StorageClass object created in 6.1 (Conditionally Mandatory) Managing a StorageClass. The value must be the same as the name of the StorageClass in dataSource.
spec.dataSource. name	Name of the source PVC object.	-
spec.volumeMod e	Volume mode.	This parameter is optional. The value can be Filesystem or Block. The default value is Filesystem. The value must be the same as the volumeMode value of the source PVC object. This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.

Parameter	Description	Remarks
spec.accessMode s	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteOncePod.
		If the volume mode is Block, LUN volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteMany, and ReadWriteOncePod. When ReadWriteMany is set, the Pod service must ensure data consistency.
		File system volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteMany, and ReadWriteOncePod.
		NOTICE The ReadWriteOncePod access mode is an alpha feature in Kubernetes v1.22+. Therefore, before using it, you need to enable the ReadWriteOncePod feature for the cluster. For details, see 9.7 Enabling the ReadWriteOncePod Feature Gate.
spec.resources.re quests.storage	Size of the volume to be created.	The value must be greater than or equal to the size of the source PVC. The value format is ***Gi. The unit is GiB.

Step 4 Run the following command to create a PVC based on the .yaml file.

kubectl create -f clone.yaml

----End

6.3.4 (Optional) Creating a PVC Using a Snapshot

Perform this operation when you want to create a PVC for an existing snapshot on Kubernetes.

Prerequisites

A snapshot has been created. For details, see **6.5 (Optional) Managing a Snapshot**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** restore.yaml command to create a file named restore.yaml. # vi restore.yaml
- **Step 3** Configure the *restore.yaml* file. The **examples**/*<Kubernetes version*>/restore.yaml template file is as follows. Set related parameters based on the site requirements and save the file in yaml format. For details, see Table 6-12.

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
name: ***
spec:
 storageClassName: ***
 dataSource:
  name: ***
  kind: VolumeSnapshot
  apiGroup: snapshot.storage.k8s.io
 volumeMode: Filesystem
 accessModes:
  - ReadWriteOnce
 resources:
  requests:
   storage: ***Gi
```

Table 6-12 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a new PVC object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.storageClass Name	Name of the StorageClass object.	Enter the name of the StorageClass object created in 6.1 (Conditionally Mandatory) Managing a StorageClass. The value must be the same as the name of the StorageClass of the original PVC in dataSource.
spec.dataSource.n ame	Name of the source VolumeSnapshot object.	-

Parameter	Description	Remarks
spec.volumeMode	Volume mode.	This parameter is optional. The value can be Filesystem or Block. The default value is Filesystem. The value must be the same as the volumeMode value of the source PVC object. This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage.
spec.accessModes	Access mode of the volume.	If the volume mode is Filesystem, LUN volumes support ReadWriteOnce, ReadOnlyMany, and ReadWriteOncePod. If the volume mode is Block, LUN volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteMany, and ReadWriteMany is set, the Pod service must ensure data consistency. File system volumes support ReadWriteOnce, ReadOnlyMany, ReadWriteOnce, ReadOnlyMany, ReadWriteOncePod. NOTICE The ReadWriteOncePod access mode is an alpha feature in Kubernetes v1.22+. Therefore, before using it, you need to enable the ReadWriteOncePod feature for the cluster. For details, see 9.7 Enabling the ReadWriteOncePod Feature Gate.
spec.resources.req uests.storage	Size of the volume to be created.	The value must be greater than or equal to the size of the source VolumeSnapshot. The value format is ***Gi. The unit is GiB.

Step 4 Run the following command to create a PVC based on the .yaml file.

kubectl create -f restore.yaml

----End

6.3.5 Deleting a PVC

This section describes how to delete a PVC.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query PVCs in the cluster.

```
# kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWX mysc 12s
```

∩ NOTE

Before deleting a PVC, if the PVC is in the **Pending** state, you are not advised to directly delete the PVC. To delete the PVC, see **10.7 Before a PVC Is Deleted, the PVC Is in the Pending State**.

Step 3 Run the following command to delete a PVC. For example, delete the PVC named *mypvc*.

```
# kubectl delete pvc mypvc persistentvolumeclaim "mypvc" deleted
```

Step 4 Run the following command to query PVCs in the cluster. If the command output does not contain the name of the PVC you want to delete, it is successfully deleted.

```
# kubectl get pvc
No resources found in default namespace.
```

----End

6.4 Managing a Pod

6.4.1 Creating a Pod

This section describes how to create a Pod. A Pod is an original storage pool or storage function set. It function as the container of virtual volumes, which means only the storage container allocates storage space to virtual volumes. This operation enables you to quickly obtain specified storage resources.

Restrictions

You are advised to create or delete a maximum of 100 Pods in a batch.

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Run the **vi** *pod.yaml* command to create a file named *pod.yaml*.

 # vi pod.yaml
- **Step 3** Press **I** or **Insert** to enter the editing mode and enter the following information in the *pod.yaml* file. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - If **volumeMode** is set to **Filesystem** in **Step 3**, the following shows a template of the *pod.yaml* file. You can also refer to the **examples/pod-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-13**.

```
kind: Pod
apiVersion: v1
metadata:
name: "mypod"
spec:
containers:
- name: "mycontainer"
image: "***"
volumeMounts:
- name: mypv
mountPath: "/mnt/path/in/container"
volumes:
- name: mypv
persistentVolumeClaim:
claimName: "mypvc"
```

Table 6-13 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a Pod object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.nam e	User-defined container name.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.imag e	Container image.	Set this parameter based on the site requirements.
spec.containers.imag e.volumeMounts.mo untPath	Volume mount path in the container.	-
spec.volumes.persist entVolumeClaim.clai mName	Name of the PVC object.	Enter the name of the PVC object created in 6.3.1 Creating a PVC.

• If **volumeMode** is set to **Block** in **Step 3**, the following shows a template of the *pod.yaml* file. You can also refer to the **examples/pod-rbd-example.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-14**. kind: Pod

```
apiVersion: v1
metadata:
name: "mypod"
spec:
containers:
- name: "mycontainer"
image: "***"
volumeDevices:
- name: mypv
devicePath: "/dev/xvda"
volumes:
- name: mypv
persistentVolumeClaim:
claimName: "mypvc"
```

Table 6-14 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a Pod object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.nam e	User-defined container name.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.containers.imag e	Container image.	Set this parameter based on the site requirements.
spec.containers.imag e.volumeDevices.devi cePath	Volume device path in the container.	-
spec.volumes.persiste ntVolumeClaim.claim Name	Name of the PVC object.	Enter the name of the PVC object created in 6.3.1 Creating a PVC .

Step 4 Run the following command to create a Pod based on the .yaml file.

kubectl create -f pod.yaml

Step 5 Run the following command to view the information about the created Pod.

```
# kubectl get pod
NAME READY STATUS RESTARTS AGE
mypod 1/1 Running 0 37s
```

----End

6.4.2 Deleting a Pod

This section describes how to delete a Pod.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query Pods in the cluster.

```
# kubectl get pod
NAME READY STATUS RESTARTS AGE
mypod 1/1 Running 0 14h
```

Step 3 Run the following command to delete a Pod. For example, delete the Pod named *mypod*.

```
# kubectl delete pod mypod pod "mypod" deleted
```

Step 4 Run the following command to query Pods in the cluster. If the command output does not contain the name of the Pod you want to delete, it is successfully deleted.

```
# kubectl get pod
No resources found in default namespace.
```

----End

6.5 (Optional) Managing a Snapshot

6.5.1 Installing the Snapshot-Dependent Component Service

Procedure for Installation Using Helm

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl api-resources** | **grep snapshot** | **awk '{print \$1}'** command to view the installation details of snapshot-related resource services.

If the following information is displayed, the snapshot-dependent component service has been installed. In this case, skip this section. If any of the following services in the command output is not displayed, go to the next step to install it. # kubectl api-resources | grep snapshot | awk '{print \$1}' volumesnapshotclasses volumesnapshotcontents volumesnapshots

Step 3 Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details, see **3.3 Uploading the Components in the Software Package**.

Step 4 Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.

cp values.yaml values.yaml.bak

Step 5 Run the **vi** *values.yaml* command to open the file, and change the value of **snapshot.enable** to **true**. After the change is complete, press **Esc** and enter **:wq!** to save the change.

vi values.yaml

Step 6 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

Procedure for Manual Installation

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl api-resources** | **grep snapshot** | **awk** '**{print \$1}**' command to view the installation details of snapshot-related resource services.

If the following information is displayed, the snapshot-dependent component service has been installed. In this case, skip this section. If any of the following services in the command output is not displayed, go to the next step to install it. # kubectl api-resources | grep snapshot | awk '{print \$1}' volumesnapshotclasses volumesnapshotcontents

- Step 3 Copy the **deploy** directory in the Kubernetes CSI component package to any directory. For details about the component package path, see 3.3 Uploading the Components in the Software Package.
- **Step 4** Go to the **deploy** directory and run the following command to install the snapshot-dependent component service.

kubectl apply -f huawei-csi-snapshot-crd.yaml

----End

volumesnapshots

6.5.2 Managing a VolumeSnapshotClass

6.5.2.1 Creating a VolumeSnapshotClass

This section describes how to create a VolumeSnapshotClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *volume-snapshot-class.yaml* command to create a file named *volume-snapshot-class.yaml*.

vi volume-snapshot-class.yaml

Step 3 Press I or Insert to enter the editing mode and enter the following information in the *volume-snapshot-class.yaml* file. After the modification is complete, press Esc and enter :wq! to save the modification.

The following shows a template of the *volume-snapshot-class.yaml* file. You can also refer to the **examples/volume-snapshot-class.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-15**.

apiVersion: snapshot.storage.k8s.io/v1beta1 kind: VolumeSnapshotClass metadata: name: mysnapclass driver: csi.huawei.com deletionPolicy: Delete

Table 6-15 Parameter description

Parameter	Description	Remarks
metadata.nam e	User-defined name of a VolumeSnapshotClass object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
driver	driver identifier.	The value is fixed to csi.huawei.com.
deletionPolicy	Handles the VolumeSnapshotContent policy when a VolumeSnapshot is deleted.	This parameter is mandatory. The value can be Delete or Retain .

Step 4 Run the following command to create a VolumeSnapshotClass based on the .yaml file.

kubectl create -f volume-snapshot-class.yaml

Step 5 Run the following command to view the information about the created VolumeSnapshotClass.

kubectl get volumesnapshotclass NAME DRIVER DELETIONPOLICY AGE mysnapclass csi.huawei.com Delete 25s

----End

6.5.2.2 Deleting a VolumeSnapshotClass

This section describes how to delete a VolumeSnapshotClass.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query VolumeSnapshotClasses in the cluster.

```
# kubectl get volumesnapshotclass
NAME DRIVER DELETIONPOLICY AGE
mysnapclass csi.huawei.com Delete 52s
```

Step 3 Run the following command to delete a VolumeSnapshotClass. For example, delete the VolumeSnapshotClass named *mysnapclass*.

```
# kubectl delete volumesnapshotclass mysnapclass volumesnapshotclass.snapshot.storage.k8s.io "mysnapclass" deleted
```

Step 4 Run the following command to query VolumeSnapshotClasses in the cluster. If the command output does not contain the name of the VolumeSnapshotClass you want to delete, it is successfully deleted.

```
# kubectl get volumesnapshotclass
No resources found in default namespace.
```

----End

6.5.3 Managing a VolumeSnapshot

6.5.3.1 Creating a VolumeSnapshot

This section describes how to create a VolumeSnapshot.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *volume-snapshot.yaml* command to create a file named *volume-snapshot.yaml*.

vi volume-snapshot.yaml

Step 3 Press **I** or **Insert** to enter the editing mode and enter the following information in the *volume-snapshot.yaml* file. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

The following shows a template of the *volume-snapshot.yaml* file. You can also refer to the **examples/volume-snapshot.yaml** example file in the software package. Set related parameters based on the site requirements and save the file in yaml format. For details, see **Table 6-16**.

```
apiVersion: snapshot.storage.k8s.io/v1beta1 kind: VolumeSnapshot metadata: name: mysnapshot spec: volumeSnapshotClassName: mysnapclass
```

source: persistentVolumeClaimName: mypvc

Table 6-16 Parameter description

Parameter	Description	Remarks
metadata.name	User-defined name of a VolumeSnapshot object.	Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.
spec.volumeSnapsho tClassName	Name of the VolumeSnapshotClass object.	Enter the name of the VolumeSnapshotClass object created in 6.5.2.1 Creating a VolumeSnapshotClass.
spec.source.persisten tVolumeClaimName	Name of the source PVC object.	Enter the name of the PVC object created in 6.3.1 Creating a PVC.

- **Step 4** Run the following command to create a VolumeSnapshot based on the .yaml file. # kubectl create -f volume-snapshot.yaml
- **Step 5** Run the following command to view the information about the created VolumeSnapshot.

kubectl get volumesnapshot
NAME READYTOUSE SOURCEPVC SOURCESNAPSHOTCONTENT RESTORESIZE
SNAPSHOTCLASS SNAPSHOTCONTENT CREATIONTIME AGE
mysnapshot true mypvc 100Gi mysnapclass
snapcontent-1009af0a-24c2-4435-861c-516224503f2d <invalid> 78s

----End

6.5.3.2 Deleting a VolumeSnapshot

This section describes how to delete a VolumeSnapshot.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query VolumeSnapshots in the cluster.

kubectl get volumesnapshot
NAME READYTOUSE SOURCEPVC SOURCESNAPSHOTCONTENT RESTORESIZE
SNAPSHOTCLASS SNAPSHOTCONTENT CREATIONTIME AGE
mysnapshot true mypvc 100Gi mysnapclass
snapcontent-1009af0a-24c2-4435-861c-516224503f2d <invalid> 78s

Step 3 Run the following command to delete a VolumeSnapshot. For example, delete the VolumeSnapshot named *mysnapshot*.

kubectl delete volumesnapshot mysnapshot volumesnapshot.snapshot.storage.k8s.io "mysnapshot" deleted

Step 4 Run the following command to query VolumeSnapshots in the cluster. If the command output does not contain the name of the VolumeSnapshot you want to delete, it is successfully deleted.

kubectl get volumesnapshot No resources found in default namespace.

----End

7 Advanced Features

This chapter describes how to configure advanced features of Huawei storage.

- 7.1 Configuring Multiple Backends
- 7.2 Creating a PVC for a Specified Backend
- 7.3 Creating a PVC for a Specified Storage Pool
- 7.4 Configuring ALUA
- 7.5 Configuring Storage Topology Awareness
- 7.6 Advanced Features of Enterprise Storage
- 7.7 Advanced Features of Distributed Storage

7.1 Configuring Multiple Backends

Huawei CSI supports multiple backends. Perform this operation when you want to configure multiple backends.

7.1.1 Configuring Multiple Backends Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.
 - # cp values.yaml values.yaml.bak

Step 4 Run the **vi** *values.yaml* command to open the file and configure multiple backends as required. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
backends:
 storage: "oceanstor-san"
name: "***"
    - "https://*.*.*:8088"
  pools:
  parameters:
    protocol: "iscsi"
    portals:
      - "*.*.*.
 - storage: "oceanstor-nas"
  name: "***"
  urls:
    - "https://*.*.*:8088"
  pools:
   parameters:
    protocol: "nfs"
    portals:
     - "*.*.*"
```

Step 5 Run the helm upgrade helm-huawei-csi ./ -n huawei-csi -f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

```
# helm upgrade helm-huawei-csi / -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

7.1.2 Manually Configuring Multiple Backends

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Configure the **huawei-csi-configmap.yaml** file. The following shows a template of the **huawei-csi-configmap.yaml** file. Set related parameters based on the site requirements and save the file in yaml format.

Multiple backends are separated by commas (,). For details about each backend, see **4.2.1 Connecting to Enterprise Storage** or **4.2.2 Connecting to Distributed Storage**.

```
kind: ConfigMap
apiVersion: v1
```

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the kubectl get configmap -n huawei-csi | grep huawei-csi-configmap command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

Step 5 Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.

----End

7.2 Creating a PVC for a Specified Backend

When multiple backends are configured, you can perform the following operations to create a PVC for a specified backend.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify parameters in the following fields. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - Add the **backend** configuration item under **parameters**.
 - The value of metadata.name is the user-defined name of a StorageClass object.

Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.

• The value of **parameters.backend** is the name of a backend in **huawei-csi-configmap.yaml**.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata:

```
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
backend: "***"
```

Step 3 Run the following command to create a StorageClass based on the .yaml file.

```
# kubectl create -f StorageClass.yaml
```

----End

7.3 Creating a PVC for a Specified Storage Pool

When multiple storage pools are configured, you can perform the following operations to create a PVC for a specified storage pool.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify parameters in the following fields. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - The value of **metadata.name** is the user-defined name of a StorageClass object.

Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit.

- Add the pool configuration item under parameters.
- The value of **pool** is the name of a storage pool in **huawei-csi-configmap.yaml**.

□ NOTE

The volume to be created using the StorageClass will be created in the specified storage pool. The existing PVC will not change the storage pool information.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
pool: "***"
```

Step 3 Run the following command to create a StorageClass based on the .yaml file.

```
# kubectl create -f StorageClass.yaml
```

----End

7.4 Configuring ALUA

7.4.1 Configuring ALUA Using Helm

This section describes how to configure ALUA using Helm.

7.4.1.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3 Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- Step 3 Back up the values.yaml file used during CSI installation. If the values.yaml file used during the last installation cannot be found, run the helm get values helm-huawei-csi -n huawei-csi -a command to query the file.

 # cp values.yaml values.yaml.bak
- **Step 4** Run the **vi** *values.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-1**.

```
backends:
 - storage: "oceanstor-san"
  name: "storage"
  urls:
   - "https://*.*.*.*:8088"
   - "https://*.*.*:8088"
  pools:
   - "***"
  parameters:
   protocol: "iscsi"
   portals:
     - "*.*.*"
   AI UA:
     HostName1:
      MULTIPATHTYPE: "*"
      FAILOVERMODE: "*"
      SPECIALMODETYPE: "*"
      PATHTYPE: "*
     HostName2:
      MULTIPATHTYPE: "*"
      FAILOVERMODE: "*"
      SPECIALMODETYPE: "*"
      PATHTYPE: "*"
```

Table 7-1 ALUA parameter description

Parameter	Description	Remarks
HostName	The value of HostName is the host name of a worker node, for example, HostName1 and HostName2 .	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
MULTIPATHTYPE	Multipathing type. The value can be: • 0: default • 1: third-party multipathing	-
FAILOVERMODE	Initiator switchover mode. The value can be: • 0: early-version ALUA • 1: common ALUA • 2: ALUA not used • 3: special ALUA	This parameter needs to be specified only when third-party multipathing is used. All OceanStor V5 models do not support early-version ALUA.
SPECIALMODETYPE	Special mode type of the initiator. The value can be: • 0: special mode 0 • 1: special mode 1 • 2: special mode 2 • 3: special mode 3	This parameter needs to be specified only when the initiator switchover mode is special ALUA.
PATHTYPE	Initiator path type. The value can be: • 0: preferred path • 1: non-preferred path	This parameter needs to be specified only when third-party multipathing is used.

□ NOTE

- The ALUA configuration may vary according to the OS. Visit https://support.huawei.com/enterprise/en/index.html, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS and configure ALUA based on the recommended configurations in the guide.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- The value of **HostName** is a regular expression. For details about how to configure it, see **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used. When you configure **HostName**, the number of host connections is limited. For details about the limitation, see **Specifications Query** and search for **Maximum number of iSCSI connections per controller enclosure**. If the number of host connections is less than or equal to the specifications, you are advised to use the general configuration. If the number of host connections is greater than the specifications, you are advised to use the common configuration.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- In the general configuration, if you need to exactly match a host, refer to example 3 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- OceanStor V3/V5 and OceanStor Dorado V3 use this configuration mode. For details about related parameters, see Table 7-2.

Table 7-2 Recommended ALUA parameter configurations for OceanStor V3/V5 and OceanStor Dorado V3

Scenario	Host Type	Whether the Storage Has the Preferred Path	Recommended ALUA Configuration
HyperMetro storage		Yes	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="0"
		No	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="1"
	SUSE/Debian host	Yes	ALUA="1" FAILOVERMODE="1" PATHTYPE="0"

Scenario	Host Type	Whether the Storage Has the Preferred Path	Recommended ALUA Configuration
		No	ALUA="1" FAILOVERMODE="1" PATHTYPE="1"
Non- HyperMetro storage	CentOS/RHEL host	N/A	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="0"
	SUSE/Debian host	N/A	ALUA="1" FAILOVERMODE="1" PATHTYPE="0"

Step 5 Run the helm upgrade helm-huawei-csi./ -n huawei-csi-f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi

Release "helm-huawei-csi" has been upgraded. Happy Helming!

NAME: helm-huawei-csi

LAST DEPLOYED: Thu Jun 9 07:58:15 2022

NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

7.4.1.2 Configuring ALUA for OceanStor Dorado 6.x Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.

cp values.yaml values.yaml.bak

Step 4 Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-3**.

```
- storage: "oceanstor-san"
 name: "storage"
 urls:
  - "https://*.*.*:8088"
- "https://*.*.*:8088"
 pools:
  - "***"
 parameters:
  protocol: "iscsi"
  portals:
  ALUA:
    HostName1:
     accessMode: "*"
     hyperMetroPathOptimized: "*"
    HostName2:
     accessMode: "*"
     hyperMetroPathOptimized: "*"
```

Table 7-3 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node, for example, HostName1 and HostName2 .	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
accessMode	Host access mode. The value can be: • 0: balanced mode • 1: asymmetric mode	-
hyperMetroPathOp- timized	Whether the path of the host on the current storage array is preferred in HyperMetro scenarios. The value can be:	This parameter needs to be specified only when the host access mode is set to asymmetric.
	• 1: yes	
	• 0 : no	

□ NOTE

- The ALUA configuration may vary according to the OS. Visit https://support.huawei.com/enterprise/en/index.html, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS and configure ALUA based on the recommended configurations in the guide.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- The value of **HostName** is a regular expression. For details about how to configure it, see **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used. When you configure **HostName**, the number of host connections is limited. For details about the limitation, see **Specifications Query** and search for **Maximum number of iSCSI connections per controller enclosure**. If the number of host connections is less than or equal to the specifications, you are advised to use the general configuration. If the number of host connections is greater than the specifications, you are advised to use the common configuration.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- In the general configuration, if you need to exactly match a host, refer to example
 3 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- If a host uses only OceanStor Dorado 6.x all-flash storage, see **Table 7-4** for detailed parameters.
- If you add OceanStor Dorado 6.x all-flash storage to a host that uses OceanStor converged storage, see **Table 7-5** for detailed parameters.

Table 7-4 Recommended ALUA parameter configurations for OceanStor Dorado 6.*x* all-flash storage

Scenario	Host Type	Host Access Mode	Recommended ALUA Configuration
HyperMetro storage	CentOS/ RHEL/SUSE/ Debian host	Load balancing mode	ALUA not required
		Asymmetric mode + Storage with the preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the non-preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="0"
Non- HyperMetro storage	CentOS/ RHEL/SUSE/ Debian host	N/A	ALUA not required

Scenario	Host Type	Host Access Mode	Recommended ALUA Configuration
HyperMet ro storage	CentOS/RHEL/SUSE/ Debian host	Load balancing mode	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the non- preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="0"
Non- HyperMet ro storage	CentOS/RHEL/SUSE/ Debian host	N/A	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"

Table 7-5 Recommended ALUA parameter configurations for hybrid OceanStor V3/V5, OceanStor Dorado V3, and OceanStor Dorado 6.*x*

Step 5 Run the helm upgrade helm-huawei-csi ./ -n huawei-csi -f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

7.4.1.3 Configuring ALUA for Distributed Storage Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any

directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.

- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.
 - # cp values.yaml values.yaml.bak
- **Step 4** Run the **vi** *values.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-6**.

```
backends:
 - storage: "oceanstor-san"
  name: "storage"
  urls:
   - "https://*.*.*.*:8088"
   - "https://*.*.*:8088"
  pools:
   _ "***"
  parameters:
   protocol: "iscsi"
    portals:
    ALUA:
     HostName1:
      switchoverMode: "*"
      pathType: "*"
     HostName2:
      switchoverMode: "*"
      pathType: "*"
```

Table 7-6 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node, for example, HostName1 and HostName2 .	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
switchoverMode	Switchover mode. The value can be:	-
	• Disable_alua : disables ALUA.	
	• Enable_alua: enables ALUA.	

Parameter	Description	Remarks
pathType	Path type. The value can be:	-
	optimal_path: preferred path	
	non_optimal_path: non-preferred path	

■ NOTE

- Only the iSCSI scenario of distributed storage is supported.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- When the value of **HostName** is a regular expression, configure it by referring to **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- In the general configuration, if you need to exactly match a host, refer to example
 3 in 11.3 Example ALUA Configuration Policy of Distributed Storage.

Step 5 Run the helm upgrade helm-huawei-csi ./ -n huawei-csi -f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

7.4.2 Manually Configuring ALUA

This section describes how to manually configure ALUA.

7.4.2.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3

This section describes how to configure ALUA if multipathing is used during the connection to block storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see **4.2.1 Connecting to Enterprise Storage** or **4.2.2 Connecting to Distributed Storage**.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-7**.

```
{
    "backends": [
    {
        "storage": "oceanstor-san",
        ...
        "parameters": {..., "ALUA": {"<HostName>": {"MULTIPATHTYPE": "*", "FAILOVERMODE": "*",
"SPECIALMODETYPE": "*", "PATHTYPE": "*"}, "<HostName>": {...}}}
    }
}
```

Table 7-7 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node.	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
MULTIPATHTYPE	Multipathing type. The value can be: • 0: default • 1: third-party multipathing	-

Parameter	Description	Remarks
FAILOVERMODE	Initiator switchover mode. The value can be: • 0: early-version ALUA • 1: common ALUA • 2: ALUA not used • 3: special ALUA	This parameter needs to be specified only when third-party multipathing is used. All OceanStor V5 models do not support early-version ALUA.
SPECIALMODETYPE	Special mode type of the initiator. The value can be: • 0: special mode 0 • 1: special mode 1 • 2: special mode 2 • 3: special mode 3	This parameter needs to be specified only when the initiator switchover mode is special ALUA.
PATHTYPE	Initiator path type. The value can be: • 0: preferred path • 1: non-preferred path	This parameter needs to be specified only when third-party multipathing is used.

□ NOTE

- The ALUA configuration may vary according to the OS. Visit https://support.huawei.com/enterprise/en/index.html, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS and configure ALUA based on the recommended configurations in the guide.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- The value of **HostName** is a regular expression. For details about how to configure it, see **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used. When you configure **HostName**, the number of host connections is limited. For details about the limitation, see **Specifications Query** and search for **Maximum number of iSCSI connections per controller enclosure**. If the number of host connections is less than or equal to the specifications, you are advised to use the general configuration. If the number of host connections is greater than the specifications, you are advised to use the common configuration.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- In the general configuration, if you need to exactly match a host, refer to example 3 in 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3.
- OceanStor V3/V5 and OceanStor Dorado V3 use this configuration mode. For details about related parameters, see Table 7-8.

Table 7-8 Recommended ALUA parameter configurations for OceanStor V3/V5 and OceanStor Dorado V3

Scenario	Host Type	Whether the Storage Has the Preferred Path	Recommended ALUA Configuration
HyperMetro storage		Yes	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="0"
		No	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="1"
	SUSE/Debian host	Yes	ALUA="1" FAILOVERMODE="1" PATHTYPE="0"

Scenario	Host Type	Whether the Storage Has the Preferred Path	Recommended ALUA Configuration
		No	ALUA="1" FAILOVERMODE="1" PATHTYPE="1"
Non- HyperMetro storage	CentOS/RHEL host	N/A	ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0" PATHTYPE="0"
	SUSE/Debian host	N/A	ALUA="1" FAILOVERMODE="1" PATHTYPE="0"

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n huawei-csi** | **grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

Step 5 Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.

----End

7.4.2.2 Configuring ALUA for OceanStor Dorado 6.x

This section describes how to configure ALUA if multipathing is used during the connection to block storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.2.1 Connecting to Enterprise Storage or 4.2.2 Connecting to Distributed Storage.

Add ALUA parameters under the parameters section. For details, see Table 7-9.

```
[
"backends": [
-{
```

```
"storage": "oceanstor-san",
...
"parameters": {..., "ALUA": {"<HostName>": {"accessMode": "*", "hyperMetroPathOptimized": "*"},
"<HostName>": {...}}}
}
]
]
```

Table 7-9 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node.	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
accessMode	Host access mode. The value can be:	
	• 0 : balanced mode	
	• 1: asymmetric mode	
hyperMetroPathOp- timized	Whether the path of the host on the current storage array is preferred in HyperMetro scenarios. The value can be:	This parameter needs to be specified only when the host access mode is set to asymmetric.
	• 1 : yes	
	• 0 : no	

□ NOTE

- The ALUA configuration may vary according to the OS. Visit https://support.huawei.com/enterprise/en/index.html, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS and configure ALUA based on the recommended configurations in the guide.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- The value of **HostName** is a regular expression. For details about how to configure it, see **Regular expression**.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used. When you configure **HostName**, the number of host connections is limited. For details about the limitation, see **Specifications Query** and search for **Maximum number of iSCSI connections per controller enclosure**. If the number of host connections is less than or equal to the specifications, you are advised to use the general configuration. If the number of host connections is greater than the specifications, you are advised to use the common configuration.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- In the general configuration, if you need to exactly match a host, refer to example
 3 in 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x.
- If a host uses only OceanStor Dorado 6.x all-flash storage, see **Table 7-10** for detailed parameters.
- If you add OceanStor Dorado 6.x all-flash storage to a host that uses OceanStor converged storage, see **Table 7-11** for detailed parameters.

Table 7-10 Recommended ALUA parameter configurations for OceanStor Dorado 6.*x* all-flash storage

Scenario	Host Type	Host Access Mode	Recommended ALUA Configuration
HyperMetro storage	CentOS/ RHEL/SUSE/	Load balancing mode	ALUA not required
	Debian host	Asymmetric mode + Storage with the preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the non-preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="0"
Non- HyperMetro storage	CentOS/ RHEL/SUSE/ Debian host	N/A	ALUA not required

Scenario	Host Type	Host Access Mode	Recommended ALUA Configuration
HyperMet ro storage	CentOS/RHEL/SUSE/ Debian host	Load balancing mode	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"
		Asymmetric mode + Storage with the non- preferred path	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="0"
Non- HyperMet ro storage	CentOS/RHEL/SUSE/ Debian host	N/A	ACCESSMODE="1" HYPERMETROPATHOPTI- MIZED="1"

Table 7-11 Recommended ALUA parameter configurations for hybrid OceanStor V3/V5. OceanStor Dorado V3. and OceanStor Dorado 6.x

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

Step 5 Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.

----End

7.4.2.3 Configuring ALUA for Distributed Storage

This section describes how to configure ALUA if multipathing is used during the connection to block storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see **4.2.1 Connecting to Enterprise Storage** or **4.2.2 Connecting to Distributed Storage**.

Add ALUA parameters under the **parameters** section. For details, see **Table 7-12**.

```
{
  "backends": [
    {
        "storage": "fusionstorage-san",
        ...
        "parameters": {..., "ALUA": {"<HostName>": {"switchoverMode": "*", "pathType": "*"},
        "<HostName>": {...}}}
    }
}
```

Table 7-12 ALUA parameter description

Parameter	Description	Remarks
<hostname></hostname>	The value of HostName is the host name of a worker node.	The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. For details about the configuration rules and priorities, see the following note.
switchoverMode	Switchover mode. The value can be: • Disable_alua: disables ALUA. • Enable_alua: enables ALUA.	
pathType	Path type. The value can be: optimal_path: preferred path non_optimal_path: non-preferred path	

- Only the iSCSI scenario of distributed storage is supported.
- A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.
- When the value of HostName is a regular expression, configure it by referring to Regular expression.

When **HostName** is set to *, the common configuration is used and takes effect on hosts with any name. When **HostName** is set to another value, the general configuration is used.

Configuration policy rules:

- Priority: General host name configuration > Common host name configuration. For details, see example 1 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- In the general configuration, use the first ALUA section that meets the configuration policy. For details, see example 2 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- In the general configuration, if you need to exactly match a host, refer to example 3 in 11.3 Example ALUA Configuration Policy of Distributed Storage.
- **Step 3** Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

Step 5 Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.

----End

7.5 Configuring Storage Topology Awareness

In the Kubernetes cluster, resources can be scheduled and provisioned based on the topology labels of nodes and the topology capabilities supported by storage backends.

Prerequisites

You need to configure topology labels on worker nodes in the cluster. The method is as follows:

- 1. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- 2. Run the **kubectl get node** command to view information about worker nodes in the current cluster.

```
# kubectl get node
NAME STATUS ROLES AGE VERSION
node01 Ready controlplane,etcd,worker 42d v1.22.3
node02 Ready worker 42d v1.22.3
node03 Ready worker 42d v1.22.3
```

3. Run the **kubectl label node** <*nodename*> **topology.kubernetes.io/** <*key*>=*<value*> command to configure a topology label for a worker node. In

the preceding command, <nodename> indicates the name of a worker node. For details about the **key** and **value** parameters, see **Table 7-13**.

kubectl label node node01 topology.kubernetes.io/zone=ChengDu node/node01 labeled

Table 7-13 Parameter description

Parameter	Description	Remarks
<key></key>	Unique identifier of a topology label.	The value can be zone , region , or protocol . < <i>protocol</i> >.
		<pre><pre><pre><pre><pre><pre><pre>can be set to iscsi, nfs, fc, or roce.</pre></pre></pre></pre></pre></pre></pre>
<value></value>	Value of a topology label.	If key is set to zone or region , value is a user-defined parameter.
		If key is set to protocol. < <i>protocol></i> , value is fixed to csi.huawei.com .

■ NOTE

- A topology label must start with **topology.kubernetes.io**. Topology label examples:
 - Example 1: topology.kubernetes.io/region=China-west
 - Example 2: topology.kubernetes.io/zone=ChengDu
 - Example 3: topology.kubernetes.io/protocol.iscsi=csi.huawei.com
 - Example 4: topology.kubernetes.io/protocol.fc=csi.huawei.com
- A key in a topology label on a node can have only one value.
- If multiple protocols are configured in a topology label on a node, when you select a backend, the backend needs to meet only one of the protocols.
- If both the region and the zone are configured in a topology label on a node, when you select a backend, the backend must meet both of them.
- 4. Run the kubectl get nodes -o=json path='{range .items[*]}
 [{.metadata.name}, {.metadata.labels}]{"\n"}{end}' | grep --color
 "topology.kubernetes.io" command to view the label information about all worker nodes in the current cluster.

kubectl get nodes -o=jsonpath='{range .items[*]}{{.metadata.name}, {.metadata.labels}}{{"\n"}{end}'
| grep --color "topology.kubernetes.io"
[node01, {"beta.kubernetes.io/arch":"amd64","beta.kubernetes.io/os":"linux","kubernetes.io/
arch":"amd64","kubernetes.io/hostname":"node01","kubernetes.io/os":"linux","node-role.kubernetes.io/
controlplane":"true","node-role.kubernetes.io/etcd":"true","node-role.kubernetes.io/
worker":"true","topology.kubernetes.io/zone":"ChengDu"}]

7.5.1 Configuring Storage Topology Awareness Using Helm

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see 3.3 Uploading the Components in the Software Package.
- Step 3 Back up the values.yaml file used during CSI installation. If the values.yaml file used during the last installation cannot be found, run the **helm get values** helmhuawei-csi -n huawei-csi -a command to query the file.

cp values.yaml values.yaml.bak

Step 4 Run the **vi** values.yaml command to open the file and configure multiple backends as required. The following is an example. After the modification is complete, press **Esc** and enter :wq! to save the modification.

```
backends:
 - storage: "oceanstor-nas"
  name: "storage"
  urls:
   - "https://*.*.*:8088"
   - "https://*.*.*:8088"
  pools:
   - "***"
   _ "***
  parameters:
   protocol: "nfs"
    portals:
    supportedTopologies:
     - {"topology.kubernetes.io/region": "China-west", "topology.kubernetes.io/zone": "ChengDu"}
     - {"topology.kubernetes.io/region": "China-south", "topology.kubernetes.io/zone": "ShenZhen"}
```

Step 5 Run the helm upgrade helm-huawei-csi ./ -n huawei-csi -f values.yaml command to upgrade the Helm chart. The upgrade command will update the **Deployment**, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, helm-huawei-csi indicates the custom chart name and huawei-csi indicates the custom namespace.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

Step 6 Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press I or Insert to enter the editing mode and add related parameters in the .yaml file. For details about the parameters, see Table 7-14. After the modification is complete, press **Esc** and enter :wq! to save the modification.

Add the following configuration items to the StorageClass.yaml file.

Example 1: Configure zone and region information in the StorageClass.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: example-storageclass
provisioner: csi.huawei.com
parameters:
volumeType: lun
```

allocType: thin

volumeBindingMode: WaitForFirstConsumer allowedTopologies:

- matchLabelExpressions:

key: topology.kubernetes.io/zone values:

- ChengDu

key: topology.kubernetes.io/region values:

- China-west

• Example 2: Configure protocol information in the StorageClass.

kind: StorageClass

apiVersion: storage.k8s.io/v1

metadata:

name: protocol-example-storageclass

provisioner: csi.huawei.com

parameters: volumeType: lun

allocType: thin volumeBindingMode: WaitForFirstConsumer

allowedTopologies:
- matchLabelExpressions:

- key: topology.kubernetes.io/protocol.iscsi values:
- csi.huawei.com

Table 7-14 Parameter description

Parameter	Description	Remarks
volumeBindin gMode	PersistentVolume binding mode, used to control the time when PersistentVolume resources are dynamically allocated and bound.	You can set this parameter to WaitForFirstConsumer or Immediate. WaitForFirstConsumer: indicates that the binding and allocation of the PersistentVolume are delayed until a Pod that uses the PVC is created. Immediate: The PersistentVolume is bound and allocated immediately after a PVC is created.
allowedTopol ogies.matchLa belExpression s	La information label,	key: This parameter can be set to topology.kubernetes.io/zone or topology.kubernetes.io/region. topology.kubernetes.io/ protocol. <protocol>: <protocol> indicates the protocol type and can be iscsi, fc, or nfs.</protocol></protocol>
	Pods cannot be created. Both key and value must be configured in a fixed format.	value: If key is topology.kubernetes.io/zone or topology.kubernetes.io/region, value must be the same as the topology label set in the prerequisites. If key is topology.kubernetes.io/protocol. <pre>cprotocol></pre> , value is fixed to csi.huawei.com.

- **Step 7** Run the following command to create a StorageClass based on the .yaml file.

 # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC with the topology capability. For details, see **6.3.1 Creating a PVC**.
- **Step 9** Use the PVC to create a Pod. For details, see **6.4.1 Creating a Pod**.

----End

7.5.2 Manually Configuring Storage Topology Awareness

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see 4.2.1 Connecting to Enterprise Storage or 4.2.2 Connecting to Distributed Storage.

Add the **supportedTopologies** field under the **backends** section in the *huawei-csi-configmap.yaml* file to configure the topology information supported by each backend. The following is a backend example.

```
{
    "backends":[
    {
        "storage": "oceanstor-san",
        ...
        "parameters": {"protocol": "iscsi", "portals": ["192.168.125.22", "192.168.125.23"]},
        "supportedTopologies": [
        {"topology.kubernetes.io/region": "China-west", "topology.kubernetes.io/zone": "ChengDu"},
        {"topology.kubernetes.io/region": "China-south", "topology.kubernetes.io/zone": "ShenZhen"}]
    }
}
```

MOTE

- **supportedTopologies** is a list. Each element in the list is a dictionary.
- Only topology.kubernetes.io/region or topology.kubernetes.io/zone can be configured for each element in the list. The parameter value must be the same as the topology label set in the prerequisites. (topology.kubernetes.io/protocol.col>
 does not need to be configured.)
- **Step 3** Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n huawei-csi | grep huawei-csi-configmap
huawei-csi-configmap 1 5s
```

- **Step 5** Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.
- **Step 6** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press I or Insert to enter the editing mode and add related parameters in the .yaml file. For details about the parameters, see Table 7-15. After the modification is complete, press **Esc** and enter :wq! to save the modification.

Add the following configuration items to the StorageClass.yaml file.

Example 1: Configure zone and region information in the StorageClass.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: example-storageclass provisioner: csi.huawei.com parameters: volumeType: lun allocType: thin

volumeBindingMode: WaitForFirstConsumer

allowedTopologies:

- matchLabelExpressions:
- key: topology.kubernetes.io/zone
- values:
- ChengDu
- key: topology.kubernetes.io/region values:
- China-west
- Example 2: Configure protocol information in the StorageClass.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: protocol-example-storageclass provisioner: csi.huawei.com parameters: volumeType: lun allocType: thin

volumeBindingMode: WaitForFirstConsumer

allowedTopologies:

- matchLabelExpressions:
- key: topology.kubernetes.io/protocol.iscsi values:
- csi.huawei.com

Table 7-15 Parameter description

Parameter	Description	Remarks
volumeBindin gMode	PersistentVolume binding mode, used to control the time when PersistentVolume resources are dynamically allocated and bound.	You can set this parameter to WaitForFirstConsumer or Immediate. WaitForFirstConsumer: indicates that the binding and allocation of the PersistentVolume are delayed until a Pod that uses the PVC is created. Immediate: The PersistentVolume is bound and allocated immediately after a PVC is created.

Parameter	Description	Remarks
allowedTopol ogies.matchLa belExpression s	Topology information label, which is used to filter CSI backends and Kubernetes nodes. If the matching fails, PVCs or Pods cannot be created. Both key and value must be configured in a fixed format.	key: This parameter can be set to topology.kubernetes.io/zone or topology.kubernetes.io/region. topology.kubernetes.io/ protocol. <protocol>: <protocol> indicates the protocol type and can be iscsi, fc, or nfs. value: If key is topology.kubernetes.io/zone or topology.kubernetes.io/region, value must be the same as the topology label set in the prerequisites. If key is topology.kubernetes.io/protocol.<pre>/protocol>, value</pre> is fixed to csi.huawei.com.</protocol></protocol>

- **Step 7** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC with the topology capability. For details, see **6.3.1 Creating a PVC**.
- **Step 9** Use the PVC to create a Pod. For details, see **6.4.1 Creating a Pod**.

----End

7.6 Advanced Features of Enterprise Storage

7.6.1 Configuring QoS

This section describes how to create a LUN/file system volume that supports QoS.

Precautions

- The QoS feature is not a standard feature of Kubernetes and is customized by storage vendors.
- A QoS policy can be specified only when a StorageClass is created. Once the QoS policy is created, it cannot be modified because the StorageClass cannot be modified on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

- The value of **metadata.name** is the user-defined name of a StorageClass object. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-).
- Add the qos configuration item under parameters. For other information about parameters, see 6.1.1.1 Creating a LUN StorageClass or 6.1.1.2 Creating a File System StorageClass.
- The value of the **qos** section is JSON character strings in dictionary format. A character string is enclosed by single quotation marks and the dictionary key by double quotation marks.
 - For details about the parameters of OceanStor V3/OceanStor V5 series storage devices, see Table 7-16.
 - For details about the parameters of OceanStor Dorado V3 series storage devices, see Table 7-17.
 - For details about the parameters of OceanStor Dorado 6.x/OceanStor 6.x series storage devices, see Table 7-18.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
qos: '{"IOTYPE": 2, "MINIOPS": 1000}'
```

Table 7-16 QoS parameters supported by OceanStor V3/OceanStor V5

Parameter	Description	Remarks
IOTYPE	Read/write type.	This parameter is optional. If it is not specified, the default value of the storage backend is used. For details, see related storage documents. The value can be: • 0: read I/O • 1: write I/O • 2: read and write I/Os
MAXBANDWIDTH	Maximum bandwidth. This is a restriction policy parameter.	The value is an integer greater than 0, expressed in MB/s.
MINBANDWIDTH	Minimum bandwidth. This is a protection policy parameter.	The value is an integer greater than 0, expressed in MB/s.
MAXIOPS	Maximum IOPS. This is a restriction policy parameter.	The value is an integer greater than 0
MINIOPS	Minimum IOPS. This is a protection policy parameter.	The value is an integer greater than 0

Parameter	Description	Remarks
LATENCY	Maximum latency. This is a protection policy parameter.	The value is an integer greater than 0, expressed in ms.

Table 7-17 QoS parameters supported by OceanStor Dorado V3

Parameter	Description	Remarks
IOTYPE	Read/write type.	The value can be: • 2: read and write I/Os
MAXBANDWIDTH	Maximum bandwidth. This is a restriction policy parameter.	The value is an integer ranging from 1 to 999999999, expressed in MB/s.
MAXIOPS	Maximum IOPS. This is a restriction policy parameter.	The value is an integer ranging from 100 to 9999999999999999999999999999999999

Table 7-18 QoS parameters supported by OceanStor Dorado 6.x/OceanStor 6.x

Parameter	Description	Remarks
IOTYPE	Read/write type.	The value can be:
		• 2: read and write I/Os
MAXBANDWIDTH	Maximum bandwidth.	The value is an integer ranging from 1 to 999999999, expressed in MB/s.
MINBANDWIDTH	Minimum bandwidth.	The value is an integer ranging from 1 to 999999999, expressed in MB/s.
MAXIOPS	Maximum IOPS.	The value is an integer ranging from 100 to 999999999.
MINIOPS	Minimum IOPS.	The value is an integer ranging from 100 to 999999999.
LATENCY	Maximum latency.	The value can be 0.5 or 1.5 , expressed in ms.

□ NOTE

- Different protection policy parameters or restriction policy parameters can be specified at the same time. However, protection policy parameters cannot coexist with restriction policy parameters.
- vStore users do not support QoS policies.
- The QoS configuration takes effect only on the newly created PVC. QoS cannot be added automatically for PVCs with the same StorageClass name that have been provisioned.
- **Step 3** Run the following command to create a StorageClass based on the .yaml file.

 # kubectl create -f StorgeClass.yaml
- **Step 4** Use the StorageClass to create a PVC with the QoS capability. For details, see **6.3.1** Creating a PVC.

----End

7.6.2 Configuring a vStore

This section describes how to configure a vStore.

Precautions

This feature does not support SAN storage. For details about the supported models, see **Table 2-7**.

7.6.2.1 Configuring a vStore Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* -**n** *huawei-csi* -**a** command to query the file.

cp values.yaml values.yaml.bak

- **Step 4** Run the **vi** *values.yaml* command to open the file and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - If the vStore uses OceanStor V3/V5 storage, add the **vstoreName** parameter to the backend configuration and set the parameter to the vStore name of the storage device. The following is an example.

```
backends:
- storage: "oceanstor-san"
name: "***"
vstoreName: "***"
urls:
- "https://*.*.*:8088"
```

```
pools:
_ "***"
parameters:
protocol: "iscsi"
portals:
_ "****"
```

• If the vStore uses OceanStor 6.1 or OceanStor Dorado 6.x storage, configure the backend by referring to 4.1.2 Installing Huawei CSI. Note that urls indicates the logical management ports of the vStore, and pools and portals must be available storage pools and logical data ports of the current vStore respectively.

```
backends:
- storage: "oceanstor-san"
name: "***"
urls:
- "https://*.*.**:8088"
pools:
- "***"
parameters:
protocol: "iscsi"
portals:
- "* * * * "
```

Step 5 Run the helm upgrade helm-huawei-csi./ -n huawei-csi-f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

7.6.2.2 Manually Configuring a vStore

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Multiple backends are separated by commas (,). For details about each backend, see **4.2.1 Connecting to Enterprise Storage**.

• If the vStore uses OceanStor V3/V5 storage, add the **vstoreName** parameter to the backend configuration and set the parameter to the vStore name of the storage device.

```
{
    "backends": [
    {
        ...
        "vstoreName": "***"
    }
    ]
}
```

• If the vStore uses OceanStor 6.1 or OceanStor Dorado 6.x storage, configure the backend by referring to 4.2.1 Connecting to Enterprise Storage. Note that urls indicates the logical management ports of the vStore, and pools and portals must be available storage pools and logical data ports of the current vStore respectively.

■ NOTE

After configuring **huawei-csi-configmap.yaml**, restart huawei-csi-controller and huawei-csi-node. Otherwise, the configuration does not take effect.

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 Run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

```
# kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s
```

Step 5 Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services**.

∩ NOTE

- When starting the huawei-csi services, enter the user name and password of the storage device vStore entered in Step 3 in 4.2.3 Starting huawei-csi Services.
- For details about the storage models supported by this feature, Table 2-7.

----End

7.6.3 Configuring NAS HyperMetro

Perform this operation when you want to configure NAS HyperMetro.

Precautions

For details about the resource objects that support NAS HyperMetro and the feature description, see **Table 7-19**.

Table 7-19 Feature description

Resource Object	Operation	Supporte d	Remarks
PVC	Creation	Yes	This feature can be used together with other features (except remote replication), such as QoS.
	Deletion	Yes	

Resource Object	Operation	Supporte d	Remarks
	Capacity expansion	Yes	-
	Synchronizing a HyperMetro pair	No	Storage supports these operations. Because Kubernetes cannot detect HyperMetro pairs, Kubernetes does not support these operations.
	Pausing a HyperMetro pair	No	
	Preferred site switchover	No	
Pod	Creation	Yes	Primary and secondary file systems can be mounted to a Pod at the same time.
	Deletion	Yes	
VolumeSnapsh ot	Creation	Yes	Snapshots can be operated only for primary storage file systems.
	Deletion	Yes	

7.6.3.1 Prerequisites

Configuring the Storage HyperMetro Relationship

Before configuring NAS HyperMetro, you need to configure the HyperMetro relationship between two storage devices, including the remote device, HyperMetro domain, and the like. The HyperMetro domain of the file system can only work in HyperMetro mode. For details about the configuration operation, see the product documentation of the corresponding storage model.

Configuring HyperMetro vStore Information

Before configuring NAS HyperMetro, you need to create vStores, vStore pairs, vStore users, and corresponding logical ports for the two storage devices. For details, see the product documentation of the desired storage model. This section uses OceanStor Dorado 6.1.3 as an example.

Step 1 Create a vStore on the primary storage device in the HyperMetro domain.

- 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
- 2. Choose Services > vStore Service > vStores.
- 3. Click Create, specify Name (for example, vStoreA), and click OK.
- **Step 2** Create a vStore on the secondary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the secondary storage device in the HyperMetro domain.
 - 2. Choose Services > vStore Service > vStores.

- 3. Click Create, specify Name (for example, vStoreB), and click OK.
- **Step 3** Create a HyperMetro vStore pair on the primary storage device in the HyperMetro domain.
 - Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Data Protection > Protection Entities > vStores > vStores.
 - Select the vStore for which you want to create a HyperMetro vStore pair (for example, vStoreA created in Step 1) and click Create HyperMetro vStore Pair.
 - 4. Set **Pair Creation** to **Manual**, set **Remote vStore** to the vStore of the secondary storage device (for example, **vStoreB** created in **Step 2**), set other parameters based on site requirements, and click **OK**.
- **Step 4** Create a logical management port for the local storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose **Services** > **Network** > **Logical Ports**.
 - Click Create, specify Name (for example, manageA), set Role to Management, Owning vStore to vStoreA, and Owning Site to Local device, set other parameters based on site requirements, and click OK.
- **Step 5** Create a logical management port for the remote storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - Choose Services > Network > Logical Ports.
 - Click Create, specify Name (for example, manageB), set Role to Management, Owning vStore to vStoreA, and Owning Site to Remote device, set other parameters based on site requirements, and click OK.
 - 4. On DeviceManager of the secondary storage device in the HyperMetro domain, check whether the logical port is correct.
 - For example, the **Home Port** parameter value of the logical port and that of the logical port corresponding to the primary storage device in the HyperMetro domain are on the same network plane.
- **Step 6** Create a logical service port for the local storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > Network > Logical Ports.
 - Click Create, specify Name (for example, dataA), set Role to Data, Data Protocol to NFS, Owning vStore to vStoreA, and Owning Site to Local device, set other parameters based on site requirements, and click OK.
- **Step 7** Create a logical service port for the remote storage device on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.

- 2. Choose Services > Network > Logical Ports.
- 3. Click **Create**, specify **Name** (for example, **dataB**), set **Role** to **Data**, **Data Protocol** to **NFS**, **Owning vStore** to **vStoreA**, and **Owning Site** to **Remote device**, set other parameters based on site requirements, and click **OK**.
- 4. On DeviceManager of the secondary storage device in the HyperMetro domain, check whether the logical port is correct.
 - For example, the **Home Port** parameter value of the logical port and that of the logical port corresponding to the primary storage device in the HyperMetro domain are on the same network plane.
- **Step 8** Create a vStore user on the primary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the primary storage device in the HyperMetro domain.
 - 2. Choose Services > vStore Service > vStores.
 - 3. Click the name of vStoreA, click User Management, and click Create. Set Type to Local user and Role to vStore administrator, set other parameters based on site requirements, and click OK.
- **Step 9** Create a vStore user on the secondary storage device in the HyperMetro domain.
 - 1. Log in to DeviceManager of the secondary storage device in the HyperMetro domain.
 - Choose Services > vStore Service > vStores.
 - 3. Click the name of **vStoreB**, click **User Management**, and click **Create**. Set **Type** to **Local user** and **Role** to **vStore administrator**, set other parameters based on site requirements, and click **OK**.

----End

7.6.3.2 Configuring NAS HyperMetro Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.
 - # cp values.yaml values.yaml.bak
- Step 4 Run the vi values.yaml command to open the file and add or modify HyperMetrorelated parameters under backends. For details about the parameters, see Table 7-20. After the modification is complete, press Esc and enter :wq! to save the modification.

An array of storages with the access info backends:

- storage: "oceanstor-nas"

```
name: "hyperMetro1"
  - "https://*.*.*.*:8088"
pools:
vstoreName: "vStore1"
metrovStorePairID: "f09838237b93c000"
metroBackend: "hyperMetro2"
parameters:
 protocol: "nfs"
 portals:
   - "*.*.*.*"
- storage: "oceanstor-nas"
name: "hyperMetro2"
urls:
  - "https://*.*.*.8088"
pools:
vstoreName: "vStore1"
metrovStorePairID: "f09838237b93c000"
metroBackend: "hyperMetro1"
parameters:
 protocol: "nfs"
 portals:
   - "*.*.*"
```

Table 7-20 HyperMetro configuration items of a back-end storage device

Configurati on Item	Format	Description	Remarks
vstoreName	String	vStore name. This parameter is conditionally mandatory (mandatory for OceanStor V3/V5).	Only vStore users support NAS HyperMetro. For details about how to configure vStore users, see 7.6.3.1 Prerequisites .
metrovStor ePairID	String	ID of the HyperMetro vStore pair to which a vStore belongs. This parameter is mandatory.	For example, the parameter of OceanStor Dorado 6.x/ OceanStor 6.1 is displayed as ID on DeviceManager.
metroBacke nd	String	Name of a peer end in HyperMetro. The two backends form a HyperMetro relationship. This parameter is mandatory.	The peer end of hyperMetro1 is hyperMetro2, and the peer end of hyperMetro2 is hyperMetro1.

Step 5 Run the **helm upgrade** *helm-huawei-csi*./ **-n** *huawei-csi* **-f** *values.yaml* command to upgrade the Helm chart. The upgrade command will update the **Deployment**,

daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

Step 6 Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and add the **hyperMetro** parameter under **parameters** in the .yaml file. For details about the parameters, see **Table 7-21**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: "***" provisioner: "csi.huawei.com" parameters: ... volumeType: fs hyperMetro: "true"

Table 7-21 Parameter description

Parameter	Description	Remarks
parameters.hyperMetro	Whether a HyperMetro volume is to be created.	If this parameter is set to true , a HyperMetro volume is to be created. If this parameter is not set or set to false , no HyperMetro volume is to be created.

- **Step 7** Run the following command to create a StorageClass based on the .yaml file.

 # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC with the NAS HyperMetro capability. For details, see **6.3.1 Creating a PVC**.

----End

7.6.3.3 Manually Configuring NAS HyperMetro

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

Step 2 Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

In the **backends** section of the *huawei-csi-configmap.yaml* file, add two backends with a HyperMetro relationship. For details about the configuration items for each backend, see **4.2.1 Connecting to Enterprise Storage**. You need to add some additional configuration parameters in the HyperMetro scenario. For details, see **Table 7-22**.

```
kind: ConfigMap
apiVersion: v1
metadata:
 name: huawei-csi-configmap
 namespace: huawei-csi
data:
 csi.json: |
      "backends": [
         {
            "storage": "oceanstor-nas",
            "name": "hyperMetro1",
            "urls": ["https://192.168.125.20:8088", "https://192.168.125.21:8088"],
            "pools": ["storagepool01", "storagepool02"],
"parameters": {"protocol": "nfs", "portals": ["192.168.125.22"]},
"vstoreName": "vStore1",
            "metrovStorePairID": "f09838237b93c000",
            "metroBackend": "hyperMetro2"
            "storage": "oceanstor-nas",
            "name": "hyperMetro2",
            "urls": ["https://192.168.125.24:8088", "https://192.168.125.25:8088"],
            "pools": ["storagepool01", "storagepool02"],
"parameters": {"protocol": "nfs", "portals": ["192.168.125.26"]},
            "vstoreName": "vStore1",
            "metrovStorePairID": "f09838237b93c000",
            "metroBackend": "hyperMetro1"
      ]
  }
```

Table 7-22 HyperMetro configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
vstoreName	String	vStore name. This parameter is conditionally mandatory (mandatory for OceanStor V3/V5).	Only vStore users support NAS HyperMetro. For details about how to configure vStore users, see 7.6.3.1 Prerequisites.
metrovStorePa irID	String	ID of the HyperMetro vStore pair to which a vStore belongs. This parameter is mandatory.	For example, the parameter of OceanStor Dorado 6.x/ OceanStor 6.1 is displayed as ID on DeviceManager.

Configuration Item	Format	Description	Remarks
metroBackend	String	Name of a peer end in HyperMetro. The two backends form a HyperMetro relationship. This parameter is mandatory.	The peer end of hyperMetro1 is hyperMetro2, and the peer end of hyperMetro2 is hyperMetro1.

Step 3 Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.

kubectl create -f huawei-csi-configmap.yaml

Step 4 After the creation is complete, run the **kubectl get configmap -n huawei-csi | grep huawei-csi-configmap** command to check whether the creation is successful. If the following information is displayed, the creation is successful.

kubectl get configmap -n huawei-csi | grep huawei-csi-configmap huawei-csi-configmap 1 5s

- **Step 5** Start huawei-csi services. For details, see **4.2.3 Starting huawei-csi Services** (use the vStore user created in **7.6.3.1 Prerequisites** for configuration).
- **Step 6** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and add the **hyperMetro** parameter under **parameters** in the .yaml file. For details about the parameters, see **Table 7-23**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata: name: "***" provisioner: "csi.huawei.com" parameters: ... volumeType: fs hyperMetro: "true"

Table 7-23 Parameter description

Parameter	Description	Remarks
parameters.hyperMetro	Whether a HyperMetro volume is to be created.	If this parameter is set to true , a HyperMetro volume is to be created. If this parameter is not set or set to false , no HyperMetro volume is to be created.

Step 7 Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorgeClass.yaml

Step 8 Use the StorageClass to create a PVC with the NAS HyperMetro capability. For details, see **6.3.1 Creating a PVC**.

----End

7.6.4 Configuring an Application Type

This section describes how to create a LUN/file system volume that supports different application types.

Precautions

- The application type feature is not a standard feature of Kubernetes and is customized by storage vendors.
- An application type can be specified only when a PVC is created.
- A created PVC cannot be modified on Kubernetes.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add the applicationType configuration item under parameters. The value of applicationType is a character string. For details, see Table 7-24. For other information about parameters, see 6.1.1.1 Creating a LUN StorageClass or 6.1.1.2 Creating a File System StorageClass.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
volumeType: "***"
applicationType: "***"
```

Table 7-24 Parameter description of applicationType

Parameter	Description	Remarks
parameters.applic ationType	Application type name on the storage device. The value is a character string.	If the value of volumeType is lun , log in to DeviceManager and choose Services > Block Service > LUN Groups > LUNs > Create to obtain the application type name.
		If the value of volumeType is fs , log in to DeviceManager and choose Services > File Service > File Systems > Create to obtain the application type name.

□ NOTE

This feature applies only to OceanStor Dorado 6.x series storage systems.

- **Step 3** Run the following command to create a StorageClass based on the .yaml file. # kubectl create -f StorgeClass.yaml
- **Step 4** Use the StorageClass to create a PVC with the application type capability. For details, see **6.3.1 Creating a PVC**.

----End

7.7 Advanced Features of Distributed Storage

7.7.1 Configuring QoS

This section describes how to create a LUN volume that supports QoS.

Precautions

- The QoS feature is not a standard feature of Kubernetes and is customized by storage vendors.
- A QoS policy can be specified only when a StorageClass is created. Once the QoS policy is created, it cannot be modified because the StorageClass cannot be modified on Kubernetes.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.
 - Add the qos configuration item under parameters. For other information about parameters, see 6.1.1.1 Creating a LUN StorageClass.
 - The value of the qos section is JSON character strings in dictionary format. A
 character string is enclosed by single quotation marks and the dictionary key
 by double quotation marks. For details about the parameters, see Table 7-25.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
qos: '{"maxMBPS": 999, "maxIOPS": 999}'
```

Parameter	Description	Remarks
maxMBPS	Maximum bandwidth.	This parameter is mandatory. The value is an integer greater than 0, expressed in MB/s.
maxIOPS	Maximum IOPS.	This parameter is mandatory. The value is an integer greater than 0

Table 7-25 Parameters in qos

Step 3 Run the following command to create a PVC based on the .yaml file.

kubectl create -f StorageClass.yaml

----End

7.7.2 Configuring a Soft Quota

This section describes how to create a PVC that supports soft quotas.

Precautions

- This feature is supported only by OceanStor Pacific series 8.1.0 and later versions.
- This feature can be configured only when a storage pool of the file system type is connected.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - Add the storageQuota configuration item under parameters. For other information about parameters, see 6.1.1.2 Creating a File System StorageClass.
 - The value of the storageQuota section is JSON character strings in dictionary format. A character string is enclosed by single quotation marks and the dictionary key by double quotation marks. For details about the parameters, see Table 7-26.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
volumeType: "fs"
...
storageQuota: '{"spaceQuota": "softQuota", "gracePeriod": 100}'
```

	-	
Parameter	Description	Remarks
spaceQuota	File quota type.	This parameter is mandatory. Only softQuota or hardQuota can be configured.
gracePeriod	Grace period allowed when the soft quota is configured.	This parameter is conditionally optional only when spaceQuota is set to softQuota .
		The value is an integer ranging from 0 to 4294967294.

Table 7-26 Parameters in storageQuota

- **Step 3** Run the following command to create a PVC based on the .yaml file.

 # kubectl create -f StorageClass.yaml
- **Step 4** Configure a StorageClass in the PVC according to **6.3.1 Creating a PVC** to finish creating a PVC.

----End

7.7.3 Configuring an Account

This section describes how to configure a backend for a specified account and create a PVC for the backend.

Precautions

This feature supports only distributed NAS series storage.

7.7.3.1 Configuring an Account Using Helm

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- Step 3 Back up the values.yaml file used during CSI installation. If the values.yaml file used during the last installation cannot be found, run the helm get values helm-huawei-csi -n huawei-csi -a command to query the file.

 # cp values.yaml values.yaml.bak
- Step 4 Run the vi values.yaml command to open the configuration file and add the accountName parameter under the backends section. For details about the parameter, see Table 7-27. Ensure that portals is set to the logical ports owned by the account. The following is an example. After the modification is complete, press Esc and enter :wq! to save the modification.

```
backends:
- storage: "fusionstorage-nas"
name: "***"

accountName: "***"
urls:
- "https://*.*.*.*:28443"
pools:
- "***"
parameters:
protocol: "nfs"
portals:
- "*.*.*"
```

Table 7-27 Parameter description of accountName

Parameter	Description	Remarks
backends.accou ntName	Account name to be specified. The value is a character string.	• If accountName is set to system or this parameter is not specified, the PVC created using the backend will be created under the system account (user name: system).
		 If accountName is set to another user, the PVC created using the backend will be created only under the user.

Step 5 Run the helm upgrade helm-huawei-csi./ -n huawei-csi -f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

7.7.3.2 Manually Configuring an Account

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *huawei-csi-configmap.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

For details about the backend configuration, see **4.2.2.3 Connecting to Distributed Storage NAS over NFS**. Add the **accountName** configuration item to the **backends** section in the **huawei-csi-configmap.yaml** file. **accountName** indicates the account name to be specified and its value is a character string. For details about the parameter, see **Table 7-28**. **portals** must be set to the logical ports owned by the account.

```
{
  "backends":[
  {
    "storage": "fusionstorage-nas",
    "name": "***",
    ...
  "accountName": "***",
    "parameters": {"protocol": "nfs", "portals": ["*.*.*.*"]}
  }
}
```

□ NOTE

portals must be the IP addresses under the configured account. To view its value, log in to DeviceManager and choose **Resources** > **Access** > **Service Network**.

Table 7-28 Parameter description of accountName

Parameter	Description	Remarks
backends.accou ntName	Account name to be specified. The value is a character string.	• If accountName is set to system or this parameter is not specified, the PVC created using the backend will be created under the system account (user name: system).
		 If accountName is set to another user, the PVC created using the backend will be created only under the user.

- **Step 3** Run the **kubectl create -f** *huawei-csi-configmap.yaml* command to create *huawei-csi-configmap*.
 - # kubectl create -f huawei-csi-configmap.yaml
- **Step 4** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 5** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 6** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add the **backend** configuration item under **parameters**. The value of **backend** is a character string and is the backend name of the account specified in **Step 2**. For other information about **parameters**, see **6.1.1.2 Creating a File System StorageClass**. For details about how to create a PVC for a specified backend, see **7.2 Creating a PVC for a Specified Backend**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
parameters:
...
volumeType: "***"
backend: "***"
```

- **Step 7** Run the following command to create a StorageClass based on the .yaml file.

 # kubectl create -f StorgeClass.yaml
- **Step 8** Use the StorageClass to create a PVC for a specified user. For details, see **6.3.1** Creating a PVC.

----End

7.7.4 Configuring Mount Parameters in the DPC Scenario

This section describes how to configure mount parameters in the DPC scenario.

Precautions

This feature supports only distributed NAS series storage that uses the DPC protocol.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.
 - Add the mountOptions configuration item. For other information about StorageClass.yaml, see 6.1.1.1 Creating a LUN StorageClass or 6.1.1.2 Creating a File System StorageClass.
 - The value of **mountOptions** is in list format. For details about the parameters, see **Table 7-29**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: "***"
provisioner: "csi.huawei.com"
mountOptions:
- ***
- ***
parameters:
```

Table 7-29 Parameters in mountOptions

Parameter	Description	Remarks
acl	The namespace supports the ACL function, and the DPC client supports POSIX ACL, NFSv4 ACL, and NT ACL authentication.	The descriptions of acl, aclonlyposix, cnflush, and cflush are for reference only. For details about the parameters, see OceanStor Pacific Series Product Documentation and choose Configuration > Basic Service Configuration Guide for File > Configuring Basic Services (DPC Scenario) > Accessing a DPC Share on a Client > Step 2.
aclonlyposi x	The namespace supports POSIX ACL, and the DPC client supports POSIX ACL authentication. The following protocols support POSIX ACL: DPC, NFSv3, and HDFS. If NFSv4 ACL or NT ACL is used, the DPC client cannot identify the ACL of this type. As a result, the ACL of this type does not take effect.	If aclonlyposix and acl are used together, only acl takes effect. That is, the namespace supports the ACL function.

Parameter	Description	Remarks
cnflush	Asynchronous disk flushing mode. That is, data is not flushed to disks immediately when files in the namespace are closed.	-
	Asynchronous flushing mode: When a file is closed, data in the cache is not flushed to storage media in synchronous mode. Instead, data is written from the cache to the storage media in asynchronous flushing mode. After the write service is complete, data is flushed from the cache to disks periodically based on the flushing period. In a multi-client scenario, if concurrent operations are performed on the same file, the file size update is affected by the disk flushing period. That is, the file size is updated only after the disk flushing is complete. Generally, the update is completed within several seconds. Synchronous I/Os are not affected by the disk flushing period.	
cflush	Synchronous disk flushing mode. That is, data is flushed to disks immediately when files in the namespace are closed.	By default, the synchronous disk flushing mode is used.

Step 3 Run the following command to create a PVC based on the .yaml file.

kubectl create -f StorageClass.yaml

----End

8 Uninstalling CSI

Perform this operation when you want to uninstall Huawei CSI.

- 8.1 Uninstalling huawei-csi Using Helm
- 8.2 Manually Uninstalling huawei-csi
- 8.3 (Optional) Uninstalling the Snapshot-Dependent Component Service

8.1 Uninstalling huawei-csi Using Helm

This section describes how to uninstall the Huawei CSI plug-in installed using Helm. For details about how to uninstall the Huawei CSI plug-in manually installed, see **8.2 Manually Uninstalling huawei-csi**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **helm uninstall** *helm-release-name* **-n** *huawei-csi* command to uninstall Huawei CSI. This command will delete the huawei-csi-controller, huawei-csi-node, huawei-csi-configmap, and RBAC configurations, but will not delete the huawei-csi-secret and huawei-csi container images.

helm uninstall helm-huawei-csi -n huawei-csi release "helm-huawei-csi" uninstalled

Step 3 Run the **helm list -n** *huawei-csi* command to check whether the service is successfully uninstalled. If the command output is empty, the service is successfully uninstalled.

In the preceding command, *huawei-csi* indicates the namespace where the chart is located.

helm list -n huawei-csi NAME NAMESPACE REVISION UPDATED STATUS CHART APP VERSION

Step 4 (Optional) Uninstall the snapshot-dependent component service. For details, see **8.3** (Optional) Uninstalling the Snapshot-Dependent Component Service.

----End

8.2 Manually Uninstalling huawei-csi

This section describes how to manually uninstall Huawei CSI. For details about how to uninstall the snapshot-dependent component service, see 8.3 (Optional) Uninstalling the Snapshot-Dependent Component Service.

8.2.1 Uninstalling the huawei-csi-node Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete daemonset huawei-csi-node -n** *huawei-csi* command to uninstall the huawei-csi-node service. Replace *huawei-csi* with the actual namespace.
 - # kubectl delete daemonset huawei-csi-node -n huawei-csi
- **Step 3** Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.
 - # kubectl get daemonset huawei-csi-node -n huawei-csi

----End

8.2.2 Uninstalling the huawei-csi-controller Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete deployment huawei-csi-controller -n** *huawei-csi* command to uninstall the huawei-csi-controller service. Replace *huawei-csi* with the actual namespace.
 - # kubectl delete deployment huawei-csi-controller -n huawei-csi
- **Step 3** Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.
 - # kubectl get deployment huawei-csi-controller -n huawei-csi

----End

8.2.3 Deleting the huawei-csi-configmap Object

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete configmap** *huawei-csi-configmap* **-n** *huawei-csi* command to delete the **configmap** object. *huawei-csi-configmap* is the name of the **configmap** object, and *huawei-csi* is the namespace where the object is located.

kubectl delete configmap huawei-csi-configmap -n huawei-csi

Step 3 Run the following command to check whether the object is successfully deleted. If **NotFound** is displayed, the object is successfully deleted.

kubectl get configmap huawei-csi-configmap -n huawei-csi

----End

8.2.4 Deleting the huawei-csi-secret Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete secret** *huawei-csi-secret* **-n** *huawei-csi* command to delete the **secret** object. *huawei-csi-secret* is the name of the **secret** object, and *huawei-csi* is the namespace where the **secret** object is located.

kubectl delete secret huawei-csi-secret -n huawei-csi

Step 3 Run the following command to check whether the **secret** object is successfully deleted. If **NotFound** is displayed in the command output, the **huawei-csi-secret** object is successfully deleted.

kubectl get secret huawei-csi-secret -n huawei-csi Error from server (NotFound): secrets "huawei-csi-secret" not found

----End

8.2.5 Deleting the RBAC Permission

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Delete the RBAC permission.
 - If the huawei-csi version is later than 2.2.15, run the following command to delete the permission. -n indicates the namespace. Change it based on site requirements.

kubectl -n huawei-csi -l provisioner=csi.huawei.com delete ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding

- If the huawei-csi version is 2.2.15 or earlier, perform the following operations to delete the permission.
 - Run the following command to create a file named delete-huawei-csi-rbac.sh. -n indicates the namespace. Change it based on site requirements.

cat <<EOF > delete-huawei-csi-rbac.sh kubectl delete ServiceAccount huawei-csi-controller -n huawei-csi kubectl delete ServiceAccount huawei-csi-node -n huawei-csi kubectl delete ClusterRole huawei-csi-attacher-runner -n huawei-csi kubectl delete ClusterRole huawei-csi-driver-registrar-runner -n huawei-csi kubectl delete ClusterRole huawei-csi-provisioner-runner -n huawei-csi kubectl delete ClusterRole huawei-csi-resizer-runner -n huawei-csi kubectl delete ClusterRole huawei-csi-snapshotter-runner -n huawei-csi kubectl delete ClusterRole snapshot-controller-runner -n huawei-csi kubectl delete ClusterRoleBinding huawei-csi-attacher-role -n huawei-csi kubectl delete ClusterRoleBinding huawei-csi-driver-registrar-role -n huawei-csi kubectl delete ClusterRoleBinding huawei-csi-provisioner-role -n huawei-csi kubectl delete ClusterRoleBinding huawei-csi-resizer-role -n huawei-csi kubectl delete ClusterRoleBinding huawei-csi-snapshotter-role -n huawei-csi kubectl delete ClusterRoleBinding snapshot-controller-role -n huawei-csi kubectl delete Role huawei-csi-resizer-cfg -n huawei-csi kubectl delete Role huawei-csi-snapshotter-leaderelection -n huawei-csi kubectl delete Role snapshot-controller-leaderelection -n huawei-csi kubectl delete RoleBinding huawei-csi-resizer-role-cfg -n huawei-csi kubectl delete RoleBinding huawei-csi-snapshotter-leaderelection -n huawei-csi kubectl delete RoleBinding snapshot-controller-leaderelection -n huawei-csi kubectl delete RoleBinding snapshot-controller-leaderelection -n huawei-csi

FOF

b. Run the following command to delete the RBAC permission. If the **NotFound** error is reported, ignore it.

sh delete-huawei-csi-rbac.sh

Step 3 Check whether the RBAC permission has been deleted.

If the huawei-csi version is later than 2.2.15, run the following command. -n indicates the namespace. Change it based on site requirements. If No resources found is displayed, the permission is successfully deleted.

kubectl -n huawei-csi -l provisioner=csi.huawei.com get ServiceAccount,role,rolebinding,ClusterRole,ClusterRoleBinding

- If the huawei-csi version is 2.2.15 or earlier, perform the following operations to check whether the RBAC permission is successfully deleted.
 - Run the following command to create a file named check-huawei-csi-rbac.sh. -n indicates the namespace. Change it based on site requirements.

cat <<EOF > check-huawei-csi-rbac.sh
kubectl get ServiceAccount -n huawei-csi | grep huawei-csi
kubectl get ClusterRole -n huawei-csi | grep huawei-csi
kubectl get ClusterRoleBinding -n huawei-csi | grep huawei-csi
kubectl get Role -n huawei-csi | grep huawei-csi
kubectl get RoleBinding -n huawei-csi | grep huawei-csi
kubectl get RoleBinding -n huawei-csi | grep huawei-csi
kubectl get ClusterRole snapshot-controller-runner -n huawei-csi --ignore-not-found=true
kubectl get ClusterRoleBinding snapshot-controller-role -n huawei-csi --ignore-not-found=true
kubectl get Role snapshot-controller-leaderelection -n huawei-csi --ignore-not-found=true
kubectl get RoleBinding snapshot-controller-leaderelection -n huawei-csi --ignore-not-found=true
EOF

b. Run the following command. If no command output is displayed, the RBAC permission has been successfully deleted.

sh check-huawei-csi-rbac.sh

----End

8.2.6 Deleting the Image of the Earlier Version

To delete the **huawei-csi** image from the cluster, you need to perform the deletion operation on all worker nodes.

To delete the image from a single node, perform the following steps.

Prerequisites

The container service that depends on the image has been stopped. Otherwise, the image cannot be deleted.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to a worker node through the management IP address.
- **Step 2** Run the following command to view all existing versions.
 - If docker is used, run the **docker image ls | grep huawei-csi** command.

docker image ls | grep huawei-csi
REPOSITORY TAG IMAGE ID CREATED SIZE
huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
huawei-csi 3.0.0 14b854dba227 2 weeks ago 79.6MB

• If containerd is used, run the **crictl image ls | grep huawei-csi** command.

crictl image ls | grep huawei-csi
REPOSITORY TAG IMAGE ID CREATED SIZE
docker.io/library/huawei-csi 2.2.15 b30b3a8b5959 2 weeks ago 79.7MB
docker.io/library/huawei-csi 3.0.0 14b854dba227 2 weeks ago 79.6MB

Step 3 Run the following command to delete the image of the earlier version:

- If docker is used, run the **docker rmi** <*REPOSITORY*>:<*TAG*> command. # docker rmi huawei-csi:2.2.15
- If containerd is used, run the **crictl rmi** <*REPOSITORY>*:<*TAG>* command. # crictl rmi huawei-csi:2.2.15
- **Step 4** Run the following command again to check whether the image is successfully deleted. If the target version is not displayed, the image of the version is successfully deleted.
 - If docker is used, run the docker image ls | grep huawei-csi command.
 # docker image ls | grep huawei-csi
 huawei-csi
 3.0.0
 14b854dba227
 10 minutes ago 80MB
 - If containerd is used, run the **crictl image ls | grep huawei-csi** command.
 # crictl image ls | grep huawei-csi
 docker.io/library/huawei-csi 3.0.0 14b854dba2273 93.1MB

----End

8.3 (Optional) Uninstalling the Snapshot-Dependent Component Service

NOTICE

- Do not uninstall the snapshot-dependent component service when snapshots exist. Otherwise, Kubernetes will automatically delete all user snapshots and they cannot be restored. Exercise caution when performing this operation. For details, see **Delete a CustomResourceDefinition**.
- Do not uninstall the snapshot-dependent component service during the CSI upgrade.

Scenario Description

1. Currently, only Huawei CSI is available in the Kubernetes cluster, and Huawei CSI is no longer used.

2. Before the uninstallation, use Huawei CSI to clear VolumeSnapshots in the Kubernetes cluster. For details, see **6.5.3.2 Deleting a VolumeSnapshot**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to uninstall the snapshot-dependent component service.

kubectl delete crd volumesnapshotclasses.snapshot.storage.k8s.io volumesnapshotcontents.snapshot.storage.k8s.io volumesnapshots.snapshot.storage.k8s.io

Step 3 Run the following command to check whether the service is successfully uninstalled.

If the command output is empty, the uninstallation is successful. # kubectl get crd | grep snapshot.storage.k8s.io

----End

9 Common Operations

- 9.1 Updating the User Name or Password of a Storage Device Configured on CSI
- 9.2 Updating the configmap Object of huawei-csi
- 9.3 Adding a Backend for huawei-csi
- 9.4 Updating the huawei-csi-controller Service
- 9.5 Updating the huawei-csi-node Service
- 9.6 Modifying the Log Output Mode
- 9.7 Enabling the ReadWriteOncePod Feature Gate
- 9.8 Configuring Access to the Kubernetes Cluster as a Non-root User

9.1 Updating the User Name or Password of a Storage Device Configured on CSI

When the user name or password of a storage device changes, you need to update the configuration information on CSI. Otherwise, huawei-csi services cannot work properly.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **chmod +x secretUpdate** command to grant the execute permission on the secretUpdate tool.

chmod +x secretUpdate

Step 3 Run the ./secretUpdate namespace command to run the secretUpdate tool and enter the ID of the backend to be configured as prompted. If **Configured** is **false**, the backend is not configured. If **Configured** is **true**, the backend is configured.

```
2 true strage-backend-02 [https://192.168.125.26:8088]
3 true strage-backend-03 [https://192.168.125.27:8088]
4 true strage-backend-04 [https://192.168.125.28:8088]
5 true strage-backend-05 [https://192.168.125.29:28443]
6 true strage-backend-06 [https://192.168.125.30:28443]
Please enter the backend number to configure (Enter 'exit' to exit):3
```

Step 4 Enter the user name and password as prompted to update the **secret** object.

Step 5 After the configuration is complete, enter **exit** to exit and save the configuration.

```
Please enter the backend number to configure (Enter 'exit' to exit): exit
Saving configuration. Please wait......
The configuration is saved successfully.
```

- **Step 6** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 7** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 8** Run the **kubectl get pod -A | grep huawei** command to check whether the services are restarted successfully.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
huawei-csi huawei-csi-node-g6f7z 3/3 Running 0 14s
```

----End

9.2 Updating the configmap Object of huawei-csi

Perform this operation when you want to add a storage pool to an existing backend or change an existing service IP address.

9.2.1 Updating the configmap Object Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.

cp values.yaml values.yaml.bak

- **Step 4** Run the **vi** *values.yaml* command to open the file, and change the parameter values based on the upgrade requirements. After the change is complete, press **Esc** and enter :**wq!** to save the change. You can modify all parameters provided in **Table 4-1**. Exercise caution when modifying backend parameters.
- Step 5 Run the helm upgrade helm-huawei-csi./ -n huawei-csi-f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

```
# helm upgrade helm-huawei-csi / -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

vi values.vaml

9.2.2 Manually Updating the configmap Object

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the kubectl edit configmap huawei-csi-configmap -n huawei-csi command, press I or Insert to enter the editing mode, and modify related parameters. The iSCSI networking is used as an example. For details about the parameters, see Table 9-1. After the modification is complete, press Esc and enter :wq! to save the modification.

```
kind: ConfigMap
apiVersion: v1
metadata:
 name: huawei-csi-configmap
 namespace: huawei-csi
data:
 csi.json: |
      "backends": [
         {
            "storage": "oceanstor-san",
            "name": "storage",
             "urls": ["https://192.168.125.20:8088", "https://192.168.125.21:8088"],
             "pools": ["storagepool01", "storagepool02"],
"parameters": {"protocol": "iscsi", "portals": ["192.168.125.22", "192.168.125.23"]}
         }
      ]
  }
```

Table 9-1 Description of configuration items

Configuration Item	Format	Description	Remarks
data."csi.json".b ackends	List	List of back-end storage devices to be connected. This parameter is mandatory.	The number of back-end storage devices is not limited. For details about the fields that can be configured for a single back-end storage device, see Table 9-2.

Table 9-2 Configuration items of a back-end storage device

Configuration Item	Format	Description	Remarks
storage	String	Type of the storage device to be connected. This parameter is mandatory.	In the scenario where the distributed storage SAN is connected, the value is fixed to fusionstoragesan.
name	String	Storage backend name.	 User-defined character string. The value can contain uppercase letters, lowercase letters, digits, and hyphens (-). This parameter cannot be modified.
urls	List	Management URL of the storage device to be connected. This parameter is mandatory.	One or more management URLs of the same storage device are supported. Use commas (,) to separate multiple management URLs. Currently, only IPv4 addresses are supported. Example: https:// 192.168.125.20:8088 NOTE A storage device has multiple controllers, and each controller has a management URL. Therefore, a storage device may have multiple management URLs.

Configuration Item	Format	Description	Remarks
pools	List	Name of a storage pool used on the storage device to be connected. This parameter is mandatory.	 One or more storage pools on the same storage device are supported. Use commas (,) to separate multiple storage pools. Currently, only storage pools can be added. You can log in to DeviceManager to obtain the storage pools that support the block storage service.
parameters	Dictionary	Variable parameters in scenarios where iSCSI is used.	In scenarios where iSCSI is used, set the protocol parameter to a fixed value: iscsi . Set the portals parameter to the iSCSI service IP addresses of the storage device. Use commas (,) to separate multiple iSCSI service IP addresses. You can log in to DeviceManager to obtain the iSCSI service IP addresses. Take OceanStor Dorado 6. <i>x</i> series as an example. On DeviceManager, choose Services > Network > Logical Ports and obtain the IP address whose data protocol is iSCSI. (For other series, see the corresponding operation description.)

- Step 3 If the storage, name, or urls parameter is modified, you need to update the user name or password of the storage device. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.
- **Step 4** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 5** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=huawei-csi | kubectl replace --force -f -

Step 6 Run the **kubectl get pod -A | grep huawei** command to check whether the services are restarted successfully.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
huawei-csi huawei-csi-node-g6f7z 3/3 Running 0 14s
----End
```

9.3 Adding a Backend for huawei-csi

Perform this operation when you want to add a storage device or a storage pool as an independent backend.

9.3.1 Adding a Backend Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- Step 3 Back up the values.yaml file used during CSI installation. If the values.yaml file used during the last installation cannot be found, run the helm get values helm-huawei-csi -n huawei-csi -a command to query the file.

 # cp values.yaml values.yaml.bak
- **Step 4** Run the **vi** *values.yaml* command to open the configuration file and add backend configurations under **backends**. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
backends:
 - storage: "oceanstor-san"
  name: "***"
  urls:
    - "https://*.*.*:8088"
  pools:
  parameters:
   protocol: "iscsi"
    portals:
     - "*.*.*.*"
 - storage: "oceanstor-nas"
  name: "***"
  urls:
    - "https://*.*.*:8088"
  pools:
  parameters:
    protocol: "nfs"
    portals:
```

Step 5 Run the helm upgrade helm-huawei-csi./ -n huawei-csi-f values.yaml command to upgrade the Helm chart. The upgrade command will update the Deployment, daemonset, and RBAC resources, but will not update the secret resource. If the backend information changes, you must manually update secret. For details, see

9.1 Updating the User Name or Password of a Storage Device Configured on CSI.

In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

9.3.2 Manually Adding a Backend

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Configure multiple backends. For details, see **7.1 Configuring Multiple Backends**.
- Step 3 Configure accounts for the new backends. For details, see 9.1 Updating the User Name or Password of a Storage Device Configured on CSI.
- **Step 4** Run the following command to restart the huawei-csi-controller service.

 # kubectl get deployment huawei-csi-controller -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 5** Run the following command to restart the huawei-csi-node service.

 # kubectl get daemonset huawei-csi-node -o yaml -n=huawei-csi | kubectl replace --force -f -
- **Step 6** Run the **kubectl get pod -A | grep huawei** command to check whether the services are restarted successfully.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
huawei-csi huawei-csi-node-g6f7z 3/3 Running 0 14s
```

----End

9.4 Updating the huawei-csi-controller Service

Perform this operation when you need to update the huawei-csi-controller service, for example, adding the snapshot or the capacity expansion function.

9.4.1 Updating the controller Service Using Helm

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any

directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.

Step 3 Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.

cp values.yaml values.yaml.bak

Step 4 Run the **vi** *values.yaml* command to open the file and modify controller parameters. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
kubernetes:
 namespace: huawei-csi
 # The image name and tag for the attacher, provisioner and registrar sidecars. These must match the
appropriate Kubernetes version.
 sidecar:
  csiAttacher: k8s.gcr.io/sig-storage/csi-attacher:v3.3.0
  csiProvisioner: k8s.gcr.io/sig-storage/csi-provisioner:v3.0.0
  csiResizer: k8s.gcr.io/sig-storage/csi-resizer:v1.3.0
  livenessProbe: k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
  csiSnapshotter: k8s.gcr.io/sig-storage/csi-snapshotter:v4.2.1
  snapshotController: k8s.gcr.io/sig-storage/snapshot-controller:v4.2.1
 # The image name and tag for the Huawei CSI Service container
 # Replace the appropriate tag name
 huaweiCSIService: huawei-csi:3.0.0
# The CSI driver parameter configuration
csi driver:
 driverName: csi.huawei.com
 endpoint: /csi/csi.sock
 backendUpdateInterval: 60
 controllerLogging:
  module: file
  level: info
  fileDir: /var/log/huawei
  fileSize: 20M
  maxBackups: 9
huaweiCsiController:
 replicas: 1 # Default number of controller replicas
# Default image pull policy for sidecar container images
sidecarImagePullPolicy: "IfNotPresent"
# Default image pull policy for Huawei plugin container images
huaweiImagePullPolicy: "IfNotPresent"
# Flag to enable or disable snapshot (Optional)
snapshot:
 enable: true
# Flag to enable or disable resize (Optional)
resizer:
 enable: true
```

Step 5 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

```
# helm upgrade helm-huawei-csi // -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
```

LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

9.4.2 Manually Updating the controller Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-controller service. For details, see **8.2.2 Uninstalling the** huawei-csi-controller Service.
- Step 3 Delete the RBAC permission. For details, see 8.2.5 Deleting the RBAC Permission.
- **Step 4** Create the RBAC permission. For details, see **Step 4**.
- **Step 5** Start the controller service. For details, see **Step 5**.
- **Step 6** After the huawei-csi service is deployed, run the **kubectl get pod -A | grep huawei-csi-controller** command to check whether the service is started.

```
# kubectl get pod -A | grep huawei-csi-controller
huawei-csi huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s
```

----End

9.5 Updating the huawei-csi-node Service

Perform this operation when you need to update the huawei-csi-node service.

9.5.1 Updating the node Service Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.
 - # cp values.yaml values.yaml.bak
- **Step 4** Run the **vi** *values.yaml* command to open the file and modify node parameters. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

kubernetes: namespace: huawei-csi

```
images:
 # The image name and tag for the attacher, provisioner and registrar sidecars. These must match the
appropriate Kubernetes version.
 sidecar:
  livenessProbe: k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
  registrar: k8s.gcr.io/sig-storage/csi-node-driver-registrar:v2.3.0
 # The image name and tag for the Huawei CSI Service container
 # Replace the appropriate tag name
 huaweiCSIService: huawei-csi:3.0.0
# The CSI driver parameter configuration
csi_driver:
 driverName: csi.huawei.com
 endpoint: /csi/csi.sock
 connectorThreads: 4
 volumeUseMultipath: true # Flag to enable or disable volume multipath access
 scsiMultipathType: DM-multipath #Required, if volume-use-multipath is set to TRUE
 nvmeMultipathType: HW-UltraPath-NVMe #Required, if volume-use-multipath is set to TRUE
 scanVolumeTimeout: 3
 backendUpdateInterval: 60
 nodeLogging:
  module: file
  level: info
  fileDir: /var/log/huawei
  fileSize: 20M
  maxBackups: 9
# Default image pull policy for sidecar container images
sidecarImagePullPolicy: "IfNotPresent"
# Default image pull policy for Huawei plugin container images
huaweilmagePullPolicy: "IfNotPresent"
```

Step 5 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

9.5.2 Manually Updating the node Service

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-node service. For details, see **8.2.1 Uninstalling the** huawei-csi-node Service.
- **Step 3** Install the new huawei-csi-node service. For details, see **4.2.3 Starting huawei-csi Services**.

Step 4 After the huawei-csi service is deployed, run the **kubectl get pod -A | grep huawei-csi-node** command to check whether the service is started.

```
# kubectl get pod -A | grep huawei-csi-node
huawei-csi huawei-csi-node-g6f7z 3/3 Running 0 14s
```

----End

9.6 Modifying the Log Output Mode

huawei-csi supports two log output modes: **file** and **console**. **file** indicates that logs are output to the fixed directory (**/var/log/huawei**), and **console** indicates that logs are output to the standard directory of the container. You can set the log output mode as required. The default mode is **file**.

9.6.1 Modifying the Log Output Mode of the controller or node Service Using Helm

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* -**n** *huawei-csi* -**a** command to query the file.

cp values.yaml values.yaml.bak

Step 4 Run the **vi** *values.yaml* command to open the file and modify controller or node log parameters. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
# The CSI driver parameter configuration
csi_driver:
controllerLogging:
module: file
level: info
fileDir: /var/log/huawei
fileSize: 20M
maxBackups: 9
nodeLogging:
module: file
level: info
fileDir: /var/log/huawei
fileSize: 20M
maxBackups: 9
```

Step 5 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
```

NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

9.6.2 Manually Modifying the Log Output Mode of the huawei-csi-controller Service

Perform this operation when you want to set the log output mode of the huawei-csi-controller service.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-controller service. For details, see **8.2.2 Uninstalling the** huawei-csi-controller Service.
- **Step 3** Go to the **deploy** directory. For details about the component package path, see **3.3 Uploading the Components in the Software Package**.
- **Step 4** Run the **vi** *huawei-csi-controller.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and modify the following parameters. For details, see **Table 9-3**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

args:

- "--endpoint=\$(CSI_ENDPOINT)"
- "--controller"
- "--containerized"
- "--driver-name=csi.huawei.com"
- "--loggingModule=file"
- "--logLevel=info"
- "--logFileDir=/var/log/huawei"
- "--logFileSize=20M"
- "--maxBackups=9"

Table 9-3 Description of log output parameters

Configuration Item	Description	Remarks
loggingModule	huawei-csi log output mode.	The value can be file or console . The default value is file .
logLevel	huawei-csi log output level.	Supported levels are debug , info , warning , error , and fatal . The default level is info .
logFileDir	huawei-csi log directory in file output mode.	This parameter is available only when loggingModule is set to file . The default log directory is /var/log/huawei .

Configuration Item	Description	Remarks
logFileSize	Size of a single huawei-csi log file in file output mode.	This parameter is available only when loggingModule is set to file . The default log file size is 20 MiB.
maxBackups	Maximum number of huawei-csi log file backups in file output mode.	This parameter is available only when loggingModule is set to file . The default number of log file backups is 9.

Step 5 Run the following command to start the controller service.

kubectl create -f huawei-csi-controller-snapshot-v1.yaml

Step 6 After the huawei-csi service is deployed, run the **kubectl get pod -A -o wide** | **grep huawei** command to check whether the service is started.

```
# kubectl get pod -A -o wide | grep huawei
huawei-csi huawei-csi-controller-b59577886-qqzm8 7/7 Running 0 18h 10.244.1.67
node <none> <none>
```

- **Step 7** View the logs of the huawei-csi-controller service.
 - If **loggingModule** is set to **file**, log in to the node, go to the log directory specified by **logFileDir**, and run the following command to view the log of huawei-csi-controller.

tail -f huawei-csi-controller

 If loggingModule is set to console, run the following command to view the log of huawei-csi-controller.

kubectl logs *huawei-csi-controller* -c huawei-csi-driver -n huawei-csi

----End

9.6.3 Manually Modifying the Log Output Mode of the huawei-csi-node Service

Perform this operation when you want to set the log output mode of the huawei-csi-node service.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Uninstall the huawei-csi-node service. For details, see **8.2.1 Uninstalling the** huawei-csi-node Service.
- Step 3 Run the vi huawei-csi-node.yaml command to modify the .yaml file. Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press Esc and enter :wq! to save the modification. Compile the huawei-csi-node.yaml file. For details, see sample file deploy/huawei-csi-node.yaml in the software package. For details about the parameters, see Table 9-4.

args:
 - "--endpoint=/csi/csi.sock"

- "--containerized"
- "--driver-name=csi.huawei.com"
- "--volume-use-multipath=false"
- "--loggingModule=file"
- "--logLevel=info"
- "--logFileDir=/var/log/huawei"
- "--logFileSize=20M"
- "--maxBackups=9"

Table 9-4 Description of log output parameters

Configuration Item	Description	Remarks
loggingModule	huawei-csi log output mode.	The value can be file or console . The default value is file .
logLevel	huawei-csi log output level.	Supported levels are debug , info , warning , error , and fatal . The default level is info .
logFileDir	huawei-csi log directory in file output mode.	This parameter is available only when loggingModule is set to file . The default log directory is /var/log/huawei .
logFileSize	Size of a single huawei-csi log file in file output mode.	This parameter is available only when loggingModule is set to file . The default log file size is 20 MiB.
maxBackups	Maximum number of huawei-csi log file backups in file output mode.	This parameter is available only when loggingModule is set to file . The default number of log file backups is 9.

Step 4 Run the following command to start the node service.

kubectl create -f huawei-csi-node.yaml

Step 5 After the huawei-csi service is deployed, run the **kubectl get pod -A -o wide** | **grep huawei-csi-node** command to check whether the service is started.

```
# kubectl get pod -A | grep huawei-csi-node
huawei-csi huawei-csi-node-4sfwr 3/3 Running 0 18h 10.244.1.68 node
<none> <none>
```

- **Step 6** View the logs of the huawei-csi-node service.
 - If **loggingModule** is set to **file**, log in to the node, go to the log directory specified by **logFileDir**, and run the following command to view the log of huawei-csi-node.

tail -f huawei-csi-node

• If **loggingModule** is set to **console**, run the following command to view the log of huawei-csi-node.

kubectl logs *huawei-csi-node* -c huawei-csi-driver -n huawei-csi

----End

9.7 Enabling the ReadWriteOncePod Feature Gate

The ReadWriteOnce access mode is the fourth access mode introduced by Kubernetes v1.22 for PVs and PVCs. If you create a Pod using a PVC in ReadWriteOncePod access mode, Kubernetes ensures that the Pod is the only Pod in the cluster that can read or write the PVC.

The ReadWriteOncePod access mode is an alpha feature in Kubernetes v1.22/1.23/1.24. Therefore, you need to enable the ReadWriteOncePod feature in **feature-gates** of kube-apiserver, kube-scheduler, and kubelet before using the access mode.

Procedure

Step 1 Enable the ReadWriteOncePod feature gate for kube-apiserver.

- 1. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Run the vi /etc/kubernetes/manifests/kube-apiserver.yaml command, press
 I or Insert to enter the editing mode, and add --featuregates=ReadWriteOncePod=true to the kube-apiserver container. After the
 modification is complete, press Esc and enter:wq! to save the modification.




Ⅲ NOTE

After the editing is complete, Kubernetes will automatically apply the updates.

Step 2 Enable the ReadWriteOncePod feature gate for kube-scheduler.

- 1. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Run the vi /etc/kubernetes/manifests/kube-scheduler.yaml command, press I or Insert to enter the editing mode, and add --featuregates=ReadWriteOncePod=true to the kube-scheduler container. After the modification is complete, press Esc and enter:wg! to save the modification.



After the editing is complete, Kubernetes will automatically apply the updates.

Step 3 Enable the ReadWriteOncePod feature gate for kubelet.

NOTICE

The dynamic Kubelet configuration function is not used since v1.22 and deleted in v1.24. Therefore, you need to perform the following operations on kubelet on each worker node in the cluster.

- 1. Use a remote access tool, such as PuTTY, to log in to any worker node in the Kubernetes cluster through the management IP address.
- Run the vi /var/lib/kubelet/config.yaml command, press I or Insert to enter the editing state, and add ReadWriteOncePod: true to the feature-gates field of the KubeletConfiguration object. If the feature-gates field does not exist, add it at the same time. After the modification is complete, press Esc and enter:wq! to save the modification.

apiVersion: kubelet.config.k8s.io/v1beta1 featureGates:

ReadWriteOncePod: true

The default path of the kubelet configuration file is /var/lib/kubelet/config.yaml. Enter the path based on site requirements.

3. After the configuration is complete, run the **systemctl restart kubelet** command to restart kubelet.

----End

9.8 Configuring Access to the Kubernetes Cluster as a Non-root User

Procedure

Step 1 Copy the authentication file of the Kubernetes cluster and modify /etc/ kubernetes/admin.conf to be the actual authentication file.

\$ mkdir -p \$HOME/.kube \$ sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config

Step 2 Change the user and user group of the authentication file.

\$ sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config

Step 3 Configure the **KUBECONFIG** environment variable of the current user. The following uses Ubuntu 20.04 as an example.

\$ echo "export KUBECONFIG=\$HOME/.kube/config" >> ~/.bashrc
\$ source ~/.bashrc

----End

10_{FAQ}

10.1 Viewing Log Information

10.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernetes Platform Is Set Up for the First Time

10.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported

10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters

10.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName

10.6 When a PVC Is Created, the PVC Is in the Pending State

10.7 Before a PVC Is Deleted, the PVC Is in the Pending State

10.8 When a Pod Is Created, the Pod Is in the ContainerCreating State

10.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created

10.1 Viewing Log Information

Viewing Logs Generated When the secret Object Is Configured

Step 1 Run the **cd /var/log/huawei** command to go to the log directory.

cd /var/log/huawei

Step 2 Run the following command to view the logs of huawei-csi-install.

vi huawei-csi-install

----Fnd

Viewing Logs of the huawei-csi-controller Service

Step 1 Run the following command to obtain the node where huawei-csi-controller is located.

```
# kubectl get pod -A -o wide | grep huawei
huawei-csi huawei-csi-controller-695b84b4d8-tg64l 7/7 Running 0 14s <host1-ip>
<host1-name> <none> <none>
```

- **Step 2** Use a remote access tool, such as PuTTY, to log in to the huawei-csi-controller node in the Kubernetes cluster through the management IP address.
- **Step 3** Run the **cd /var/log/huawei** command to go to the log directory. # cd /var/log/huawei
- **Step 4** Run the following command to view the customized output logs of the container.

 # vi huawei-csi-controller
- **Step 5** Run the **cd /var/log/containers** command to go to the container directory. # cd /var/log/containers
- **Step 6** Run the following command to view the standard output logs of the container.

 # vi huawei-csi-controller-<name>_huawei-csi_huawei-csi-driver-<contrainer-id>.log

----End

Viewing Logs of the huawei-csi-node Service

Step 1 Run the following command to obtain the node where huawei-csi-node is located. # kubectl get pod -A -o wide | grep huawei

huawei-csi huawei-csi-node-g6f7z 3/3 **Running** 0 14s <host2-ip> <host2-name> <none>

- **Step 2** Use a remote access tool, such as PuTTY, to log in to the huawei-csi-node node in the Kubernetes cluster through the management IP address.
- **Step 3** Run the **cd /var/log/huawei** command to go to the log directory.

 # cd /var/log/huawei
- **Step 4** Run the following command to view the customized output logs of the container.

 # vi huawei-csi-node
- **Step 5** Run the **cd /var/log/containers** command to go to the container directory. # cd /var/log/containers
- **Step 6** Run the following command to view the standard output logs of the container.

 # vi huawei-csi-node# vi huawei-csi-node-

----End

10.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernetes Platform Is Set Up for the First Time

Symptom

When you create a Pod, error Cannot connect ISCSI portal *.*.*: libkmod: kmod_module_insert_module: could not find module by name='iscsi_tcp' is reported in the /var/log/huawei-csi-node log.

Root Cause Analysis

The iscsi_tcp service may be stopped after the Kubernetes platform is set up and the iscsi service is installed. You can run the **lsmod | grep iscsi | grep iscsi_tcp** command to check whether the service is stopped.

Solution or Workaround

Run the following command to manually load the iscsi_tcp service.

```
# modprobe iscsi_tcp
# lsmod | grep iscsi | grep iscsi_tcp
iscsi_tcp 18333 6
libiscsi_tcp 25146 1 iscsi_tcp
```

10.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported

Symptom

The huawei-csi-node service cannot be started. When you run the **kubectl describe daemonset huawei-csi-node -n huawei-csi** command, error message "/var/lib/iscsi is not a directory" is reported.

Root Cause Analysis

The /var/lib/iscsi directory does not exist in the huawei-csi-node container.

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the following command to delete the huawei-csi-node service (huawei-csi-node.yaml is the configuration file in Step 6 in 4.2.3 Starting huawei-csi-Services).

```
# kubectl delete -f huawei-csi-node.yaml
```

- **Step 3** Run the following command to view the huawei-csi-node service. If no command output is displayed, the deletion is complete.
 - # kubectl get pod -A | grep huawei-csi-node
- Step 4 Run the vi huawei-csi-node.yaml command to modify the .yaml file. Press I or Insert to enter the editing mode, set path in huawei-csi-node.yaml > volumes > iscsi-dir > hostPath to /var/lib/iscsi and delete the type line. After the modification is complete, press Esc and enter :wq! to save the modification. Compile the huawei-csi-node.yaml file. For details, see sample file deploy/huawei-csi-node.yaml in the software package.

Step 5 Run the following command to start the node service.

kubectl create -f huawei-csi-node.yaml

Step 6 After the huawei-csi service is deployed, run the **kubectl get pod -A | grep huawei-csi-node** command to check whether the service is started.

kubectl get pod -A | grep huawei-csi-node huawei-csi huawei-csi-node-q6f7z 3/3 Running 0 14s

----End

10.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters

Symptom

A Pod is running on worker node A, and an external block device is mounted to the Pod through CSI. After worker node A is powered off abnormally, the Kubernetes platform detects that the node is faulty and switches the Pod to worker node B. After worker node A recovers, the drive letters on worker node A change from normal to faulty.

Environment Configuration

Kubernetes version: 1.18 or later

Storage type: block storage

Root Cause Analysis

After worker node A recovers, Kubernetes initiates an unmapping operation on the storage, but does not initiate a drive letter removal operation on the host. After Kubernetes completes the unmapping, residual drive letters exist on worker node A

Solution or Workaround

Currently, you can only manually clear the residual drive letters on the host. Alternatively, restart the host again and use the disk scanning mechanism during the host restart to clear the residual drive letters. The specific method is as follows:

Step 1 Check the residual drive letters on the host.

1. Run the **multipath -ll** command to check whether a DM multipathing device with abnormal multipathing status exists.

As shown in the following figure, the path status is **failed faulty running**, the corresponding DM multipathing device is **dm-12**, and the associated SCSI disks are **sdi** and **sdj**. If multiple paths are configured, multiple SCSI disks exist. Record these SCSI disks.

multipath -ll mpathb (3618cf24100f8f457014a764c000001f6) dm-12 HUAWEI "XSG1

- If yes, go to Step 1.2.
- If no, no further action is required.
- 2. Check whether the residual DM multipathing device is readable.

Run the **dd if=/dev/***dm-xx* **of=/dev/null count=1 bs=1M iflag=direct** command.

dm-xx indicates the device ID obtained in **Step 1.1**.

If the returned result is **Input/output error** and the read data is **0 bytes (0 B) copied**, the device is unreadable.

```
#dd if=/dev/dm-12 of=/dev/null count=1 bs=1M iflag=direct
dd: error reading '/dev/dm-12': Input/output error
0+0 records in
0+0 records out
0 bytes (0 B) copied, 0.0236862 s, 0.0 kB/s
```

- If yes, record the residual dm-xx device and associated disk IDs (for details, see Step 1.1) and perform the clearing operation.
- If the command execution is suspended, go to Step 1.3.
- If other cases, contact technical support engineers.
- 3. Log in to the node again in another window.
 - Run the following command to view the suspended process.
 # ps -ef | grep dm-12 | grep -w dd
 root 21725 9748 0 10:33 pts/10 00:00:00 dd if=/dev/dm-12 of=/dev/null count=1 bs=10M
 iflag=direct
 - b. Kill the pid. # kill -9 pid
 - c. Record the residual *dm-xx* device and associated disk IDs (for details, see **Step 1.1**) and perform the clearing operation.

Step 2 Clear the residual drive letters on the host.

1. Run the **multipath -f /dev/***dm-** command to delete residual multipathing aggregation device information according to the DM multipathing device obtained in **Step 1**.

```
# multipath -f /dev/dm-12
```

If an error is reported, contact technical support engineers.

2. Run the following command to clear the residual SCSI disks according to the drive letters of the residual disks obtained in the troubleshooting method. echo 1 > /sys/block/xxxx/device/delete

When multiple paths are configured, clear the residual disks based on the drive letters. The residual paths are **sdi** and **sdj**.

```
# echo 1 > /sys/block/sdi/device/delete
# echo 1 > /sys/block/sdj/device/delete
```

If an error is reported, contact technical support engineers.

3. Check whether the DM multipathing device and SCSI disk information has been cleared.

Run the **multipath** -ll, ls -l /sys/block/, and ls -l /dev/disk/by-id/ commands in sequence to query the path and disk information. If the residual dm-12 device and SCSI disks **sdi** and **sdj** are cleared, the clearing is complete.

```
# multipath -ll
mpathb (3618cf24100f8f457014a764c000001f6) dm-3 HUAWEI ,XSG1
size=100G features='0' hwhandler='0' wp=rw
-+- policy='service-time 0' prio=-1 status=active
|- 39:0:0:1
               sdd 8:48 active ready running
 `- 38:0:0:1
               sde 8:64 active ready running
mpathn (3618cf24100f8f457315a764c000001f6) dm-5 HUAWEI ,XSG1
size=100G features='0' hwhandler='0' wp=rw
-+- policy='service-time 0' prio=-1 status=active
|- 39:0:0:2
               sdc 8:32 active ready running
  - 38:0:0:2
               sdb 8:16 active ready running
# ls -l /sys/block/
total 0
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-0 -> ../devices/virtual/block/dm-0
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-1 -> ../devices/virtual/block/dm-1
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-2 -> ../devices/virtual/block/dm-2
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-3 -> ../devices/virtual/block/dm-3
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdb -> ../devices/platform/host35/session2/
target35:0:0/35:0:0:1/block/sdb
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdc -> ../devices/platform/host34/
target34:65535:5692/34:65535:5692:0/block/sdc
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdd -> ../devices/platform/host39/session6/
target39:0:0/39:0:0:1/block/sdd
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sde -> ../devices/platform/host38/session5/
target38:0:0/38:0:0:1/block/sde
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdh -> ../devices/platform/host39/session6/
target39:0:0/39:0:0:3/block/sdh
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdi -> ../devices/platform/host38/session5/target38:0:0/38:0:0:3/
block/sdi
ls -l /dev/disk/by-id/
total 0
lrwxrwxrwx 1 root root 10 Aug 11 19:57 dm-name-mpathb -> ../../dm-3
lrwxrwxrwx 1 root root 10 Aug 11 19:58 dm-name-mpathn -> ../../dm-5
lrwxrwxrwx 1 root root 10 Aug 11 19:57 dm-uuid-mpath-3618cf24100f8f457014a764c000001f6 -> ../../
lrwxrwxrwx 1 root root 10 Aug 11 19:58 dm-uuid-mpath-3618cf24100f8f457315a764c000001f6 -> ../../
dm-5
lrwxrwxrwx 1 root root 9 Aug 11 19:57 scsi-3618cf24100f8f457014a764c000001f6 -> ../../sdd
lrwxrwxrwx 1 root root 9 Aug 11 19:57 scsi-3618cf24100f8f45712345678000103e8 -> ../../sdi
lrwxrwxrwx 1 root root 9 Aug 3 15:17 scsi-3648435a10058805278654321ffffffff -> ../../sdb
lrwxrwxrwx 1 root root 9 Aug 2 14:49 scsi-368886030000020aff44cc0d060c987f1 -> ../../sdc
lrwxrwxrwx 1 root root 9 Aug 11 19:57 wwn-0x618cf24100f8f457014a764c000001f6 -> ../../sdd
lrwxrwxrwx 1 root root 9 Aug 11 19:57 wwn-0x618cf24100f8f45712345678000103e8 -> ../../sdi
lrwxrwxrwx 1 root root 9 Aug 3 15:17 wwn-0x648435a10058805278654321ffffffff -> ../../sdb
lrwxrwxrwx 1 root root 9 Aug 2 14:49 wwn-0x68886030000020aff44cc0d060c987f1 -> ../../sdc
```

----End

10.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName

Symptom

The huawei-csi services (huawei-csi-controller or huawei-csi-node) cannot be started. After the **kubectl get pod -A | grep huawei** command is executed, the command output shows that the service status is **InvalidImageName**.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-fd5f97768-qlldc 6/7 InvalidImageName 0 16s
huawei-csi huawei-csi-node-25txd 2/3 InvalidImageName 0 15s
```

Root Cause Analysis

In the .yaml configuration files of the controller and node, the Huawei CSI image version number is incorrect. For example:

```
...
- name: huawei-csi-driver
image: huawei-csi:3.0.0
...
```

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to modify the configuration file of the huawei-csi-node service. Press I or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

kubectl edit daemonset huawei-csi-node -o yaml -n=huawei-csi

□ NOTE

 In the image configuration item under huawei-csi-driver in the sample .yaml file, huawei-csi:*.** must be replaced with <Name>:<Version> of the created Huawei CSI image.

containers:

- name: huawei-csi-driver image: huawei-csi:3.0.0
- **Step 3** Run the following command to modify the configuration file of the huawei-csi-controller service: Press I or **Insert** to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

kubectl edit deployment huawei-csi-controller -o yaml -n=huawei-csi

∩ NOTE

• In the **image** configuration item under **huawei-csi-driver** in the sample .yaml file, huawei-csi:*.** must be replaced with <Name>:<Version> of the created Huawei CSI image.

containers:

- name: huawei-csi-driver image: huawei-csi:3.0.0
- **Step 4** Wait until the huawei-csi-node and huawei-csi-controller services are started.
- **Step 5** Run the following command to check whether the huawei-csi services are started.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-58799449cf-zvhmv 7/7 Running 0 2m29s
huawei-csi huawei-csi-node-7fxh6 3/3 Running 0 12m
```

----End

10.6 When a PVC Is Created, the PVC Is in the Pending State

Symptom

A PVC is created. After a period of time, the PVC is still in the **Pending** state.

Root Cause Analysis

Cause 1: A StorageClass with the specified name is not created in advance. As a result, Kubernetes cannot find the specified StorageClass name when a PVC is created.

Cause 2: The storage pool capability does not match the StorageClass capability. As a result, huawei-csi fails to select a storage pool.

Cause 3: An error code (for example, 50331651) is returned by a RESTful interface of the storage. As a result, huawei-csi fails to create a PVC.

Cause 4: The storage does not return a response within the timeout period set by huawei-csi. As a result, huawei-csi returns a timeout error to Kubernetes.

Cause 5: Other causes.

Solution or Workaround

When a PVC is created, if the PVC is in the **Pending** state, you need to take different measures according to the following causes.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to view details about the PVC.

 # kubectl describe pvc mypvc
- **Step 3** Perform the corresponding operation according to the **Events** information in the detailed PVC information.
 - If the PVC is in the **Pending** state due to cause 1, perform the following steps.

Events:							
Type	Reason	Age	From	Message			
Warning ProvisioningFailed 0s (x15 over 3m24s) persistentvolume-controller							
storageclass.storage.k8s.io " <i>mysc</i> " not found							

- a. Delete the PVC. For details, see **6.3.5** Deleting a PVC.
- b. Create a StorageClass. For details, see **6.1.1 Creating a StorageClass**.
- c. Create a PVC. For details, see 6.3.1 Creating a PVC.
- If the PVC is in the **Pending** state due to cause 2, perform the following steps.

Events: Type From	Reason	Age	Message
Normal	Provisioning	63s (x3 over	64s) csi.huawei.com_huawei-csi-controller-b59577886-
qqzm8_5	8533e4a-884c-4	4c7f-92c3-6e8a7b3	27515 External provisioner is provisioning volume for

claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = **failed to select pool**, the capability filter failed, error: failed to select pool, the final filter field: **replication**, parameters map[allocType:thin replication:True size: 1099511627776 volumeType:lun]. please check your storage class

- a. Delete the PVC. For details, see **6.3.5** Deleting a PVC.
- b. Delete the StorageClass. For details, see **6.1.2 Deleting a StorageClass**.
- c. Modify the **StorageClass.yaml** file based on the **Events** information.
- d. Create a StorageClass. For details, see 6.1.1 Creating a StorageClass.
- e. Create a PVC. For details, see **6.3.1 Creating a PVC**.
- If the PVC is in the **Pending** state due to cause 3, contact Huawei engineers.

Events: Type From	Reason	Age	Message

Normal Provisioning 63s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 62s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = Create volume map[ALLOCTYPE:1 CAPACITY:20 DESCRIPTION:Created from Kubernetes CSI NAME:pvc-63ebfda5-4cf0-458e-83bd-ecc PARENTID:0] error: 50331651

• If the PVC is in the **Pending** state due to cause 4, perform the following steps.

Type From	Reason	Age		Message

Normal Provisioning 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = context deadline exceeded (Client.Timeout exceeded while awaiting headers)

- a. Wait for 10 minutes and check the PVC details again by referring to this section.
- b. If it is still in the **Pending** state, contact Huawei engineers.
- If the PVC is in the **Pending** state due to cause 5, contact Huawei engineers.

----End

10.7 Before a PVC Is Deleted, the PVC Is in the Pending State

Symptom

Before a PVC is deleted, the PVC is in the **Pending** state.

Root Cause Analysis

Cause 1: A StorageClass with the specified name is not created in advance. As a result, Kubernetes cannot find the specified StorageClass name when a PVC is created.

Cause 2: The storage pool capability does not match the StorageClass capability. As a result, huawei-csi fails to select a storage pool.

Cause 3: An error code (for example, 50331651) is returned by a RESTful interface of the storage. As a result, huawei-csi fails to create a PVC.

Cause 4: The storage does not return a response within the timeout period set by huawei-csi. As a result, huawei-csi returns a timeout error to Kubernetes.

Cause 5: Other causes.

Solution or Workaround

To delete a PVC in the **Pending** state, you need to take different measures according to the following causes.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to view details about the PVC. # kubectl describe pvc mypvc
- **Step 3** Perform the corresponding operation according to the **Events** information in the detailed PVC information.
 - If the PVC is in the **Pending** state due to cause 1, run the **kubectl delete pvc** *mypvc* command to delete the PVC.

Events:							
Type	Reason	Age	From	Message			
Warning ProvisioningFailed 0s (x15 over 3m24s) persistentvolume-controller							
storageclass.storage.k8s.io " <i>mysc</i> " not found							

• If the PVC is in the **Pending** state due to cause 2, run the **kubectl delete pvc** *mypvc* command to delete the PVC.

Events:						
Type	Reason	Age				
From				Message		
Norma	l Provisioning	63s (:	x3 over 64s) o	csi.huawei.com_huawei-csi-controller-b59577886-		
qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for						
claim "d	efault/mypvc"					
Warnin	g Provisioning	Failed 63s	(x3 over 64s)	csi.huawei.com_huawei-csi-controller-b59577886-		
qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass						
"mysc": rpc error: code = Internal desc = failed to select pool , the capability filter failed, error: failed						
to select	pool, the final	filter field:	<i>replication</i> , pa	arameters map[allocType:thin replication:True size:		
1099511627776 volumeType:lun]. please check your storage class						

• If the PVC is in the **Pending** state due to cause 3, run the **kubectl delete pvc** *mypvc* command to delete the PVC.

Events: Type From	Reason	Age	Message
	l Provisioning	63s (x4 over 68s) csi.huawei.com	_huawei-csi-controller-b59577886-
	58533e4a-884c-4c	7f-92c3-6e8a7b327515 External prov	isioner is provisioning volume for

claim "default/mypvc"

Warning ProvisioningFailed 62s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = Create volume map[ALLOCTYPE:1 CAPACITY:20 DESCRIPTION:Created from Kubernetes CSI NAME:pvc-63ebfda5-4cf0-458e-83bd-ecc PARENTID:0] error: 50331651

• If the PVC is in the **Pending** state due to cause 4, contact Huawei engineers.

```
Events:

Type Reason Age

From Message

---- Mormal Provisioning 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = context deadline exceeded (Client.Timeout exceeded while awaiting headers)
```

• If the PVC is in the **Pending** state due to cause 5, contact Huawei engineers.

----Fnd

10.8 When a Pod Is Created, the Pod Is in the ContainerCreating State

Symptom

A Pod is created. After a period of time, the Pod is still in the **ContainerCreating** state. Check the log information (for details, see **10.1 Viewing Log Information**). The error message "Fibre Channel volume device not found" is displayed.

Root Cause Analysis

This problem occurs because residual disks exist on the host node. As a result, disks fail to be found when a Pod is created next time.

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query information about the node where the Pod resides.

```
# kubectl get pod -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE
READINESS GATES
mypod 0/1 ContainerCreating 0 51s 10.244.1.224 node1 <none>
```

- **Step 3** Delete the Pod. For details, see **6.4.2 Deleting a Pod**.
- **Step 4** Use a remote access tool, such as PuTTY, to log in to the *node1* node in the Kubernetes cluster through the management IP address. *node1* indicates the node queried in **Step 2**.
- **Step 5** Clear the residual drive letters. For details, see **Solution or Workaround**.

----End

10.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created

Symptom

When a Pod is being created, the Pod is in the **ContainerCreating** state for a long time. Check the huawei-csi-node log (for details, see **10.1 Viewing Log Information**). No Pod creation information is recorded in the huawei-csi-node log. After the **kubectl get volumeattachment** command is executed, the name of the PV used by the Pod is not displayed in the **PV** column. After a long period of time (more than ten minutes), the Pod is normally created and the Pod status changes to **Running**.

Root Cause Analysis

The kube-controller-manager component of Kubernetes is abnormal.

Solution or Workaround

Contact container platform engineers to rectify the fault.

11 Appendix

- 11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3
- 11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x
- 11.3 Example ALUA Configuration Policy of Distributed Storage

11.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3

Example 1: The configuration file content is as follows:

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

According to the configuration policy rules in **7.4.2.1 Configuring ALUA for OceanStor V3/V5 and OceanStor Dorado V3**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. A matches the beginning of a character string, and \$ matches the end of a character string.

11.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x

Example 1: The configuration file content is as follows:

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2.2 Configuring ALUA for OceanStor Dorado 6.x**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2.2 Configuring ALUA for OceanStor Dorado 6.x**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

According to the configuration policy rules in **7.4.2.2 Configuring ALUA for OceanStor Dorado 6.x**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.

11.3 Example ALUA Configuration Policy of Distributed Storage

Example 1: The configuration file content is as follows:

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2.3 Configuring ALUA for Distributed Storage**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **7.4.2.3 Configuring ALUA for Distributed Storage**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

According to the configuration policy rules in **7.4.2.3 Configuring ALUA for Distributed Storage**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.