eSDK Huawei Storage Kubernetes CSI Plugins V4.2.1

User Guide

Issue 01

Date 2023-10-27





Copyright © Huawei Technologies Co., Ltd. 2023. All rights reserved.

No part of this document may be reproduced or transmitted in any form or by any means without prior written consent of Huawei Technologies Co., Ltd.

Trademarks and Permissions

HUAWEI and other Huawei trademarks are trademarks of Huawei Technologies Co., Ltd. All other trademarks and trade names mentioned in this document are the property of their respective holders.

Notice

The purchased products, services and features are stipulated by the contract made between Huawei and the customer. All or part of the products, services and features described in this document may not be within the purchase scope or the usage scope. Unless otherwise specified in the contract, all statements, information, and recommendations in this document are provided "AS IS" without warranties, guarantees or representations of any kind, either express or implied.

The information in this document is subject to change without notice. Every effort has been made in the preparation of this document to ensure accuracy of the contents, but all statements, information, and recommendations in this document do not constitute a warranty of any kind, express or implied.

Huawei Technologies Co., Ltd.

Address: Huawei Industrial Base

Bantian, Longgang Shenzhen 518129

People's Republic of China

Website: https://e.huawei.com

Security Declaration

Product Life Cycle

Huawei's regulations on product life cycle are subject to the Product End of Life Policy. For details about the policy, see the following website: https://support.huawei.com/ecolumnsweb/en/warranty-policy

Vulnerability

Huawei's regulations on product vulnerability management are subject to "Vul. Response Process". For details about the policy, see the following website: https://www.huawei.com/en/psirt/vul-response-process

For enterprise customers who need to obtain vulnerability information, visit: https://securitybulletin.huawei.com/enterprise/en/security-advisory

Preconfigured Digital Certificate

Huawei has released the Huawei Preset Digital Certificate Disclaimer for the preconfigured digital certificates delivered with devices. For details about the disclaimer, visit the following website:https://support.huawei.com/enterprise/en/bulletins-service/ENEWS2000015789

Life Cycle of Product Documentation

Huawei released the Huawei Product Documentation Lifecycle Policy for after-sales customer documentation. For details about this policy, see the website of Huawei's official website.: https://support.huawei.com/enterprise/en/bulletins-website/ENEWS2000017761

About This Document

Intended Audience

This document is intended for:

- Technical support engineers
- O&M engineers
- Engineers with basic knowledge of storage and Kubernetes

Symbol Conventions

The symbols that may be found in this document are defined as follows.

| Symbol | Description | |
|------------------|---|--|
| ▲ DANGER | Indicates a hazard with a high level of risk which, if not avoided, will result in death or serious injury. | |
| ⚠ WARNING | Indicates a hazard with a medium level of risk which, if not avoided, could result in death or serious injury. | |
| ⚠ CAUTION | Indicates a hazard with a low level of risk which, if not avoided, could result in minor or moderate injury. | |
| NOTICE | Indicates a potentially hazardous situation which, if not avoided, could result in equipment damage, data loss, performance deterioration, or unanticipated results. NOTICE is used to address practices not related to personal injury. | |
| ◯ NOTE | Supplements the important information in the main text. NOTE is used to address information not related to personal injury, equipment damage, and environment deterioration. | |

Change History

| Issue | Date | Description |
|-------|------------|---|
| 01 | 2023-10-27 | This issue is the first official release. |

Contents

| About This Document | iii |
|---|-----|
| 1 Overview | 1 |
| 2 Compatibility and Features | 3 |
| 2.1 Kubernetes and OS Compatibility | |
| 2.2 Kubernetes Feature Matrix | 5 |
| 2.3 Compatibility with Huawei Enterprise Storage | 6 |
| 2.4 Compatibility with Huawei Distributed Storage | g |
| 3 Installation Preparations | 12 |
| 3.1 Prerequisites | 12 |
| 3.2 Downloading Huawei CSI Software Package | 13 |
| 3.3 Uploading a Huawei CSI Image | 14 |
| 3.3.1 Uploading an Image to the Image Repository | 14 |
| 3.3.2 Importing an Image to All Nodes | 15 |
| 3.4 Checking the Images on Which CSI Depends | 16 |
| 3.5 Checking Volume Snapshot-Dependent Components | 17 |
| 3.6 Checking the Host Multipathing Configuration | 19 |
| 3.7 Checking the Accounts on Huawei Storage | 20 |
| 3.8 Checking the Status of Host-Dependent Software | |
| 3.9 Communication Matrix | 21 |
| 4 Installation and Deployment | 22 |
| 4.1 Installing Huawei CSI Using Helm | 22 |
| 4.1.1 Preparing the values.yaml File | 22 |
| 4.1.1.1 images Parameters | 22 |
| 4.1.1.2 controller Parameters | 23 |
| 4.1.1.3 node Parameters | 26 |
| 4.1.1.4 csiDriver Parameters | 28 |
| 4.1.1.5 Other Parameters | 35 |
| 4.1.2 Installing Huawei CSI | 39 |
| 4.1.2.1 Installing Huawei CSI on Kubernetes, OpenShift, and Tanzu | |
| 4.1.2.2 Installing Huawei CSI on CCE/CCE Agile | |
| 4.1.2.2.1 Creating a Helm Installation Package | |
| 4.1.2.2.2 Installing Huawei CSI | 44 |

| 4.2 Manually Installing Huawei CSI | 46 |
|---|----|
| 5 Uninstalling Huawei CSI | 48 |
| 5.1 Uninstalling Huawei CSI Using Helm | 48 |
| 5.1.1 Uninstalling Huawei CSI on Kubernetes, OpenShift, and Tanzu | 48 |
| 5.1.2 Uninstalling Huawei CSI on CCE/CCE Agile | 49 |
| 5.1.3 Deleting CSI-Dependent Component Services | 49 |
| 5.1.3.1 Deleting the huawei-csi-host-info Object | 50 |
| 5.1.3.2 Deleting a Webhook Resource | 50 |
| 5.1.3.3 Uninstalling the Snapshot-Dependent Component Service | 50 |
| 5.2 Manually Uninstalling Huawei CSI | 51 |
| 5.2.1 Uninstalling the huawei-csi-node Service | 51 |
| 5.2.2 Uninstalling the huawei-csi-controller Service | 52 |
| 5.2.3 Uninstalling the csidriver Object | 52 |
| 5.2.4 Deleting the RBAC Permission | 52 |
| 5.2.5 Uninstalling Other Resources | 53 |
| 6 Upgrade/Rollback Operations | 54 |
| 6.1 Upgrading or Rolling Back Huawei CSI Using Helm | 54 |
| 6.1.1 Upgrading Huawei CSI | 54 |
| 6.1.1.1 Upgrading Huawei CSI on Kubernetes, OpenShift, and Tanzu | 54 |
| 6.1.2 Rolling Back Huawei CSI | 55 |
| 6.1.2.1 Rolling Back Huawei CSI on Kubernetes, OpenShift, and Tanzu | 55 |
| 6.2 Manual Upgrade/Rollback | 56 |
| 6.2.1 Upgrading Huawei CSI | 56 |
| 6.2.2 Rolling Back Huawei CSI | 57 |
| 7 Storage Backend Management | 58 |
| 7.1 Adding a Storage Backend | 58 |
| 7.1.1 Preparing the Storage Backend Configuration File | 59 |
| 7.1.2 Creating a Storage Backend | |
| 7.2 Querying a Storage Backend | 71 |
| 7.3 Updating a Storage Backend | |
| 7.4 Deleting a Storage Backend | 72 |
| 7.5 (Optional) Adding a Certificate to a Storage Backend | 73 |
| 7.5.1 Creating a Certificate for a Storage Backend | |
| 7.6 (Optional) Querying a Storage Backend Certificate | 74 |
| 7.7 (Optional) Updating a Storage Backend Certificate | 74 |
| 7.8 (Optional) Deleting a Storage Backend Certificate | 75 |
| 8 Using Huawei CSI | 77 |
| 8.1 Managing a PV/PVC | 77 |
| 8.1.1 Creating a PVC | 78 |
| 8.1.1.1 Dynamic Volume Provisioning | 78 |
| 8.1.1.1.1 Configuring a StorageClass | 78 |

| 8.1.1.1.2 Configuring a PVC | 98 |
|--|----------|
| 8.1.1.2 Manage Volume Provisioning | 102 |
| 8.1.1.2.1 Configuring a StorageClass | 103 |
| 8.1.1.2.2 Configuring a PVC | 114 |
| 8.1.1.3 Static Volume Provisioning | 121 |
| 8.1.1.3.1 Configuring a PV | 121 |
| 8.1.1.3.2 Configuring a PVC | 128 |
| 8.1.2 Expanding the Capacity of a PVC | 132 |
| 8.1.3 Cloning a PVC | 134 |
| 8.1.4 Creating a PVC Using a Snapshot | 134 |
| 8.2 Creating a VolumeSnapshot | 135 |
| 8.2.1 Checking Information About Volume Snapshot-dependent Components | 136 |
| 8.2.2 Configuring a VolumeSnapshotClass | 136 |
| 8.2.3 Configuring a VolumeSnapshot | 137 |
| 9 Advanced Features | 140 |
| 9.1 Configuring ALUA | |
| 9.1.1 Configuring ALUA Using Helm | 140 |
| 9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend | |
| 9.1.1.2 Configuring ALUA Parameters for a Distributed Storage Backend | 144 |
| 9.2 Configuring Storage Topology Awareness | 146 |
| 9.2.1 Configuring Storage Topology Awareness Using Helm | 147 |
| 10 Common Operations | 151 |
| 10.1 Collecting Logs | |
| 10.2 Updating the huawei-csi-controller Service | |
| 10.3 Updating the huawei-csi-node Service | 155 |
| 10.4 Modifying the Log Output Mode | 157 |
| 10.5 Enabling the ReadWriteOncePod Feature Gate | |
| 10.6 Configuring Access to the Kubernetes Cluster as a Non-root User | 160 |
| 11 FAQ | 162 |
| 11.1 How Do I View Huawei CSI Logs? | |
| 11.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Ku Platform Is Set Up for the First Time | bernetes |
| 11.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a dire | ectory" |
| 11.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete Source Host Where the Pod Resides Has Residual Drive Letters | but the |
| 11.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName | |
| 11.6 When a PVC Is Created, the PVC Is in the Pending State | |
| 11.7 Before a PVC Is Deleted, the PVC Is in the Pending StateState | |
| 11.8 When a Pod Is Created, the Pod Is in the ContainerCreating StateState | |
| 11.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created | |

| 11.10 A Pod Fails to Be Created and the Log Shows That the Execution of the mount Command Time Out | |
|---|-------|
| 11.11 A Pod Fails to Be Created and the Log Shows That the mount Command Fails to Be Executed | 174 |
| 11.12 After a Pod Fails to Be Created or kubelet Is Restarted, Logs Show That the Mount Point Alrea | dy |
| Exists | 175 |
| 11.13 How Do I Download a Container Image to the Local PC? | 175 |
| 11.14 How Do I Obtain CSI Version Information? | 176 |
| 11.15 Common Problems and Solutions for Interconnecting with the Tanzu Kubernetes Cluster | 177 |
| 11.15.1 A Pod Cannot Be Created Because the PSP Permission Is Not Created | 177 |
| 11.15.2 Changing the Mount Point of a Host | 178 |
| 11.15.3 Changing the Default Port of the livenessprobe Container | 178 |
| 11.16 Common Problems and Solutions for Using the Tanzu Kubernetes Cluster | 179 |
| 11.16.1 Failed to Create an Ephemeral Volume | 179 |
| 11.17 Failed to Expand the Capacity of a Generic Ephemeral Volume | 180 |
| 11.18 Failed to Expand the PVC Capacity Because the Target Capacity Exceeds the Storage Pool Capa | icity |
| | |
| 11.19 A webhook Fails to Be Called When the oceanctl Tool Is Used to Manage Backends | 181 |
| 11.20 An Account Is Locked After the Password Is Updated on the Storage Device | 181 |
| 12 Appendix | . 183 |
| 12.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3 | 183 |
| 12.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x | 184 |
| 12.3 Example ALUA Configuration Policy of Distributed Storage | 185 |
| 12.4 Installing Helm 3 | |
| 12.5 Configuring Custom Permissions | 187 |

1 Overview

Container Storage Interface (CSI) is an industry standard used to expose block and file storage systems to container workloads on container orchestration systems (COs) such as Kubernetes. Huawei CSI plug-in is used to communicate with Huawei enterprise storage and distributed storage products and provide storage services for Kubernetes container workloads. It is a mandatory plug-in used by Huawei enterprise storage and distributed storage in the Kubernetes environment.

Kubernetes uses a series of officially maintained sidecar components to register and listen to Kubernetes object resources and call CSI Driver when necessary. Huawei CSI Driver implements the call initiated by sidecar on Huawei storage, for example, creating a **Persistent Volume (PV)** is to create a LUN or file system on Huawei storage. The following figure shows the overall structure of Kubernetes, Huawei CSI, and Huawei storage.

Kubernetes Cluster Kubernetes core External component **Huawei CSI** (K8s CSI Team) Master Driver registrar CSI Identity External provisioner Node Node External attacher CSI Controller External resizer Kubelet Kubelet CSI Node External snapshotter Huawei Storage

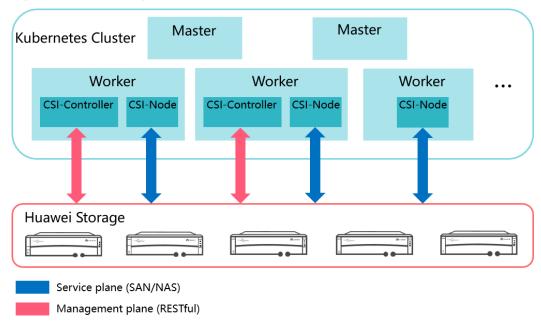
Figure 1-1 CSI overall architecture

Huawei CSI consists of two components: CSI Controller and CSI Node.

- CSI Controller: one or more Pods running in Deployment mode. It is used to interact with Huawei storage using RESTful. Therefore, the node running the CSI Controller component must be connected to the management plane network of the storage.
- CSI Node: a Pod that runs on Kubernetes worker nodes in DaemonSet mode. It is used to mount and unmount a LUN/file system provided by Huawei storage on worker nodes. Therefore, the node running the CSI Node component must be connected to the service plane network of the storage.

The following figure shows the deployment model of Huawei CSI.

Figure 1-2 CSI deployment model



This document describes how to install, deploy, and use the Huawei CSI V4.2.1 plug-in.

2 Compatibility and Features

This chapter describes the container management platforms, operating systems (OSs), and multipathing software supported by Huawei CSI plug-in, as well as the features and functions provided by the CSI plug-in when working with Huawei storage.

- 2.1 Kubernetes and OS Compatibility
- 2.2 Kubernetes Feature Matrix
- 2.3 Compatibility with Huawei Enterprise Storage
- 2.4 Compatibility with Huawei Distributed Storage

2.1 Kubernetes and OS Compatibility

Huawei CSI plug-in supports the following container management platforms.

Table 2-1 Supported container management platforms

| Container Management Platform | Version |
|--------------------------------------|---|
| Kubernetes | 1.16, 1.18 to 1.27 |
| Red Hat OpenShift Container Platform | 4.6 EUS, 4.7, 4.8, 4.9, 4.10, 4.11, 4.12, 4.13 |
| Tanzu Kubernetes | TKGI 1.14.1 |
| CCE Agile | 22.3.2 |
| CCE | 22.9.5 |

NOTICE

- For details about how to connect Huawei CSI to Red Hat OpenShift, see eSDK
 Enterprise Storage Plugins User Guide (Kubernetes CSI for Red Hat OpenShift).
- The connection between Huawei CSI and Tanzu Kubernetes supports only the NAS scenario. For related FAQ, see 11.15 Common Problems and Solutions for Interconnecting with the Tanzu Kubernetes Cluster.

The following table lists the OSs and multipathing software supported by the Huawei CSI plug-in.

Table 2-2 Supported host OSs and multipathing software versions

| OS Name | OS Version | Native DM- Multipath Version | Huawei UltraPath version |
|-----------------------------|---|---|---|
| CentOS x86_64 | 7.6, 7.7, 7.9 | Delivered with the OS, supporting FC/iSCSI | UltraPath 31.1.0, supporting FC/iSCSI |
| CentOS x86_64 | 8.2, 8.4 | Delivered with the OS, supporting FC/iSCSI | UltraPath 31.1.0, supporting FC/iSCSI UltraPath-NVMe 31.1.RC8, supporting NVMe over RoCE/NVMe over FC |
| CentOS ARM | 7.6 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| Rocky Linux x86_64 | 8.6 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| SUSE 15 x86_64 | SP2, SP3 | Delivered with the OS, supporting FC/ iSCSI | UltraPath 31.1.0, supporting FC/iSCSI UltraPath-NVMe 31.1.RC8, supporting NVMe over RoCE/NVMe over FC |
| Red Hat CoreOS x86_64 | 4.6, 4.7, 4.8, 4.9, 4.10, 4.11, 4.12, 4.13 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| Ubuntu x86_64 | 18.04, 20.04, 22.04 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| Kylin x86_64 | V10 SP1, V10 SP2, 7.6 | Delivered with the OS, supporting FC/iSCSI | UltraPath 31.2.0, supporting FC/iSCSI ¹ |

| OS Name | OS Version | Native DM- Multipath Version | Huawei UltraPath version |
|-------------------|--------------------------|--|--------------------------|
| Kylin ARM | V10 SP1, V10 SP2 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| Debian x86_64 | 9, 11 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| EulerOS x86_64 | V2R9, V2R10, V2R11 | Delivered with the OS, supporting FC/iSCSI | Not supported |
| EulerOS ARM | V2R10 | Delivered with the OS, supporting FC/iSCSI | Not supported |

Note 1: Only Kylin x86_64 V10 SP2 supports UltraPath 31.2.0.

□ NOTE

For DM-Multipath 0.7, some virtual devices may not be displayed in the command output after the **multipathd show maps** command is executed. Therefore, you are advised to use version 0.8 or later.

You can query the DM-Multipath version in either of the following ways:

- If the rpm package is used, run the **rpm -qa | grep device-mapper** command.
- If the deb package is used, run the **dpkg -l | grep multipath** command.

2.2 Kubernetes Feature Matrix

This section describes the features of different Kubernetes versions supported by Huawei CSI.

Table 2-3 Kubernetes versions and supported features

| Feature | V1.16 | V1.18 | V1.19 | V1.20 | V1.21+ |
|-------------------------------------|-------|-------|-------|-------|----------|
| Static Provisioning | √ | √ | √ | √ | ✓ |
| Dynamic Provisioning | √ | √ | √ | √ | √ |
| Manage Provisioning ¹ | √ | √ | √ | √ | √ |
| Expand Persistent Volume | √ | √ | √ | √ | √ |
| Create VolumeSnapshot | х | √ | √ | √ | √ |

| Feature | V1.16 | V1.18 | V1.19 | V1.20 | V1.21+ |
|---------------------------------|-------|-------|-------|-------|----------|
| Restore VolumeSnapshot | x | √ | √ | √ | √ |
| Delete VolumeSnapshot | x | √ | √ | √ | √ |
| Clone Persistent Volume | х | √ | √ | √ | √ |
| Raw Block Volume | √ | √ | √ | √ | √ |
| Topology | √ | √ | √ | √ | √ |
| Generic Ephemeral Volumes | × | × | × | × | √ |

• Note 1: Manage Provisioning is a volume management feature customized by Huawei CSI. This feature allows existing storage resources to be managed by Kubernetes. A storage resource cannot be managed for multiple times.

2.3 Compatibility with Huawei Enterprise Storage

Huawei CSI plug-in is compatible with Huawei OceanStor series all-flash storage and hybrid flash storage. The following table lists the supported storage versions.

Table 2-4 Supported Huawei enterprise storage

| Storage Product | Version |
|---------------------|-----------------------------------|
| OceanStor V3 | V300R006 |
| OceanStor V5 | V500R007, V500R007 Kunpeng |
| OceanStor Dorado V3 | V300R002 |
| OceanStor V6 | 6.1.3, 6.1.5, 6.1.6 |
| OceanStor Dorado V6 | 6.1.0, 6.1.2, 6.1.3, 6.1.5, 6.1.6 |

Huawei CSI plug-in supports the following features for Huawei enterprise storage.

Table 2-5 Features supported by Huawei enterprise storage and constraints

| Feature | OceanSt or V3 | OceanStor V5 | OceanStor Dorado V3 | OceanStor V6 | OceanStor Dorado V6 |
|---|---|--|--|---|--|
| Static Provisionin g | SAN: FC/ iSCSI ² NAS: NFS | SAN: FC/ iSCSI ² NAS: NFS | SAN: FC/ iSCSI ² | SAN: FC/ iSCSI/NVMe over RoCE/ NVMe over | SAN: FC/ iSCSI/NVMe over RoCE/ NVMe over |
| Dynamic Provisionin g | 3 | 3 | | FC ³ NAS: NFS 3/4.0/4.1 | FC ³ NAS: NFS 3/4.0/4.1 ⁴ |
| Manage Provisionin g ¹ | | | | 3,4.0,4.1 | 3/4.0/4.1 |
| Expand Persistent Volume | Volumes cr mode are s | | amic Provisior | ing or Manage | Provisioning |
| Create VolumeSna pshot | Volumes created in Dynamic Provisioning or Manage Provisioning mode are supported. | | | | e Provisioning |
| Delete VolumeSna pshot | Supporte d | Supported | Supported | Supported | Supported |
| Restore VolumeSna pshot | Supporte d | Supported | Supported | SAN: supported NAS: supported only in 6.1.5 and later versions | SAN: supported NAS: supported only in 6.1.5 and later versions |
| Clone Persistent Volume | Non-HyperMetro volumes created in Dynamic Provisioning or Manage Provisioning mode are supported. | | | SAN: supports HyperMetro v created in Dy Provisioning of Provisioning of NAS: Only 6.1 versions supp HyperMetro v created in Dy Provisioning of Provisioning of | volumes namic or Manage mode. I.5 and later ort non- volumes namic or Manage |
| Raw Block Volume | Only SAN volumes are supporte d. | Only SAN volumes are supported. | Only SAN volumes are supported. | Only SAN volumes are supported. | Only SAN volumes are supported. |

| Feature | OceanSt or V3 | OceanStor V5 | OceanStor Dorado V3 | OceanStor V6 | OceanStor Dorado V6 |
|---------------------------------|--|------------------|------------------------|--|------------------------|
| Topology | Supporte d | Supported | Supported | Supported | Supported |
| Generic Ephemeral Volumes | Supporte d | Supported | Supported | Supported | Supported |
| Access Mode | RWO/ROX/RWOP: supported by all types of volumes. RWOP is supported only by Kubernetes 1.22 and later versions. RWX: supported only by Raw Block volumes and NFS volumes. | | | | S. |
| QoS | Supporte d ⁵ | Supported 5 | Supported | Supported | Supported |
| Application type | N/A | N/A | N/A | Supported | Supported |
| Volume HyperMetr o | Not supporte d | Not supported | N/A | Only NAS volumes are supported. | |
| Storage multi- tenant | Only NAS volumes are supported. | | N/A | Only NAS volumes are supported. ⁶ | |

- Note 1: Manage Provisioning is a volume management feature customized by Huawei CSI. This feature allows existing storage resources to be managed by Kubernetes. A storage resource cannot be managed for multiple times.
- Note 2: If the user's container platform is deployed in a virtualization environment, only iSCSI networking is supported.
- Note 3: If NVMe over RoCE or NVMe over FC is used, the version of the nvmecli tool on worker nodes must be 1.9 or later. To query the version, run the nvme version command.
- Note 4: Only OceanStor Dorado V6 6.1.0 and later versions support NFS. Only OceanStor Dorado V6 6.1.3 and later versions support NFS 4.1.
- Note 5: Only system users can configure QoS.
- Note 6: Only OceanStor Dorado V6 6.1.3 and later versions support multitenant.

Huawei CSI plug-in supports the following Dtree features for Huawei enterprise storage.

Table 2-6 Features supported by Dtree

| Feature | Supported |
|---------------------|--------------|
| Static Provisioning | \checkmark |

| Feature | Supported |
|--------------------------|---|
| Dynamic Provisioning | √ |
| Expand Persistent Volume | √ |
| Access Mode | √ (RWX/RWO/ROX/RWOP: Kubernetes 1.22 or later supports RWOP.) |
| Multi-tenancy | √ |
| Create VolumeSnapshot | X |
| Delete VolumeSnapshot | X |
| Restore VolumeSnapshot | X |
| Clone Persistent Volume | X |
| QoS | X |
| Volume HyperMetro | X |
| Application type | X |

Table 2-7 Huawei storage versions supported by Dtree

| Storage Product | Version | |
|---------------------|-----------------------------------|--|
| OceanStor Dorado V6 | 6.1.0, 6.1.2, 6.1.3, 6.1.5, 6.1.6 | |
| OceanStor Dorado V3 | Not supported | |

2.4 Compatibility with Huawei Distributed Storage

Huawei CSI plug-in is compatible with Huawei OceanStor series distributed storage systems. The following table lists the supported storage versions.

Table 2-8 Supported Huawei distributed storage

| Storage Product | Version | |
|--------------------------|-----------------------------------|--|
| FusionStorage Block | 8.0.0, 8.0.1 | |
| OceanStor Pacific series | 8.1.0, 8.1.1, 8.1.2, 8.1.3, 8.1.5 | |

Huawei CSI plug-in supports the following features for Huawei distributed storage.

Table 2-9 Features supported by Huawei distributed storage and constraints

| Feature | FusionStorage Block | OceanStor Pacific Series |
|-------------------------------------|--|---|
| Static Provisioning | SAN: iSCSI/SCSI | SAN: iSCSI/SCSI |
| Dynamic | | NAS: DPC ² /NFS 3/4.1 ³ |
| Provisioning | | |
| Manage Provisioning ¹ | | |
| Expand Persistent Volume | Volumes created in Dynamic I Provisioning mode are suppor | 5 5 |
| Create VolumeSnapsho t | SAN volumes created in Dyna Provisioning mode are suppor | · · · · · · · · · · · · · · · · · · · |
| Delete VolumeSnapsho t | Supported | Only SAN volume snapshots are supported. |
| Restore VolumeSnapsho t | Supported | Only SAN volume snapshots are supported. |
| Clone Persistent Volume | SAN volumes created in Dyna Provisioning mode are suppor | · · · · · · · · · · · · · · · · · · · |
| Raw Block Volume | Only SAN volumes are supported. | Only SAN volumes are supported. |
| Topology | Supported | Supported |
| Generic Ephemeral Volumes | Supported | Supported |
| Access Mode | RWO/ROX/RWOP: supported by all types of volumes. RWOP is supported only by Kubernetes 1.22 and later versions. RWX: supported only by Raw Block volumes and NFS volumes. | |
| QoS | Supported | Supported |
| Soft and hard quotas | Not supported | Only NAS volumes are supported. |
| Storage multi- tenant | Not supported | Only NAS volumes are supported. |

- Note 1: Manage Provisioning is a volume management feature customized by Huawei CSI. This feature allows existing storage resources to be managed by Kubernetes. A storage resource cannot be managed for multiple times.
- Note 2: Only OceanStor Pacific series 8.1.2 and later versions support DPC. For details about whether the OSs supported by Huawei CSI support DPC, see the compatibility document of the corresponding product version.
- Note 3: Only OceanStor Pacific series 8.1.2 and later versions support NFS 4.1.

3 Installation Preparations

This chapter describes the preparations for the installation.

- 3.1 Prerequisites
- 3.2 Downloading Huawei CSI Software Package
- 3.3 Uploading a Huawei CSI Image
- 3.4 Checking the Images on Which CSI Depends
- 3.5 Checking Volume Snapshot-Dependent Components
- 3.6 Checking the Host Multipathing Configuration
- 3.7 Checking the Accounts on Huawei Storage
- 3.8 Checking the Status of Host-Dependent Software
- 3.9 Communication Matrix

3.1 Prerequisites

Before performing the operations described in this chapter, ensure that the following conditions are met:

- A container management platform has been deployed and is running properly, and its compatibility meets the requirements described in 2.1 Kubernetes and OS Compatibility.
- (Mandatory for enterprise storage) Initial configuration for interconnecting
 with Huawei enterprise storage has been completed, including storage pool
 division and port configuration. The version of the storage product meets the
 requirements in 2.3 Compatibility with Huawei Enterprise Storage.
- (Mandatory for distributed storage) Initial configuration for interconnecting
 with Huawei distributed storage has been completed, including storage pool
 division and port configuration. The version of the storage product meets the
 requirements in 2.4 Compatibility with Huawei Distributed Storage.
- If a multipathing network is used, ensure that multipathing software has been installed on all worker nodes. For details, see **Table 2-2**.
- The connectivity between Huawei storage and the container platform host has been configured. For example, the worker node running huawei-csi-

controller communicates properly with the management IP address of the storage device to be connected, and the worker node running huawei-csi-node communicates properly with the service IP address of the storage device to be connected. In iSCSI scenarios, the **ping** command can be used to verify the connectivity.

- Software clients required by the corresponding protocol, such as iSCSI and NFS clients, have been installed on all worker nodes of the container cluster.
- Ensure that the language of the operating system is English.

3.2 Downloading Huawei CSI Software Package

This section describes how to download the software package and the component structure of the software package.

- **Step 1** Open a browser and enter https://github.com/Huawei/eSDK_K8S_Plugin/releases in the address box.
- **Step 2** Download the software package of the 4.2.1 version based on the storage type and CPU architecture.

□ NOTE

Software package naming rule: Plug-in name (eSDK_Huawei_Storage_Kubernetes_CSI_Plugin) + Version number + CPU architecture

Step 3 Decompress the downloaded software package. The following table shows the component structure of the software package.

Table 3-1 Component description

| Component | Description | | |
|---|---|--|--|
| image/huawei-csi-v4.2.1- arch.tar | huawei-csi-driver image. <i>arch</i> is x86 or arm . | | |
| image/storage-backend- controller-v4.2.1- <i>arch</i> .tar | Back-end management controller image. <i>arch</i> is x86 or arm . | | |
| image/storage-backend- sidecar-v4.2.1- <i>arch</i> .tar | Back-end management sidecar image. <i>arch</i> is x86 or arm . | | |
| bin/ | Binary file used by an image provided by Huawei. | | |
| bin/oceanctl | Command line tool provided by Huawei, which can be used to manage storage backends. | | |
| helm/ | Helm project used to deploy Huawei CSI. | | |
| manual/ | Used to manually install and deploy Huawei CSI. | | |
| examples/ | .yaml sample file used during CSI use. | | |
| examples/backend | .yaml sample file used to create a storage backend. | | |

| Component | Description |
|-----------|---|
| tools/ | Script used to upload images when no image repository is available. |

----End

3.3 Uploading a Huawei CSI Image

Huawei provides the **huawei-csi** image for users. For details about how to obtain the image file, see **3.2 Downloading Huawei CSI Software Package**.

To use the CSI image on the container management platform, you need to import the CSI image to the cluster in advance using either of the following methods:

- (Recommended) Use Docker to upload the CSI image to the image repository.
- Use the image upload script to import the CSI image to all nodes where Huawei CSI needs to be deployed.

3.3.1 Uploading an Image to the Image Repository

The installation of Huawei CSI depends on the following three image files provided by Huawei. Import and upload the image files in sequence. For details about how to obtain the image files, see **3.2 Downloading Huawei CSI Software Package**.

- huawei-csi-v4.2.1-arch.tar
- storage-backend-controller-v4.2.1-arch.tar
- storage-backend-sidecar-v4.2.1-*arch*.tar

Prerequisites

A Linux host with Docker installed is available, and the host can access the image repository.

Procedure

Step 1 Run the **docker load -i huawei-csi.tar** command to import the CSI image to the current node.

docker load -i huawei-csi.tar Loaded image: huawei-csi:4.2.1

Step 2 Run the docker tag huawei-csi:4.2.1 repo.huawei.com/huawei-csi:4.2.1 command to add the image repository address to the image tag. repo.huawei.com indicates the image repository address.

docker tag huawei-csi:4.2.1 repo.huawei.com/huawei-csi:4.2.1

Step 3 Run the docker push repo.huawei.com/huawei-csi:4.2.1 command to upload the CSI image to the image repository. repo.huawei.com indicates the image repository address.

docker push repo.huawei.com/huawei-csi:4.2.1

----End

NOTICE

- You can also use containerd to import and upload the images.
- For details about how to import and upload images to the CCE/CCE Agile platform, see the user manual of the platform.

3.3.2 Importing an Image to All Nodes

If the image has been uploaded to the image repository, skip this section.

Prerequisites

- The host where the image is located can communicate with all hosts to which the image is to be imported using SSH.
- The **expect**, **sshpass**, and **scp** software packages have been installed on the host where the image is located.

Procedure

Step 1 Run the **vi** *worker-list.txt* command to create and open the **worker-list.txt** configuration file.

vi worker-list.txt

Step 2 Configure the **worker-list.txt** file. The template of the **worker-list.txt** file is as follows. Press **I** or **Insert** to enter the editing mode and add node information. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

ip 192.168.128.16 192.168.128.17

- **Step 3** Upload and import an image.
 - If containerd is used as the container runtime, run the ./containerd-upload.sh worker-list.txt huawei-csi.tar command, enter the user name and password as prompted, and upload and import an image.
 # ./containerd-upload.sh worker-list.txt huawei-csi.tar
 - If Docker is used as the container runtime, run the ./docker-upload.sh worker-list.txt huawei-csi.tar command, enter the user name and password as prompted, and upload and import an image.

 # ./docker-upload.sh worker-list.txt huawei-csi.tar
- **Step 4** Check whether the image is successfully imported.
 - 1. If **All images are uploaded successfully** is displayed at the end of the script, the image has been successfully imported to all nodes.

 All images are uploaded successfully
 - 2. If the following information is displayed at the end of the script, the image fails to be imported to the nodes in the list.

List of nodes to which the image fails to be imported: 192.168.128.16 192.168.128.17

----End

3.4 Checking the Images on Which CSI Depends

The installation of Huawei CSI depends on the images listed in the following table. If all worker nodes in the cluster have been connected to the Internet and can pull images online, skip this section. If nodes in the cluster cannot connect to the Internet, download the corresponding image file based on the Kubernetes version and upload it to the image repository or import it to all worker nodes in the Kubernetes cluster.

The huawei-csi-controller service depends on the following sidecar images: livenessprobe, csi-provisioner, csi-attacher, csi-resizer, csi-snapshotter, and snapshot-controller.

The huawei-csi-node service depends on the following sidecar images: livenessprobe, csi-node-driver-registrar, and huawei-csi-driver.

For details about the functions and details of each image, see the following table.

Table 3-2 Images on which Huawei CSI depends

| Container Name | Container Image | K8s Version Requirement s | Feature Description |
|-----------------------------------|---|---------------------------------|---|
| livenesspro be | k8s.gcr.io/sig- storage/ livenessprobe:v2.5. 0 | v1.13+ | Monitors the health status of CSI and reports it to Kubernetes so that Kubernetes can automatically detect CSI program problems and restart the Pod to rectify the problems. |
| csi-resizer | k8s.gcr.io/sig- storage/csi- resizer:v1.4.0 | v1.13+ | Calls CSI to provide more storage space for a PVC when expanding the capacity of the PVC. |
| csi-node- driver- registrar | k8s.gcr.io/sig- storage/csi-node- driver- registrar:v2.3.0 | v1.13+ | Obtains CSI information and registers a node with kubelet using the plug-in registration mechanism of kubelet so that Kubernetes can detect the connection between the node and Huawei storage. |
| csi- snapshotter | k8s.gcr.io/sig- storage/csi- snapshotter:v4.2.1 | v1.17+ | Calls CSI to create or delete a snapshot on the storage system when creating or deleting a VolumeSnapshot. |

| Container Name | Container Image | K8s Version Requirement s | Feature Description |
|-------------------------|---|---------------------------------|--|
| snapshot- controller | k8s.gcr.io/sig- storage/snapshot- controller:v4.2.1 | v1.17+ | Listens to the VolumeSnapshot and VolumeSnapshotContent objects in the Kubernetes API and triggers csi- snapshotter to create a snapshot on the storage system when creating or deleting a VolumeSnapshot. |
| csi- provisioner | k8s.gcr.io/sig- storage/csi- provisioner:v3.0.0 | v1.17+ | Calls the CSI Controller service to create a LUN or file system on the |
| | quay.io/k8scsi/csi- provisioner:v1.4.0 | v1.16.x | storage system as a PV and bind the PV to a PVC when creating a PVC. |
| | | | Calls the CSI Controller service to unbind a PV from a PVC and delete the LUN or file system corresponding to the PV when deleting a PVC. |
| csi-attacher | k8s.gcr.io/sig- storage/csi- attacher:v3.4.0 | v1.17+ | Calls the CSI Controller service to perform the "Publish/Unpublish Volume" |
| | quay.io/k8scsi/csi- attacher:v1.2.1 | v.1.16.x | operation when creating or deleting a Pod. |

■ NOTE

For details about how to download container images to the local host, see 11.13 How Do I Download a Container Image to the Local PC?.

3.5 Checking Volume Snapshot-Dependent Components

This section describes how to check the volume snapshot-dependent components in the cluster.



Kubernetes earlier than v1.17.0 does not support the snapshot function. If the snapshot CRD is deployed, the cluster may be faulty. Therefore, if Huawei CSI is deployed on Kubernetes earlier than v1.17.0, perform the check according to **Kubernetes Earlier Than v1.17.0**.

Kubernetes Earlier Than v1.17.0

If the Kubernetes version is earlier than v1.17.0, the cluster may be faulty during snapshot deployment. Perform the following steps to delete the snapshot CRD installation file.

Step 1 Run the following command to check the Kubernetes version. In the following example, the Kubernetes version is v1.16.0.

```
# kubectl get node
NAME STATUS ROLES AGE VERSION
test-master Ready master 311d v1.16.0
test-node Ready <none> 311d v1.16.0
```

Step 2 Go to the /helm/esdk/crds/snapshot-crds directory and run the following command to delete the snapshot CRD installation file. For details about the component package path, see Table 3-1.

rm -rf ./huawei-csi-snapshot-crd-v1.yaml

----End

Kubernetes v1.17.0 or Later

If you need to use volume snapshots and features associated with volume snapshots in the container environment, perform the following steps to check whether volume snapshot-dependent components have been deployed in your environment and check the api-versions information about volume snapshots.

Step 1 Run the following command to view the installation details of snapshot-related resource services.

```
# kubectl api-resources | grep snapshot | awk '{print $1}' volumesnapshotclasses volumesnapshotcontents volumesnapshots
```

- If the preceding command output is displayed, the snapshot-dependent component services have been installed. In this case, go to **Step 2** and check the api-versions information.
- If any of the services in the preceding command output is not displayed, go to **Step 3** and install the snapshot-dependent component service.
- **Step 2** Run the following command to query the api-versions information about volume snapshots.

```
# kubectl api-versions | grep "snapshot.storage.k8s.io" snapshot.storage.k8s.io/v1 snapshot.storage.k8s.io/v1beta1
```

- If the preceding command output is displayed, the snapshot-dependent component services support v1 and v1beta1. In this case, skip this section.
- If any of the services in the preceding command output is not displayed, go to the next step to install it.

Step 3 Go to the /helm/esdk/crds/snapshot-crds directory and run the following command to install the snapshot-dependent component services. For details about the component package path, see Table 3-1.

kubectl apply -f huawei-csi-snapshot-crd-v1.yaml --validate=false Warning: resource volumesnapshots/mysnapshot is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.

volumesnapshot.snapshot.storage.k8s.io/mysnapshot configured

----End

After the installation is complete, you can run the command in **Step 1** to check the installation details of snapshot-related resource services.

3.6 Checking the Host Multipathing Configuration

If you plan to use the FC/iSCSI/NVMe over RoCE/NVMe over FC protocol to access Huawei storage in a container environment, you are advised to use host multipathing software to enhance the link redundancy and performance of the host and storage. If you do not want to use the software, skip this section.

For details about the OSs and multipathing software supported by Huawei CSI, see **Table 2-2**.

□ NOTE

- If you want to use the FC/iSCSI protocol to connect to Huawei storage, you are advised to use native DM-Multipath provided by the OS.
- If you want to use the NVMe over RoCE/NVMe over FC protocol to connect to Huawei storage, you are advised to use Huawei-developed UltraPath-NVMe.
- If you want to use the SCSI protocol to connect to Huawei storage, disable DM-Multipath provided by the OS.

Prerequisites

Multipathing software has been correctly installed on a host.

- If you use native DM-Multipath provided by the OS, contact your host or OS provider to obtain the documents and software packages required for the installation.
- If you use Huawei-developed UltraPath or UltraPath-NVMe, contact Huawei engineers to obtain the UltraPath or UltraPath-NVMe documents and software packages. For details about the software package versions, see Table 2-2.

Procedure

- Step 1 If you use the iSCSI/FC protocol to connect to Huawei enterprise storage, configure and check host multipathing by referring to Configuring Multipathing > Non-HyperMetro Scenarios in OceanStor Dorado 6.x and OceanStor 6.x Host Connectivity Guide for Red Hat.
- **Step 2** If you use the NVMe over RoCE/NVMe over FC protocol to connect to Huawei enterprise storage, configure and check host multipathing by referring to

Configuring Multipathing > Non-HyperMetro Scenarios > UltraPath in OceanStor Dorado 6.x and OceanStor 6.x Host Connectivity Guide for Red Hat.

- **Step 3** If you use iSCSI to connect to Huawei distributed storage, configure and check host multipathing by referring to **Configuring Multipathing for an Application Server** in *FusionStorage 8.0.1 Block Storage Basic Service Configuration Guide*.
- **Step 4** If you use the native multipathing software provided by the OS, check whether the **/etc/multipath.conf** file contains the following configuration item.

```
defaults {
    user_friendly_names yes
    find_multipaths no
}
```

If the configuration item does not exist, add it to the beginning of the **/etc/multipath.conf** file.

□ NOTE

For details about the functions of the **user_friendly_names** and **find_multipaths** parameters, see **dm_multipath/config_file_defaults**.

----End

3.7 Checking the Accounts on Huawei Storage

After Huawei storage is connected to the container platform, Huawei CSI needs to manage storage resources on Huawei storage based on service requirements, such as creating and mapping volumes. In this case, Huawei CSI needs to use the accounts created on Huawei storage to communicate with Huawei storage. The following table lists the accounts required for different storage devices.

Table 3-3 Account requirements for connecting storage to CSI

| Storage Type | User Type | Role | Level | Туре |
|--|-------------|---|-------------------|------------|
| OceanStor V3/V5 | System user | Administrat or | Administrat or | Local user |
| | vStore user | vStore administrat or | Administrat or | Local user |
| OceanStor Dorado V3 | System user | Administrat or | Administrat or | Local user |
| OceanStor 6.1.3, 6.1.5, 6.1.6 | System user | Administrat or/User- defined role ¹ | N/A | Local user |
| OceanStor Dorado 6.1.0, 6.1.2, 6.1.3, 6.1.5, 6.1.6 | System user | Administrat or/User- defined role ¹ | N/A | Local user |

| Storage Type | User Type | Role | Level | Туре |
|--------------------------|-------------|-----------------------------|-------|------------|
| | vStore user | vStore administrat or | N/A | Local user |
| OceanStor Pacific series | System user | Administrat or | N/A | Local user |

 Note 1: If a user-defined role is used, you need to configure permissions for the role. For details about how to configure the minimum permissions, see 12.5 Configuring Custom Permissions.

3.8 Checking the Status of Host-Dependent Software

This section describes how to check whether the status of host-dependent software on worker nodes in a cluster is normal. In this example, the host OS is CentOS 7.9 x86_64.

- Check the status of the iSCSI client.
 # systemctl status iscsi iscsid
- Check the status of the NFS client. # systemctl status rpcbind
- Check the status of DM-Multipath. # systemctl status multipathd.socket multipathd
- Check the status of UltraPath.
 # systemctl status nxup
- Check the status of UltraPath-NVMe.
 # systemctl status upudev upService_plus

3.9 Communication Matrix

Contact Huawei engineers to obtain the document from Huawei support website. For details about enterprise storage, see eSDK Enterprise Storage Plugins 2.5.0. For details about distributed storage, see eSDK Cloud Storage Plugins 2.5.RC4.

4 Installation and Deployment

This chapter describes how to install Huawei CSI on the Kubernetes, OpenShift, Tanzu, CCE, and CCE Agile platforms.

4.1 Installing Huawei CSI Using Helm

4.2 Manually Installing Huawei CSI

4.1 Installing Huawei CSI Using Helm

4.1.1 Preparing the values.yaml File

When using Helm to install CSI, you need to prepare the **values.yaml** file of the Helm project based on the features required during deployment. Huawei CSI provides the **values.yaml** template file in the **helm/esdk** directory of the software package.

This section describes the configuration items in the **values.yaml** file and backend configuration examples in typical scenarios.

4.1.1.1 images Parameters

The images parameters in the **values.yaml** file are used to configure the component image information on which Huawei CSI depends during running. Set the following parameters:

Table 4-1 images parameters

| Parameter | Description | Mandatory | Default Value |
|--------------------------------------|---|-----------|-----------------------------------|
| images.huaw eiCSIService | huawei-csi image. | Yes | huawei-csi:4.2.1 |
| images.stora geBackendSi decar | Huawei back-end management sidecar image. | Yes | storage-backend- sidecar:4.2.1 |

| Parameter | Description | Mandatory | Default Value |
|---|--|-----------|---|
| images.stora geBackendC ontroller | Huawei back-end management controller image. | Yes | storage-backend- controller:4.2.1 |
| images.sidec ar.livenessPro be | livenessprobe sidecar image. | Yes | k8s.gcr.io/sig- storage/ livenessprobe:v2.5 |
| images.sidec ar.provisioner | csi-provisioner sidecar image. | Yes | k8s.gcr.io/sig- storage/csi- provisioner:v3.0.0 |
| images.sidec ar.attacher | csi-attacher sidecar image. | Yes | k8s.gcr.io/sig- storage/csi- attacher:v3.4.0 |
| images.sidec ar.resizer | csi-resizer sidecar image. | Yes | k8s.gcr.io/sig- storage/csi- resizer:v1.4.0 |
| images.sidec ar.snapshotte r | csi-snapshotter sidecar image. | Yes | k8s.gcr.io/sig- storage/csi- snapshotter:v4.2.1 |
| images.sidec ar.snapshotC ontroller | snapshot-controller sidecar image. | Yes | k8s.gcr.io/sig- storage/snapshot- controller:v4.2.1 |
| images.sidec ar.registrar | csi-node-driver-registrar sidecar image. | Yes | k8s.gcr.io/sig- storage/csi-node- driver- registrar:v2.3.0 |

NOTICE

- For details about the values of huaweiCSIService, storageBackendSidecar, and storageBackendController, see 3.3 Uploading a Huawei CSI Image. Use the name and version of the finally generated image.
- For details about other sidecar image parameters, see 3.4 Checking the Images on Which CSI Depends. Use the name and version of the finally uploaded image.

4.1.1.2 controller Parameters

The controller parameters are used to configure the huawei-csi-controller component.

Table 4-2 controller parameters

| Paramete r | Description | Mandator y | Default Value | Remarks |
|------------------------------------|--|---------------|------------------|---------|
| controller. controller Count | Number of huawei- csi-controller component copies. | Yes | 1 | - |

| Paramete r | Description | Mandator y | Default Value | Remarks |
|-------------------------------|---|---------------|------------------|--|
| controller. volumeNa mePrefix | PV name prefix. The default value is pvc , that is, the name of a created PV is pvc - <uuid>. The prefix must comply with the naming rules of a DNS subdomain name, and the total length of the PV name cannot exceed 253 characters. </uuid> | No | pvc | The corresponding provisioner parameter name isvolume-name-prefix. For details, see Configuring the PV Name Prefix. If the connected backend is OceanStor V3/V5 SAN storage, it is recommended that the prefix contain a maximum of 5 characters. If the connected backend is OceanStor V3/V5 NAS storage, the prefix can contain only lowercase letters, hyphens (-), and digits. If the connected backend is OceanStor Dorado V6 or converged V6 storage, the prefix can contain only lowercase letters, hyphens (-), and digits. If the connected backend is OceanStor Dorado V6 or converged V6 storage, the prefix can contain only lowercase letters, hyphens (-), and digits. If the connected backend is OceanStor Pacific series storage, the |

| Paramete r | Description | Mandator y | Default Value | Remarks |
|-------------------------------------|---|---------------|------------------|---|
| | | | | prefix can contain a maximum of 95 characters. |
| controller. webhookP ort | Port used by the webhook service. | Yes | 4433 | If a port conflict occurs, change the port number to one that does not conflict with the port number. |
| controller. snapshot. enabled | Whether to enable the snapshot feature. | Yes | true | If you want to use snapshot-related functions, enable this feature. The Kubernetes version must be |
| | | | | later than v1.17. |
| controller. resizer.ena bled | Whether to enable the capacity expansion feature. | Yes | true | The Kubernetes version must be later than v1.16. |
| controller. nodeSelec tor | Node selector of huawei-csi-controller. After this parameter is set, huawei-csi-controller will be scheduled only to a node with the label. | No | - | For details about the node selector, see Assign Pods to Nodes. |
| controller. toleration s | Taint toleration of huawei-csi-controller. After this parameter is set, huawei-csi-controller can tolerate taints on a node. | No | - | For details about taints and tolerations, see Taints and Tolerations. |

□ NOTE

If **controller.snapshot.enabled** is set to **true**, when the **helm install** command is executed, the system automatically reads the volume snapshot CRD in the **helm/crd** directory and installs the snapshot CRD resource.

4.1.1.3 node Parameters

The node parameters are used to configure the huawei-csi-node component.

Table 4-3 node parameters

| Parameter | Description | Mandato ry | Default Value | Remarks |
|----------------------------|--|---------------|---|---|
| node.maxVolum esPerNode | Maximum number of volumes provisioned by Huawei CSI that can be used by a node. If this parameter is not specified or is set to 0, the number is unlimited. If nodeName is specified during Pod creation, this configuration will be ignored. | No | 100 | For details, see Volume Limits. |
| node.nodeSelect or | Node selector of huawei-csi-node. After this parameter is set, huawei-csi-node will be scheduled only to a node with the label. | No | - | For details about the node selector, see Assign Pods to Nodes. |
| node.tolerations | Taint toleration of huawei-csi-node. After this parameter is set, huawei-csi-node can tolerate taints on a node. | No | - key: "node.kubernetes.i o/memory- pressure" operator: "Exists" effect: "NoExecute" - key: "node.kubernetes.i o/disk-pressure" operator: "Exists" effect: "NoExecute" - key: "node.kubernetes.i o/network- unavailable" operator: "Exists" effect: "NoExecute" | For details about taints and tolerations, see Taints and Tolerations. |

4.1.1.4 csiDriver Parameters

The csiDriver parameters include the basic configurations for running Huawei CSI, such as Huawei driver name and multipathing type.

Table 4-4 csiDriver parameters

| Parameter | Description | Mandator y | Default Value | Remarks |
|--------------------------------------|---|---------------|----------------|--|
| csiDriver.driv erName | Registered driver name. | Yes | csi.huawei.com | Use the default value. For the CCE Agile platform, modify this field. For example, csi.oceanstorcom. |
| csiDriver.end point | Communicatio n endpoint. | Yes | /csi/csi.sock | Use the default value. |
| csiDriver.con nectorThread s | Maximum number of disks that can be concurrently scanned/ detached. The value is an integer ranging from 1 to 10. | Yes | 4 | A larger value indicates that more concurrent disk scanning and detaching operations are performed on a single node at the same time. When DM-Multipath is used, a large number of concurrent requests may cause unknown problems and affect the overall time. |
| csiDriver.volu meUseMultip ath | Whether to use multipathing software. The value is a Boolean value. | Yes | true | It is strongly recommended that multipathing software be enabled to enhance the redundancy and performance of storage links. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|-------------------------------------|--|---|-----------------------|---|
| csiDriver.scsi MultipathTyp e | Multipathing software used when the storage protocol is fc or iscsi . The following parameter values can be configured: • DM-multipath • HW-UltraPath • HW-UltraPath-NVMe | Mandator y when volumeUs eMultipat h is set to true. | DM-multipath | The DM - multipath value is recommended. |
| csiDriver.nvm eMultipathTy pe | Multipathing software used when the storage protocol is roce or fc-nvme. Only HW-UltraPath-NVMe is supported. | Mandator y when volumeUs eMultipat h is set to true. | HW-UltraPath- NVMe | |
| csiDriver.scan VolumeTime out | Timeout interval for waiting for multipathing aggregation when DM-Multipath is used on the host. The value ranges from 1 to 600 seconds. | Yes | 3 | - |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--------------------------------------|--|---------------|---------------|--|
| csiDriver.exec CommandTi meout | Timeout interval for running commands on the host. | Yes | 30 | In scenarios such as mounting and capacity expansion, the CSI plug-in needs to run some host commands, for example, running the mount command to mount a file system. This parameter is used to control the timeout interval for running a single command. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|------------------------------|--|---|---------------|---------|
| csi_driver.allP athOnline | Whether to check whether the number of paths aggregated by DM-Multipath is equal to the actual number of online paths. The following parameter values can be configured: • true: The drive letter mounting condition is met only when the number of paths aggregated by DM-Multipath is equal to the actual number of online paths. • false: By default, the number of paths aggregated by DM-Multipath is not checked. As long as virtual drive letters are generated upon aggregatio n, the drive letter | This parameter is mandator y when csi_driver. scsiMulti pathType is set to DM-multipath . | false | |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--|---|---------------|---------------|--|
| | mounting condition is met. | | | |
| csiDriver.ena bleLabel | Whether Kubernetes resource object labels can be delivered to storage. | Yes | false | Only centralized storage 6.1.7 and later versions support this function. Distributed storage, dtrees, and static PVs do not support this function. |
| csiDriver.back endUpdateIn terval | Interval for updating backend capabilities. The value ranges from 60 to 600 seconds. | Yes | 60 | - |
| csiDriver.cont rollerLogging .module | Record type of the controller log. The following parameter values can be configured: • file • console | Yes | file | When the value is file , logs are retained in the specified directory of the node. When the Pod where CSI is located is destroyed, logs are still retained. When the value is console , logs are retained in the temporary space of the Pod where CSI is located. When the Pod where CSI is located is destroyed, the logs are also destroyed. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--|---|---------------|-----------------|---|
| csiDriver.cont rollerLogging .level | Output level of the controller log. The following parameter values can be configured: debug info warning error fatal | Yes | info | |
| csiDriver.cont rollerLogging .fileDir | Directory of the controller log in file output mode. | Yes | /var/log/huawei | Ensure that the directory has sufficient space for storing logs. It is recommended that the space be greater than or equal to 200 MB. |
| csiDriver.cont rollerLogging .fileSize | Size of a single controller log file in file output mode. | Yes | 20M | - |
| csiDriver.cont rollerLogging .maxBackups | Maximum number of controller log file backups in file output mode. | Yes | 9 | - |

| Parameter | Description | Mandator y | Default Value | Remarks |
|---------------------------------------|---|---------------|-----------------|--|
| csiDriver.nod eLogging.mo dule | Record type of the node log. The following parameter values can be configured: • file • console | Yes | file | When the value is file , logs are retained in the specified directory of the node. When the Pod where CSI is located is destroyed, logs are still retained. When the value is console , logs are retained in the temporary space of the Pod where CSI is located. When the Pod where CSI is located. When the Pod where CSI is located is destroyed, the logs are also destroyed. |
| csiDriver.nod eLogging.lev el | Output level of the node log. The following parameter values can be configured: • debug • info • warning • error • fatal | Yes | info | - |
| csiDriver.nod eLogging.file Dir | Directory of the node log in file output mode. | Yes | /var/log/huawei | Ensure that the directory has sufficient space for storing logs. It is recommended that the space be greater than or equal to 200 MB. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--|--|---------------|---------------|---------|
| csiDriver.nod eLogging.file Size | Size of a single node log file in file output mode. | Yes | 20M | - |
| csiDriver.nod eLogging.ma xBackups | Maximum number of node log file backups in file output mode. | Yes | 9 | - |

♠ CAUTION

If Huawei CSI has been deployed in your container environment, ensure that the value of **csiDriver.driverName** is the same as that configured during previous deployment. Otherwise, existing volumes or snapshots provisioned by Huawei CSI in the system cannot be managed by the newly deployed Huawei CSI.

4.1.1.5 Other Parameters

Other parameters include some features of the CSI plug-in or the policies for obtaining images.

Table 4-5 Other parameters

| Parameter | Description | Mandator y | Default Value | Remarks |
|--------------------------|--|---------------|------------------|---------|
| kubernetes.n amespace | Kubernetes namespace where Huawei CSI is running, which can be customized. The name must consist of lowercase letters, digits, and hyphens (-), for example, my-name and 123- abc. | No | huawei- csi | - |

| Parameter | Description | Mandator y | Default Value | Remarks |
|------------------------------|---|---------------|----------------------|---|
| kubeletConfi gDir | Working directory of kubelet. | Yes | /var/lib/ kubelet | Use the default value. |
| | | | | For the Tanzu platform, change the value of this field to /var/ vcap/data/ kubelet. |
| | | | | • For the CCE Agile platform, change the value of this field to /mnt/ paas/ kubernetes/ kubelet. |
| sidecarImage PullPolicy | Pull policy of the sidecar image. | Yes | IfNotPres ent | - |
| huaweilmag ePullPolicy | Pull policy of the huawei-csi image. | Yes | IfNotPres ent | - |
| CSIDriverObj ect.isCreate | Whether to create the CSIDriver object. | Yes | false | The CSIDriver feature is a GA version in Kubernetes v1.18. Therefore, to use this feature, the Kubernetes version must be later than v1.18. If the Kubernetes version is earlier than v1.18, set this parameter to false. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--|---|---------------|------------------|---|
| CSIDriverObj ect.attachRe quired | Whether the CSI plug-in skips the attach operation. The following parameter values can be configured: • true: The attach operation is required. • false: The attach operation is skipped. | Yes | true | The attachRequired parameter can be configured in Kubernetes v1.18. If CSIDriverObjec t.isCreate is set to true and attachRequired is set to false, the huawei-csi plug-in will not deploy the csi-attacher sidecar. If NAS storage is used, this parameter can be set to false. If SAN storage is used, set this parameter to true. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|---------------------------------------|---|---------------|------------------|---|
| CSIDriverObj ect.fsGroupP olicy | Whether the ownership and permissions of a basic volume can be changed before the volume is mounted. The following parameter values can be configured: ReadWriteOnce WithFSType: The volume ownership and permission can be changed only when fsType is specified and accessModes of the volume contains ReadWriteOnce. File: Kubernetes can use fsGroup to change the permissions and ownership of a volume to match fsGroup requested by a user in the Pod security policy, regardless of fsGroup or accessModes. None: A volume is mounted without any change. null: The fsGroupPolicy parameter is not set. | No | null | The fsGroupPolicy parameter can be configured in Kubernetes v1.20, and takes effect only when CSIDriverObjec t.isCreate is set to true. This feature is a Beta version in Kubernetes v1.20 but a GA version in Kubernetes v1.23. Therefore, the Kubernetes version must be later than v1.20. |
| leaderElectio n.leaseDurati on | Leader duration. | No | 8s | This parameter takes effect only in the multicontroller scenario. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--------------------------------------|---------------------------------------|---------------|------------------|---|
| leaderElectio n.renewDead line | Time for the leader to be re-elected. | No | 6s | This parameter takes effect only in the multicontroller scenario. |
| leaderElectio n.retryPeriod | Leader election retry time. | No | 2s | This parameter takes effect only in the multicontroller scenario. |

NOTICE

Ensure that the namespace entered in **kubernetes.namespace** exists on Kubernetes. If the namespace does not exist, run the following command to create it. In this example, the namespace for running Huawei CSI is **huawei-csi**. # kubectl create namespace *huawei-csi*

4.1.2 Installing Huawei CSI

This section describes how to install Huawei CSI using Helm 3.

NOTICE

- Huawei CSI can be installed as the root user or a non-root user. When installing
 Huawei CSI as a non-root user, ensure that the current user can access the API
 Server of the Kubernetes cluster. For details about how to configure access to
 the Kubernetes cluster as a non-root user, see 10.6 Configuring Access to the
 Kubernetes Cluster as a Non-root User.
- Huawei CSI must be run as the root user.

Helm is a software package management tool in the Kubernetes ecosystem. Similar to Ubuntu APT, CentOS YUM, or Python pip, Helm manages Kubernetes application resources.

You can use Helm to package, distribute, install, upgrade, and roll back Kubernetes applications in a unified manner.

- For details about how to obtain and install Helm, see https://helm.sh/docs/intro/install/.
- For details about the mapping between Helm and Kubernetes versions, see https://helm.sh/docs/topics/version_skew/.

When installing huawei-csi-controller, Helm deploys the following components in the workloads of the Deployment type in the specified namespace:

- huawei-csi-driver: Huawei CSI driver.
- storage-backend-controller: Huawei backend management controller, used to manage storageBackendClaim resources.
- storage-backend-sidecar: used to manage storageBackendContent resources.
- Kubernetes External Provisioner: used to provide volumes.
- Kubernetes External Attacher: used to attach volumes.
- (Optional) Kubernetes External Snapshotter: used to provide snapshot support (installed as CRD).
- (Optional) Kubernetes External Snapshot Controller: used to control volume snapshots.
- Kubernetes External Resizer: used to expand the capacity of volumes.
- Kubernetes External liveness-probe: used to determine the health status of a Pod.

When installing huawei-csi-node, Helm deploys the following components in the workloads of the DaemonSet type in the specified namespace:

- huawei-csi-driver: Huawei CSI driver.
- Kubernetes Node Registrar: used to process driver registration.
- liveness-probe: used to determine the health status of a Pod.

4.1.2.1 Installing Huawei CSI on Kubernetes, OpenShift, and Tanzu

This section describes how to install Huawei CSI on the Kubernetes, OpenShift, and Tanzu platforms.

Prerequisites

- A Huawei CSI image has been created and uploaded to the image repository or imported to all nodes by following the instructions provided in 3.3 Uploading a Huawei CSI Image.
- The component images on which Huawei CSI installation and running depend have been uploaded to the image repository or imported to all nodes. For details, see 3.4 Checking the Images on Which CSI Depends.
- The volume snapshot component CRD on which the running of Huawei CSI depends has been installed. For details, see 3.5 Checking Volume Snapshot-Dependent Components.
- If you want to use multipathing to connect to Huawei storage, ensure that multipathing software has been installed on all compute nodes. For details, see 3.6 Checking the Host Multipathing Configuration.
- Helm 3 has been installed on the container management platform.
- The values.yaml file required for installing CSI has been prepared. For details, see 4.1.1 Preparing the values.yaml File.
- All worker nodes of the cluster communicate properly with the service IP address of the storage device to be connected. In iSCSI scenarios, the ping command can be used to verify the connectivity.
- Software clients required by the corresponding protocol, such as iSCSI and NFS clients, have been installed on all worker nodes of the cluster.

 The accounts required for connecting to Huawei CSI have been created on the Huawei storage to be connected. For details, see 3.7 Checking the Accounts on Huawei Storage.

Installation Preparations

For the OpenShift platform, run the following commands to create the **SecurityContextConstraints** resource.

• Run the **vi** helm_scc.yaml command to create a **SecurityContextConstraints** file. In the following command output, **huawei-csi** indicates the created namespace. Replace it based on site requirements.

```
# vi helm_scc.yaml
apiVersion: security.openshift.io/v1
kind: SecurityContextConstraints
name: helm-scc
allowHostDirVolumePlugin: true
allowHostIPC: true
allowHostNetwork: true
allowHostPID: true
allowHostPorts: true
allowPrivilegeEscalation: true
allowPrivilegedContainer: true
defaultAddCapabilities:
- SYS_ADMIN
runAsUser:
type: RunAsAny
seLinuxContext:
type: RunAsAny
fsGroup:
type: RunAsAny
users:
- system:serviceaccount:huawei-csi:huawei-csi-controller
- system:serviceaccount:huawei-csi:huawei-csi-node
```

 Run the oc create -f helm_scc.yaml command to create a SecurityContextConstraints file.

oc create -f helm_scc.yaml

On the Tanzu platform, run the following commands to configure the **use** permission of **podsecuritypolicies** and the **kubelet** installation directory.

• Run the vi psp-use.yaml command to create a file named psp-use.yaml

```
# vi psp-use.yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
name: huawei-csi-psp-role
- apiGroups: ['policy']
 resources: ['podsecuritypolicies']
verbs: ['use']
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBindina
metadata:
name: huawei-csi-psp-role-cfg
roleRef:
 kind: ClusterRole
 name: huawei-csi-psp-role
 apiGroup: rbac.authorization.k8s.io
subjects:
- kind: Group
apiGroup: rbac.authorization.k8s.io
```

```
name: system:serviceaccounts:huawei-csi
- kind: Group
apiGroup: rbac.authorization.k8s.io
name: system:serviceaccounts:default
```

 Run the kubectl create -f psp-use.yaml command to create the PSP permission.

kubectl create -f psp-use.yaml

Go to the helm/esdk directory in the installation package, run the vi values.yaml command to open the configuration file, modify the file, and save the file. For details about the installation package directory, see Table 3-1.

```
# vi values.yaml

# Specify kubelet config dir path.

# kubernetes and openshift is usually /var/lib/kubelet

# Tanzu is usually /var/vcap/data/kubelet

# CCE is usually /mnt/paas/kubernetes/kubelet
kubeletConfigDir: /var/vcap/data/kubelet
```

Deployment Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the cluster through the management IP address.
- **Step 2** Copy the **helm** directory in the Kubernetes CSI component package to any directory on the master node. For details about the Helm tool path, see **Table 3-1**.
- **Step 3** Go to the **helm/esdk** working directory.

 # cd helm/esdk
- **Step 4** Run the **kubectl apply -f** *crds/backend/* command to update the CRD.

 # kubectl apply -f crds/backend/
 customresourcedefinition.apiextensions.k8s.io/storagebackendclaims.xuanwu.huawei.io configured
 customresourcedefinition.apiextensions.k8s.io/storagebackendcontents.xuanwu.huawei.io configured
- Step 5 Ensure that snapshot-dependent components have been checked by following the instructions provided in 3.5 Checking Volume Snapshot-Dependent Components. Then run the helm install helm-huawei-csi./ -n huawei-csi -- create-namespace command to install Huawei CSI.

In the preceding command, *helm-huawei-csi* indicates the custom Helm chart name, ./ indicates that the Helm project in the current directory is used, and *huawei-csi* indicates the custom Helm chart namespace.

```
# helm install helm-huawei-csi ./ -n huawei-csi --create-namespace

NAME: helm-huawei-csi
LAST DEPLOYED: Wed Jun 8 11:50:28 2022

NAMESPACE: huawei-csi
STATUS: deployed

REVISION: 1

TEST SUITE: None
```

Step 6 After the huawei-csi service is deployed, run the **kubectl get pod -n** *huawei-csi* command to check whether the service is started. On the OpenShift platform, run the **oc get pod -n** *huawei-csi* command to check whether the service is started.

```
# kubectl get pod -n huawei-csi
NAME READY STATUS RESTARTS AGE
huawei-csi-controller-6dfcc4b79f-9vjtq 9/9 Running 0 24m
huawei-csi-controller-6dfcc4b79f-csphc 9/9 Running 0 24m
huawei-csi-node-g6f4k 3/3 Running 0 20m
huawei-csi-node-tqs87 3/3 Running 0 20m
```

4.1.2.2 Installing Huawei CSI on CCE/CCE Agile

This section describes how to install Huawei CSI on the CCE/CCE Agile platform.

Prerequisites

- A Huawei CSI image has been created and uploaded to the image repository or imported to all nodes by following the instructions provided in 3.3 Uploading a Huawei CSI Image.
- The component images on which Huawei CSI installation and running depend have been uploaded to the image repository or imported to all nodes. For details, see 3.4 Checking the Images on Which CSI Depends.
- The volume snapshot component CRD on which the running of Huawei CSI depends has been installed. For details, see 3.5 Checking Volume Snapshot-Dependent Components.
- If you want to use multipathing to connect to Huawei storage, ensure that multipathing software has been installed on all compute nodes. For details, see 3.6 Checking the Host Multipathing Configuration.
- Helm 3 has been installed on the container management platform.
- The **values.yaml** file required for installing CSI has been prepared. For details, see **4.1.1 Preparing the values.yaml File**.
- All worker nodes of the cluster communicate properly with the service IP address of the storage device to be connected. In iSCSI scenarios, the ping command can be used to verify the connectivity.
- Software clients required by the corresponding protocol, such as iSCSI and NFS clients, have been installed on all worker nodes of the cluster.
- The accounts required for connecting to Huawei CSI have been created on the Huawei storage to be connected. For details, see 3.7 Checking the Accounts on Huawei Storage.

4.1.2.2.1 Creating a Helm Installation Package

To install Huawei CSI on the CCE/CCE Agile platform, you need to create a Helm installation package. This section describes how to create a Helm installation package.

Prerequisites

- Helm 3 has been installed on a node server.
- The **values.yaml** file required for installing CSI has been prepared. For details, see **4.1.1 Preparing the values.yaml File**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any node where Helm is deployed through the management IP address.
- **Step 2** Copy the **helm** directory in the Huawei CSI component package to any directory on the node. For details about the Helm tool path, see **Table 3-1**.
- **Step 3** Go to the **helm** working directory.

cd helm/

Step 4 Modify the **kubeletConfigDir** and **csiDriver.driverName** parameters in the **helm/esdk/values.yaml** file.

- # vi ./esdk/values.yaml
- # Specify kubelet config dir path.
- # kubernetes and openshift is usually /var/lib/kubelet
- # Tanzu is usually /var/vcap/data/kubelet
- # CCE is usually /mnt/paas/kubernetes/kubelet

kubeletConfigDir: /mnt/paas/kubernetes/kubelet

The CSI driver parameter configuration

csiDriver:

- # Driver name, it is strongly recommended not to modify this parameter
- # The CCE platform needs to modify this parameter, e.g. csi.oceanstor.com

driverName: csi.oceanstor.com

Step 5 Run the **helm package ./esdk/ -d ./** command to create a Helm installation package. This command will generate the installation package to the current path.

helm package ./esdk/ -d ./

Successfully packaged chart and saved it to: esdk-1.0.0.tgz

----End

4.1.2.2.2 Installing Huawei CSI

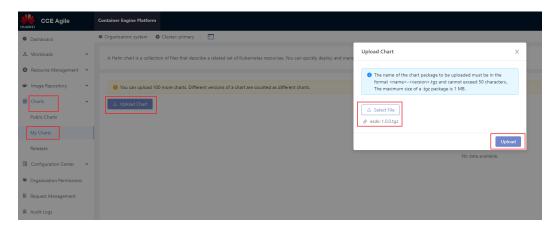
This section describes how to install Huawei CSI on the CCE/CCE Agile platform. The following uses CCE Agile v22.3.2 as an example.

Prerequisites

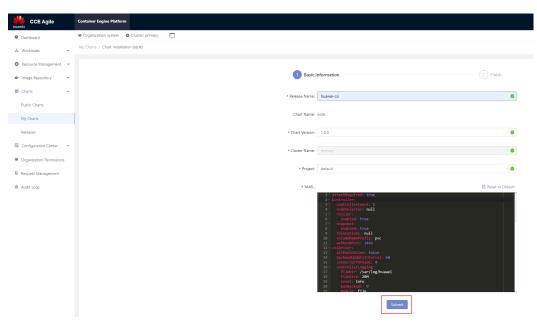
- Installation preparations have been checked according to **Prerequisites**.
- A Huawei CSI Helm installation package has been created. For details, see
 4.1.2.2.1 Creating a Helm Installation Package.

Procedure

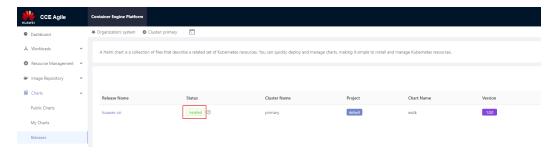
- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node where the CCE Agile platform is deployed through the management IP address.
- **Step 2** Run the **kubectl create namespace** *huawei-csi* command to create a namespace for deploying Huawei CSI. *huawei-csi* indicates the custom namespace.
 - # kubectl create namespace huawei-csi
- **Step 3** Export the Helm installation package. For details, see **4.1.2.2.1 Creating a Helm Installation Package**.
- **Step 4** On the home page, choose **Charts** > **My Charts** > **Upload Chart**. The **Upload Chart** dialog box is displayed. Import the exported Helm installation package to the CCE Agile platform.



Step 5 After the installation package is uploaded, choose **Charts > My Charts**. On the **My Charts** page that is displayed, choose **Install > Submit**. The chart release name can be customized.



Step 6 On the home page, choose **Charts** > **Releases** and select the project specified during installation (for example, **default** in the following figure). After the installation is successful, **Installed** is displayed in the **Status** column.



4.2 Manually Installing Huawei CSI

This section describes how to manually install Huawei CSI.

■ NOTE

Currently, only the Kubernetes platform supports manual installation of Huawei CSI.

Prerequisites

- A Huawei CSI image has been created and uploaded to the image repository or imported to all nodes by following the instructions provided in 3.3 Uploading a Huawei CSI Image.
- The component images on which Huawei CSI installation and running depend have been uploaded to the image repository or imported to all nodes. For details, see 3.4 Checking the Images on Which CSI Depends.
- The volume snapshot component CRD on which the running of Huawei CSI depends has been installed. For details, see 3.5 Checking Volume Snapshot-Dependent Components.
- If you want to use multipathing to connect to Huawei storage, ensure that multipathing software has been installed on all compute nodes. For details, see 3.6 Checking the Host Multipathing Configuration.
- All worker nodes of the cluster communicate properly with the service IP address of the storage device to be connected. In iSCSI scenarios, the ping command can be used to verify the connectivity.
- Software clients required by the corresponding protocol, such as iSCSI and NFS clients, have been installed on all worker nodes of the cluster.
- The accounts required for connecting to Huawei CSI have been created on the Huawei storage to be connected. For details, see 3.7 Checking the Accounts on Huawei Storage.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the cluster through the management IP address.
- **Step 2** Copy the **manual** directory in the Kubernetes CSI component package to any directory on the master node.
- **Step 3** Go to the **manual/esdk** working directory. For details about the path, see **Table 3-1**.

cd manual/esdk

Step 4 Run the **kubectl apply -f** ./crds/backend/ command to install the storage backend CRD.

kubectl apply -f ./crds/backend/ customresourcedefinition.apiextensions.k8s.io/storagebackendclaims.xuanwu.huawei.io configured customresourcedefinition.apiextensions.k8s.io/storagebackendcontents.xuanwu.huawei.io configured

Step 5 Run the **kubectl apply -f** ./crds/resourcetopologies command to install the topology resource CRD.

kubectl apply -f ./crds/resourcetopologies customresourcedefinition.apiextensions.k8s.io/resourcetopologies.xuanwu.huawei.io configured

Step 6 Run the **kubectl apply -f** ./crds/snapshot-crds command to install the snapshot CRD.

kubectl apply -f ./crds/snapshot-crds customresourcedefinition.apiextensions.k8s.io/volumesnapshotclasses.snapshot.storage.k8s.io configured customresourcedefinition.apiextensions.k8s.io/volumesnapshotcontents.snapshot.storage.k8s.io configured customresourcedefinition.apiextensions.k8s.io/volumesnapshots.snapshot.storage.k8s.io configured

Step 7 Run the **kubectl create -f** ./deploy command to install and deploy CSI.

kubectl create -f ./deploy csidriver.storage.k8s.io/csi.huawei.com created clusterrole.rbac.authorization.k8s.io/storage-backend-controller-role created clusterrole.rbac.authorization.k8s.io/storage-backend-sidecar-role created clusterrolebinding.rbac.authorization.k8s.io/storage-backend-controller-binding created clusterrolebinding.rbac.authorization.k8s.io/storage-backend-sidecar-binding created serviceaccount/huawei-csi-controller created rolebinding.rbac.authorization.k8s.io/huawei-csi-attacher-role-cfg created role.rbac.authorization.k8s.io/huawei-csi-attacher-cfg created rolebinding.rbac.authorization.k8s.io/huawei-csi-provisioner-role-cfg created role.rbac.authorization.k8s.io/huawei-csi-provisioner-cfg created clusterrolebinding.rbac.authorization.k8s.io/huawei-csi-provisioner-role created clusterrole.rbac.authorization.k8s.io/huawei-csi-provisioner-runner created clusterrolebinding.rbac.authorization.k8s.io/huawei-csi-attacher-role created clusterrole.rbac.authorization.k8s.io/huawei-csi-attacher-runner created rolebinding.rbac.authorization.k8s.io/huawei-csi-csi-resizer-role-cfg created role.rbac.authorization.k8s.io/huawei-csi-resizer-cfg created clusterrolebinding.rbac.authorization.k8s.io/huawei-csi-csi-resizer-role created clusterrole.rbac.authorization.k8s.io/huawei-csi-resizer-runner created clusterrole.rbac.authorization.k8s.io/huawei-csi-controller-runner created clusterrolebinding.rbac.authorization.k8s.io/huawei-csi-controller-role created deployment.apps/huawei-csi-controller created service/huawei-csi-controller created serviceaccount/huawei-csi-node created clusterrolebinding.rbac.authorization.k8s.io/huawei-csi-driver-registrar-role created clusterrole.rbac.authorization.k8s.io/huawei-csi-driver-registrar-runner created clusterrole.rbac.authorization.k8s.io/huawei-csi-node-runner created clusterrolebinding.rbac.authorization.k8s.io/huawei-csi-node-role created daemonset.apps/huawei-csi-node created

Step 8 After the service is deployed in **Step 7**, run the **kubectl get pod -n** *huawei-csi* command to check whether the service is started.

kubectl get pod -n huawei-csi
NAME READY STATUS RESTARTS AGE
huawei-csi-controller-6dfcc4b79f-9vjtq 9/9 Running 0 24m
huawei-csi-node-g6f4k 3/3 Running 0 20m
huawei-csi-node-tqs87 3/3 Running 0 20m

5 Uninstalling Huawei CSI

- 5.1 Uninstalling Huawei CSI Using Helm
- 5.2 Manually Uninstalling Huawei CSI

5.1 Uninstalling Huawei CSI Using Helm

This chapter describes how to uninstall Huawei CSI. The uninstallation method varies according to the installation mode.



If you do not uninstall Huawei CSI for the purpose of an upgrade, ensure that all resources (such as PV, PVC, snapshot, and storage backend resources) provisioned by Huawei CSI have been cleared on your container platform before uninstalling Huawei CSI. Otherwise, once you uninstall Huawei CSI, these resources cannot be automatically scheduled, managed, or cleared.

5.1.1 Uninstalling Huawei CSI on Kubernetes, OpenShift, and Tanzu

This section describes how to uninstall Huawei CSI on the Kubernetes, OpenShift, and Tanzu platforms.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the helm uninstall helm-huawei-csi -n huawei-csi command to uninstall Huawei CSI. In the command, helm-huawei-csi indicates the custom Helm chart name and huawei-csi indicates the namespace where the Helm chart resides. This command will delete the huawei-csi-controller, huawei-csi-node, and RBAC resources of Huawei CSI.

helm uninstall helm-huawei-csi -n huawei-csi release "helm-huawei-csi" uninstalled

After the deletion command is executed, you need to check whether the uninstallation is successful.

helm list -n huawei-csi

NAME NAMESPACE REVISION UPDATED STATUS CHART APP VERSION

In the preceding command, *huawei-csi* indicates the namespace where the chart is located.

If the command output is empty, the service is successfully deleted.

- **Step 3** Delete the huawei-csi-host-info object. For details, see **5.1.3.1 Deleting the huawei-csi-host-info Object**.
- **Step 4** Delete the webhook resource. For details, see **5.1.3.2 Deleting a Webhook Resource**.
- **Step 5** (Optional) Uninstall the snapshot-dependent component service. For details, see **5.1.3.3 Uninstalling the Snapshot-Dependent Component Service**.

----End

5.1.2 Uninstalling Huawei CSI on CCE/CCE Agile

This section describes how to uninstall Huawei CSI on the CCE/CCE Agile platform. The following uses CCE Agile v22.3.2 as an example.

Procedure

- **Step 1** Log in to the CCE Agile platform.
- **Step 2** On the home page, choose **Charts** > **Releases**. The **Releases** page is displayed.
- **Step 3** Select a Huawei CSI release and click **Uninstall**. In the displayed dialog box, click **OK**.



- **Step 4** Delete the huawei-csi-host-info object. For details, see **5.1.3.1 Deleting the** huawei-csi-host-info Object.
- **Step 5** Delete the webhook resource. For details, see **5.1.3.2 Deleting a Webhook Resource**.
- **Step 6** (Optional) Uninstall the snapshot-dependent component service. For details, see **5.1.3.3 Uninstalling the Snapshot-Dependent Component Service**.

----End

5.1.3 Deleting CSI-Dependent Component Services

This section describes how to delete the CSI-dependent component services.

5.1.3.1 Deleting the huawei-csi-host-info Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete secret** *huawei-csi-host-info* **-n** *huawei-csi* command to delete the **secret** object. *huawei-csi-host-info* is the name of the **secret** object, and *huawei-csi* is the namespace where the **secret** object is located.

kubectl delete secret huawei-csi-host-info -n huawei-csi

Step 3 Run the following command to check whether the **secret** object is successfully deleted. If **NotFound** is displayed in the command output, the **huawei-csi-host-info** object is successfully deleted.

kubectl get secret huawei-csi-host-info -n huawei-csi Error from server (NotFound): secrets "huawei-csi-host-info" not found

----End

5.1.3.2 Deleting a Webhook Resource

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query the webhook-dependent component service.

kubectl get validatingwebhookconfigurations.admissionregistration.k8s.io NAME WEBHOOKS AGE storage-backend-controller.xuanwu.huawei.io 1 12d

Step 3 Run the following command to uninstall the webhook-dependent component service.

kubectl delete validatingwebhookconfigurations.admissionregistration.k8s.io storage-backend-controller.xuanwu.huawei.io

Step 4 Run the following command to check whether the service is successfully uninstalled. If the command output is empty, the uninstallation is successful.

 ${\it\#}\ kubectl\ get\ validating we bhook configurations. admission registration. k8s. io\ storage-backend-controller. xuanwu.huawei.io$

----End

5.1.3.3 Uninstalling the Snapshot-Dependent Component Service

CAUTION

- Do not uninstall the snapshot-dependent component service when snapshots exist. Otherwise, Kubernetes will automatically delete all user snapshots and they cannot be restored. Exercise caution when performing this operation. For details, see **Delete a CustomResourceDefinition**.
- Do not uninstall the snapshot-dependent component service during the CSI upgrade.

Scenario Description

- Currently, Huawei CSI uses the snapshot feature.
- Currently, only Huawei CSI is available in the Kubernetes cluster, and Huawei CSI is no longer used.
- Before the uninstallation, ensure that no VolumeSnapshot resource managed by Huawei CSI exists in the Kubernetes cluster.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to uninstall the snapshot-dependent component service.

kubectl delete crd volumesnapshotclasses.snapshot.storage.k8s.io volumesnapshotcontents.snapshot.storage.k8s.io volumesnapshots.snapshot.storage.k8s.io

Step 3 Run the following command to check whether the service is successfully uninstalled.

If the command output is empty, the uninstallation is successful. # kubectl get crd | grep snapshot.storage.k8s.io

----End

5.2 Manually Uninstalling Huawei CSI

This section describes how to manually uninstall Huawei CSI.



If you do not uninstall Huawei CSI for the purpose of an upgrade, ensure that all resources (such as PV, PVC, snapshot, and storage backend resources) provisioned by Huawei CSI have been cleared on your container platform before uninstalling Huawei CSI. Otherwise, once you uninstall Huawei CSI, these resources cannot be automatically scheduled, managed, or cleared.

5.2.1 Uninstalling the huawei-csi-node Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete daemonset huawei-csi-node -n** *huawei-csi* command to uninstall the huawei-csi-node service. Replace *huawei-csi* with the namespace where Huawei CSI is located.

kubectl delete daemonset huawei-csi-node -n huawei-csi

Step 3 Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get daemonset huawei-csi-node -n huawei-csi

----End

5.2.2 Uninstalling the huawei-csi-controller Service

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete deployment huawei-csi-controller -n** *huawei-csi* command to uninstall the huawei-csi-controller service. Replace *huawei-csi* with the namespace where Huawei CSI is located.
 - # kubectl delete deployment huawei-csi-controller -n huawei-csi
- **Step 3** Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get deployment huawei-csi-controller -n huawei-csi

----End

5.2.3 Uninstalling the csidriver Object

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **kubectl delete csidriver csi.huawei.com** command to uninstall the csidriver object.
 - # kubectl delete csidriver csi.huawei.com
- **Step 3** Run the following command to check whether the service is successfully uninstalled. If **NotFound** is displayed, the service is successfully uninstalled.

kubectl get csidriver csi.huawei.com

----End

5.2.4 Deleting the RBAC Permission

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Delete the RBAC permission.

kubectl -n huawei-csi -l provisioner=csi.huawei.com delete ServiceAccount,Service,role,rolebinding,ClusterRole,ClusterRoleBinding

5.2.5 Uninstalling Other Resources

Procedure

- **Step 1** Delete the huawei-csi-host-info object. For details, see **5.1.3.1 Deleting the** huawei-csi-host-info Object.
- **Step 2** Delete the webhook resource. For details, see **5.1.3.2 Deleting a Webhook Resource**.
- **Step 3** (Optional) Uninstall the snapshot-dependent component service. For details, see **5.1.3.3 Uninstalling the Snapshot-Dependent Component Service**.

6 Upgrade/Rollback Operations

This chapter describes how to upgrade or roll back Huawei CSI.

- 6.1 Upgrading or Rolling Back Huawei CSI Using Helm
- 6.2 Manual Upgrade/Rollback

6.1 Upgrading or Rolling Back Huawei CSI Using Helm

6.1.1 Upgrading Huawei CSI

This section describes how to upgrade Huawei CSI.

During the upgrade or rollback, the existing resources such as PVCs, snapshots, and Pods will run properly and will not affect your service access.

<u>^</u> CAUTION

- Some CSI 2.x versions are unavailable now. If the upgrade fails, CSI may fail to be rolled back to a version which is unavailable now.
- During the upgrade or rollback, you cannot use Huawei CSI to create new resources or mount or unmount an existing PVC.
- During the upgrade or rollback, do not uninstall the snapshot-dependent component service.

6.1.1.1 Upgrading Huawei CSI on Kubernetes, OpenShift, and Tanzu

Upgrading CSI from 2.x or 3.x to 4.2.1

To upgrade CSI from 2.x or 3.x to 4.2.1, perform the following operations:

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- Step 2 Run the kubectl get cm huawei-csi-configmap -n huawei-csi -o json > configmap.json command to back up the backend information to the configmap.json file. For the OpenShift platform, replace kubectl with oc. # kubectl get cm huawei-csi-configmap -n huawei-csi -o json > configmap.json
- Step 3 Uninstall CSI. For details, see 5.1 Uninstalling Huawei CSI Using Helm.
- **Step 4** Install CSI of the current version. For details, see **4.1.2 Installing Huawei CSI**.
- Step 5 Install the backend information backed up in Step 2 according to 7.1 Adding a Storage Backend.

----End

Upgrading CSI from 4.x to 4.2.1

To upgrade CSI from 4.x to 4.2.1, perform the following operations:

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Uninstall CSI. For details, see 5.1 Uninstalling Huawei CSI Using Helm.
- **Step 3** Install CSI of the current version. For details, see **4.1.2 Installing Huawei CSI**.

----End

6.1.2 Rolling Back Huawei CSI

If CSI fails to be upgraded from 2.x or 3.x to 4.2.1 and needs to be rolled back, uninstall CSI by referring to 5.1 Uninstalling Huawei CSI Using Helm and then download and install CSI of the source version.

<u>A</u> CAUTION

- During the upgrade or rollback, the existing resources such as PVCs, snapshots, and Pods will run properly and will not affect your service access.
- During the upgrade or rollback, you cannot use Huawei CSI to create new resources or mount or unmount an existing PVC.
- During the upgrade or rollback, do not uninstall the snapshot-dependent component service.

6.1.2.1 Rolling Back Huawei CSI on Kubernetes, OpenShift, and Tanzu

Prerequisites

You have downloaded the CSI software package of the source version.

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- Step 2 Uninstall CSI. For details, see 5.1 Uninstalling Huawei CSI Using Helm.
- **Step 3** Reinstall CSI of the original version by referring to the user guide in the CSI installation package of the original version.

----End

6.2 Manual Upgrade/Rollback

6.2.1 Upgrading Huawei CSI

This section describes how to manually upgrade Huawei CSI.

During the upgrade or rollback, the existing resources such as PVCs, snapshots, and Pods will run properly and will not affect your service access.

♠ CAUTION

- Some CSI 2.x versions are unavailable now. If the upgrade fails, CSI may fail to be rolled back to a version which is unavailable now.
- During the upgrade or rollback, you cannot use Huawei CSI to create new resources or mount or unmount an existing PVC.
- During the upgrade or rollback, do not uninstall the snapshot-dependent component service.

Upgrading CSI from 2.x or 3.x to 4.2.1

To upgrade CSI from 2.x or 3.x to 4.2.1, perform the following operations:

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Run the kubectl get cm huawei-csi-configmap -n huawei-csi -o json > configmap.json command to back up the backend information to the configmap.json file. For the OpenShift platform, replace kubectl with oc.
 - # kubectl get cm huawei-csi-configmap -n huawei-csi -o json > configmap.json
- **Step 3** Uninstall CSI. For details, see **5.2 Manually Uninstalling Huawei CSI**.
- **Step 4** Install CSI of the current version. For details, see **4.2 Manually Installing Huawei** CSI.
- Step 5 Install the backend information backed up in Step 2 according to 7.1 Adding a Storage Backend.

----End

Upgrading CSI from 4.x to 4.2.1

To upgrade CSI from 4.x to 4.2.1, perform the following operations:

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Uninstall CSI. For details, see **5.2 Manually Uninstalling Huawei CSI**.
- **Step 3** Install CSI of the current version. For details, see **4.2 Manually Installing Huawei CSI**.

----End

6.2.2 Rolling Back Huawei CSI

Uninstall CSI by referring to **5.2 Manually Uninstalling Huawei CSI**, and then download and install CSI of the source version.

<u>A</u> CAUTION

- During the upgrade or rollback, the existing resources such as PVCs, snapshots, and Pods will run properly and will not affect your service access.
- During the upgrade or rollback, you cannot use Huawei CSI to create new resources or mount or unmount an existing PVC.
- During the upgrade or rollback, do not uninstall the snapshot-dependent component service.

Prerequisites

You have downloaded the CSI software package of the source version.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Uninstall CSI. For details, see 5.2 Manually Uninstalling Huawei CSI.
- **Step 3** Reinstall CSI of the source version. For details, see **4.2 Manually Installing Huawei CSI**.

Storage Backend Management

Backend is an abstract concept of Huawei storage resources. Each Huawei storage device can abstract multiple backend resources using features such as tenants, storage pools, and protocols. Each backend exists independently and defines Huawei storage information required for providing persistent volumes for Kubernetes clusters.

This section describes how to use the oceanctl tool to manage storage backends, including creating, querying, updating, and deleting backends. For details about the help information of the oceanctl tool, run the **oceanctl** --help command.

Prerequisites

- You have obtained the oceanctl tool, copied the oceanctl tool to the environment directory, for example, /usr/local/bin, and obtained the execute permission. The oceanctl tool is stored in /bin/oceanctl of the software package.
- The oceanctl tool must be executed on a master node.
- huawei-csi is the default namespace used by oceanctl to create a backend.
- 7.1 Adding a Storage Backend
- 7.2 Querying a Storage Backend
- 7.3 Updating a Storage Backend
- 7.4 Deleting a Storage Backend
- 7.5 (Optional) Adding a Certificate to a Storage Backend
- 7.6 (Optional) Querying a Storage Backend Certificate
- 7.7 (Optional) Updating a Storage Backend Certificate
- 7.8 (Optional) Deleting a Storage Backend Certificate

7.1 Adding a Storage Backend

This section describes how to create a storage backend. Currently, you can create a backend based on the configured backend yaml file or the exported configmap.json file.

If you create a backend by adding a backend yaml file, configure the backend file by referring to **7.1.1 Preparing the Storage Backend Configuration File**.

If the exported configmap.json file exists, create a storage backend by referring to **7.1.2 Creating a Storage Backend**.

7.1.1 Preparing the Storage Backend Configuration File

□ NOTE

- Before configuring NAS HyperMetro, you need to configure the HyperMetro relationship between two storage devices, including the remote device, HyperMetro domain, and the like. The HyperMetro domain of the file system can only work in HyperMetro mode. For details about the configuration operation, see the product documentation of the corresponding storage model.
- The accounts for connecting to NAS HyperMetro backends must be the administrator accounts of the storage vStores.

Backend Configuration File Description

An example template of the backend configuration file is **/examples/backend/backend.yaml**. The following table lists the parameters.

Table 7-1 backend parameters

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|---------------|---|---------------|-------------------|--|
| storage | If enterprise storage provides SAN, set this parameter to oceanstor-san. If enterprise storage provides NAS, set this parameter to oceanstor-nas. If enterprise storage provides NAS of the Dtree type, set this parameter to oceanstor-dtree. If distributed storage provides SAN, set this parameter to fusionstorage-san. If distributed storage provides SAN, set this parameter to fusionstorage-san. If distributed storage provides NAS, set this parameter to fusionstorage-nas. | Yes | oceanst or-nas | One backend can provide only one storage service. If a single Huawei storage system can provide both SAN and NAS storage services, you can configure multiple backends and use different storage service types for each backend. |
| name | Storage backend name. The value can contain a maximum of 63 characters, including lowercase letters, digits, and hyphens (-). It must start with a letter or digit. | Yes | - | Ensure that the storage backend name is unique. |
| namespa ce | Namespace. | No | - | The storage backend must be in the same namespace as Huawei CSI. |

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|-----------------|--|------------------------------------|-------------------|--|
| vstoreNa me | vStore name on the storage side. This parameter needs to be specified when the connected backend is OceanStor V3/V5 and resources need to be provisioned under a specified vStore. | Conditio nally mandato ry | - | This parameter needs to be specified only when the backend is OceanStor V3/V5 and vStores need to be supported. |
| account Name | Account name on the storage side. This parameter is mandatory when OceanStor Pacific series NAS is connected and NAS resources need to be provisioned under a specified account. | Conditio nally mandato ry | - | This parameter needs to be specified only when the backend is OceanStor Pacific series NAS and accounts need to be supported. |
| urls | Management URLs of storage device. The value format is a list. The value can be a domain name or an IP address + port number. Only IPv4 addresses are supported. | Yes | - | If the connected backend is OceanStor V6 or OceanStor Dorado V6 and resources need to be provisioned under a specified vStore, set this parameter to the URL of the logical management port of the vStore. |
| pools | Storage pools of storage devices. The value format is a list. | Conditio nally mandato ry | - | This parameter is optional when storage is set to oceanstor-dtree. |

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|-----------------------------------|--|---------------|-------------------|--|
| er paramet ers.proto col | Storage protocol. The value is a character string. • iscsi • fc • roce • fc-nvme • nfs • dpc • scsi | Yes | | If the value is set to iscsi, ensure that an iSCSI client has been installed on the connected compute node. If the value is set to nfs, ensure that an NFS client tool has been installed on the connected compute node. If the value is set to fc-nvme or roce, ensure that the nvme-cli tool of version 1.9 or later has been installed on the connected compute node. If the value is set to dpc, ensure that DPC has been installed on the connected compute node and the node has been added as a DPC compute node on the storage device to be connected. If the value is set to scsi, ensure that a distributed storage VBS client |
| | | | | has been installed on the connected compute node. |

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|-------------------------------|--|------------------------------------|-------------------|---|
| paramet ers.porta ls | Service access port. Nodes will use this port to read and write storage resources. The value format is a list. Multiple ports can be configured if the protocol is iscsi or roce. Only one port can be configured if the protocol is nfs. Service ports do not need to be configured if the protocol is fc, fc-nvme, or dpc. If the protocol is scsi, the port is in dictionary format where the key indicates the host name and the value indicates the IP address (only IPv4 addresses are supported). | Conditio nally mandato ry | _ | If a vStore or account is used to connect to a backend, portals must be set to the logical port information of the vStore or account. If nfs is used, the value can be a domain name. |
| paramet ers.ALUA | ALUA configuration of the storage backend. If the worker node uses the native multipathing software provided by the OS and ALUA is enabled, you need to configure this parameter. | Conditio nally mandato ry | - | If ALUA is enabled for the host multipathing software, ensure that the backend ALUA configuration is the same as that of the host ALUA configuration. For details about the ALUA configuration, see 9.1.1 Configuring ALUA Using Helm. |
| paramet ers.paren tname | Name of a file system on the current storage device. Dtree is created in the file system. This parameter is mandatory when storage is set to oceanstor-dtree. | Conditio nally mandato ry | - | Query the name on the File Systems page of DeviceManager. |

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|-----------------------------|---|------------------------------------|-------------------|---|
| metrovSt orePairID | HyperMetro vStore pair ID. This parameter is mandatory when a PV to be created on the storage side needs to support the NAS HyperMetro feature. In this case, you need to enter the ID of the HyperMetro vStore pair to which the PV to be created belongs. | Conditio nally mandato ry | - | You can query the HyperMetro vStore pair ID on DeviceManager. |
| metroBa ckend | Backend name of the HyperMetro peer. The value is a character string. This parameter is mandatory when a PV to be created on the storage side needs to support the NAS HyperMetro feature. In this case, you need to enter the name of the other backend to form a HyperMetro pair with the current backend. | Conditio nally mandato ry | | The names of the two backends in the pair must be entered. After the two backends form a HyperMetro relationship, they cannot form a HyperMetro relationship with other backends. |
| supporte dTopolog ies | Storage topology awareness configuration. The parameter format is JSON of the list type. | Conditio nally mandato ry | - | This parameter is mandatory if storage topology awareness is enabled. For details, see 9.2.1 Configuring Storage Topology Awareness Using Helm. |
| maxClien tThreads | Maximum number of concurrent connections to a storage backend. | No | 30 | If this parameter is not specified, the default maximum number of connections is 30. |

For details about how to configure **backend** in typical scenarios, see the following examples:

- Configuring a Storage Backend of the iSCSI Type
- Configuring a Storage Backend of the FC Type
- Configuring a Storage Backend of the NVMe over RoCE Type
- Configuring a Storage Backend of the NVMe over FC Type
- Configuring a Storage Backend of the NFS Type
- Configuring a Storage Backend of the SCSI Type
- Configuring a Storage Backend of the DPC Type
- Configuring Storage Backends of the Dtree Type
- Configuring Storage Backends of the HyperMetro Type

Configuring a Storage Backend of the iSCSI Type

The following is an example of the backend configuration file of the iSCSI type for enterprise storage:

The following is an example of the backend configuration file of the iSCSI type for distributed storage:

```
storage: "fusionstorage-san"
name: "pacific-iscsi-125"
namespace: "huawei-csi"
urls:
- "https://192.168.129.125:8088"
- "https://192.168.129.126:8088"
pools:
- "StoragePool001"
parameters:
protocol: "iscsi"
portals:
- "192.168.128.122"
- "192.168.128.123"
maxClientThreads: "30"
```

Configuring a Storage Backend of the FC Type

The following is an example of the backend configuration file of the FC type for enterprise storage:

```
storage: "oceanstor-san"
name: "fc-155"
namespace: "huawei-csi"
urls:
    - "https://192.168.129.155:8088"
    - "https://192.168.129.156:8088"
pools:
    - "StoragePool001"
```

```
parameters:
protocol: "fc"
maxClientThreads: "30"
```

Configuring a Storage Backend of the NVMe over RoCE Type

The following is an example of the backend configuration file of the NVMe over RoCE type for enterprise storage:

Configuring a Storage Backend of the NVMe over FC Type

The following is an example of the backend configuration file of the NVMe over FC type for enterprise storage:

```
storage: "oceanstor-san"
name: "fc-nvme-155"
namespace: "huawei-csi"
urls:
- "https://192.168.129.155:8088"
- "https://192.168.129.156:8088"
pools:
- "StoragePool001"
parameters:
protocol: "fc-nvme"
maxClientThreads: "30"
```

Configuring a Storage Backend of the NFS Type

The following is an example of the backend configuration file of the NFS type for enterprise storage:

The following is an example of the backend configuration file of the NFS type for distributed storage:

```
storage: "fusionstorage-nas"
name: "nfs-126"
```

```
namespace: "huawei-csi"
urls:
- "https://192.168.129.125:8088"
- "https://192.168.129.126:8088"
pools:
- "StoragePool001"
parameters:
protocol: "nfs"
portals:
- "192.168.128.123"
maxClientThreads: "30"
```

Configuring a Storage Backend of the SCSI Type

The following is an example of the backend configuration file of the SCSI type for distributed storage:

```
storage: "fusionstorage-san"
name: "scsi-155"
namespace: "huawei-csi"
urls:
    - "https://192.168.129.155:8088"
pools:
    - "StoragePool001"
parameters:
protocol: "scsi"
portals:
    - {"hostname01": "192.168.125.21"}
    - {"hostname02": "192.168.125.22"}
maxClientThreads: "30"
```

Configuring a Storage Backend of the DPC Type

The following is an example of the backend configuration file of the DPC type for distributed storage:

```
storage: "fusionstorage-nas"
name: "dpc-155"
namespace: "huawei-csi"
urls:
- "https://192.168.129.155:8088"
- "https://192.168.129.156:8088"
pools:
- "StoragePool001"
parameters:
protocol: "dpc"
maxClientThreads: "30"
```

Configuring Storage Backends of the Dtree Type

The following is an example of the backend configuration file of the Dtree type for enterprise storage:

```
storage: "oceanstor-dtree"
name: "nfs-dtree"
namespace: "huawei-csi"
urls:
- "https://192.168.129.155:8088"
parameters:
protocol: "nfs"
parentname: "parent-filesystem"
portals:
- "192.168.128.155"
maxClientThreads: "30"
```

Configuring Storage Backends of the HyperMetro Type

CSI allows you to provision HyperMetro volumes of the NFS type on the storage side when connecting with OceanStor V6 or OceanStor Dorado V6. When configuring storage backends that work in HyperMetro mode, you need to create two configuration files and create backends one by one.

This example shows how to configure backends of the HyperMetro type for Huawei OceanStor V6 or OceanStor Dorado V6. First, create local storage backend configuration file **nfs-hypermetro-155.yaml**.

```
storage: "oceanstor-nas"
name: "nfs-hypermetro-155"
namespace: "huawei-csi"
urls:
- "https://192.168.129.155:8088"
- "https://192.168.129.156:8088"
pools:
- "StoragePool001"
metrovStorePairID: "f09838237b93c000"
metroBackend: "nfs-hypermetro-157"
parameters:
protocol: "nfs"
portals:
- "192.168.129.155"
maxClientThreads: "30"
```

After the local backend is created, create remote storage backend configuration file **nfs-hypermetro-157.yaml**.

```
storage: "oceanstor-nas"
name: "nfs-hypermetro-157"
namespace: "huawei-csi"
urls:
- "https://192.168.129.157:8088"
- "https://192.168.129.158:8088"
pools:
- "StoragePool001"
metrovStorePairID: "f09838237b93c000"
metroBackend: "nfs-hypermetro-155"
parameters:
protocol: "nfs"
portals:
- "192.168.129.157"
maxClientThreads: "30"
```

7.1.2 Creating a Storage Backend

◯ NOTE

- When oceanctl is used to create a storage backend, the entered account and key information is stored in the Secret object. It is recommended that the customer container platform encrypt the Secret object based on the suggestions of the supplier or K8s community. For details about how to encrypt the Secret object in the K8s community, see Enable Encryption at Rest.
- 2. When a backend is created using a .json file, the backend name of an earlier version may contain uppercase letters or underscores (_). In this case, the old name is remapped to a new name. The mapping process automatically occurs and does not affect the original functions. For example, ABC_123 is mapped to abc-123-fd68e. The mapping rules are as follows:
 - Uppercase letters are converted to lowercase letters.
 - An underscore (_) is converted to a hyphen (-).
 - A 5-digit hash code is added to the end.

Command Description

 Run the following command to obtain the help information about creating a backend.

oceanctl create backend -h

- Run the following command to create a storage backend based on the specified yaml file.
 - oceanctl create backend -f /path/to/backend.yaml -i yaml
- Run the following command to create a storage backend based on the specified json file. The **huawei-csi-configmap** file can be exported only in ison format.

oceanctl create backend -f /path/to/configmap.json -i json

- Run the following command to create a storage backend in the specified namespace.
 - oceanctl create backend -f /path/to/backend.yaml -i yaml -n <namespace>
- Run the following command to create a storage backend and ignore the storage backend name verification, for example, uppercase letters and underscores (_).

oceanctl create backend -f /path/to/backend.yaml -i yaml --not-validate-name

 Run the following command to create a storage backend and specify provisioner. csi.oceanstor.com is the driver name specified during installation. For details, see Step 4.

| \bigcap | NOTE |
|-----------|------|
| | |

This command is used only when a backend is created on the CCE/CCE Agile platform. oceanctl create backend -f /path/to/backend.yaml -i yaml --provisioner=csi.oceanstor.com

Example of Creating a Backend

Step 1 Prepare the backend configuration file, for example, **backend.yaml**. For details, see **7.1.1 Preparing the Storage Backend Configuration File**. To create multiple backends, separate them with ---.

storage: "oceanstor-san" name: "backend-1" namespace: "huawei-csi"

```
- "https://192.168.129.157:8088"
pools:
- "StoragePool001"
parameters:
 protocol: "roce"
 portals:
 - "10.10.30.20"
  - "10.10.30.21"
maxClientThreads: "30"
storage: "oceanstor-san"
name: "backend-2"
namespace: "huawei-csi"
urls:
- "https://192.168.129.158:8088"
pools:
- "StoragePool001"
parameters:
 protocol: "roce"
 portals:
  - "10.10.30.20"
  - "10.10.30.21"
maxClientThreads: "30"
```

Step 2 Run the following command to create a storage backend.

Step 3 Enter the serial number of the backend to be created and enter the account and password.

Step 4 Check the storage backend creation result.

----End

7.2 Querying a Storage Backend

Command Description

 Run the following command to obtain the help information about querying a backend.

oceanctl get backend -h

 Run the following command to query a single storage backend in the default namespace.

oceanctl get backend <backend-name>

• Run the following command to query all storage backends in the specified namespace.

oceanctl get backend -n <namespace>

 Run the following command to format the output. Currently, json, yaml, and wide are supported.

oceanctl get backend <backend-name> -o json

7.3 Updating a Storage Backend

NOTICE

- Currently, only the password in the storage backend information can be updated.
- If the backend account password is updated on the storage device, the CSI plug-in will retry due to login failures. As a result, the account may be locked. If the account is locked, change the password by referring to 11.20 An Account Is Locked After the Password Is Updated on the Storage Device.

Command Description

• Run the following command to obtain the help information about updating a backend.

oceanctl update backend -h

• Run the following command to update the specified storage backend in the default namespace.

oceanctl update backend <backend-name> --password

• Run the following command to update a storage backend in the specified namespace.

oceanctl update backend <backend-name> -n <namespace> --password

Example of Updating a Backend

Step 1 Run the following command to obtain the help information about updating a storage backend.

oceanctl update backend -h Update a backend for Ocean Storage in Kubernetes

Usage:

oceanctl update backend <name> [flags]

```
Examples:
# Update backend account information in default(huawei-csi) namespace
oceanctl update backend <name> --password

# Update backend account information in specified namespace
oceanctl update backend <name> -n namespace --password

Flags:
-h, --help help for backend
-n, --namespace string namespace of resources
--password Update account password
```

Step 2 Run the following command to update a storage backend.

```
# oceanctl update backend backend-1 --password
Please enter this backend user name:admin
Please enter this backend password:
backend/backend-1 updated
```

----End

7.4 Deleting a Storage Backend

Command Description

 Run the following command to obtain the help information about deleting a backend.

oceanctl delete backend -h

 Run the following command to delete the specified storage backend in the default namespace.

oceanctl delete backend <backend-name>

• Run the following command to delete all storage backends in the default namespace.

oceanctl delete backend --all

• Run the following command to delete a storage backend in the specified namespace.

oceanctl delete backend <backend-name...> -n <namespace>

Example of Deleting a Backend

Step 1 Run the following command to obtain information about a storage backend.

Step 2 Run the following command to delete the specified storage backend.

```
# oceanctl delete backend backend-1 backend/backend-1 deleted
```

Step 3 Check the deletion result. If **not found** is displayed, the deletion is successful.

oceanctl get backend backend-1 Error from server (NotFound): backend "backend-1" not found

----End

7.5 (Optional) Adding a Certificate to a Storage Backend

This section describes how to create a certificate for a storage backend. If certificate verification is required for logging in to the storage, you can add a certificate by referring to this section. Currently, you can create a certificate for a storage backend based on the specified .crt or .pem file.

NOTICE

Before creating a certificate for a storage backend, import the prepared certificate to the storage array.

7.5.1 Creating a Certificate for a Storage Backend

Prerequisites

A certificate has been created. Take OceanStor Dorado V6 as an example. For details about how to create a certificate, **click here**.

Command Description

• Run the following command to obtain the help information about querying a certificate.

oceanctl create cert -h

- Run the following command to create a certificate for a single storage backend in the default namespace based on the specified .crt certificate file. oceanctl create cert <name> -f /path/to/cert.crt -b <backend-name>
- Run the following command to create a certificate for a single storage backend in the specified namespace based on the specified .crt certificate file. oceanctl create cert <name> -f /path/to/cert.crt -b <backend-name> -n <namespace>
- Run the following command to create a certificate for a single storage backend in the specified namespace based on the specified .pem certificate file.

oceanctl create cert <name> -f /path/to/cert.pem -b <backend-name> -n <namespace>

Example of Creating a Certificate

- **Step 1** Prepare a certificate file in advance, for example, **cert.crt**.
- **Step 2** Run the following command to obtain information about a storage backend.

Step 3 Run the following command to create a certificate for the specified storage backend.

```
# oceanctl create cert cert-1 -b backend-1 -f /path/to/cert.crt cert/cert-1 created
```

Step 4 Check the certificate creation result.

```
# oceanctl get cert -b backend-1
+------+
| NAMESPACE | NAME | BOUNDBACKEND |
+-----+
| huawei-csi | cert-1 | backend-1 |
+------+
```

----End

7.6 (Optional) Querying a Storage Backend Certificate

Command Description

- Run the following command to obtain the help information about querying a certificate.
 oceanctl get cert -h
- Run the following command to query the certificate of a specified storage backend in the default namespace.
 oceanctl get cert -b <backend-name>
- Run the following command to query the certificate of a specified storage backend in the specified namespace.
 oceanctl get cert -b <backend-name> -n <namespace>

7.7 (Optional) Updating a Storage Backend Certificate

Before updating a certificate, prepare a new certificate file and update the storage backend certificate by following the instructions provided in this section. If the certificate is no longer used, delete the certificate from the storage backend by referring to 7.8 (Optional) Deleting a Storage Backend Certificate.

Command Description

- Run the following command to obtain the help information about updating a certificate.
 - oceanctl update cert -h
- Run the following command to update a certificate for a specified storage backend in the default namespace based on the specified .crt certificate file. oceanctl update cert -b <backend-name> -f /path/to/cert.crt

- Run the following command to update a certificate for a specified storage backend in the specified namespace based on the specified .crt certificate file. oceanctl update cert -b <backend-name> -n <namespace> -f /path/to/cert.crt
- Run the following command to update a certificate for a specified storage backend in the specified namespace based on the specified .pem certificate file.

oceanctl update cert -b <backend-name> -n <namespace> -f /path/to/cert.pem

Example of Updating a Certificate

Step 1 Run the following command to obtain information about a storage backend.

Step 2 Run the following command to check whether the specified storage backend has a certificate.

```
# oceanctl get cert -b backend-1
+------+
| NAMESPACE | NAME | BOUNDBACKEND |
+------+
| huawei-csi | cert-1 | backend-1 |
+------+
```

Step 3 Run the following command to update the certificate of the specified storage backend.

```
# oceanctl update cert -b backend-1 -f /path/to/cert.crt
cert/cert-1 updated
```

----End

7.8 (Optional) Deleting a Storage Backend Certificate

Command Description

 Run the following command to obtain the help information about deleting a certificate.

oceanctl delete cert -h

- Run the following command to delete the certificate of a specified storage backend in the default namespace.
 oceanctl delete cert -b <backend-name>
- Run the following command to delete the certificate of a specified storage backend in the specified namespace.
 oceanctl delete cert -b <backend-name> -n <namespace>

Example of Deleting a Certificate

Step 1 Run the following command to obtain information about a storage backend.

Step 2 Run the following command to obtain information about the certificate of the specified storage backend.

```
# oceanctl get cert -b backend-1
+-----+
| NAMESPACE | NAME | BOUNDBACKEND |
+------+
| huawei-csi | cert-1 | backend-1 |
+------+
```

Step 3 Run the following command to delete the certificate of the specified storage backend.

```
# oceanctl delete cert -b backend-1 cert/cert-1 deleted
```

Step 4 Check the deletion result. If **no cert found** is displayed, the deletion is successful.

```
# oceanctl get cert -b backend-1
Error from server (NotFound): no cert found on backend backend-1 in huawei-csi namespace
```

----End

8 Using Huawei CSI

This chapter describes how to use Huawei CSI to manage the lifecycle of PVs and snapshots.

NOTICE

When block volumes are mapped, Huawei CSI automatically creates associated objects, such as hosts, host groups, and LUN groups, as well as mapping views. If these objects are manually created on the storage, the mapping logic of Huawei CSI will be affected. Therefore, ensure that these objects are deleted before mapping volumes using Huawei CSI.

8.1 Managing a PV/PVC

8.2 Creating a VolumeSnapshot

8.1 Managing a PV/PVC

Based on service requirements, files in containers need to be persistently stored on disks. When the containers are re-built or re-allocated to new nodes, the persistent data can still be used.

To persistently store data on storage devices, you need to use the **PersistentVolume (PV)** and **PersistentVolumeClaim (PVC)** when provisioning containers.

- PV: a piece of storage in the Kubernetes cluster that has been provisioned by an administrator or dynamically provisioned using a **StorageClass**.
- PVC: a request for storage by a user. A PVC consumes PV resources. A PVC can request specific size and access modes. For example, a PV can be mounted in ReadWriteOnce, ReadOnlyMany, or ReadWriteMany mode. For details, see Access Modes.

This section describes how to use Huawei CSI to create, expand the capacity of, and clone a PV/PVC, as well as create a PVC using a snapshot.

8.1.1 Creating a PVC

Huawei CSI allows storage resources (LUNs or file systems) to be created on Huawei storage and provided for containers based on user settings. For details about the supported features, see **Table 2-5** or **Table 2-9**.

A PVC can be created in dynamic volume provisioning or static volume provisioning mode.

- Dynamic volume provisioning does not require a PV to be created in advance. Huawei CSI automatically creates resources required by a PV on storage devices based on a StorageClass. In addition, you can create a PV when creating a PVC.
- Static volume provisioning requires the administrator to create required resources on a storage device in advance and use existing resources by creating a PV. In addition, you can specify the associated PV when creating a PVC.

8.1.1.1 Dynamic Volume Provisioning

Dynamic volume provisioning allows storage volumes to be created on demand. Dynamic volume provisioning depends on the StorageClass objects. The cluster administrator can define multiple StorageClass objects as required and specify a StorageClass that meets service requirements when declaring a PV or PVC. When applying for resources from Huawei storage devices, Huawei CSI creates storage resources that meet service requirements based on the preset StorageClass.

To implement dynamic volume provisioning, perform the following steps:

- Configuring a StorageClass
- Configuring a PVC

8.1.1.1.1 Configuring a StorageClass

A **StorageClass** provides administrators with methods to describe a storage "class". Different types may map to a different group of capability definitions. Kubernetes cluster users can dynamically provision volumes based on a StorageClass.

A StorageClass supports the following parameters.

If SAN storage is used, refer to example file /examples/sc-lun.yaml. If NAS storage is used, refer to example file /examples/sc-fs.yaml.

Table 8-1 StorageClass configuration parameters

| Parameter | Description | Mandator y | Default Value | Remarks |
|--------------------------|---|---------------|--------------------|---|
| metadata.n ame | User-defined name of a StorageClass object. | Yes | - | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |
| provisioner | Name of the provisioner. | Yes | csi.hua wei.com | Set this parameter to the driver name set during Huawei CSI installation. The value is the same as that of driverName in the values.yaml file. |
| reclaimPolic y | Reclamation policy. The following types are supported: • Delete: Resources are automatically reclaimed. • Retain: Resources are manually reclaimed. | No | Delete | Delete: When a PV/PVC is deleted, resources on the storage device are also deleted. Retain: When a PV/PVC is deleted, resources on the storage device are not deleted. |
| allowVolum eExpansion | Whether to allow volume expansion. If this parameter is set to true , the capacity of the PV that uses the StorageClass can be expanded. | No | false | This function can only be used to expand PV capacity but cannot be used to reduce PV capacity. The PV capacity expansion function is supported in Kubernetes 1.14 (alpha) and later versions. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|------------------------|--|---------------|------------------|---|
| parameters. backend | Name of the backend where the resource to be created is located. | No | - | If this parameter is not set, Huawei CSI will randomly select a backend that meets the capacity requirements to create resources. You are advised to specify a backend to |
| | | | | ensure that the created resource is located on the expected backend. |
| parameters. pool | Name of the storage resource pool where the resource to be created is located. If this parameter is set, parameters.back end must also be specified. | No | - | If this parameter is not set, Huawei CSI will randomly select a storage pool that meets the capacity requirements from the selected backend to create resources. You are advised to specify a storage pool to ensure that the created resource is located in the expected storage pool. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|---------------------------|--|---------------|------------------|---|
| parameters. volumeType | Type of the volume to be created. The following types are supported: • lun: A LUN is provisioned on the storage side. • fs: A file system is provisioned on the storage side. • dtree: A volume of the Dtree type is provisioned on the storage side. | Yes | | If NAS storage is used, this parameter must be set to fs. If SAN storage is used, this parameter must be set to lun. If NAS storage of the Dtree type is used, this parameter must be set to dtree. |
| parameters. allocType | Allocation type of the volume to be created. The following types are supported: • thin: Not all required space is allocated during creation. Instead, the space is dynamically allocated based on the usage. • thick: All required space is allocated during creation. | No | - | If this parameter is not specified, thin will be used. Not all required space is allocated during creation. Instead, the space is dynamically allocated based on the usage. OceanStor Dorado V3/V6 does not support thick. FusionStorage Block and OceanStor Pacific series do not support thick. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|---------------------------|--|------------------------------------|------------------|---|
| parameters. fsType | Type of a host file system. The supported types are: • ext2 • ext3 • ext4 • xfs | No | ext4 | This parameter is valid only when volumeType of a StorageClass is set to lun and volumeMode of a PVC is set to Filesystem. |
| parameters. authClient | IP address of the NFS client that can access the volume. This parameter is mandatory when volumeType is set to fs. You can enter the client host name (a full domain name is recommended), client IP address, or client IP address segment. | Conditiona lly mandator y | | The asterisk (*) can be used to indicate any client. If you are not sure about the IP address of the access client, you are advised to use the asterisk (*) to prevent the client access from being rejected by the storage system. If the client host name is used, you are advised to use the full domain name. The IP addresses can be IPv4 addresses, IPv6 addresses, or a combination of IPv4 and IPv6 addresses. You can enter multiple host names, IP address segments and separate them with semicolons (;) or spaces or by pressing Enter. Example: 192.168.0.10;192.16 8.0.0/24;myserver1.t est |

| Parameter | Description | Mandator y | Default Value | Remarks |
|------------------------------------|--|---------------|------------------|--|
| parameters. cloneSpeed | Cloning speed. The value ranges from 1 to 4. | No | 3 | 4 indicates the highest speed. This parameter is available when you clone a PVC or create a PVC using a snapshot. For details, see 8.1.3 Cloning a PVC or 8.1.4 Creating a PVC Using a Snapshot. |
| parameters. applicationT ype | Application type name for creating a LUN or NAS when the backend is OceanStor Dorado V6. | No | - | If the value of volumeType is lun, log in to DeviceManager and choose Services > Block Service > LUN Groups > LUNs > Create to obtain the application type name. If the value of volumeType is fs, log in to DeviceManager and choose Services > File Service > File Systems > Create to obtain the application type name. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|---------------------------------|--|---------------|------------------|--|
| parameters. qos | LUN/NAS QoS settings of the PV on the storage side. The value of the parameter is JSON character strings in dictionary format. A character string is enclosed by single quotation marks and the dictionary key by double quotation marks. Example: '{"maxMBPS": 999, "maxIOPS": 999}' | No | - | For details about the supported QoS configurations, see Table 8-2. |
| parameters. storageQuot a | Quota of a PV on the storage device. This parameter is valid only when NAS is used for connecting to OceanStor Pacific series storage. The value of the parameter is JSON character strings in dictionary format. A character string is enclosed by single quotation marks and the dictionary key by double quotation marks. Example: '{"spaceQuota": "softQuota", "gracePeriod": 100}' | No | | For details about the supported quota configurations, see Table 8-3. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|-----------------------------|--|------------------------------------|------------------|---|
| parameters. hyperMetro | Whether a HyperMetro volume is to be created. This parameter needs to be configured when the backend is of the HyperMetro type. "true": The created volume is a HyperMetro volume. "false": The created volume is a common volume. | Conditiona lly mandator y | false | Set this parameter when the used backend is a HyperMetro backend and a HyperMetro volume needs to be provisioned. |
| parameters. fsPermission | Permission on the directory mounted to a container. | No | - | For details about the configuration format, refer to the Linux permission settings, for example, 777 and 755. All SAN storage devices are supported. Only the following NAS storage devices are supported: OceanStor Dorado V6, OceanStor V6, and OceanStor Pacific series 8.1.2 and later versions. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|---------------------------|---|---------------|------------------|--------------------------------|
| parameters. rootSquash | Controls the root permission of the client. The value can be: root_squash: The client cannot access the storage system as user root. If a client accesses the storage system as user root, the client will be mapped as an anonymous user. no_root_squas h: A client can access the storage system as user root and has the permission of user root. | No | - | Only NAS storage is supported. |
| parameters. allSquash | Whether to retain the user ID (UID) and group ID (GID) of a shared directory. The value can be: • all_squash: The UID and GID of the shared directory are mapped to anonymous users. • no_all_squash: The UID and GID of the shared directory are retained. | No | 1 | Only NAS storage is supported. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--|---|---------------|------------------|---|
| parameters. snapshotDir ectoryVisibil ity | Whether the snapshot directory is visible. The value can be: • visible: The snapshot directory is visible. • invisible: The snapshot directory is invisible. | No | - | Only NAS storage is supported. |
| parameters. reservedSna pshotSpace Ratio | Configures reserved snapshot space. Value type: character string Value range: 0 to 50 | No | - | OceanStor Dorado 6.1.5+ and OceanStor 6.1.5+ NAS storage devices are supported. |
| parameters. description | Configures the description of the created file system or LUN. Value type: character string The value contains 0 to 255 characters. | No | - | Only enterprise storage file systems and LUNs are supported. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|--------------------------|---|---------------|------------------|---|
| mountOptio ns.nfsvers | NFS mount option on the host. The following mount option is supported: nfsvers: protocol version for NFS mounting. The value can be 3, 4, 4.0, or 4.1. | No | | This parameter is optional after the -o parameter when the mount command is executed on the host. The value is in list format. If the NFS version is specified for mounting, NFS 3, 4.0, and 4.1 protocols are supported (the protocol must be supported and enabled on storage devices). If nfsvers is set to 4, the latest protocol version NFS 4 may be used for mounting due to different OS configurations, for example, 4.1. If protocol 4.0 is required, you are advised to set nfsvers to 4.0. |
| mountOptio ns.acl | The DPC namespace supports the ACL function. The DPC client supports POSIX ACL, NFSv4 ACL, and NT ACL authentication. | No | - | The descriptions of acl, aclonlyposix, cnflush, and cflush are for reference only. For details about the parameters, see OceanStor Pacific Series Product Documentation and choose Configuration > Basic Service Configuration Guide for File > Configuring Basic Services (DPC Scenario) > Accessing a DPC Share on a Client > Step 2. |

| Parameter | Description | Mandator y | Default Value | Remarks |
|-----------------------------------|--|---------------|------------------|---|
| mountOptio ns.aclonlypo six | The DPC namespace supports POSIX ACL, and the DPC client supports POSIX ACL authentication. The following protocols support POSIX ACL: DPC, NFSv3, and HDFS. If NFSv4 ACL or NT ACL is used, the DPC client cannot identify the ACL of this type. As a result, the ACL of this | No | - | If aclonlyposix and acl are used together, only acl takes effect. That is, the namespace supports the ACL function. |
| | type does not take effect. | | | |

| Parameter | Description | Mandator y | Default Value | Remarks |
|-------------------------|--|---------------|------------------|--|
| mountOptio | Asynchronous disk flushing mode. That is, data is not flushed to disks immediately when files in the namespace are closed. | No | | Asynchronous flushing mode: When a file is closed, data in the cache is not flushed to storage media in synchronous mode. Instead, data is written from the cache to the storage media in asynchronous flushing mode. After the write service is complete, data is flushed from the cache to disks periodically based on the flushing period. In a multi-client scenario, if concurrent operations are performed on the same file, the file size update is affected by the disk flushing period. That is, the file size is updated only after the disk flushing is complete. Generally, the update is completed within several seconds. Synchronous I/Os are not affected by the disk flushing period. |
| mountOptio ns.cflush | Synchronous disk flushing mode. That is, data is flushed to disks immediately when files in the namespace are closed. | No | - | By default, the synchronous disk flushing mode is used. |

Table 8-2 Supported QoS configurations

| Storage Type | Parameter | Description | Remarks |
|-------------------------------------|------------------|---|---|
| OceanStor V3/ OceanStor V5 | IOTYPE | Read/write type. | This parameter is optional. If it is not specified, the default value of the storage backend is used. For details, see related storage documents. The value can be: • 0: read I/O • 1: write I/O • 2: read and write I/Os |
| | MAXBAND WIDTH | Maximum bandwidth. This is a restriction policy parameter. | The value is an integer greater than 0, expressed in MB/s. |
| | MINBAND WIDTH | Minimum bandwidth. This is a protection policy parameter. | The value is an integer greater than 0, expressed in MB/s. |
| | MAXIOPS | Maximum IOPS. This is a restriction policy parameter. | The value is an integer greater than 0. |
| | MINIOPS | Minimum IOPS. This is a protection policy parameter. | The value is an integer greater than 0. |
| | LATENCY | Maximum latency. This is a protection policy parameter. | The value is an integer greater than 0, expressed in ms. |
| OceanStor Dorado V3 | IOTYPE | Read/write type. | The value can be: • 2: read and write I/Os |
| | MAXBAND WIDTH | Maximum bandwidth. This is a restriction policy parameter. | The value is an integer ranging from 1 to 999999999, expressed in MB/s. |
| | MAXIOPS | Maximum IOPS. This is a restriction policy parameter. | The value is an integer ranging from 100 to 999999999. |

| Storage Type | Parameter | Description | Remarks |
|--|------------------|---|---|
| OceanStor Dorado V6/ OceanStor V6 | IOTYPE | Read/write type. | The value can be: • 2: read and write I/Os |
| | MAXBAND WIDTH | Maximum bandwidth. This is a restriction policy parameter. | The value is an integer ranging from 1 to 999999999, expressed in MB/s. |
| | MINBAND WIDTH | Minimum bandwidth. This is a protection policy parameter. | The value is an integer ranging from 1 to 999999999, expressed in MB/s. |
| | MAXIOPS | Maximum IOPS. This is a restriction policy parameter. | The value is an integer ranging from 100 to 999999999. |
| | MINIOPS | Minimum IOPS. This is a protection policy parameter. | The value is an integer ranging from 100 to 999999999. |
| | LATENCY | Maximum latency. This is a protection policy parameter. | The value can be 0.5 or 1.5 , expressed in ms. |
| FusionStor age/ OceanStor Pacific series | maxMBPS | Maximum bandwidth. This is a restriction policy parameter. | This parameter is mandatory. The value is an integer greater than 0, expressed in MB/s. For details about the maximum value, see the actual limit of the storage device. For example, the maximum value of OceanStor Pacific NAS is 1073741824. |
| | maxIOPS | Maximum IOPS. This is a restriction policy parameter. | This parameter is mandatory. The value is an integer greater than 0. For details about the maximum value, see the actual limit of the storage device. For example, the maximum value of OceanStor Pacific NAS is 1073741824000. |

□ NOTE

- vStore users of OceanStor V3/OceanStor V5 cannot configure QoS policies.
- The QoS configuration takes effect only on the newly created PVC. QoS cannot be added automatically for PVCs with the same StorageClass name that have been provisioned.

Table 8-3 Supported quota configurations

| Parameter | Description | Remarks |
|-------------|---|---|
| spaceQuota | File quota type. | This parameter is mandatory. Only softQuota or hardQuota can be configured. |
| gracePeriod | Grace period allowed when the soft quota is configured. | This parameter is conditionally optional only when spaceQuota is set to softQuota . |
| | | The value is an integer ranging from 0 to 4294967294. |

For details about how to configure a StorageClass in typical scenarios, see the following examples:

- Setting the Backend and Storage Pool in a StorageClass
- Setting the NFS Access Mode in a StorageClass
- Setting a Dtree Type in a StorageClass
- Setting the Local File System Access Mode in a StorageClass
- Setting the DPC Access Mode in a StorageClass
- Setting an Application Type in a StorageClass
- Setting a Soft Quota in a StorageClass
- Setting HyperMetro in a StorageClass
- Setting the Permission on a Mount Directory in a StorageClass
- Setting QoS in a StorageClass
- Configuring a StorageClass on the CCE/CCE Agile Platform

Setting the Backend and Storage Pool in a StorageClass

If multiple Huawei backends are configured in a Kubernetes cluster or a Huawei backend provides multiple storage pools, you are advised to configure the specified backend and storage pool information in the StorageClass. This prevents Huawei CSI from randomly selecting backends and storage pools and ensures that the storage device where the volume resides complies with the plan.

For details about how to set the backend and storage pool for SAN storage, see the following configuration example.

kind: StorageClass apiVersion: storage.k8s.io/v1 metadata:

```
name: mysc
provisioner: csi.huawei.com
allowVolumeExpansion: true
parameters:
backend: "iscsi_san_181"
pool: "pool001"
volumeType: lun
allocType: thin
```

For details about how to set the backend and storage pool for NAS storage, see the following configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
allowVolumeExpansion: true
parameters:
backend: "iscsi_nas_181"
pool: "pool001"
volumeType: fs
allocType: thin
authClient: "*"
```

Setting the NFS Access Mode in a StorageClass

When a container uses an NFS file system as a storage resource, refer to the following configuration example. In this example, NFS version 4.1 is specified for mounting.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_nas_181
pool: pool001
volumeType: fs
allocType: thin
authClient: "192.168.0.10;192.168.0.0/24;myserver1.test" #use * for all client
mountOptions:
- nfsvers=4.1
```

Setting a Dtree Type in a StorageClass

When a container uses a Dtree as a storage resource, refer to the following configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_dtree
volumeType: dtree
allocType: thin
authClient: "*"
mountOptions:
- nfsvers=4.1
```

Setting the Local File System Access Mode in a StorageClass

If a container uses a LUN of enterprise storage or distributed storage as a storage resource and a file system needs to be formatted as a local file system, refer to the following example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: iscsi_lun_181
pool: pool001
volumeType: lun
allocType: thin
fsType: xfs
```

Setting the DPC Access Mode in a StorageClass

If a container uses OceanStor Pacific series storage and the storage supports DPC-based access, you can configure mounting parameters for DPC-based access in the StorageClass. In this example, **acl** is used as the authentication parameter for mounting, and **cnflush** is used to set the asynchronous disk flushing mode.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_dpc_101
pool: pool001
volumeType: fs
allocType: thin
authClient: "*" #use * for all client
mountOptions:
- acl
- cnflush
```

Setting an Application Type in a StorageClass

When a container uses a LUN of OceanStor Dorado V6 as the storage, if the default application type of the storage cannot meet the I/O model requirements of some services (for example, the container provides the database OLAP service), you can configure an application type in the StorageClass to improve storage performance. For details about the application types to be used, see the product documentation of the corresponding storage product.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: iscsi_lun_181
pool: pool001
volumeType: lun
allocType: thin
fsType: xfs
applicationType: Oracle_OLAP
```

Setting a Soft Quota in a StorageClass

If a container uses a file system of OceanStor Pacific series as the storage, you can configure a soft quota in the StorageClass. The following is a configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_pacific_101
pool: pool001
volumeType: fs
allocType: fs
allocType: thin
authClient: "*"
storageQuota: '{"spaceQuota": "softQuota", "gracePeriod": 100}'
mountOptions:
- nfsvers=3
```

Setting QoS in a StorageClass

When containers use enterprise storage or distributed storage as storage resources, you can set QoS for the storage resources used by containers to ensure that the storage read and write operations of these containers meet certain service levels.

Storage devices of different models or versions support different QoS settings. For details about how to find the configuration items of the corresponding storage devices, see **Table 8-2**. In this example, the backend is OceanStor Dorado V6. For other storage devices, refer to this example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: iscsi_qos_181
pool: pool001
volumeType: lun
allocType: thin
fsType: xfs
qos: '{"IOTYPE": 2, "MINIOPS": 1000}'
```

Setting HyperMetro in a StorageClass

When a container uses an NFS HyperMetro file system as a storage resource, refer to the following configuration example. In this example, the used backend supports HyperMetro, and **hyperMetro** is set to **true**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_hypermetro_dorado_181
pool: pool001
volumeType: fs
hyperMetro: "true"
allocType: thin
authClient: "*"
```

NOTICE

Before provisioning a NAS HyperMetro volume, you need to configure the HyperMetro relationship between two storage devices, including the remote device, HyperMetro domain, and the like. The HyperMetro domain of the file system can only work in HyperMetro mode. For details about the configuration operation, see the product documentation of the corresponding storage model.

Setting the Permission on a Mount Directory in a StorageClass

To modify the permission on a mount directory in a container, you can configure the directory permission in a StorageClass. The following is a configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
allowVolumeExpansion: true
parameters:
volumeType: fs
allocType: thin
authClient: "*"
fsPermission: "777"
rootSquash: "no_root_squash" # This parameter supports only NAS storage.
allSquash: "no_all_squash" # This parameter supports only NAS storage.
```

After the StorageClass configuration is complete, perform the following steps to create a StorageClass.

Step 1 Run the following command to create a StorageClass based on the .yaml file.

```
# kubectl create -f mysc.yaml
storageclass.storage.k8s.io/mysc created
```

Step 2 Run the following command to view the information about the created StorageClass.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION AGE
mysc csi.huawei.com Delete Immediate false 34s
```

After creating a StorageClass, you can use the StorageClass to create a PV or PVC.

----End

CAUTION

Pay attention to the following when using a StorageClass:

- Do not delete a StorageClass that is being used by a PV. Otherwise, the StorageClass information of the PV will be missing. As a result, an error occurs during mounting on the host.
- Modifications to a StorageClass do not take effect on existing PVs. You need to
 delete these PVs and create them again using the modified StorageClass to
 apply the modified parameters.

Configuring a StorageClass on the CCE/CCE Agile Platform

Create a StorageClass of the NAS type on the CCE/CCE Agile platform. The following is a configuration example. The value of **provisioner** must be the same as that of **driverName** in the **values.yaml** file.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
annotations:
storageclass.kubernetes.io/storageType: file
provisioner: csi.oceanstor.com
allowVolumeExpansion: true
parameters:
volumeType: fs
allocType: thin
authClient: "*"
```

Create a StorageClass of the Block type on the CCE/CCE Agile platform. The following is a configuration example. The value of **provisioner** must be the same as that of **driverName** in the **values.yaml** file.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
annotations:
storageclass.kubernetes.io/storageType: block
provisioner: csi.oceanstor.com
allowVolumeExpansion: true
parameters:
volumeType: lun
allocType: thin
```

8.1.1.1.2 Configuring a PVC

After configuring a StorageClass, you can use the StorageClass to configure a PVC. For details about the PVC configuration template, see example file **pvc*.yaml** in the **examples** directory in Huawei CSI software package.

Table 8-4 Parameters in the pvc*.yaml file

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|-------------------|------------------------------------|---------------|-------------------|--|
| metadat a.name | User-defined name of a PVC object. | Yes | - | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|---|--|---------------|-------------------|--|
| spec.volu meMode | Volume mode. This parameter is optional. When LUN volumes are used, the following types are supported: • Filesystem: local file system. • Block: raw device. | No | Filesyst em | This parameter takes effect when a PV is mounted. The default value is Filesystem. • Filesystem indicates that a container accesses a PV using a local file system. The local file system type is specified by the fsType field in the specified StorageClass. Storage of the Dtree type also uses this parameter. • Block indicates that a PV is accessed in raw volume mode. |
| spec.stor ageClass Name | Name of the StorageClass object. | Yes | - | Name of the StorageClass object required by services. |
| spec.reso urces.req uests.sto rage | Size of the volume to be created. The format is ***Gi and the unit is GiB. The size must be an integer multiple of 512 bytes. | Yes | 10Gi | The PVC capacity depends on storage specifications and host specifications. For example, OceanStor Dorado 6.1.2 or OceanStor Pacific series 8.1.0 is connected to CentOS 7. If ext4 file systems are used, see Table 8-5. If XFS file systems are used, see Table 8-6. If NFS or raw devices are used, the capacity must meet the specifications of the used Huawei storage device model and version. If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications. |

| Paramet er | Description | Mandat ory | Defaul t Value | Remarks |
|----------------------|---|---------------|-----------------------|---|
| spec.acce ssModes | Access mode of the volume. RWO (ReadWriteOnc e): A volume can be mounted to a node in read/write mode. This mode also allows multiple Pods running on the same node to access the volume. ROX (ReadOnlyMany): A volume can be mounted to multiple nodes in read-only mode. RWX (ReadWriteMan y): A volume can be mounted to multiple nodes in read/write mode. RWOP (ReadWriteOnc ePod): A volume can only be mounted to a single Pod in read/write mode. Kubernetes 1.22 and later versions support this feature. | Yes | ReadW riteOnc e | RWO/ROX/RWOP: supported by all types of volumes. RWOP is supported only by Kubernetes 1.22 and later versions. Check whether this feature is enabled for your Kubernetes cluster by referring to 10.5 Enabling the ReadWriteOncePod Feature Gate. RWX: supported by volumes whose volumeMode is set to Block or NFS. |

Table 8-5 ext4 capacity specifications

| Storage Type | Storage Specification s | ext4 Specifications | CSI Specifications |
|--------------------------------|-------------------------------|------------------------|-----------------------|
| OceanStor Dorado 6.1.2 | 512 Ki to 256 Ti | 50 Ti | 512 Ki to 50 Ti |
| OceanStor Pacific series 8.1.0 | 64 Mi to 512 Ti | 50 Ti | 64 Mi to 50 Ti |

Table 8-6 XFS capacity specifications

| Storage Type | Storage Specifications | XFS Specifications | CSI Specifications |
|--------------------------------|---------------------------|-----------------------|-----------------------|
| OceanStor Dorado 6.1.2 | 512 Ki to 256 Ti | 500 Ti | 512 Ki to 500 Ti |
| OceanStor Pacific series 8.1.0 | 64 Mi to 512 Ti | 500 Ti | 64 Mi to 500 Ti |

Step 1 Based on service requirements, modify specific parameters by referring to the description in this section and the PVC configuration file example to generate the PVC configuration file to be created, for example, the **mypvc.yaml** file in this example.

kind: PersistentVolumeClaim apiVersion: v1 metadata: name: mypvc spec: accessModes: - ReadWriteOnce volumeMode: Filesystem storageClassName: mysc resources: requests: storage: 100Gi

Step 2 Run the following command to create a PVC using the configuration file.

kubectl create -f mypvc.yaml persistentvolumeclaim/mypvc created

Step 3 After a period of time, run the following command to view the information about the created PVC. If the PVC status is **Bound**, the PVC has been created and can be used by a Pod.

```
# kubectl get pvc mypvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWO mysc 12s
```

! CAUTION

- After the PVC is created, if the PVC is in the Pending state after a long time (for example, one minute), refer to 11.6 When a PVC is Created, the PVC is in the Pending State.
- You are advised to create or delete a maximum of 100 PVCs in a batch.

----End

After a PVC is created, you can use the PVC to create a Pod. The following is a simple example of using a PVC. In this example, the created Pod uses the newly created *mypvc*.

```
apiVersion: apps/v1
kind: Deployment
metadata:
name: nginx-deployment
 selector:
  matchLabels:
   app: nginx
 replicas: 2
 template:
  metadata:
   labels:
    app: nginx
  spec:
   containers:
    - image: nginx:alpine
     name: container-0
     volumeMounts:
     - mountPath: /tmp
      name: pvc-mypvc
   restartPolicy: Always
   volumes:
    - name: pvc-mypvc
     persistentVolumeClaim:
      claimName: mypvc
                                        # name of PVC
```

8.1.1.2 Manage Volume Provisioning

Manage Volume Provisioning allows administrators to use resources created on storage as PVs and supports features of dynamic volumes, such as capacity expansion, snapshot, and clone. This is a custom capability of Huawei CSI. This feature applies to the following scenarios:

- In the reconstruction containerized applications, existing storage volumes need to be used.
- The Kubernetes cluster is rebuilt.
- Storage data is migrated in disaster recovery (DR) scenarios.

Prerequisites

- You have registered the storage where the volume to be managed resides with CSI.
- You have logged in to the storage device to obtain the name and capacity of the volume to be managed.

8.1.1.2.1 Configuring a StorageClass

A **StorageClass** provides administrators with methods to describe a storage "class". Different types may map to a different group of capability definitions. Kubernetes cluster users can dynamically provision volumes based on a StorageClass.

A StorageClass supports the following parameters.

If SAN storage is used, refer to example file **/examples/sc-lun.yaml**. If NAS storage is used, refer to example file **/examples/sc-fs.yaml**.

Table 8-7 StorageClass configuration parameters

| Parameter | Description | Mandato ry | Default Value | Remarks |
|-------------------|---|---------------|--------------------|---|
| metadata.na me | User-defined name of a StorageClass object. | Yes | - | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |
| provisioner | Name of the provisioner. | Yes | csi.huaw ei.com | Set this parameter to the driver name set during Huawei CSI installation. The value is the same as that of driverName in the values.yaml file. |
| reclaimPolic y | Reclamation policy. The following types are supported: • Delete: Resources are automatically reclaimed. • Retain: Resources are manually reclaimed. | Yes | - | Delete: When a PV/PVC is deleted, resources on the storage device are also deleted. Retain: When a PV/PVC is deleted, resources on the storage device are not deleted. |

| Parameter | Description | Mandato ry | Default Value | Remarks |
|---------------------------|---|---------------|------------------|---|
| allowVolum eExpansion | Whether to allow volume expansion. If this parameter is set to true , the capacity of the PV that uses the StorageClass can be expanded. | No | false | This function can only be used to expand PV capacity but cannot be used to reduce PV capacity. The PV capacity expansion function is supported in Kubernetes 1.14 (alpha) and later versions. |
| parameters. backend | Name of the backend where the resource to be created is located. | No | - | If this parameter is not set, Huawei CSI will randomly select a backend that meets the capacity requirements to create resources. You are advised to specify a backend to ensure that the created resource is located on the expected backend. |
| parameters. volumeType | Type of the volume to be created. The following types are supported: • lun: A LUN is provisioned on the storage side. • fs: A file system is provisioned on the storage side. | Yes | - | If NAS storage is used, this parameter must be set to fs. If SAN storage is used, this parameter must be set to lun. |

| Parameter | Description | Mandato ry | Default Value | Remarks |
|-----------------------------|---|---------------|------------------|---|
| parameters.f sType | Type of a host file system. The supported types are: • ext2 • ext3 • ext4 • xfs | No | ext4 | This parameter is valid only when volumeType of a StorageClass is set to lun and volumeMode of a PVC is set to Filesystem. |
| parameters.f sPermission | Permission on the directory mounted to a container. | No | - | For details about the configuration format, refer to the Linux permission settings, for example, 777 and 755. This parameter is available when volumeType is set to lun. |
| mountOptio ns.nfsvers | NFS mount option on the host. The following mount option is supported: nfsvers: protocol version for NFS mounting. The value can be 3, 4, 4.0, or 4.1. | No | | This parameter is optional after the -o parameter when the mount command is executed on the host. The value is in list format. If the NFS version is specified for mounting, NFS 3, 4.0, and 4.1 protocols are supported (the protocol must be supported and enabled on storage devices). If nfsvers is set to 4, the latest protocol version NFS 4 may be used for mounting due to different OS configurations, for example, 4.1. If protocol 4.0 is required, you are advised to set nfsvers to 4.0. |

| Parameter | Description | Mandato ry | Default Value | Remarks |
|-----------------------------------|---|---------------|------------------|---|
| mountOptio ns.acl | The DPC namespace supports the ACL function. The DPC client supports POSIX ACL, NFSv4 ACL, and NT ACL authentication. | No | | The descriptions of acl, aclonlyposix, cnflush, and cflush are for reference only. For details about the parameters, see OceanStor Pacific Series Product Documentation and choose Configuration > Basic Service Configuration Guide for File > Configuring Basic Services (DPC Scenario) > Accessing a DPC Share on a Client > Step 2. |
| mountOptio ns.aclonlypo six | The DPC namespace supports POSIX ACL, and the DPC client supports POSIX ACL authentication. The following protocols support POSIX ACL: DPC, NFSv3, and HDFS. If NFSv4 ACL or NT ACL is used, the DPC client cannot identify the ACL of this type. As a result, the ACL of this type does not take effect. | No | - | If aclonlyposix and acl are used together, only acl takes effect. That is, the namespace supports the ACL function. |

| Parameter | Description | Mandato ry | Default Value | Remarks |
|-------------------------|--|---------------|------------------|--|
| mountOptio | Asynchronous disk flushing mode. That is, data is not flushed to disks immediately when files in the namespace are closed. | No | - | Asynchronous flushing mode: When a file is closed, data in the cache is not flushed to storage media in synchronous mode. Instead, data is written from the cache to the storage media in asynchronous flushing mode. After the write service is complete, data is flushed from the cache to disks periodically based on the flushing period. In a multi-client scenario, if concurrent operations are performed on the same file, the file size update is affected by the disk flushing period. That is, the file size is updated only after the disk flushing is complete. Generally, the update is completed within several seconds. Synchronous I/Os are not affected by the disk flushing period. |
| mountOptio ns.cflush | Synchronous disk flushing mode. That is, data is flushed to disks immediately when files in the namespace are closed. | No | - | By default, the synchronous disk flushing mode is used. |

Table 8-8 Supported QoS configurations

| Storage Type | Paramete r | Description | Remarks |
|---------------------------------------|-------------------|--|---|
| OceanSt or V3/ OceanSt or V5 | or V3/ OceanSt | | This parameter is optional. If it is not specified, the default value of the storage backend is used. For details, see related storage documents. The value can be: • 0: read I/O • 1: write I/O • 2: read and write I/Os |
| | MAXBAN DWIDTH | Maximum bandwidth. This is a restriction policy parameter. | The value is an integer greater than 0, expressed in MB/s. |
| | MINBAND WIDTH | Minimum bandwidth. This is a protection policy parameter. | The value is an integer greater than 0, expressed in MB/s. |
| | MAXIOPS | Maximum IOPS. This is a restriction policy parameter. | The value is an integer greater than 0. |
| | MINIOPS | Minimum IOPS. This is a protection policy parameter. | The value is an integer greater than 0. |
| | LATENCY | Maximum latency. This is a protection policy parameter. | The value is an integer greater than 0, expressed in ms. |
| OceanSt or Dorado | IOTYPE | Read/write type. | The value can be: • 2: read and write I/Os |
| V3 | MAXBAN DWIDTH | Maximum bandwidth. This is a restriction policy parameter. | The value is an integer ranging from 1 to 999999999, expressed in MB/s. |
| | MAXIOPS | Maximum IOPS. This is a restriction policy parameter. | The value is an integer ranging from 100 to 9999999999999999999999999999999999 |
| OceanSt or Dorado | IOTYPE | Read/write type. | The value can be: • 2: read and write I/Os |
| V6/ OceanSt or V6 | MAXBAN DWIDTH | Maximum bandwidth. This is a restriction policy parameter. | The value is an integer ranging from 1 to 999999999, expressed in MB/s. |

| Storage Type | Paramete r | Description | Remarks |
|-------------------------------|------------------|--|---|
| | MINBAND WIDTH | Minimum bandwidth. This is a protection policy parameter. | The value is an integer ranging from 1 to 999999999, expressed in MB/s. |
| | MAXIOPS | Maximum IOPS. This is a restriction policy parameter. | The value is an integer ranging from 100 to 9999999999999999999999999999999999 |
| | MINIOPS | Minimum IOPS. This is a protection policy parameter. | The value is an integer ranging from 100 to 9999999999999999999999999999999999 |
| | LATENCY | Maximum latency. This is a protection policy parameter. | The value can be 0.5 or 1.5 , expressed in ms. |
| FusionSt orage/ OceanSt | maxMBPS | Maximum bandwidth. This is a restriction policy parameter. | This parameter is mandatory. The value is an integer greater than 0, expressed in MB/s. |
| or Pacific series | maxIOPS | Maximum IOPS. This is a restriction policy parameter. | This parameter is mandatory. The value is an integer greater than 0. |

□ NOTE

- vStore users of OceanStor V3/OceanStor V5 cannot configure QoS policies.
- The QoS configuration takes effect only on the newly created PVC. QoS cannot be added automatically for PVCs with the same StorageClass name that have been provisioned.

Table 8-9 Supported quota configurations

| Parameter | Description | Remarks |
|-------------|---|--|
| spaceQuota | File quota type. | This parameter is mandatory. Only softQuota or hardQuota can be configured. |
| gracePeriod | Grace period allowed when the soft quota is configured. | This parameter is conditionally optional only when spaceQuota is set to softQuota. |
| | | The value is an integer ranging from 0 to 4294967294. |

For details about how to configure a StorageClass in typical scenarios, see the following examples:

- Setting the Backend and Storage Pool in a StorageClass
- Setting the NFS Access Mode in a StorageClass
- Setting the Local File System Access Mode in a StorageClass
- Setting the DPC Access Mode in a StorageClass
- Setting an Application Type in a StorageClass
- Setting a Soft Quota in a StorageClass
- Setting HyperMetro in a StorageClass
- Setting the Permission on a Mount Directory in a StorageClass
- Setting QoS in a StorageClass

Setting the Backend and Storage Pool in a StorageClass

If multiple Huawei backends are configured in a Kubernetes cluster or a Huawei backend provides multiple storage pools, you are advised to configure the specified backend and storage pool information in the StorageClass. This prevents Huawei CSI from randomly selecting backends and storage pools and ensures that the storage device where the volume resides complies with the plan.

For details about how to set the backend and storage pool for SAN storage, see the following configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
allowVolumeExpansion: true
parameters:
backend: "iscsi_san_181"
pool: "pool001"
volumeType: lun
allocType: thin
```

For details about how to set the backend and storage pool for NAS storage, see the following configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
allowVolumeExpansion: true
parameters:
backend: "iscsi_nas_181"
pool: "pool001"
volumeType: fs
allocType: thin
authClient: "*"
```

Setting the NFS Access Mode in a StorageClass

When a container uses an NFS file system as a storage resource, refer to the following configuration example. In this example, NFS version 4.1 is specified for mounting.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_nas_181
pool: pool:pool001
volumeType: fs
allocType: thin
authClient: "192.168.0.10;192.168.0.0/24;myserver1.test" #use * for all client
mountOptions:
- nfsvers=4.1
```

Setting the Local File System Access Mode in a StorageClass

If a container uses a LUN of enterprise storage or distributed storage as a storage resource and a file system needs to be formatted as a local file system, refer to the following example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: iscsi_lun_181
pool: pool001
volumeType: lun
allocType: thin
fsType: xfs
```

Setting the DPC Access Mode in a StorageClass

If a container uses OceanStor Pacific series storage and the storage supports DPC-based access, you can configure mounting parameters for DPC-based access in the StorageClass. In this example, **acl** is used as the authentication parameter for mounting, and **cnflush** is used to set the asynchronous disk flushing mode.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_dpc_101
pool: pool001
volumeType: fs
allocType: thin
authClient: "*" #use * for all client
mountOptions:
- acl
- cnflush
```

Setting an Application Type in a StorageClass

When a container uses a LUN of OceanStor Dorado V6 as the storage, if the default application type of the storage cannot meet the I/O model requirements of some services (for example, the container provides the database OLAP service), you can configure an application type in the StorageClass to improve storage performance. For details about the application types to be used, see the product documentation of the corresponding storage product.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: iscsi_lun_181
pool: pool001
volumeType: lun
allocType: thin
fsType: xfs
applicationType: Oracle_OLAP
```

Setting a Soft Quota in a StorageClass

If a container uses a file system of OceanStor Pacific series as the storage, you can configure a soft quota in the StorageClass. The following is a configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_pacific_101
pool: pool001
volumeType: fs
allocType: thin
authClient: "*"
storageQuota: '{"spaceQuota": "softQuota", "gracePeriod": 100}'
mountOptions:
- nfsvers=3
```

Setting QoS in a StorageClass

When containers use enterprise storage or distributed storage as storage resources, you can set QoS for the storage resources used by containers to ensure that the storage read and write operations of these containers meet certain service levels.

Storage devices of different models or versions support different QoS settings. For details about how to find the configuration items of the corresponding storage devices, see **Table 8-8**. In this example, the backend is OceanStor Dorado V6. For other storage devices, refer to this example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: iscsi_qos_181
pool: pool001
volumeType: lun
allocType: thin
fsType: xfs
qos: '{"IOTYPE": 2, "MINIOPS": 1000}'
```

Setting HyperMetro in a StorageClass

When a container uses an NFS HyperMetro file system as a storage resource, refer to the following configuration example. In this example, the used backend supports HyperMetro, and **hyperMetro** is set to **true**.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
parameters:
backend: nfs_hypermetro_dorado_181
pool: pool001
volumeType: fs
hyperMetro: "true"
allocType: thin
authClient: "*"
```

NOTICE

Before provisioning a NAS HyperMetro volume, you need to configure the HyperMetro relationship between two storage devices, including the remote device, HyperMetro domain, and the like. The HyperMetro domain of the file system can only work in HyperMetro mode. For details about the configuration operation, see the product documentation of the corresponding storage model.

Setting the Permission on a Mount Directory in a StorageClass

To modify the permission on a mount directory in a container, you can configure the directory permission in a StorageClass. The following is a configuration example.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: mysc
provisioner: csi.huawei.com
allowVolumeExpansion: true
parameters:
volumeType: fs
allocType: thin
authClient: "*"
fsPermission: "777"
rootSquash: "no_root_squash" # This parameter supports only NAS storage.
allSquash: "no_all_squash" # This parameter supports only NAS storage.
```

After the StorageClass configuration is complete, perform the following steps to create a StorageClass.

Step 1 Run the following command to create a StorageClass based on the .yaml file.

```
# kubectl create -f mysc.yaml
storageclass.storage.k8s.io/mysc created
```

Step 2 Run the following command to view the information about the created StorageClass.

```
# kubectl get sc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE ALLOWVOLUMEEXPANSION AGE
mysc csi.huawei.com Delete Immediate false 34s
```

After creating a StorageClass, you can use the StorageClass to create a PV or PVC.

----End

CAUTION

Pay attention to the following when using a StorageClass:

- Do not delete a StorageClass that is being used by a PV. Otherwise, the StorageClass information of the PV will be missing. As a result, an error occurs during mounting on the host.
- Modifications to a StorageClass do not take effect on existing PVs. You need to delete these PVs and create them again using the modified StorageClass to apply the modified parameters.

8.1.1.2.2 Configuring a PVC

After configuring a StorageClass, you can use the StorageClass to configure a PVC. For details about the PVC configuration template, see example file **pvc-manager.yaml** in the **examples** directory in Huawei CSI software package.

Table 8-10 Parameters in the pvc-manager.yaml file

| Paramete r | Description | Mandato ry | Default Value | Remarks |
|------------------------------|---|---------------|---|---|
| metadata. annotatio ns | PVC object annotations. Set the following parameters: • Driver name/ manageVolum eName: volume name on the storage. • Driver name/ manageBacke ndName: name of the backend to which the volume belongs. | Yes | csi.huaw ei.com/ manage VolumeN ame: * csi.huaw ei.com/ manage Backend Name: * | For details about how to obtain Driver name, see 4.1.1.4 csiDriver Parameters. Driver name manageVolumeName: name of an existing volume on the storage. Only English characters are supported. Driver name manageBackendName: name of the storage backend in CSI. You can run the oceanctl get backend nhuawei-csi command to obtain the backend name. |

| Paramete r | Description | Mandato ry | Default Value | Remarks |
|---------------------|--|---------------|------------------|--|
| metadata. labels | PVC object labels. | Yes | - | Format: provisioner: <i>Driver name specified during installation</i> |
| | | | | Example: provisioner: csi.huawei.com |
| | | | | This parameter takes effect when a PVC is created. It is used to listen to PVC resources and obtain information about metadata.annotation s. |
| metadata. name | User-defined name of a PVC object. | Yes | - | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |

| Paramete r | Description | Mandato ry | Default Value | Remarks |
|---------------------|--|---------------|------------------|---|
| spec.volu meMode | Volume mode. This parameter is optional. When LUN volumes are used, the following types are supported: • Filesystem: local file system. • Block: raw device. | No | Filesyste m | This parameter takes effect when a PV is mounted. • Filesystem indicates that a container accesses a PV using a local file system. The local file system type is specified by the fsType field in the specified StorageClass. • Block indicates that a PV is accessed in raw volume mode. |

| Paramete r | Description | Mandato ry | Default Value | Remarks |
|-------------------------------|--|---------------|------------------|---|
| r | NOTE This parameter takes effect when a PV is mounted. The use method of this parameter must be the same as that of the managed volume. If a volume is used as a raw volume before being managed, volumeMode must be set to Block. If a volume is used in ext2, ext3, or ext4 mode before being managed, volumeMode must be set to Filesystem and fsType in the | ry | Value | |
| | StorageClass must be set to ext2, ext3, or ext4. If a volume is used in XFS mode before being managed, volumeMode must be set to Filesystem and fsType in the StorageClass must be set to xfs. | | | |
| spec.stora geClassN ame | Name of the StorageClass object. | Yes | - | The configuration of the StorageClass must be the same as that of the managed volume. |

| Paramete r | Description | Mandato ry | Default Value | Remarks |
|---|---|---------------|------------------|---|
| spec.reso urces.requ ests.stora ge | Size of the volume to be created. The format is ***Gi and the unit is GiB. The size must be an integer multiple of 512 bytes. | Yes | _ | The PVC capacity depends on storage specifications and host specifications. For example, OceanStor Dorado 6.1.2 or OceanStor Pacific series 8.1.0 is connected to CentOS 7. If ext4 file systems are used, see Table 8-11. If XFS file systems are used, see Table 8-12. If NFS or raw devices are used, the capacity must meet the specifications of the used Huawei storage device model and version. If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications. |

| Paramete r | Description | Mandato ry | Default Value | Remarks |
|----------------------|--|---------------|------------------|---|
| spec.acces sModes | Access mode of the volume. RWO (ReadWriteOnce): A volume can be mounted to a node in read/write mode. This mode also allows multiple Pods running on the same node to access the volume. ROX (ReadOnlyMany): A volume can be mounted to multiple nodes in read-only mode. RWX (ReadWriteMany): A volume can be mounted to multiple nodes in read/write mode. RWOP (ReadWriteOncePod): A volume can only be mounted to a single Pod in read/write mode. Kubernetes 1.22 and later versions support this feature. | Yes | ReadWriteOnce | RWO/ROX/RWOP: supported by all types of volumes. RWOP is supported only by Kubernetes 1.22 and later versions. Check whether this feature is enabled for your Kubernetes cluster by referring to 10.5 Enabling the ReadWriteOncePod Feature Gate. RWX: supported by volumes whose volumeMode is set to Block or NFS. |

Table 8-11 ext4 capacity specifications

| Storage Type | Storage Specification s | ext4 Specifications | CSI Specifications |
|--------------------------------|-------------------------------|------------------------|-----------------------|
| OceanStor Dorado 6.1.2 | 512 Ki to 256 Ti | 50 Ti | 512 Ki to 50 Ti |
| OceanStor Pacific series 8.1.0 | 64 Mi to 512 Ti | 50 Ti | 64 Mi to 50 Ti |

Table 8-12 XFS capacity specifications

| Storage Type | Storage Specifications | XFS Specifications | CSI Specifications |
|--------------------------------|---------------------------|-----------------------|-----------------------|
| OceanStor Dorado 6.1.2 | 512 Ki to 256 Ti | 500 Ti | 512 Ki to 500 Ti |
| OceanStor Pacific series 8.1.0 | 64 Mi to 512 Ti | 500 Ti | 64 Mi to 500 Ti |

Step 1 Based on service requirements, modify specific parameters by referring to the description in this section and the PVC configuration file example to generate the PVC configuration file to be created, for example, the **mypvc.yaml** file in this example.

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
 name: mypvc
 annotations:
  csi.huawei.com/manageVolumeName: *
  csi.huawei.com/manageBackendName: *
 labels:
  provisioner: csi.huawei.com
spec:
 accessModes:
  - ReadWriteOnce
volumeMode: Filesystem
storageClassName: mysc
resources:
 requests:
  storage: 100Gi
```

Step 2 Run the following command to create a PVC using the configuration file.

```
# kubectl create -f mypvc.yaml persistentvolumeclaim/mypvc created
```

Step 3 After a period of time, run the following command to view the information about the created PVC. If the PVC status is **Bound**, the PVC has been created and can be used by a Pod.

```
# kubectl get pvc mypvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWO mysc 12s
```

! CAUTION

- After the PVC is created, if the PVC is in the Pending state after a long time (for example, one minute), refer to 11.6 When a PVC is Created, the PVC is in the Pending State.
- You are advised to create or delete a maximum of 100 PVCs in a batch.

----End

After a PVC is created, you can use the PVC to create a Pod. The following is a simple example of using a PVC. In this example, the created Pod uses the newly created *mypvc*.

```
apiVersion: apps/v1
kind: Deployment
metadata:
name: nginx-deployment
spec:
 selector:
  matchl abels:
   app: nginx
 replicas: 2
 template:
  metadata:
   labels:
     app: nginx
  spec:
   containers:
   - image: nginx:alpine
     name: container-0
     volumeMounts:
     - mountPath: /tmp
      name: pvc-mypvc
   restartPolicy: Always
   volumes:
    - name: pvc-mypvc
     persistentVolumeClaim:
      claimName: mypvc
                                        # name of PVC
```

8.1.1.3 Static Volume Provisioning

Static volume provisioning allows administrators to use a resource created on the storage side as a PV for containers in the cluster.

To implement static volume provisioning, perform the following steps:

- Configuring a PV
- Configuring a PVC

8.1.1.3.1 Configuring a PV

If static volume provisioning is used, you do not need to configure a StorageClass. Instead, you can directly create a PV. Before creating a PV, you need to configure a PV. The following example is a configuration file for static volume provisioning.

```
kind: PersistentVolume
apiVersion: v1
metadata:
name: mypv
spec:
volumeMode: Filesystem
```

storage: 100Gi

storageClassName: ""
accessModes:
- ReadWriteOnce
csi:
driver: csi.huawei.com
volumeHandle: iscsi-dorado-181.lun0001
fsType: xfs
capacity:

As shown in the preceding example, in the configuration file for static volume provisioning, **storageClassName** must be set to "". Otherwise, Kubernetes will use the default StorageClass. For details about other parameters, see **Table 8-13**.

 Table 8-13 Static volume provisioning parameters

| Paramet er | Description | Mandato ry | Defaul t Value | Remarks |
|-------------------------------|--|---------------|----------------------|--|
| metadat a.name | User-defined name of a PVC object. | Yes | - | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |
| spec.volu meMode | Volume mode. This parameter is optional. When LUN volumes are used, the following types are supported: • Filesystem: local file system. • Block: raw device. | No | Filesyst | This parameter takes effect when a PV is mounted. The default value is Filesystem. • Filesystem indicates that a container accesses a PV using a local file system. The local file system type is specified by the fsType field in the specified StorageClass. • Block indicates that a PV is accessed in raw volume mode. |
| spec.stor ageClass Name | Name of the StorageClass object. This parameter is mandatory. | Yes | - | Set the parameter to an empty string, that is, enter "". |

| Paramet er | Description | Mandato ry | Defaul t Value | Remarks |
|----------------------|--|---------------|----------------------|---|
| spec.acce ssModes | Access mode of the volume. RWO (ReadWriteOnc e): A volume can be mounted to a node in read/write mode. This mode also allows multiple Pods running on the same node to access the volume. ROX (ReadOnlyMan y): A volume can be mounted to multiple nodes in read-only mode. RWX (ReadWriteMan y): A volume can be mounted to multiple nodes in read/write mode. RWOP (ReadWriteOnc ePod): A volume can only be mounted to a single Pod in read/write mode. Kubernetes 1.22 and later versions support this feature. | Yes | ReadW riteOn ce | RWO/ROX/RWOP: supported by all types of volumes. RWOP is supported only by Kubernetes 1.22 and later versions. Check whether this feature is enabled for your Kubernetes cluster by referring to 10.5 Enabling the ReadWriteOncePod Feature Gate. RWX: supported by volumes whose volumeMode is set to Block or NFS. |

| Paramet er | Description | Mandato ry | Defaul t Value | Remarks |
|-------------------------------|--|---------------|------------------------|--|
| spec.csi.d river | CSI driver name. | Yes | csi.hua wei.co m | Set this parameter to the driver name set during Huawei CSI installation. |
| spec.csi.v olumeHa ndle | Unique identifier of a storage resource. This parameter is mandatory. Format: <backendname>.< volume-name></backendname> | Yes | | The value of this parameter consists of the following parts: • <backered name="">: indicates the name of the backend where the volume resides. You can run the following command to obtain the configured backend information. • ceanct get backend • <volume-name>: indicates the name of a resource (LUN/file system) on the storage. You can obtain the value from DeviceManager.</volume-name></backered> |
| spec.csi.fs Type | Type of a host file system. This parameter is optional. The supported types are: • ext2 • ext3 • ext4 • xfs | No | - | If this parameter is not set, the default value ext4 is used. This parameter is available only when volumeMode is set to Filesystem . |

| Paramet er | Description | Mandato ry | Defaul t Value | Remarks |
|-----------------------------------|---|---------------|----------------------|---|
| spec.capa city.stora ge | Volume size. | Yes | 100Gi | Ensure that the size is the same as that of the corresponding resource on the storage. Kubernetes will not invoke CSI to check whether the value of this parameter is correct. Therefore, the PV can be successfully created even if its capacity is inconsistent with that of the corresponding resource on the storage. |
| spec.mou ntOption s.nfsvers | NFS mount option on the host. The following mount option is supported: nfsvers: protocol version for NFS mounting. The value can be 3, 4, 4.0, or 4.1. | No | | This parameter is optional after the -o parameter when the mount command is executed on the host. The value is in list format. If the NFS version is specified for mounting, NFS 3, 4.0, and 4.1 protocols are supported (the protocol must be supported and enabled on storage devices). If nfsvers is set to 4, the latest protocol version NFS 4 may be used for mounting due to different OS configurations, for example, 4.1. If protocol 4.0 is required, you are advised to set nfsvers to 4.0. |

| Paramet er | Description | Mandato ry | Defaul t Value | Remarks |
|--|---|---------------|----------------------|---|
| spec.mou ntOption s.acl | The DPC namespace supports the ACL function. The DPC client supports POSIX ACL, NFSv4 ACL, and NT ACL authentication. | No | | The descriptions of acl, aclonlyposix, cnflush, and cflush are for reference only. For details about the parameters, see OceanStor Pacific Series Product Documentation and choose Configuration > Basic Service Configuration Guide for File > Configuring Basic Services (DPC Scenario) > Accessing a DPC Share on a Client > Step 2. |
| spec.mou ntOption s.aclonly posix | The DPC namespace supports POSIX ACL, and the DPC client supports POSIX ACL authentication. The following protocols support POSIX ACL: DPC, NFSv3, and HDFS. If NFSv4 ACL or NT ACL is used, the DPC client cannot identify the ACL of this type. As a result, the ACL of this type does not take effect. | No | - | If aclonlyposix and acl are used together, only acl takes effect. That is, the namespace supports the ACL function. |

| Paramet er | Description | Mandato ry | Defaul t Value | Remarks |
|-----------------------------------|--|---------------|----------------------|--|
| spec.mou ntOption s.cnflush | Asynchronous disk flushing mode. That is, data is not flushed to disks immediately when files in the namespace are closed. | No | | Asynchronous flushing mode: When a file is closed, data in the cache is not flushed to storage media in synchronous mode. Instead, data is written from the cache to the storage media in asynchronous flushing mode. After the write service is complete, data is flushed from the cache to disks periodically based on the flushing period. In a multi-client scenario, if concurrent operations are performed on the same file, the file size update is affected by the disk flushing period. That is, the file size is updated only after the disk flushing is complete. Generally, the update is completed within several seconds. Synchronous I/Os are not affected by the disk flushing period. |
| spec.mou ntOption s.cflush | Synchronous disk flushing mode. That is, data is flushed to disks immediately when files in the namespace are closed. | No | - | By default, the synchronous disk flushing mode is used. |

Prerequisites

- A storage resource, such as a LUN or file system, required by the PV to be created exists on the storage device. If the storage resource is a file system, you also need to create the share and client information of the file system.
- You have configured the PV configuration file by referring to Table 8-13.

Procedure

Step 1 Run the following command to create a PV based on the prepared .yaml file.

```
# kubectl create -f mypv.yaml
persistentvolume/mypv created
```

Step 2 After a period of time, run the following command to view the information about the created PV. If the PV status is **Available**, the PV is successfully created.

```
# kubectl get pv
NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS
REASON AGE
mypv 100Gi RWO Retain Available 4s
```

----End

8.1.1.3.2 Configuring a PVC

After a PV is created in static volume provisioning mode, you can create a PVC based on the PV for containers. The following example is a PVC configuration file for static volume provisioning.

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
name: mypvc
spec:
accessModes:
- ReadWriteOnce
volumeMode: Filesystem
resources:
requests:
storage: 100Gi
volumeName: mypv
```

As shown in the preceding example, set the **volumeName** parameter in the PVC configuration file to the PV created in static volume provisioning mode. For details about the parameters, see **Table 8-14**.

Table 8-14 PVC parameters

| Param eter | Description | Mandat ory | Defaul t Value | Remarks |
|-----------------------|------------------------------------|---------------|-------------------|--|
| metada ta.nam e | User-defined name of a PVC object. | Yes | - | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |

| Param eter | Description | Mandat ory | Defaul t Value | Remarks |
|--------------------------|--|---------------|-----------------------|--|
| spec.ac cessMo des | Access mode of the volume. RWO (ReadWriteOnce): A volume can be mounted to a node in read/write mode. This mode also allows multiple Pods running on the same node to access the volume. ROX (ReadOnlyMany): A volume can be mounted to multiple nodes in read-only mode. RWX (ReadWriteMany): A volume can be mounted to multiple nodes in read/write mode. RWOP (ReadWriteOnceP od): A volume can only be mounted to a single Pod in read/write mode. Kubernetes 1.22 and later versions support this feature. | Yes | ReadW riteOnc e | RWO/ROX/RWOP: supported by all types of volumes. RWOP is supported only by Kubernetes 1.22 and later versions. Check whether this feature is enabled for your Kubernetes cluster by referring to 10.5 Enabling the ReadWriteOncePod Feature Gate. RWX: supported by volumes whose volumeMode is set to Block or NFS. |

| Param eter | Description | Mandat ory | Defaul t Value | Remarks |
|-------------------------|--------------|---------------|-------------------|---|
| spec.vol umeMo de | Volume mode. | No | Filesyst em | This parameter is optional. The value can be Filesystem or Block. The default value is Filesystem. This parameter takes effect when a Pod is created. Filesystem indicates that a file system is created on a PVC to access the storage. Block indicates that a raw volume is used to access the storage. |

| Param eter | Description | Mandat ory | Defaul t Value | Remarks |
|---------------------------------|-----------------------------------|---------------|-------------------|---|
| spec.res ources.r equests | Size of the volume to be created. | Yes | - | Size of the volume to be created. The format is ***Gi and the unit is GiB. |
| .storag e | | | | The PVC capacity depends on storage specifications and host specifications. For example, OceanStor Dorado 6.1.2 or OceanStor Pacific series 8.1.0 is connected to CentOS 7. If ext4 file systems are used, see Table 8-5. If XFS file systems are used, see Table 8-6. If NFS or raw devices are used, the capacity must meet the specifications of the used Huawei storage device model and version. If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications. When a PVC is created using a static PV and the PVC capacity is smaller than the capacity of the bound PV, the PVC capacity is set to the capacity of the bound PV, the PVC cannot be created. |
| spec.vol umeNa me | Name of the PV object. | Yes | - | This parameter is mandatory when a PVC is created statically. |

Procedure

Step 1 Run the following command to create a PVC based on the configured .yaml file.

```
# kubectl create -f mypvc.yaml
persistentvolumeclaim/mypvc created
```

Step 2 After a period of time, run the following command to view the information about the created PVC.

```
# kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-840054d3-1d5b-4153-b73f-826f980abf9e 100Gi RWO 12s
```


- After the PVC is created, if the PVC is in the Pending state after a long time (for example, one minute), refer to 11.6 When a PVC is Created, the PVC is in the Pending State
- You are advised to create or delete a maximum of 100 PVCs in a batch.

----End

After a PVC is created, you can use the PVC to create a Pod. The following is a simple example of using a PVC. In this example, the created Pod uses the newly created *mypvc*.

```
apiVersion: apps/v1
kind: Deployment
metadata:
name: nginx-deployment
spec:
 selector:
  matchLabels:
   app: nginx
 replicas: 2
 template:
  metadata:
   labels:
     app: nginx
  spec:
   containers:
    - image: nginx:alpine
     name: container-0
     volumeMounts:
     - mountPath: /tmp
      name: pvc-mypvc
   restartPolicy: Always
    volumes:
    - name: pvc-mypvc
     persistentVolumeClaim:
      claimName: mypvc
                                         # name of PVC
```

8.1.2 Expanding the Capacity of a PVC

When the capacity of a PVC used by a container is insufficient, you need to expand the capacity of the PVC.

Prerequisites

- A PVC has been created, the backend to which it resides exists and supports capacity expansion.
- For details about the storage devices that support capacity expansion, see **Table 2-5** and **Table 2-9**. For details about the Kubernetes versions that support capacity expansion, see **Table 2-3**.

The csi-resizer service is enabled for huawei-csi-controller.

kubectl describe deploy huawei-csi-controller -n huawei-csi | grep csi-resizer csi-resizer:

Image: k8s.gcr.io/sig-storage/csi-resizer:v1.4.0

Procedure

Step 1 Run the **kubectl get pvc** *mypvc* command to query the StorageClass name of the PVC. In the preceding command, *mypvc* indicates the name of the PVC to be expanded.

kubectl get pvc mypvc
NAME STATUS VOLUME CAPACITY ACCESS MODES
STORAGECLASS AGE
mypvc Bound pvc-3383be36-537c-4cb1-8f32-a415fa6ba384 2Gi RW0
mysc 145m

Step 2 Run the **kubectl get sc** *mysc* command to check the StorageClass supports capacity expansion. In the preceding command, *mysc* indicates the name of the StorageClass to be queried.

```
# kubectl get sc mysc
NAME PROVISIONER RECLAIMPOLICY VOLUMEBINDINGMODE

ALLOWVOLUMEEXPANSION AGE

mysc csi.huawei.com Delete Immediate true 172m
```

If the value of **ALLOWVOLUMEEXPANSION** is **true**, the current StorageClass supports capacity expansion. In this case, go to **Step 4**.

Step 3 Run the following command to change the value of **allowVolumeExpansion** to **true**. In the preceding command, *mysc* indicates the name of the StorageClass to be modified.

kubectl patch sc mysc --patch '{"allowVolumeExpansion":true}'

Step 4 Run the following command to expand the capacity.

```
# kubectl patch pvc mypvc-p '{"spec":{"resources":{"requests":{"storage":"120Gi"}}}}'
```

In the preceding command, *mypvc* indicates the name of the PVC to be expanded, and *120Gi* indicates the capacity after expansion. Change the values based on the site requirements.

- The PVC capacity depends on storage specifications and host specifications. For example, OceanStor Dorado 6.1.2 or OceanStor Pacific series 8.1.0 is connected to CentOS 7. If ext4 file systems are used, see Table 8-5. If XFS file systems are used, see Table 8-6. If NFS or raw devices are used, the capacity must meet the specifications of the used Huawei storage device model and version.
- If the PVC capacity does not meet the specifications, a PVC or Pod may fail to be created due to the limitations of storage specifications or host file system specifications.
- If the capacity expansion fails because the target capacity exceeds the storage pool
 capacity, see 11.18 Failed to Expand the PVC Capacity Because the Target Capacity
 Exceeds the Storage Pool Capacity.
- **Step 5** Run the following command to check whether the capacity modification takes effect.

```
# kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS AGE
mypvc Bound pvc-3383be36-537c-4cb1-8f32-a415fa6ba384 120Gi RWO mysc 24s
```

----End

8.1.3 Cloning a PVC

This section describes how to clone a PVC.

When cloning a PVC, you need to specify the data source. The following is a simple example of cloning a PVC. In this example, **mypvc** is used as the data source and a PVC named **myclone** is created.

kind: PersistentVolumeClaim
apiVersion: v1
metadata:
name: myclone
spec:
storageClassName: mysc
dataSource:
name: mypvc
kind: PersistentVolumeClaim
volumeMode: Filesystem
accessModes:
- ReadWriteOnce
resources:
requests:
storage: 2Gi

CAUTION

- The specified storageClassName must be the same as the StorageClass of the source volume in dataSource.
- The capacity of the clone volume must be greater than or equal to that of the source volume. Equal capacity is recommended.

Prerequisites

The source PVC already exists in the system, and the backend where the source PVC resides supports cloning. For details about the storage devices that support cloning, see **Table 2-5** and **Table 2-9**. For details about the Kubernetes versions that support cloning, see **Table 2-3**.

Procedure

Step 1 Run the following command to create a PVC based on the configuration file of the clone volume.

kubectl create -f myclone.yaml persistentvolumeclaim/myclone created

----End

8.1.4 Creating a PVC Using a Snapshot

This section describes how to create a PVC using a snapshot.

When creating a PVC, you need to specify the data source. The following is a simple example of creating a PVC using a snapshot. In this example, **mysnapshot** is used as the data source and a PVC named **myrestore** is created.

apiVersion: v1 kind: PersistentVolumeClaim

metadata:
name: myrestore
spec:
storageClassName: mysc
dataSource:
name: mysnapshot
kind: VolumeSnapshot
apiGroup: snapshot.storage.k8s.io
volumeMode: Filesystem
accessModes:
- ReadWriteOnce
resources:
requests:
storage: 100Gi

CAUTION

- The specified storageClassName must be the same as the StorageClass of the snapshot source volume in dataSource.
- The capacity of the clone volume must be greater than or equal to that of the snapshot. Equal capacity is recommended.

Prerequisites

A snapshot already exists in the system, and the backend where the snapshot resides supports cloning. For details about the storage devices that support PVC creation using a snapshot, see **Table 2-5** and **Table 2-9**. For details about the Kubernetes versions that support PVC creation using a snapshot, see **Table 2-3**.

Procedure

Step 1 Run the following command to create a PVC based on the configuration file for creating a volume using a snapshot.

kubectl create -f myrestore.yaml persistentvolumeclaim/myrestore created

----End

8.2 Creating a VolumeSnapshot

In Kubernetes, a **VolumeSnapshot** is a snapshot of a volume on a storage system. The VolumeSnapshot capability provides Kubernetes users with a standard way to replicate the content of a volume at a specified point in time without creating a volume. For example, this function enables database administrators to back up the database before making changes such as editing or deleting.

This section describes how to create a VolumeSnapshot using Huawei CSI. To create a VolumeSnapshot, perform the following steps:

- Checking information about volume snapshot-dependent components
- Configuring a VolumeSnapshotClass
- Configuring a VolumeSnapshot

8.2.1 Checking Information About Volume Snapshotdependent Components

If you need to use volume snapshots and features associated with volume snapshots in the container environment, perform the operations in 3.5 Checking Volume Snapshot-Dependent Components to check whether volume snapshot-dependent components have been deployed in your environment and check the api-versions information about volume snapshots.

8.2.2 Configuring a VolumeSnapshotClass

VolumeSnapshotClass provides a way to describe the "classes" of storage when provisioning a VolumeSnapshot. Each VolumeSnapshotClass contains the **driver**, **deletionPolicy**, and **parameters** fields, which are used when a VolumeSnapshot belonging to the class needs to be dynamically provisioned.

The name of a VolumeSnapshotClass object is significant, and is how users can request a particular class. Administrators set the name and other parameters of a class when first creating VolumeSnapshotClass objects, and the objects cannot be updated once they are created.

The following is an example of a VolumeSnapshotClass used by Huawei CSI:

• If api-versions in your environment supports v1, use the following example:

apiVersion: snapshot.storage.k8s.io/v1

kind: VolumeSnapshotClass

metadata:

name: mysnapclass driver: csi.huawei.com deletionPolicy: Delete

 If api-versions in your environment supports v1beta1, use the following example:

apiVersion: snapshot.storage.k8s.io/v1beta1

kind: VolumeSnapshotClass

metadata:

name: mysnapclass driver: csi.huawei.com deletionPolicy: Delete

• If api-versions in your environment supports both v1 and v1beta1, v1 is recommended.

You can modify the parameters according to **Table 8-15**. Currently, Huawei CSI does not support user-defined parameters (**parameters**) in a VolumeSnapshotClass. Therefore, you are advised to create a VolumeSnapshotClass for all snapshots.

Table 8-15 VolumeSnapshotClass parameters

| Parameter | Description | Remarks |
|-------------------|---|--|
| metadata.n ame | User-defined name of a VolumeSnapshotCla ss object. | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. |

| Parameter | Description | Remarks |
|--------------------|---|---|
| driver | driver identifier. This parameter is mandatory. | Set this parameter to the driver name set during Huawei CSI installation. The default driver name is csi.huawei.com . |
| deletionPoli cy | Snapshot deletion policy. This parameter is mandatory. The value can be: Delete Retain | If the deletion policy is Delete, the snapshot on the storage device will be deleted together with the VolumeSnapshotContent object. If the deletion policy is Retain, the snapshot and VolumeSnapshotContent object on the storage device will be retained. |

Prerequisites

Huawei CSI supports snapshots, and the volume snapshot component CRD on which its running depends has been installed. For details about the CRD, see 3.5 Checking Volume Snapshot-Dependent Components. For details about the Kubernetes versions that support VolumeSnapshot creation, see Table 2-3.

Procedure

Step 1 Run the following command to create a VolumeSnapshotClass using the created VolumeSnapshotClass configuration file.

kubectl create -f mysnapclass.yaml volumesnapshotclass.snapshot.storage.k8s.io/mysnapclass created

Step 2 Run the following command to view the information about the created VolumeSnapshotClass.

```
# kubectl get volumesnapshotclass
NAME DRIVER DELETIONPOLICY AGE
mysnapclass csi.huawei.com Delete 25s
```

----End

8.2.3 Configuring a VolumeSnapshot

VolumeSnapshot can be provisioned in two ways: pre-provisioning and dynamic provisioning. Currently, Huawei CSI supports only dynamic provisioning. This section describes how to dynamically provision a VolumeSnapshot using Huawei CSI.

The following is an example of the VolumeSnapshot configuration file:

If api-versions in your environment supports v1, use the following example:

```
apiVersion: snapshot.storage.k8s.io/v1
kind: VolumeSnapshot
metadata:
name: mysnapshot
spec:
volumeSnapshotClassName: mysnapclass
source:
persistentVolumeClaimName: mypvc
```

• If api-versions in your environment supports v1beta1, use the following example:

apiVersion: snapshot.storage.k8s.io/v1beta1 kind: VolumeSnapshot metadata: name: mysnapshot spec: volumeSnapshotClassName: mysnapclass source: persistentVolumeClaimName: mypvc

• The api-versions information in the VolumeSnapshot must be the same as the version used for creating the VolumeSnapshotClass.

You can modify the parameters according to Table 8-16.

Table 8-16 VolumeSnapshot parameters

| Parameter | Description | Remarks | |
|---|--|--|--|
| metadata.name | User-defined name of a VolumeSnapshot object. | Take Kubernetes v1.22.1 as an example. The value can contain digits, lowercase letters, hyphens (-), and periods (.), and must start and end with a letter or digit. | |
| spec.volumeSnapshotCl assName | Name of the VolumeSnapshotCl ass object. | | |
| spec.source.persistentVo lumeClaimName | Name of the source PVC object. | Name of the source PVC of the snapshot | |

Prerequisites

- The source PVC exists, and the backend where the PVC resides supports
 VolumeSnapshot creation. For details about the storage devices that support
 VolumeSnapshot creation, see Table 2-5 and Table 2-9. For details about the
 Kubernetes versions that support VolumeSnapshot creation, see Table 2-3.
- The volume snapshot component CRD on which the running of Huawei CSI depends has been installed. For details, see 3.5 Checking Volume Snapshot-Dependent Components.
- A VolumeSnapshotClass that uses Huawei CSI exists in the system.

Procedure

Step 1 Run the following command to create a VolumeSnapshot using the created VolumeSnapshot configuration file.

kubectl create -f mysnapshot.yaml volumesnapshot.snapshot.storage.k8s.io/mysnapshot created

Step 2 Run the following command to view the information about the created VolumeSnapshot.

kubectl get volumesnapshot
NAME READYTOUSE SOURCEPVC SOURCESNAPSHOTCONTENT RESTORESIZE

SNAPSHOTCLASS SNAPSHOTCONTENT CREATIONTIME AGE mysnapshot **true** mypvc 100Gi mysnapclass snapcontent-1009af0a-24c2-4435-861c-516224503f2d <invalid> 78s

----End

9 Advanced Features

- 9.1 Configuring ALUA
- 9.2 Configuring Storage Topology Awareness

9.1 Configuring ALUA

Asymmetric Logical Unit Access (ALUA) is a model that supports access to multiple target ports. In the multipathing state, ALUA presents active/passive volumes to the host and provides a port access status switchover interface to switch over the working controllers for volumes. For example, when a volume of a controller fails, you can set the status of ports on the controller to **Unavailable**. After the host multipathing software that supports ALUA detects the status, it switches subsequent I/Os from the failed controller to the peer controller.

9.1.1 Configuring ALUA Using Helm

9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend

For details about how to configure ALUA for Huawei enterprise storage, see the host connectivity guide of the corresponding product.

The ALUA configuration may vary according to the OS. Visit **Huawei Technical Support**, enter **Host Connectivity Guide** in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS. Configure ALUA according to the actual situation and the description in the guide. Huawei CSI will apply the configuration items you set to the initiator of the host on Huawei storage.

□ NOTE

A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.

ALUA Parameters for OceanStor V3/V5 and OceanStor Dorado V3 Series

Table 9-1 lists the ALUA parameters supported by Huawei CSI for OceanStor V3/V5 and OceanStor Dorado V3 series.

Table 9-1 ALUA parameters supported by Huawei CSI for OceanStor V3/V5 and OceanStor Dorado V3 series

| Parameter | Description | Remarks |
|-------------------|---|--|
| HostName | Host name rule. This parameter is mandatory. You can use a regular expression. | The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. When HostName is set to *, the configuration takes effect on hosts with any name. For details, see Regular expression. If the host name of a compute node matches multiple ALUA configuration options, they will be sorted based on the matching accuracy and the first ALUA configuration option will be used. For details about the sorting rules, see Rules for Matching ALUA Configuration Items with Host Names. |
| MULTIPATHTY PE | Multipathing type. This parameter is mandatory. The value can be: • 0: Third-party multipathing is not used. • 1: Third-party multipathing is used. | |
| FAILOVERMO DE | Initiator switchover mode. This parameter is conditionally mandatory. The value can be: • 0: early-version ALUA • 1: common ALUA • 2: ALUA not used • 3: special ALUA | This parameter needs to be specified only when third-party multipathing is used. Configure the initiator switchover mode by referring to the connectivity guide. |

| Parameter | Description | Remarks |
|---------------------|--|--|
| SPECIALMODE TYPE | Special mode type of the initiator. This parameter is conditionally mandatory. The value can be: • 0: special mode 0 • 1: special mode 1 • 2: special mode 2 • 3: special mode 3 | This parameter needs to be specified only when the initiator switchover mode is special ALUA. Configure the special mode type of the initiator by referring to the connectivity guide. |
| PATHTYPE | Initiator path type. This parameter is conditionally mandatory. The value can be: • 0: preferred path • 1: non-preferred path | This parameter needs to be specified only when third-party multipathing is used. Configure the initiator path type by referring to the connectivity guide. |

The following uses OceanStor 18500 V5 as an example to describe how to connect to Red Hat. For details about the host connectivity guide, see *Huawei SAN*Storage Host Connectivity Guide for Red Hat.

The following ALUA configuration example is recommended in the OceanStor 18500 V5 host connectivity guide for Red Hat in non-HyperMetro storage scenarios. In this example, the OS on compute node **myhost01** in the Kubernetes cluster is RHEL 5.x, and that on other compute nodes is RHEL 7.x. According to the recommendation, the switchover mode of RHEL 5.x should be "ALUA not used", and that of RHEL 7.x should be "common ALUA".

```
storage: oceanstor-san
name: oceanstor-iscsi-155
urls:
 - https://192.168.129.155:8088
- https://192.168.129.156:8088
pools:
 - StoragePool001
parameters:
 protocol: iscsi
 portals:
  - 192.168.128.120
  - 192.168.128.121
 ALUA:
  ^myhost01$:
   MULTIPATHTYPE: 1
   FAILOVERMODE: 2
   PATHTYPE: 0
   MULTIPATHTYPE: 1
   FAILOVERMODE: 1
   PATHTYPE: 0
```

ALUA Parameters for OceanStor V6 and OceanStor Dorado V6 Series

Table 9-2 lists the ALUA parameters supported by Huawei CSI for OceanStor V6 and OceanStor Dorado V6 series.

□ NOTE

By default, the initiator host access mode of OceanStor V6 and OceanStor Dorado V6 series storage is "balanced mode". Therefore, you are not advised to configure ALUA parameters for OceanStor V6 and OceanStor Dorado V6 series storage.

Table 9-2 ALUA parameters for OceanStor V6 and OceanStor Dorado V6 series

| Parameter | Description | Remarks |
|-----------------------------|---|--|
| HostName | Host name rule. This parameter is mandatory. You can use a regular expression. | The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. When HostName is set to *, the configuration takes effect on hosts with any name. For details, see Regular expression. If the host name of a compute node matches multiple ALUA configuration options, they will be sorted based on the matching accuracy and the first ALUA configuration option will be used. For details about the sorting rules, see Rules for Matching ALUA Configuration Items with Host Names. |
| accessMode | Host access mode. This parameter is mandatory. The value can be: • 0: balanced mode • 1: asymmetric mode | The balanced mode is recommended in non-HyperMetro scenarios. Currently, Huawei CSI does not support SAN HyperMetro scenarios. Exercise caution when using the asymmetric mode. |
| hyperMetroPathO ptimized | Whether the path of the host on the current storage array is preferred in HyperMetro scenarios. The value can be: 1: yes 0: no | This parameter needs to be specified only when the host access mode is set to asymmetric. Currently, Huawei CSI does not support SAN HyperMetro scenarios. Exercise caution when using the asymmetric mode. |

The following uses OceanStor Dorado 18000 V6 as an example to describe how to connect to Red Hat. For details about the host connectivity guide, see *OceanStor Dorado 6.x and OceanStor 6.x Host Connectivity Guide for Red Hat*.

The following ALUA configuration example is recommended in the OceanStor Dorado 18000 V6 host connectivity guide for Red Hat in non-HyperMetro storage scenarios.

Rules for Matching ALUA Configuration Items with Host Names

 If the configured host name rule exactly matches the host name of the service node, the ALUA configuration item corresponding to the host name rule is used.

For example, the host name rule in configuration item 1 is * and that in configuration item 2 is **^myhost01\$**. If the host name of a compute node is **myhost01**, it exactly matches configuration item 2. In this case, Huawei CSI will apply the configuration information in configuration item 2 to the storage side.

 If the configured host name rule does not exactly match the host name of the service node, the first ALUA configuration item matched by regular expressions is used.

For example, the host name rule in configuration item 1 is **myhost0[0-9]** and that in configuration item 2 is **myhost0[5-9]**. In this case, configuration item 1 has a higher priority than configuration item 2. If the host name of a compute node is **myhost06**, both configuration items can be matched. In this case, Huawei CSI will apply the configuration information in configuration item 1 to the storage side.

9.1.1.2 Configuring ALUA Parameters for a Distributed Storage Backend

For details about how to configure ALUA for Huawei distributed storage, see the host connectivity guide of the corresponding product.

The ALUA configuration may vary according to the OS. Visit Huawei Technical Support, enter Host Connectivity Guide in the search box, and click the search button. In the search result, select the host connectivity guide for the desired OS. Configure ALUA according to the actual situation and the description in the guide. Huawei CSI will apply the configuration items you set to the initiator of the host on Huawei storage.

■ NOTE

A node with a Pod provisioned does not proactively change ALUA information. The host ALUA configuration changes only after a Pod is provisioned again to the node.

In non-HyperMetro scenarios of distributed storage, you are advised to set the switchover mode to "disable ALUA" (default value). This is because the storage system is in active/active mode and "enables ALUA" is meaningless. Therefore, you are not advised to configure ALUA parameters for distributed storage.

Table 9-3 lists the ALUA parameters supported by Huawei CSI for distributed storage.

Table 9-3 ALUA parameters for distributed storage

| Parameter | Description | Remarks |
|----------------|---|--|
| HostName | The value of HostName is the host name of a worker node, for example, HostName1 and HostName2. | The host name can be obtained by running the cat /etc/hostname command. It can be matched by using regular expressions. When HostName is set to *, the configuration takes effect on hosts with any name. For details, see Regular expression. If the host name of a compute node matches multiple ALUA configuration options, they will be sorted based on the matching accuracy and the first ALUA configuration option will be used. For details about the sorting rules, see Rules for Matching ALUA Configuration Items with Host Names. |
| switchoverMode | Switchover mode. This parameter is mandatory. The value can be: • Disable_alua: disables ALUA. • Enable_alua: enables ALUA. | In non-HyperMetro scenario, you are advised to set the switchover mode to "disable ALUA". This is because the storage system is in active/active mode and "enables ALUA" is meaningless. Currently, Huawei CSI does not support SAN HyperMetro scenarios. Exercise caution when enabling ALUA. |

| Parameter | Description | Remarks |
|-----------|---|--|
| pathType | Path type. This parameter is conditionally mandatory. The value can be: | This parameter is mandatory when the switchover mode is set to "enables ALUA". |
| | optimal_path: preferred path | |
| | non_optimal_path: non-preferred path | |

Rules for Matching ALUA Configuration Items with Host Names

 If the configured host name rule exactly matches the host name of the service node, the ALUA configuration item corresponding to the host name rule is used.

For example, the host name rule in configuration item 1 is * and that in configuration item 2 is **^myhost01\$**. If the host name of a compute node is **myhost01**, it exactly matches configuration item 2. In this case, Huawei CSI will apply the configuration information in configuration item 2 to the storage side.

• If the configured host name rule does not exactly match the host name of the service node, the first ALUA configuration item matched by regular expressions is used.

For example, the host name rule in configuration item 1 is **myhost0[0-9]** and that in configuration item 2 is **myhost0[5-9]**. In this case, configuration item 1 has a higher priority than configuration item 2. If the host name of a compute node is **myhost06**, both configuration items can be matched. In this case, Huawei CSI will apply the configuration information in configuration item 1 to the storage side.

9.2 Configuring Storage Topology Awareness

In the Kubernetes cluster, resources can be scheduled and provisioned based on the topology labels of nodes and the topology capabilities supported by storage backends.

Prerequisites

You need to configure topology labels on worker nodes in the cluster. The method is as follows:

- 1. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Run the kubectl get node command to view information about worker nodes in the current cluster.

```
# kubectl get node

NAME STATUS ROLES AGE VERSION

node01 Ready controlplane,etcd,worker 42d v1.22.3

node02 Ready worker 42d v1.22.3

node03 Ready worker 42d v1.22.3
```

3. Run the **kubectl label node** <nodename> topology.kubernetes.io/ <key>= <value> command to configure a topology label for a worker node. In the preceding command, <nodename> indicates the name of a worker node. For details about the **key** and **value** parameters, see **Table 9-4**. # kubectl label node node01 topology.kubernetes.io/zone=ChengDu node/node01 labeled

Table 9-4 Parameter description

| Paramete r | Description | Remarks |
|-----------------|--|---|
| <key></key> | Unique identifier of a topology label. | The value can be zone , region , or protocol . <pre><pre>con be set to iscsi, nfs, fc, or roce.</pre></pre> |
| <value></value> | Value of a topology label. | If key is set to zone or region , value is a user-defined parameter. |
| | | If key is set to protocol. <pre><pre><pre>// return of the content of the cont</pre></pre></pre> |

■ NOTE

- A topology label must start with **topology.kubernetes.io**. Topology label examples:
 - Example 1: topology.kubernetes.io/region=China-west
 - Example 2: topology.kubernetes.io/zone=ChengDu
 - Example 3: topology.kubernetes.io/protocol.iscsi=csi.huawei.com
 - Example 4: topology.kubernetes.io/protocol.fc=csi.huawei.com
- A key in a topology label on a node can have only one value.
- If multiple protocols are configured in a topology label on a node, when you select a backend, the backend needs to meet only one of the protocols.
- If both the region and the zone are configured in a topology label on a node, when you select a backend, the backend must meet both of them.
- 4. Run the kubectl get nodes -o=jsonpath='{range .items[*]} [{.metadata.name}, {.metadata.labels}]{"\n"}{end}' | grep --color "topology.kubernetes.io" command to view the label information about all worker nodes in the current cluster.

kubectl get nodes -o=jsonpath='{range .items[*]}[{.metadata.name}, {.metadata.labels}]{"\n"}{end}' | grep --color "topology.kubernetes.io" [node01, {"beta.kubernetes.io/arch":"amd64","beta.kubernetes.io/os":"linux","kubernetes.io/arch":"amd64","kubernetes.io/hostname":"node01","kubernetes.io/os":"linux","node-role.kubernetes.io/controlplane":"true","node-role.kubernetes.io/etcd":"true","node-role.kubernetes.io/worker":"true","topology.kubernetes.io/zone":"ChengDu"}]

9.2.1 Configuring Storage Topology Awareness Using Helm

Procedure

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **Table 3-1**.
- **Step 3** Go to the backend service configuration directory **/examples/backend/** and back up the **backend.yaml** file.

cp backend.yaml backend.yaml.bak

Step 4 Run the **vi** backend.yaml command to open the file and configure topology awareness as required. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
storage: "oceanstor-san"
name: "dorado-iscsi-155"
namespace: "huawei-csi"
urls:
 - "https://192.168.129.155:8088"
pools:
 - "StoragePool001"
parameters:
 protocol: "iscsi"
 portals:
   - "10.10.30.20"
  - "10.10.30.21"
supportedTopologies:
 - { "topology.kubernetes.io/region": "China-west", "topology.kubernetes.io/zone": "ChengDu" }
 - { "topology.kubernetes.io/region": "China-south", "topology.kubernetes.io/zone": "ShenZhen" }
maxClientThreads: "30"
```

Step 5 Run the following command to delete the storage backend to be modified. In the command, **dorado-iscsi-155** indicates the storage backend name.

```
# oceanctl delete backend dorado-iscsi-155 -n huawei-csi backend/dorado-iscsi-155 deleted
```

Step 6 Run the following command to create a storage backend.

```
# oceanctl create backend -f ../examples/backend/backend.yaml
Please enter this backend user name:admin
Please enter this backend password:
backend/dorado-iscsi-155 created
```

Step 7 Run the **vi** *StorageClass.yaml* command to modify the .yaml file. Press **I** or **Insert** to enter the editing mode and add related parameters in the .yaml file. For details about the parameters, see **Table 9-5**. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

Add the following configuration items to the StorageClass.yaml file.

• Example 1: Configure zone and region information in the StorageClass.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
name: example-storageclass
provisioner: csi.huawei.com
parameters:
volumeType: lun
allocType: thin
volumeBindingMode: WaitForFirstConsumer
allowedTopologies:
- matchLabelExpressions:
- key: topology.kubernetes.io/zone
values:
- ChengDu
- key: topology,kubernetes.io/region
```

values:

- China-west

Example 2: Configure protocol information in the StorageClass.

kind: StorageClass apiVersion: storage.k8s.io/v1

metadata:

name: protocol-example-storageclass provisioner: csi.huawei.com

parameters:

volumeType: lun

allocType: thin

volumeBindingMode: WaitForFirstConsumer

allowedTopologies:

- matchLabelExpressions:
- key: topology.kubernetes.io/protocol.iscsi values:
- csi.huawei.com

Table 9-5 Parameter description

| Parameter | Description | Remarks |
|---|---|---|
| volumeBindin gMode | PersistentVolume binding mode, used to control the time when PersistentVolume resources are dynamically allocated and | You can set this parameter to WaitForFirstConsumer or Immediate. WaitForFirstConsumer: indicates that the binding and allocation of the PersistentVolume are delayed until a Pod that uses the PVC is created. Immediate: The PersistentVolume is bound and allocated immediately after a |
| allowedTopol ogies.matchLa belExpression s | ogies.matchLa information label, which is used to | PVC is created. key: This parameter can be set to topology.kubernetes.io/zone or topology.kubernetes.io/region. topology.kubernetes.io/ protocol. <pre>cprotocol></pre> indicates the protocol type and can be iscsi, fc, or nfs. |
| | | value: If key is topology.kubernetes.io/zone or topology.kubernetes.io/region, value must be the same as the topology label set in the prerequisites. If key is topology.kubernetes.io/protocol. <pre>/protocol>, value</pre> is fixed to csi.huawei.com. |

Step 8 Run the following command to create a StorageClass based on the .yaml file.

kubectl create -f StorgeClass.yaml

Step 9 Use the StorageClass to create a PVC with the topology capability. For details, see **8.1.1.1.2 Configuring a PVC**.

----End

10 Common Operations

- 10.1 Collecting Logs
- 10.2 Updating the huawei-csi-controller Service
- 10.3 Updating the huawei-csi-node Service
- 10.4 Modifying the Log Output Mode
- 10.5 Enabling the ReadWriteOncePod Feature Gate
- 10.6 Configuring Access to the Kubernetes Cluster as a Non-root User

10.1 Collecting Logs

Performing Check Before Collection

- **Step 1** Use a remote access tool, such as PuTTY, to log in to the node where the oceanctl tool is installed in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **oceanctl version** command. The displayed version is **v4.2.1**.

```
$ oceanctl version
+------+
| OCEANCTL VERSION |
+------+
| v4.2.1 |
```

Step 3 Run the **oceanctl --help** command. The following information is displayed.

```
$ oceanctl --help
A CLI tool for Ocean Storage in Kubernetes
Usage:
 oceanctl [command]
Available Commands:
 collect collect messages in Kubernetes
 create
          Create a resource to Ocean Storage in Kubernetes
 delete
          Delete one or more resources from Ocean Storage in Kubernetes
         Get one or more resources from Ocean Storage in Kubernetes
 get
         Help about any command
 help
 update Update a resource for Ocean Storage in Kubernetes
 version Print the version of oceanctl
```

```
Flags:
-h, --help help for oceanctl

Use "oceanctl [command] --help" for more information about a command.
```

Step 4 Run the **kubectl get deploy -n** *\${NAMESPACE}* command to check whether a Pod is started properly. In the preceding command, *\${NAMESPACE}* indicates the namespace for installing CSI. **huawei-csi** is used as an example.

```
# kubectl get deploy -n huawei-csi
NAME READY UP-TO-DATE AVAILABLE AGE
huawei-csi-controller 1/1 1 1 21h
```

----End

Collecting All Logs in the CSI Namespace Using oceanctl

- **Step 1** Use a remote access tool, such as PuTTY, to log in to the node checked in **Performing Check Before Collection** through the management IP address.
- **Step 2** Run the **oceanctl collect logs -n <namespace> -a** command to collect CSI logs of all nodes where CSI containers reside in the cluster.

Step 3 Check the log package generated in the **/tmp** directory. You can run the **unzip** \$ {zip_name} -d collect_logs command to decompress the log package. In the preceding command, \${zip_name} indicates the package name.

```
# date
Wed Sep 20 02:49:24 EDT 2023
# ls
huawei-csi-2023-09-20-02:48:22-all.zip
```

----End

Collecting the Log of a Single CSI Node Using oceanctl

- **Step 1** Use a remote access tool, such as PuTTY, to log in to the node checked in **Performing Check Before Collection** through the management IP address.
- **Step 2** Run the **oceanctl collect logs -n** *<namespace>* **-N** *<nodeName>* command to collect CSI logs of all nodes where CSI containers reside in the cluster.

Step 3 Check the log package generated in the **/tmp** directory. You can run the **unzip** \$ {zip_name} -d collect_logs command to decompress the log package. In the preceding command, \${zip_name} indicates the package name.

```
# date
Thu Sep 21 04:08:47 EDT 2023
# ls
huawei-csi-2023-09-21-04:05:15-node-1.zip
```

----End

10.2 Updating the huawei-csi-controller Service

Perform this operation when you need to update the huawei-csi-controller service, for example, adding the snapshot or the capacity expansion function.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Back up the values.yaml file used during CSI installation. You can run the helm get values helm-huawei-csi -n huawei-csi -a > values.yaml.bak command to back up the file. In the preceding command, helm-huawei-csi indicates the Helm chart name defined during installation, and huawei-csi indicates the Helm chart namespace defined during installation.
 - # helm get values helm-huawei-csi -n huawei-csi -a > values.yaml.bak
- **Step 3** Go to the /helm/esdk directory, run the vi values.yaml command to open the file and modify controller parameters. The following is an example. After the modification is complete, press **Esc** and enter :wq! to save the modification. For details about the component package path, see Table 3-1.

```
kubernetes:
 namespace: huawei-csi
images:
 # Images provided by Huawei
 huaweiCSIService: huawei-csi:4.2.1
 storageBackendSidecar: storage-backend-sidecar:4.2.1
 storageBackendController: storage-backend-controller:4.2.1
 # CSI-related sidecar images provided by the Kubernetes community.
 # These must match the appropriate Kubernetes version.
 sidecar:
  attacher: k8s.gcr.io/sig-storage/csi-attacher:v3.4.0
  provisioner: k8s.gcr.io/sig-storage/csi-provisioner:v3.0.0
  resizer: k8s.gcr.io/sig-storage/csi-resizer:v1.4.0
  registrar: k8s.gcr.io/sig-storage/csi-node-driver-registrar:v2.3.0
  livenessProbe: k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
  snapshotter: k8s.gcr.io/sig-storage/csi-snapshotter:v4.2.1
  snapshotController: k8s.gcr.io/sig-storage/snapshot-controller:v4.2.1
 # The image name and tag for the Huawei CSI Service container
 # Replace the appropriate tag name
 huaweiCSIService: huawei-csi:4.2.1
# The CSI driver parameter configuration
csiDriver:
 # Driver name, it is strongly recommended not to modify this parameter
 # The CCE platform needs to modify this parameter, e.g. csi.oceanstor.com
 driverName: csi.huawei.com
 # Endpoint, it is strongly recommended not to modify this parameter
 endpoint: /csi/csi.sock
 # DR Endpoint, it is strongly recommended not to modify this parameter
 drEndpoint: /csi/dr-csi.sock
 # Maximum number of concurrent disk scans or detaches, support 1~10
 connectorThreads: 4
 # Flag to enable or disable volume multipath access, support [true, false]
 volumeUseMultipath: true
 # Multipath software used by fc/iscsi. support [DM-multipath, HW-UltraPath, HW-UltraPath-NVMe]
 scsiMultipathType: DM-multipath
```

```
# Multipath software used by roce/fc-nvme. only support [HW-UltraPath-NVMe]
 nvmeMultipathType: HW-UltraPath-NVMe
 # Timeout interval for waiting for multipath aggregation when DM-multipath is used on the host. support
1~600
 scanVolumeTimeout: 3
 # Interval for updating backend capabilities. support 60~600
 backendUpdateInterval: 60
 # Huawei-csi-controller log configuration
 controllerLogging:
  # Log record type, support [file, console]
  module: file
  # Log Level, support [debug, info, warning, error, fatal]
  level: info
  # Directory for storing logs
  fileDir: /var/log/huawei
  # Size of a single log file
  fileSize: 20M
  # Maximum number of log files that can be backed up.
  maxBackups: 9
 # Huawei-csi-node log configuration
 nodeLogging:
  # Log record type, support [file, console]
  module: file
  # Log Level, support [debug, info, warning, error, fatal]
  level: info
  # Directory for storing logs
  fileDir: /var/log/huawei
  # Size of a single log file
  fileSize: 20M
  # Maximum number of log files that can be backed up.
  maxBackups: 9
controller:
 # controllerCount: Define the number of huawei-csi controller
 # Allowed values: n, where n > 0
 # Default value: 1
 # Recommended value: 2
 controllerCount: 1
 # volumeNamePrefix: Define a prefix that is prepended to volumes.
 # THIS MUST BE ALL LOWER CASE.
 # Default value: pvc
 # Examples: "volumes", "vol"
 volumeNamePrefix: pvc
 # Port used by the webhook service. The default port is 4433.
 # You can change the port to another port that is not occupied.
 webhookPort: 4433
  # enabled: Enable/Disable volume snapshot feature
  # If the Kubernetes version is lower than 1.17, set this parameter to false.
  # true: enable volume snapshot feature(install snapshotter sidecar)
  # false: disable volume snapshot feature(do not install snapshotter sidecar)
  # Default value: None
  enabled: true
 resizer:
  # enabled: Enable/Disable volume expansion feature
  # Allowed values:
  # true: enable volume expansion feature(install resizer sidecar)
  # false: disable volume snapshot feature(do not install resizer sidecar)
  # Default value: None
  enabled: true
 # nodeSelector: Define node selection constraints for controller pods.
 # For the pod to be eligible to run on a node, the node must have each
 # of the indicated key-value pairs as labels.
```

```
# Leave as blank to consider all nodes
# Allowed values: map of key-value pairs
# Default value: None
nodeSelector:
# Uncomment if nodes you wish to use have the node-role.kubernetes.io/master taint
# node-role.kubernetes.io/master: "
# Uncomment if nodes you wish to use have the node-role.kubernetes.io/control-plane taint
# node-role.kubernetes.io/control-plane: ""
# tolerations: Define tolerations that would be applied to controller deployment
# Leave as blank to install controller on worker nodes
# Allowed values: map of key-value pairs
# Default value: None
tolerations:
# Uncomment if nodes you wish to use have the node-role.kubernetes.io/master taint
# - key: "node-role.kubernetes.io/master"
# Uncomment if nodes you wish to use have the node-role.kubernetes.io/control-plane taint
# - key: "node-role.kubernetes.io/control-plane"
    operator: "Exists"
# effect: "NoSchedule"
```

Step 4 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

10.3 Updating the huawei-csi-node Service

Perform this operation when you need to update the huawei-csi-node service.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- Step 2 Back up the values.yaml file used during CSI installation. You can run the helm get values helm-huawei-csi -n huawei-csi -a > values.yaml.bak command to back up the file. In the preceding command, helm-huawei-csi indicates the Helm chart name defined during installation, and huawei-csi indicates the Helm chart namespace defined during installation.

helm get values helm-huawei-csi -n huawei-csi -a > values.yaml.bak

Step 3 Go to the /helm/esdk directory, run the vi values.yaml command to open the file and modify node parameters. The following is an example. After the modification is complete, press Esc and enter :wq! to save the modification. For details about the component package path, see Table 3-1.

```
kubernetes:
namespace: huawei-csi
```

```
images:
 # Images provided by Huawei
 huaweiCSIService: huawei-csi:4.2.1
 storageBackendSidecar: storage-backend-sidecar:4.2.1
 storageBackendController: storage-backend-controller:4.2.1
 # CSI-related sidecar images provided by the Kubernetes community.
 # These must match the appropriate Kubernetes version.
 sidecar:
  attacher: k8s.gcr.io/sig-storage/csi-attacher:v3.4.0
  provisioner: k8s.gcr.io/sig-storage/csi-provisioner:v3.0.0
  resizer: k8s.gcr.io/sig-storage/csi-resizer:v1.4.0
  registrar: k8s.gcr.io/sig-storage/csi-node-driver-registrar:v2.3.0
  livenessProbe: k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
  snapshotter: k8s.gcr.io/sig-storage/csi-snapshotter:v4.2.1
  snapshotController: k8s.gcr.io/sig-storage/snapshot-controller:v4.2.1
 # The image name and tag for the Huawei CSI Service container
 # Replace the appropriate tag name
 huaweiCSIService: huawei-csi:4.2.1
# The CSI driver parameter configuration
csiDriver:
 # Driver name, it is strongly recommended not to modify this parameter
 # The CCE platform needs to modify this parameter, e.g. csi.oceanstor.com
 driverName: csi.huawei.com
 # Endpoint, it is strongly recommended not to modify this parameter
 endpoint: /csi/csi.sock
 # DR Endpoint, it is strongly recommended not to modify this parameter
 drEndpoint: /csi/dr-csi.sock
 # Maximum number of concurrent disk scans or detaches, support 1~10
 connectorThreads: 4
 # Flag to enable or disable volume multipath access, support [true, false]
 volumeUseMultipath: true
 # Multipath software used by fc/iscsi. support [DM-multipath, HW-UltraPath, HW-UltraPath-NVMe]
 scsiMultipathType: DM-multipath
 # Multipath software used by roce/fc-nvme. only support [HW-UltraPath-NVMe]
 nvmeMultipathType: HW-UltraPath-NVMe
 # Timeout interval for waiting for multipath aggregation when DM-multipath is used on the host. support
1~600
 scanVolumeTimeout: 3
 # Interval for updating backend capabilities. support 60~600
 backendUpdateInterval: 60
 # Huawei-csi-controller log configuration
 controllerLogging:
  # Log record type, support [file, console]
  module: file
  # Log Level, support [debug, info, warning, error, fatal]
  level: info
  # Directory for storing logs
  fileDir: /var/log/huawei
  # Size of a single log file
  fileSize: 20M
  # Maximum number of log files that can be backed up.
  maxBackups: 9
 # Huawei-csi-node log configuration
 nodeLogging:
  # Log record type, support [file, console]
  module: file
  # Log Level, support [debug, info, warning, error, fatal]
  level: info
  # Directory for storing logs
  fileDir: /var/log/huawei
  # Size of a single log file
  fileSize: 20M
  # Maximum number of log files that can be backed up.
  maxBackups: 9
```

```
node:
 # maxVolumesPerNode: Defines the maximum number of volumes that can be used by a node.
 # Uncomment if you want to limit the number of volumes that can be used in a Node.
 # maxVolumesPerNode: 100
 # nodeSelector: Define node selection constraints for node pods.
 # For the pod to be eligible to run on a node, the node must have each
 # of the indicated key-value pairs as labels.
 # Leave as blank to consider all nodes
 # Allowed values: map of key-value pairs
 # Default value: None
 nodeSelector:
 # Uncomment if nodes you wish to use have the node-role.kubernetes.io/master taint
 # node-role.kubernetes.io/master: '
 # Uncomment if nodes you wish to use have the node-role.kubernetes.io/control-plane taint
 # node-role.kubernetes.io/control-plane: ""
 # tolerations: Define tolerations that would be applied to node daemonset
 # Add/Remove tolerations as per requirement
 # Leave as blank if you wish to not apply any tolerations
 # Allowed values: map of key-value pairs
 # Default value: None
 tolerations:
  - key: "node.kubernetes.io/memory-pressure"
   operator: "Exists"
   effect: "NoExecute"
  - key: "node.kubernetes.io/disk-pressure"
   operator: "Exists"
   effect: "NoExecute"
  - key: "node.kubernetes.io/network-unavailable"
   operator: "Exists"
   effect: "NoExecute"
   - key: "node-role.kubernetes.io/control-plane"
     operator: "Exists"
     effect: "NoSchedule"
   - key: "node-role.kubernetes.io/master"
     operator: "Exists"
     effect: "NoSchedule"
```

Step 4 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

```
# helm upgrade helm-huawei-csi ./ -n huawei-csi
Release "helm-huawei-csi" has been upgraded. Happy Helming!
NAME: helm-huawei-csi
LAST DEPLOYED: Thu Jun 9 07:58:15 2022
NAMESPACE: huawei-csi
STATUS: deployed
REVISION: 2
TEST SUITE: None
```

----End

10.4 Modifying the Log Output Mode

huawei-csi supports two log output modes: **file** and **console**. **file** indicates that logs are output to the fixed directory (**/var/log/huawei**), and **console** indicates that logs are output to the standard directory of the container. You can set the log output mode as required. The default mode is **file**.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **Table 3-1**.
- **Step 3** Back up the **values.yaml** file used during CSI installation. If the **values.yaml** file used during the last installation cannot be found, run the **helm get values** *helm-huawei-csi* **-n** *huawei-csi* **-a** command to query the file.
 - # cp values.yaml values.yaml.bak
- **Step 4** Run the **vi** *values.yaml* command to open the file and modify controller or node log parameters. The following is an example. After the modification is complete, press **Esc** and enter :**wq!** to save the modification.

```
# The CSI driver parameter configuration
csiDriver:
 # Driver name, it is strongly recommended not to modify this parameter
 # The CCE platform needs to modify this parameter, e.g. csi.oceanstor.com
 driverName: csi.huawei.com
 # Endpoint, it is strongly recommended not to modify this parameter
 endpoint: /csi/csi.sock
 # DR Endpoint, it is strongly recommended not to modify this parameter
 drEndpoint: /csi/dr-csi.sock
 # Maximum number of concurrent disk scans or detaches, support 1~10
 connectorThreads: 4
 # Flag to enable or disable volume multipath access, support [true, false]
 volumeUseMultipath: true
 # Multipath software used by fc/iscsi. support [DM-multipath, HW-UltraPath, HW-UltraPath-NVMe]
 scsiMultipathType: DM-multipath
 # Multipath software used by roce/fc-nvme. only support [HW-UltraPath-NVMe]
 nvmeMultipathType: HW-UltraPath-NVMe
 # Timeout interval for waiting for multipath aggregation when DM-multipath is used on the host. support
 scanVolumeTimeout: 3
 # Interval for updating backend capabilities. support 60~600
 backendUpdateInterval: 60
 # Huawei-csi-controller log configuration
 controllerLogging:
  # Log record type, support [file, console]
  module: file
  # Log Level, support [debug, info, warning, error, fatal]
  level: info
  # Directory for storing logs
  fileDir: /var/log/huawei
  # Size of a single log file
  fileSize: 20M
  # Maximum number of log files that can be backed up.
  maxBackups: 9
 # Huawei-csi-node log configuration
 nodeLogging:
  # Log record type, support [file, console]
  module: file
  # Log Level, support [debug, info, warning, error, fatal]
  level: info
  # Directory for storing logs
  fileDir: /var/log/huawei
  # Size of a single log file
  fileSize: 20M
  # Maximum number of log files that can be backed up.
  maxBackups: 9
```

Step 5 Run the **helm upgrade** *helm-huawei-csi* ./ **-n** *huawei-csi* command to upgrade the Helm chart.

In the preceding command, *helm-huawei-csi* indicates the name of the chart to be upgraded, ./ indicates the Helm project in the current directory, and *huawei-csi* indicates the namespace where the chart is located.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

10.5 Enabling the ReadWriteOncePod Feature Gate

The ReadWriteOnce access mode is the fourth access mode introduced by Kubernetes v1.22 for PVs and PVCs. If you create a Pod using a PVC in ReadWriteOncePod access mode, Kubernetes ensures that the Pod is the only Pod in the cluster that can read or write the PVC.

The ReadWriteOncePod access mode is an alpha feature in Kubernetes v1.22/1.23/1.24. Therefore, you need to enable the ReadWriteOncePod feature in **feature-gates** of kube-apiserver, kube-scheduler, and kubelet before using the access mode.

Currently, the CCE/CCE Agile platform does not support the ReadWriteOncePod feature gate.

Procedure

- **Step 1** Enable the ReadWriteOncePod feature gate for kube-apiserver.
 - 1. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
 - 2. Run the vi /etc/kubernetes/manifests/kube-apiserver.yaml command, press I or Insert to enter the editing mode, and add --feature-gates=ReadWriteOncePod=true to the kube-apiserver container. After the modification is complete, press Esc and enter :wq! to save the modification.

...
spec:
containers:
- command:
- kube-apiserver
- --feature-gates=ReadWriteOncePod=true
...

□ NOTE

After the editing is complete, Kubernetes will automatically apply the updates.

- **Step 2** Enable the ReadWriteOncePod feature gate for kube-scheduler.
 - 1. Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

2. Run the vi /etc/kubernetes/manifests/kube-scheduler.yaml command, press I or Insert to enter the editing mode, and add --feature-gates=ReadWriteOncePod=true to the kube-scheduler container. After the modification is complete, press Esc and enter :wq! to save the modification.

... spec:
containers:
- command:
- kube-scheduler
- --feature-gates=ReadWriteOncePod=true
...

□ NOTE

After the editing is complete, Kubernetes will automatically apply the updates.

Step 3 Enable the ReadWriteOncePod feature gate for kubelet.

NOTICE

The dynamic Kubelet configuration function is not used since v1.22 and deleted in v1.24. Therefore, you need to perform the following operations on kubelet on each worker node in the cluster.

- 1. Use a remote access tool, such as PuTTY, to log in to any worker node in the Kubernetes cluster through the management IP address.
- 2. Run the vi /var/lib/kubelet/config.yaml command, press I or Insert to enter the editing state, and add ReadWriteOncePod: true to the featureGates field of the KubeletConfiguration object. If the featureGates field does not exist, add it at the same time. After the modification is complete, press Esc and enter :wq! to save the modification.

apiVersion: kubelet.config.k8s.io/v1beta1
featureGates:

ReadWriteOncePod: true
...

□ NOTE

The default path of the kubelet configuration file is /var/lib/kubelet/config.yaml. Enter the path based on site requirements.

3. After the configuration is complete, run the **systemctl restart kubelet** command to restart kubelet.

----End

10.6 Configuring Access to the Kubernetes Cluster as a Non-root User

Procedure

Step 1 Copy the authentication file of the Kubernetes cluster and modify /etc/ kubernetes/admin.conf to be the actual authentication file.

\$ mkdir -p \$HOME/.kube

\$ sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config

Step 2 Change the user and user group of the authentication file.

\$ sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config

Step 3 Configure the **KUBECONFIG** environment variable of the current user. The following uses Ubuntu 20.04 as an example.

\$ echo "export KUBECONFIG=\$HOME/.kube/config" >> ~/.bashrc
\$ source ~/.bashrc

----End

11 FAQ

- 11.1 How Do I View Huawei CSI Logs?
- 11.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernetes Platform Is Set Up for the First Time
- 11.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported
- 11.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters
- 11.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName
- 11.6 When a PVC Is Created, the PVC Is in the Pending State
- 11.7 Before a PVC Is Deleted, the PVC Is in the Pending State
- 11.8 When a Pod Is Created, the Pod Is in the ContainerCreating State
- 11.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created
- 11.10 A Pod Fails to Be Created and the Log Shows That the Execution of the mount Command Times Out
- 11.11 A Pod Fails to Be Created and the Log Shows That the mount Command Fails to Be Executed
- 11.12 After a Pod Fails to Be Created or kubelet Is Restarted, Logs Show That the Mount Point Already Exists
- 11.13 How Do I Download a Container Image to the Local PC?
- 11.14 How Do I Obtain CSI Version Information?
- 11.15 Common Problems and Solutions for Interconnecting with the Tanzu Kubernetes Cluster
- 11.16 Common Problems and Solutions for Using the Tanzu Kubernetes Cluster
- 11.17 Failed to Expand the Capacity of a Generic Ephemeral Volume

- 11.18 Failed to Expand the PVC Capacity Because the Target Capacity Exceeds the Storage Pool Capacity
- 11.19 A webhook Fails to Be Called When the oceanctl Tool Is Used to Manage Backends
- 11.20 An Account Is Locked After the Password Is Updated on the Storage Device

11.1 How Do I View Huawei CSI Logs?

Viewing Logs of the huawei-csi-controller Service

Step 1 Run the following command to obtain the node where huawei-csi-controller is located.

```
# kubectl get pod -A -o wide | grep huawei
huawei-csi huawei-csi-controller-695b84b4d8-tg64l 9/9 Running 0 14s <host1-ip>
<host1-name> <none> <none>
```

- **Step 2** Use a remote access tool, such as PuTTY, to log in to the huawei-csi-controller node in the Kubernetes cluster through the management IP address.
- **Step 3** Run the **cd /var/log/huawei** command to go to the log directory. # cd /var/log/huawei
- **Step 4** Run the following command to view the customized output logs of the container. # vi huawei-csi-controller
- **Step 5** Run the **cd /var/log/containers** command to go to the container directory.

 # cd /var/log/containers
- **Step 6** Run the following command to view the standard output logs of the container.

 # vi huawei-csi-controller-<name>_huawei-csi_huawei-csi-driver-<contrainer-id>.log
 - ----End

Viewing Logs of the huawei-csi-node Service

Step 1 Run the following command to obtain the node where huawei-csi-node is located. # kubectl get pod -A -o wide | grep huawei

huawei-csi huawei-csi-node-g6f7z 3/3 **Running** 0 14s <host2-ip> <host2-name> <none>

- **Step 2** Use a remote access tool, such as PuTTY, to log in to the huawei-csi-node node in the Kubernetes cluster through the management IP address.
- **Step 3** Run the **cd /var/log/huawei** command to go to the log directory.

 # cd /var/log/huawei
- **Step 4** Run the following command to view the customized output logs of the container.
- **Step 5** Run the **cd /var/log/containers** command to go to the container directory. # cd /var/log/containers
- **Step 6** Run the following command to view the standard output logs of the container. # vi huawei-csi-node-<name>_huawei-csi_huawei-csi-driver-<contrainer-id>.log
 - ----End

11.2 Failed to Create a Pod Because the iscsi_tcp Service Is Not Started Properly When the Kubernetes Platform Is Set Up for the First Time

Symptom

When you create a Pod, error Cannot connect ISCSI portal *.*.*: libkmod: kmod_module_insert_module: could not find module by name='iscsi_tcp' is reported in the /var/log/huawei-csi-node log.

Root Cause Analysis

The iscsi_tcp service may be stopped after the Kubernetes platform is set up and the iscsi service is installed. You can run the **lsmod | grep iscsi | grep iscsi_tcp** command to check whether the service is stopped.

Solution or Workaround

Run the following command to manually load the iscsi_tcp service.

```
# modprobe iscsi_tcp
# lsmod | grep iscsi | grep iscsi_tcp
iscsi_tcp 18333 6
libiscsi_tcp 25146 1 iscsi_tcp
```

11.3 Failed to Start the huawei-csi-node Service with Error Message "/var/lib/iscsi is not a directory" Reported

Symptom

The huawei-csi-node service cannot be started. When you run the **kubectl describe daemonset huawei-csi-node -n huawei-csi** command, error message "/var/lib/iscsi is not a directory" is reported.

Root Cause Analysis

The /var/lib/iscsi directory does not exist in the huawei-csi-node container.

Solution or Workaround

Step 1 Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.

- **Step 2** Go to the directory where the Helm project is located. If the previous Helm project cannot be found, copy the **helm** directory in the component package to any directory on the master node. For details about the component package path, see **Table 3-1**.
- **Step 3** Go to the **templates** directory and find the **huawei-csi-node.yaml** file.

cd /templates

Step 4 Run the following command to set **path** in **huawei-csi-node.yaml** > **volumes** > **iscsi-dir** > **hostPath** to **/var/lib/iscsi**, save the file, and exit.

vi huawei-csi-node.yaml

Step 5 Run the **helm upgrade** *helm-huawei-csi*./ **-n** *huawei-csi* **-f** *values.yaml* command to upgrade the Helm chart. The upgrade command will update the Deployment, DaemonSet, and RBAC resources. In the preceding command, *helm-huawei-csi* indicates the custom chart name and *huawei-csi* indicates the custom namespace.

helm upgrade helm-huawei-csi ./ -n huawei-csi Release "helm-huawei-csi" has been upgraded. Happy Helming! NAME: helm-huawei-csi LAST DEPLOYED: Thu Jun 9 07:58:15 2022 NAMESPACE: huawei-csi STATUS: deployed REVISION: 2 TEST SUITE: None

----End

11.4 After a Worker Node in the Cluster Breaks Down and Recovers, Pod Failover Is Complete but the Source Host Where the Pod Resides Has Residual Drive Letters

Symptom

A Pod is running on worker node A, and an external block device is mounted to the Pod through CSI. After worker node A is powered off abnormally, the Kubernetes platform detects that the node is faulty and switches the Pod to worker node B. After worker node A recovers, the drive letters on worker node A change from normal to faulty.

Environment Configuration

Kubernetes version: 1.18 or later

Storage type: block storage

Root Cause Analysis

After worker node A recovers, Kubernetes initiates an unmapping operation on the storage, but does not initiate a drive letter removal operation on the host. After Kubernetes completes the unmapping, residual drive letters exist on worker node A.

Solution or Workaround

Currently, you can only manually clear the residual drive letters on the host. Alternatively, restart the host again and use the disk scanning mechanism during the host restart to clear the residual drive letters. The specific method is as follows:

Step 1 Check the residual drive letters on the host.

1. Run the **multipath -ll** command to check whether a DM multipathing device with abnormal multipathing status exists.

As shown in the following figure, the path status is **failed faulty running**, the corresponding DM multipathing device is **dm-12**, and the associated SCSI disks are **sdi** and **sdj**. If multiple paths are configured, multiple SCSI disks exist. Record these SCSI disks.

- If yes, go to Step 1.2.
- If no, no further action is required.
- 2. Check whether the residual DM multipathing device is readable.

Run the **dd if=/dev/***dm-xx* **of=/dev/null count=1 bs=1M iflag=direct** command.

dm-xx indicates the device ID obtained in Step 1.1.

If the returned result is **Input/output error** and the read data is **0 bytes (0 B) copied**, the device is unreadable.

```
#dd if=/dev/dm-12 of=/dev/null count=1 bs=1M iflag=direct
dd: error reading '/dev/dm-12': Input/output error
0+0 records in
0+0 records out
0 bytes (0 B) copied, 0.0236862 s, 0.0 kB/s
```

- If yes, record the residual dm-xx device and associated disk IDs (for details, see Step 1.1) and perform the clearing operation.
- If the command execution is suspended, go to Step 1.3.
- If other cases, contact technical support engineers.
- 3. Log in to the node again in another window.
 - a. Run the following command to view the suspended process.

 # ps -ef | grep dm-12 | grep -w dd

 root 21725 9748 0 10:33 pts/10 00:00:00 dd if=/dev/dm-12 of=/dev/null count=1 bs=10M iflag=direct
 - b. Kill the pid. # kill -9 pid
 - c. Record the residual *dm-xx* device and associated disk IDs (for details, see **Step 1.1**) and perform the clearing operation.

Step 2 Clear the residual drive letters on the host.

1. Run the **multipath -f /dev/***dm-** command to delete residual multipathing aggregation device information according to the DM multipathing device obtained in **Step 1**.

```
# multipath -f /dev/dm-12
```

If an error is reported, contact technical support engineers.

2. Run the following command to clear the residual SCSI disks according to the drive letters of the residual disks obtained in **Step 1**.

```
echo 1 > /sys/block/xxxx/device/delete
```

When multiple paths are configured, clear the residual disks based on the drive letters. The residual paths are **sdi** and **sdi**.

```
# echo 1 > /sys/block/sdi/device/delete
# echo 1 > /sys/block/sdj/device/delete
```

If an error is reported, contact technical support engineers.

3. Check whether the DM multipathing device and SCSI disk information has been cleared.

Run the **multipath** -**ll**, **ls** -**l** /**sys/block**/, and **ls** -**l** /**dev/disk/by-id/** commands in sequence to query the path and disk information. If the residual **dm-12** device and SCSI disks **sdi** and **sdj** are cleared, the clearing is complete.

```
mpathb (3618cf24100f8f457014a764c000001f6) dm-3 HUAWEI ,XSG1
size=100G features='0' hwhandler='0' wp=rw
`-+- policy='service-time 0' prio=-1 status=active
|- 39:0:0:1
               sdd 8:48 active ready running
 `- 38:0:0:1
               sde 8:64 active ready running
mpathn (3618cf24100f8f457315a764c000001f6) dm-5 HUAWEI ,XSG1
size=100G features='0' hwhandler='0' wp=rw
-+- policy='service-time 0' prio=-1 status=active
|- 39:0:0:2

`- 38:0:0:2
               sdc 8:32 active ready running
               sdb 8:16 active ready running
# ls -l /sys/block/
total 0
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-0 -> ../devices/virtual/block/dm-0
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-1 -> ../devices/virtual/block/dm-1
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-2 -> ../devices/virtual/block/dm-2
lrwxrwxrwx 1 root root 0 Aug 11 19:56 dm-3 -> ../devices/virtual/block/dm-3
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdb -> ../devices/platform/host35/session2/
target35:0:0/35:0:0:1/block/sdb
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdc -> ../devices/platform/host34/
target34:65535:5692/34:65535:5692:0/block/sdc
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdd -> ../devices/platform/host39/session6/
target39:0:0/39:0:0:1/block/sdd
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sde -> ../devices/platform/host38/session5/
target38:0:0/38:0:0:1/block/sde
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdh -> ../devices/platform/host39/session6/
target39:0:0/39:0:0:3/block/sdh
lrwxrwxrwx 1 root root 0 Aug 11 19:56 sdi -> ../devices/platform/host38/session5/target38:0:0/38:0:0:3/
block/sdi
ls -l /dev/disk/by-id/
total 0
lrwxrwxrwx 1 root root 10 Aug 11 19:57 dm-name-mpathb -> ../../dm-3
lrwxrwxrwx 1 root root 10 Aug 11 19:58 dm-name-mpathn -> ../../dm-5
lrwxrwxrwx 1 root root 10 Aug 11 19:57 dm-uuid-mpath-3618cf24100f8f457014a764c000001f6 -> ../../
lrwxrwxrwx 1 root root 10 Aug 11 19:58 dm-uuid-mpath-3618cf24100f8f457315a764c000001f6 -> ../../
lrwxrwxrwx 1 root root 9 Aug 11 19:57 scsi-3618cf24100f8f457014a764c000001f6 -> ../../sdd
lrwxrwxrwx 1 root root 9 Aug 11 19:57 scsi-3618cf24100f8f45712345678000103e8 -> ../../sdi
lrwxrwxrwx 1 root root 9 Aug 3 15:17 scsi-3648435a10058805278654321ffffffff -> ../../sdb
lrwxrwxrwx 1 root root 9 Aug 2 14:49 scsi-36888603000020aff44cc0d060c987f1 -> ../../sdc
lrwxrwxrwx 1 root root 9 Aug 11 19:57 wwn-0x618cf24100f8f457014a764c000001f6 -> ../../sdd
lrwxrwxrwx 1 root root 9 Aug 11 19:57 wwn-0x618cf24100f8f45712345678000103e8 -> ../../sdi
lrwxrwxrwx 1 root root 9 Aug 3 15:17 wwn-0x648435a10058805278654321ffffffff -> ../../sdb
lrwxrwxrwx 1 root root 9 Aug 2 14:49 wwn-0x68886030000020aff44cc0d060c987f1 -> ../../sdc
```

----End

11.5 Failed to Start huawei-csi Services with the Status Displayed as InvalidImageName

Symptom

The huawei-csi services (huawei-csi-controller or huawei-csi-node) cannot be started. After the **kubectl get pod -A | grep huawei** command is executed, the command output shows that the service status is **InvalidImageName**.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-fd5f97768-qlldc 6/9 InvalidImageName 0 16s
huawei-csi huawei-csi-node-25txd 2/3 InvalidImageName 0 15s
```

Root Cause Analysis

In the .yaml configuration files of the controller and node, the Huawei CSI image version number is incorrect. For example:

```
...
- name: huawei-csi-driver
image: huawei-csi:4.2.1
...
```

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to modify the configuration file of the huawei-csi-node service. Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

kubectl edit daemonset huawei-csi-node -o yaml -n=huawei-csi

• In huawei-csi-driver in the sample .yaml file, modify image to Huawei CSI image huawei-csi:4.2.1.

containers:

- name: huawei-csi-driver image: huawei-csi:4.2.1
- **Step 3** Run the following command to modify the configuration file of the huawei-csi-controller service: Press I or Insert to enter the editing mode and modify related parameters. After the modification is complete, press **Esc** and enter :wq! to save the modification.

kubectl edit deployment huawei-csi-controller -o yaml -n=huawei-csi

 In huawei-csi-driver in the sample .yaml file, modify image to Huawei CSI image huawei-csi:4.2.1.

containers:

- name: huawei-csi-driver image: huawei-csi:4.2.1

- **Step 4** Wait until the huawei-csi-node and huawei-csi-controller services are started.
- **Step 5** Run the following command to check whether the huawei-csi services are started.

```
# kubectl get pod -A | grep huawei
huawei-csi huawei-csi-controller-58799449cf-zvhmv 9/9 Running 0 2m29s
huawei-csi huawei-csi-node-7fxh6 3/3 Running 0 12m
```

----End

11.6 When a PVC Is Created, the PVC Is in the Pending State

Symptom

A PVC is created. After a period of time, the PVC is still in the **Pending** state.

Root Cause Analysis

Cause 1: A StorageClass with the specified name is not created in advance. As a result, Kubernetes cannot find the specified StorageClass name when a PVC is created.

Cause 2: The storage pool capability does not match the StorageClass capability. As a result, huawei-csi fails to select a storage pool.

Cause 3: An error code (for example, 50331651) is returned by a RESTful interface of the storage. As a result, huawei-csi fails to create a PVC.

Cause 4: The storage does not return a response within the timeout period set by huawei-csi. As a result, huawei-csi returns a timeout error to Kubernetes.

Cause 5: Other causes.

Solution or Workaround

When a PVC is created, if the PVC is in the **Pending** state, you need to take different measures according to the following causes.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to view details about the PVC. # kubectl describe pvc mypvc
- **Step 3** Perform the corresponding operation according to the **Events** information in the detailed PVC information.
 - If the PVC is in the **Pending** state due to cause 1, perform the following steps.

```
Events:

Type Reason Age From Message
---- Warning ProvisioningFailed 0s (x15 over 3m24s) persistentvolume-controller

storageclass.storage.k8s.io "mysc" not found
```

- a. Delete the PVC.
- b. Create a StorageClass. For details, see **8.1.1.1.1 Configuring a StorageClass**.

- Create a PVC. For details, see 8.1.1.1.2 Configuring a PVC.
- If the PVC is in the **Pending** state due to cause 2, perform the following steps.

Events:
Type Reason Age
From Message

Normal Provisioning 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 64s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = **failed to select pool**, the capability filter failed, error: failed to select pool, the final filter field: **replication**, parameters map[allocType:thin replication:True size:1099511627776 volumeType:lun]. please check your storage class

- a. Delete the PVC.
- b. Delete the StorageClass.
- c. Modify the **StorageClass.yaml** file based on the **Events** information.
- d. Create a StorageClass. For details, see **8.1.1.1.1 Configuring a StorageClass**.
- e. Create a PVC. For details, see 8.1.1.1.2 Configuring a PVC.
- If the PVC is in the **Pending** state due to cause 3, contact Huawei engineers.

Events:
Type Reason Age
From Message

Normal Provisioning 63s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 62s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = Create volume map[ALLOCTYPE:1 CAPACITY:20 DESCRIPTION:Created from Kubernetes CSI NAME:pvc-63ebfda5-4cf0-458e-83bd-ecc PARENTID:0] error: 50331651

• If the PVC is in the **Pending** state due to cause 4, perform the following steps.

Events:
Type Reason Age
From Message

Normal Provisioning 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = context deadline exceeded (Client.Timeout exceeded while awaiting headers)

- a. Wait for 10 minutes and check the PVC details again by referring to this section.
- b. If it is still in the **Pending** state, contact Huawei engineers.
- If the PVC is in the **Pending** state due to cause 5, contact Huawei engineers.

----End

11.7 Before a PVC Is Deleted, the PVC Is in the Pending State

Symptom

Before a PVC is deleted, the PVC is in the **Pending** state.

Root Cause Analysis

Cause 1: A StorageClass with the specified name is not created in advance. As a result, Kubernetes cannot find the specified StorageClass name when a PVC is created.

Cause 2: The storage pool capability does not match the StorageClass capability. As a result, huawei-csi fails to select a storage pool.

Cause 3: An error code (for example, 50331651) is returned by a RESTful interface of the storage. As a result, huawei-csi fails to create a PVC.

Cause 4: The storage does not return a response within the timeout period set by huawei-csi. As a result, huawei-csi returns a timeout error to Kubernetes.

Cause 5: Other causes.

Solution or Workaround

To delete a PVC in the **Pending** state, you need to take different measures according to the following causes.

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to view details about the PVC. # kubectl describe pvc mypvc
- **Step 3** Perform the corresponding operation according to the **Events** information in the detailed PVC information.
 - If the PVC is in the **Pending** state due to cause 1, run the **kubectl delete pvc** mypvc command to delete the PVC.

| יקנייי | mypre command to detete the rive. | | | | |
|--|-----------------------------------|-----|------|---------|--|
| Events: | | | | | |
| Type | Reason | Age | From | Message | |
| | | | | | |
| Warning ProvisioningFailed 0s (x15 over 3m24s) persistentvolume-controller | | | | | |
| storageclass.storage.k8s.io " <i>mysc</i> " not found | | | | | |

• If the PVC is in the **Pending** state due to cause 2, run the **kubectl delete pvc** *mypvc* command to delete the PVC.

| Events: | | | |
|-----------|-------------------|------------------------|---|
| Type | Reason | Age | |
| From | | | Message |
| | | | |
| | | | |
| Norma | l Provisioning | 63s (x3 over 64s) c | si.huawei.com_huawei-csi-controller-b59577886- |
| qqzm8_5 | 8533e4a-884c-4c | 7f-92c3-6e8a7b327515 | External provisioner is provisioning volume for |
| claim "de | efault/mypvc" | | |
| Warnin | g ProvisioningFai | iled 63s (x3 over 64s) | csi.huawei.com_huawei-csi-controller-b59577886- |
| | - | | |

qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = **failed to select pool**, the capability filter failed, error: failed to select pool, the final filter field: *replication*, parameters map[allocType:thin replication:True size:1099511627776 volumeType:lun]. please check your storage class

If the PVC is in the **Pending** state due to cause 3, run the **kubectl delete pvc** mypvc command to delete the PVC.

Events:
Type Reason Age
From Message
---Normal Provisioning 63s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"
Warning ProvisioningFailed 62s (x4 over 68s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = Create volume map[ALLOCTYPE:1 CAPACITY:20 DESCRIPTION:Created from Kubernetes CSI NAME:pvc-63ebfda5-4cf0-458e-83bd-ecc PARENTID:0]

• If the PVC is in the **Pending** state due to cause 4, contact Huawei engineers.

| Events: Type | Reason | Age | |
|-----------------|--------|-----|---------|
| From | | | Message |
| | | | |
| | | | |

Normal Provisioning 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 External provisioner is provisioning volume for claim "default/mypvc"

Warning ProvisioningFailed 63s (x3 over 52s) csi.huawei.com_huawei-csi-controller-b59577886-qqzm8_58533e4a-884c-4c7f-92c3-6e8a7b327515 failed to provision volume with StorageClass "mysc": rpc error: code = Internal desc = context deadline exceeded (Client.Timeout exceeded while awaiting headers)

• If the PVC is in the **Pending** state due to cause 5, contact Huawei engineers.

----End

error: 50331651

11.8 When a Pod Is Created, the Pod Is in the ContainerCreating State

Symptom

A Pod is created. After a period of time, the Pod is still in the **ContainerCreating** state. Check the log information (for details, see **11.1 How Do I View Huawei CSI Logs?**). The error message "Fibre Channel volume device not found" is displayed.

Root Cause Analysis

This problem occurs because residual disks exist on the host node. As a result, disks fail to be found when a Pod is created next time.

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query information about the node where the Pod resides.

```
# kubectl get pod -o wide

NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE

READINESS GATES

mypod 0/1 ContainerCreating 0 51s 10.244.1.224 node1 <none> <none>
```

- **Step 3** Delete the Pod.
- **Step 4** Use a remote access tool, such as PuTTY, to log in to the *node1* node in the Kubernetes cluster through the management IP address. *node1* indicates the node queried in **Step 2**.
- **Step 5** Clear the residual drive letters. For details, see **Solution or Workaround**.

----End

11.9 A Pod Is in the ContainerCreating State for a Long Time When It Is Being Created

Symptom

When a Pod is being created, the Pod is in the **ContainerCreating** state for a long time. Check the huawei-csi-node log (for details, see **11.1 How Do I View Huawei CSI Logs?**). No Pod creation information is recorded in the huawei-csi-node log. After the **kubectl get volumeattachment** command is executed, the name of the PV used by the Pod is not displayed in the **PV** column. After a long period of time (more than ten minutes), the Pod is normally created and the Pod status changes to **Running**.

Root Cause Analysis

The kube-controller-manager component of Kubernetes is abnormal.

Solution or Workaround

Contact container platform engineers to rectify the fault.

11.10 A Pod Fails to Be Created and the Log Shows That the Execution of the mount Command Times Out

Symptom

When a Pod is being created, the Pod keeps in the **ContainerCreating** status. In this case, check the log information of huawei-csi-node (for details, see 11.1 How **Do I View Huawei CSI Logs?**). The log shows that the execution of the mount command times out.

Root Cause Analysis

Cause 1: The configured service IP address is disconnected. As a result, the **mount** command execution times out and fails.

Cause 2: For some operating systems, such as Kylin V10 SP1 and SP2, it takes a long time to run the **mount** command in a container using NFSv3. As a result, the

mount command may time out. The possible cause is that the value of **LimitNOFILE** of container runtime containerd is too large (1+ billion).

Cause 3: The mounting may fail due to network problems. The default mounting timeout period of CSI is 30 seconds. If the mounting still fails after 30 seconds, logs show that the execution of the **mount** command times out.

Solution or Workaround

- **Step 1** Run the **ping** command to check whether the service IP network is connected. If the ping fails, the fault is caused by cause 1. In this case, configure an available service IP address. If the ping succeeds, go to **Step 2**.
- Step 2 Go to any container where the mount command can be executed and use NFSv3 to run the mount command. If the command times out, the fault may be caused by cause 2. Run the systemctl status containerd.service command to check the configuration file path, and then run the cat /xxx/containerd.service command to check the configuration file. If the file contains LimitNOFILE=infinity or the value of LimitNOFILE is 1 billion, go to Step 3. Otherwise, contact Huawei technical support engineers.
- **Step 3** For cause 2, perform the following operations:
 - Try using NFSv4.0 or NFSv4.1.
 - Change the value of **LimitNOFILE** to a proper one by referring to **change solution provided by the community**. This solution will restart the container runtime. Evaluate the impact on services.
- **Step 4** Manually mount the file system on the host machine where the mounting fails. If the required time exceeds 30 seconds, check whether the network between the host machine and the storage node is normal. An example of the **mount** command is as follows.
 - Run the following command to create a test directory. # mkdir /tmp/test_mount
 - Run the mount command to mount the file system and observe the time consumed. The value of ip:nfs_share_path can be obtained from the huawei-csi-node log. For details, see 11.1 How Do I View Huawei CSI Logs?.
 # time mount ip:nfs_share_path /tmp/test_mount
 - After the test is complete, run the following command to unmount the file system.

umount /tmp/test_mount

----End

11.11 A Pod Fails to Be Created and the Log Shows That the mount Command Fails to Be Executed

Symptom

In NAS scenarios, when a Pod is being created, the Pod keeps in the **ContainerCreating** status. In this case, check the log information of huawei-csi-node (for details, see 11.1 How Do I View Huawei CSI Logs?). The log shows that the mount command fails to be executed.

Root Cause Analysis

The possible cause is that the NFS 4.0/4.1 protocol is not enabled on the storage side. After the NFS v4 protocol fails to be used for mounting, the host does not negotiate to use the NFS v3 protocol for mounting.

Solution or Workaround

- Enable the NFS 3/4.0/4.1 protocol on the storage side and retry the default mounting.
- Specify an available NFS protocol for mounting. For details, see 8.1.1.1.1
 Configuring a StorageClass.

11.12 After a Pod Fails to Be Created or kubelet Is Restarted, Logs Show That the Mount Point Already Exists

Symptom

When a Pod is being created, the Pod is always in the **ContainerCreating** state. Alternatively, after kubelet is restarted, logs show that the mount point already exists. Check the log information of huawei-csi-node (for details, see 11.1 How Do I View Huawei CSI Logs?). The error information is: The mount /var/lib/kubelet/pods/xxx/mount is already exist, but the source path is not /var/lib/kubelet/plugins/kubernetes.io/xxx/globalmount

Root Cause Analysis

The root cause of this problem is that Kubernetes performs repeated mounting operations.

Solution or Workaround

Run the following command to unmount the existing path. In the command, /var/lib/kubelet/pods/xxx/mount indicates the existing mount path displayed in the logs.

umount /var/lib/kubelet/pods/xxx/mount

11.13 How Do I Download a Container Image to the Local PC?

Download a Container Image Using containerd

Step 1 Run the **ctr image pull** *image:tag* command to download an image to the local PC. In the preceding command, *image:tag* indicates the image to be pulled and its tag.

ctr image pull k8s.gcr.io/sig-storage/livenessprobe:v2.5.0

Step 2 Run the **ctr image export** *image.tar image:tag* command to export the image to a file. In the preceding command, *image:tag* indicates the image to be exported, and *image.tar* indicates the name of the exported image file.

ctr image export livenessprobe.tar k8s.gcr.io/sig-storage/livenessprobe:v2.5.0

----End

Download a Container Image Using Docker

- **Step 1** Run the **docker pull** *image:tag* command to download an image to the local PC. In the preceding command, *image:tag* indicates the image to be pulled.
 - # docker pull k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
- **Step 2** Run the **docker save** *image:tag* **-o** *image.tar* command to export the image to a file. In the preceding command, *image:tag* indicates the image to be exported, and *image.tar* indicates the name of the exported image file.

docker save k8s.gcr.io/sig-storage/livenessprobe:v2.5.0 -o livenessprobe.tar

----End

Download a Container Image Using Podman

- **Step 1** Run the **podman pull** *image:tag* command to download an image to the local PC. In the preceding command, *image:tag* indicates the image to be pulled.
 - # podman pull k8s.gcr.io/sig-storage/livenessprobe:v2.5.0
- **Step 2** Run the **podman save** *image:tag* **-o** *image.tar* command to export the image to a file. In the preceding command, *image:tag* indicates the image to be exported, and *image.tar* indicates the name of the exported image file.

podman save k8s.gcr.io/sig-storage/livenessprobe:v2.5.0 -o livenessprobe.tar

----End

11.14 How Do I Obtain CSI Version Information?

This section describes how to view the CSI version.

Procedure

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the following command to query information about the node where huaweicsi-node resides.

```
# kubectl get pod -A -owide | grep huawei-csi-node
                                       READY STATUS RESTARTS
NAMESPACE
            NAME
                                                                    AGE
                                                                          ΙP
           NOMINATED NODE READINESS GATES
NODE
huawei-csi huawei-csi-node-87mss
                                         3/3
                                               Running 0
                                                                 6m41s
192.168.129.155
               node-1
                           <none>
                                        <none>
huawei-csi huawei-csi-node-xp8cc
                                         3/3
                                              Running 0
                                                                6m41s 192.168.129.156
node-2
          <none>
```

- **Step 3** Use a remote access tool, such as PuTTY, to log in to any node where huawei-csi-node resides through the node IP address.
- **Step 4** Run the following command to view the CSI version.

cat /var/lib/kubelet/plugins/csi.huawei.com/version 4.2.1

----End

11.15 Common Problems and Solutions for Interconnecting with the Tanzu Kubernetes Cluster

This section describes the common problems and solutions for interconnecting with the Tanzu Kubernetes cluster. Currently, the following problems occur during interconnection with the Tanzu Kubernetes cluster:

- A Pod cannot be created because the PSP permission is not created.
- The mount point of the host is different from that of the native Kubernetes. As a result, a volume fails to be mounted.
- The livenessprobe container port conflicts with the Tanzu vSphere port. As a result, the container restarts repeatedly.

11.15.1 A Pod Cannot Be Created Because the PSP Permission Is Not Created

Symptom

When huawei-csi-controller and huawei-csi-node are created, only the Deployment and DaemonSet resources are successfully created, and no Pod is created for the controller and node.

Root Cause Analysis

The service account used for creating resources does not have the "use" permission of the PSP policy.

Solution or Workaround

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run the **vi** *psp-use.yaml* command to create a file named **psp-use.yaml** # vi psp-use.yaml

Step 3 Configure the **psp-use.yaml** file.

```
apiVersion: rbac.authorization.k8s.io/v1 kind: ClusterRole metadata: name: huawei-csi-psp-role rules: - apiGroups: ['policy'] resources: ['podsecuritypolicies'] verbs: ['use'] --- apiVersion: rbac.authorization.k8s.io/v1 kind: ClusterRoleBinding metadata: name: huawei-csi-psp-role-cfg roleRef:
```

kind: ClusterRole
name: huawei-csi-psp-role
apiGroup: rbac.authorization.k8s.io
subjects:
- kind: Group
apiGroup: rbac.authorization.k8s.io
name: system:serviceaccounts:huawei-csi
- kind: Group
apiGroup: rbac.authorization.k8s.io
name: system:serviceaccounts:default

Step 4 Run the **kubectl create -f** *psp-use.yaml* command to create the PSP permission.

kubectl create -f psp-use.yaml

----End

11.15.2 Changing the Mount Point of a Host

Symptom

A Pod fails to be created, and error message "mount point does not exist" is recorded in Huawei CSI logs.

Root Cause Analysis

The native Kubernetes cluster in the **pods-dir** directory of huawei-csi-node is inconsistent with the Tanzu Kubernetes cluster.

Solution or Workaround

Step 1 Go to the **helm/esdk/** directory and run the **vi values.yaml** command to open the configuration file.

vi values.yaml

Step 2 Change the value of **kubeletConfigDir** to the actual installation directory of kubelet.

#
Specify kubelet config dir path.
kubernetes and openshift is usually /var/lib/kubelet
Tanzu is usually /var/vcap/data/kubelet
kubeletConfigDir: /var/vcap/data/kubelet

----End

11.15.3 Changing the Default Port of the livenessprobe Container

Symptom

The livenessprobe container of the huawei-csi-controller component keeps restarting.

Root Cause Analysis

The default port (9808) of the livenessprobe container of huawei-csi-controller conflicts with the existing vSphere CSI port of Tanzu.

Solution or Workaround

Change the default port of the livenessprobe container to an idle port.

Step 1 Go to the **helm/esdk/templates** directory and run the **vi huawei-csi-controller.yaml** command to open the controller configuration file.

vi huawei-csi-controller.yaml

Step 2 Change the default port (9808) of the livenessprobe container to an idle port.

----End

11.16 Common Problems and Solutions for Using the Tanzu Kubernetes Cluster

This section describes the common problems and solutions for using the Tanzu Kubernetes cluster.

11.16.1 Failed to Create an Ephemeral Volume

Symptom

A generic ephemeral volume fails to be created, and the error message PodSecurityPolicy: unable to admit pod: [spec.volumes[0]: Invalid value: "ephemeral": ephemeral volumes are not allowed to be used spec.volumes[0] is displayed.

Root Cause Analysis

The current PSP policy does not contain the permission to use ephemeral volumes.

Solution or Workaround

Add the permission to use ephemeral volumes to the default PSP **pks-privileged** and **pks-restricted**. The following is an example of modifying **pks-privileged**:

- **Step 1** Use a remote access tool, such as PuTTY, to log in to any master node in the Kubernetes cluster through the management IP address.
- **Step 2** Run **kubectl edit psp pks-privileged** to modify the **pks-privileged** configuration. # kubectl edit psp pks-privileged
- **Step 3** Add **ephemeral** to **spec.volumes**. The following is an example.

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
annotations:
apparmor.security.beta.kubernetes.io/allowedProfileName: '*'
seccomp.security.alpha.kubernetes.io/allowedProfileNames: '*'
creationTimestamp: "2022-10-11T08:07:00Z"
name: pks-privileged
resourceVersion: "1227763"
```

```
uid: 2f39c44a-2ce7-49fd-87ca-2c5dc3bfc0c6
spec:
allowPrivilegeEscalation: true
allowedCapabilities:
- '*'
supplementalGroups:
rule: RunAsAny
volumes:
- glusterfs
- hostPath
- iscsi
- nfs
- persistentVolumeClaim
- ephemeral
```

Step 4 Run the **kubectl get psp pks-privileged -o yaml** command to check whether the addition is successful.

kubectl get psp pks-privileged -o yaml

----End

11.17 Failed to Expand the Capacity of a Generic Ephemeral Volume

Symptom

In an environment where the Kubernetes version is earlier than 1.25, the capacity of a **generic ephemeral volume** of the LUN type fails to be expanded. The system displays a message indicating that the PV capacity has been expanded, but the PVC capacity fails to be updated.

Root Cause Analysis

This problem is caused by a Kubernetes **bug**, which has been resolved in Kubernetes 1.25.

11.18 Failed to Expand the PVC Capacity Because the Target Capacity Exceeds the Storage Pool Capacity

Symptom

In a Kubernetes environment earlier than 1.23, PVC capacity expansion fails when the target capacity exceeds the storage pool capacity.

Root Cause Analysis

This is a known issue in the Kubernetes community. For details, see **Recovering** from Failure when Expanding Volumes.

Solution or Workaround

For details, see Recovering from Failure when Expanding Volumes.

11.19 A webhook Fails to Be Called When the oceanctl Tool Is Used to Manage Backends

Symptom

After the webhook configuration is changed, for example, the value of the **webhookPort** parameter is changed, an error is reported indicating that a webhook fails to be called when the oceanctl tool is used to manage backends, as shown in the following figure.

```
root@ubuntu-master:/opt/huawei-csi/backend# oceanctl delete backend onas

Error: secret "onas" deleted
configmap "onas" deleted
configmap "onas" deleted
rror from server (InternalError): Internal error occurred: failed calling webhook "storage-backend-controller.xuanwu.huawei.i
ler.huawei-csi.svc:443/storagebackendclaim?timeout=10s": no service port 443 found for service "huawei-csi-controller"
```

Root Cause Analysis

After the webhook configuration changes, the **validatingwebhookconfiguration** resource becomes invalid.

Solution or Workaround

Step 1 Run the following command to delete the **validatingwebhookconfiguration** resource.

kubectl delete validatingwebhookconfiguration storage-backend-controller.xuanwu.huawei.io validatingwebhookconfiguration.admissionregistration.k8s.io "storage-backend-controller.xuanwu.huawei.io" deleted

Step 2 Run the following command to restart CSI Controller. Run the --replicas=* command to set the number of CSI Controller copies to be restored. In the following example, the number of copies to be restored is 1. Change it based on site requirements.

```
# kubectl scale deployment huawei-csi-controller -n huawei-csi --replicas=0 deployment.apps/huawei-csi-controller scaled # kubectl scale deployment huawei-csi-controller -n huawei-csi --replicas=1 deployment.apps/huawei-csi-controller scaled
```

Step 3 Run the following command to check whether CSI Controller is successfully started.

```
# kubectl get pod -n huawei-csi
NAME READY STATUS RESTARTS AGE
huawei-csi-controller-58d5b6b978-s2dsq 9/9 Running 0 19s
huawei-csi-node-dt6nd 3/3 Running 0 77m
```

----End

11.20 An Account Is Locked After the Password Is Updated on the Storage Device

Symptom

After a user changes the password on the storage device, the account is locked.

Root Cause Analysis

CSI uses the account and password configured on the storage device to log in to the storage device. After the account password is changed on the storage device, CSI attempts to log in to the storage device again after the login fails. Take Dorado V6 as an example. The default login policy is that an account will be locked after three consecutive password verification failures. Therefore, when CSI retries for more than three times, the account will be locked.

Solution or Workaround

- **Step 1** If the backend account is **admin**, run the following command to set the number of huawei-csi-controller service copies to 0. If an account other than **admin** is used, skip this step.
 - kubectl scale deployment huawei-csi-controller -n huawei-csi --replicas=0
- Step 2 Log in to the storage device as user admin and modify the login policy. Take Dorado V6 as an example. On DeviceManager, choose Settings > User and Security > Security Policies > Login Policy, click Modify, and disable Account Lockout.
- **Step 3** If the backend account is admin, run the following command to restore the number of CSI Controller copies using --replicas=*. In the following example, the number of copies is restored to 1. Change it based on the site requirements. If an account other than **admin** is used, skip this step.

 kubectl scale deployment huawei-csi-controller -n huawei-csi --replicas=1
- **Step 4** Use the oceanctl tool to change the storage backend password. For details about how to change the backend password, see **7.3 Updating a Storage Backend**.
- Step 5 Log in to the storage device as user admin and modify the login policy. Take Dorado V6 as an example. On DeviceManager, choose Settings > User and Security > Security Policies > Login Policy, click Modify, and enable Account Lockout.
 - ----End

12 Appendix

- 12.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3
- 12.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x
- 12.3 Example ALUA Configuration Policy of Distributed Storage
- 12.4 Installing Helm 3
- 12.5 Configuring Custom Permissions

12.1 Example ALUA Configuration Policy of OceanStor V3/V5 and OceanStor Dorado V3

Example 1: The configuration file content is as follows:

```
parameters:
ALUA:

"*":

MULTIPATHTYPE: 1

FAILOVERMODE: 3

SPECIALMODETYPE: 0

PATHTYPE: 0

node1:

MULTIPATHTYPE: 1

FAILOVERMODE: 3

SPECIALMODETYPE: 0

PATHTYPE: 1
```

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

```
parameters:
ALUA:
node[0-9]:
MULTIPATHTYPE: 1
```

```
FAILOVERMODE: 3
SPECIALMODETYPE: 0
PATHTYPE: 0
node[5-7]:
MULTIPATHTYPE: 1
FAILOVERMODE: 3
SPECIALMODETYPE: 0
PATHTYPE: 1
```

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

```
parameters:
ALUA:
node$:
MULTIPATHTYPE: 1
FAILOVERMODE: 3
SPECIALMODETYPE: 0
PATHTYPE: 0
node10$:
MULTIPATHTYPE: 1
FAILOVERMODE: 3
SPECIALMODETYPE: 0
PATHTYPE: 1
```

According to the configuration policy rules in **9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.

12.2 Example ALUA Configuration Policy of OceanStor Dorado 6.x

Example 1: The configuration file content is as follows:

```
parameters:
ALUA:

"*":

accessMode: 1
hyperMetroPathOptimized: 1
node1:
accessMode: 1
hyperMetroPathOptimized: 0
```

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

```
parameters:
ALUA:
node[0-9]:
accessMode: 1
```

```
hyperMetroPathOptimized: 1
node[5-7]:
accessMode: 1
hyperMetroPathOptimized: 0
```

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend**, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

```
parameters:
node1$:
node[0-9]:
accessMode: 1
hyperMetroPathOptimized: 1
node10$:
accessMode: 1
hyperMetroPathOptimized: 0
```

According to the configuration policy rules in **9.1.1.1 Configuring ALUA Parameters for a Huawei Enterprise Storage Backend**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.

12.3 Example ALUA Configuration Policy of Distributed Storage

Example 1: The configuration file content is as follows:

```
parameters:
ALUA:

"*":

switchoverMode: Enable_alua
pathType: optimal_path
node1:
switchoverMode: Enable_alua
pathType: non_optimal_path
```

If the host name is **node1**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in **9.1.1.2 Configuring ALUA Parameters for a Distributed Storage Backend**, the priority of the second configuration section (where **HostName** is **node1**) is higher than that of the first configuration section (where **HostName** is *).

Example 2: The configuration file content is as follows:

```
parameters:
ALUA:
node[0-9]:
switchoverMode: Enable_alua
pathType: optimal_path
node[5-7]:
switchoverMode: Enable_alua
pathType: non_optimal_path
```

If the host name is **node6**, both of the preceding ALUA configuration sections can be used to configure initiators. According to the configuration policy rules in

9.1.1.2 Configuring ALUA Parameters for a Distributed Storage Backend, select the first ALUA configuration section to configure initiators.

Example 3: The configuration file content is as follows:

```
parameters:
ALUA:
node1$:
switchoverMode: Enable_alua
pathType: optimal_path
node10$:
switchoverMode: Enable_alua
pathType: non_optimal_path
```

According to the configuration policy rules in **9.1.1.2 Configuring ALUA Parameters for a Distributed Storage Backend**: For host **node1**, select the first ALUA configuration section to configure initiators. For host **node10**, select the second ALUA configuration section to configure initiators. ^ matches the beginning of a character string, and \$ matches the end of a character string.

12.4 Installing Helm 3

This section describes how to install Helm 3.

For details, see https://helm.sh/docs/intro/install/.

Prerequisites

Ensure that the master node in the Kubernetes cluster can access the Internet.

Procedure

- **Step 1** Run the following command to download the Helm 3 installation script.

 # curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
- **Step 2** Run the following command to modify the permission on the Helm 3 installation script.

chmod 700 get_helm.sh

Step 3 Determine the Helm version to be installed based on the version mapping between Helm and Kubernetes. For details about the version mapping, see Helm Version Support Policy. Then run the following command to change the DESIRED_VERSION environment variable to the Helm version to be installed and run the installation command.

DESIRED_VERSION=v3.9.0 ./get_helm.sh

Step 4 Run the following command to check whether Helm 3 of the specified version is successfully installed.

```
# helm version version:"v3.9.0", GitCommit:"7ceeda6c585217a19a1131663d8cd1f7d641b2a7", GitTreeState:"clean", GoVersion:"go1.17.5"}
```

----End

12.5 Configuring Custom Permissions

User-defined Role Configurations

For different storage resources, refer to the following configurations:

- For NAS resources, configure the minimum permissions by referring to **Table** 12-1.
- For SAN resources, configure the minimum permissions by referring to **Table** 12-2.

□ NOTE

For details about how to configure permissions for user-defined roles, see OceanStor Dorado 6000, Dorado 18000 Series Product Documentation.

Table 12-1 Minimum permissions for NAS resources

| Permission Object | Parent Object | Read/Write Permission | Function |
|------------------------|-----------------------|--------------------------|---|
| workload_type | file_storage_service | Read-only | Queries the workload type. |
| file_system | file_storage_service | Read and write | Manages file systems. |
| fs_snapshot | file_storage_service | Read and write | Manages file system snapshots. |
| quota | file_storage_service | Read and write | Manages file system quotas. |
| nfs_service | file_storage_service | Read-only | Queries NFS services. |
| share | file_storage_service | Read and write | Manages NFS shares. |
| dtree | file_storage_service | Read and write | Manages dtrees. |
| hyper_metro_p air | hyper_metro | Read and write | Creates file system HyperMetro pairs. |
| hyper_metro_d omain | hyper_metro | Read-only | Queries information about file system HyperMetro domains. |
| remote_device | local_data_protection | Read-only | Queries remote device information. |
| storage_pool | pool | Read-only | Queries storage pool information. |

| Permission Object | Parent Object | Read/Write Permission | Function |
|----------------------|---------------------------------|--------------------------|---|
| smart_qos | resource_performanc e_tuning | Read and write | Manages SmartQoS policies. |
| system | system | Read-only | Queries storage device information (this object needs to be configured only when the owning group is the system group). |
| vstore | vstore | Read-only | Queries vStore information. |

Table 12-2 Minimum permissions for SAN resources

| Permission Object | Parent Object | Read/Write Permission | Function |
|----------------------|-----------------------|--------------------------|------------------------------------|
| remote_device | local_data_protection | Read-only | Queries remote device information. |
| hyper_clone | local_data_protection | Read and write | Manages clone pairs. |
| lun_snapshot | local_data_protection | Read and write | Manages LUN snapshots. |
| workload_type | lun | Read-only | Queries the workload type. |
| lun | lun | Read and write | Manages LUNs. |
| host | mapping_view | Read and write | Manages hosts. |
| host_group | mapping_view | Read and write | Manages host groups. |
| initiator | mapping_view | Read and write | Manages initiators. |
| lun_group | mapping_view | Read and write | Manages LUN groups. |
| mapping_view | mapping_view | Read and write | Manages mapping views. |
| target | mapping_view | Read-only | Queries iSCSI initiators. |
| port | network | Read-only | Queries logical ports. |

| Permission Object | Parent Object | Read/Write Permission | Function |
|----------------------|---------------------------------|--------------------------|---|
| storage_pool | pool | Read-only | Queries storage pool information. |
| smart_qos | resource_performanc e_tuning | Read and write | Manages SmartQoS policies. |
| system | system | Read-only | Queries storage device information (this object needs to be configured only when the owning group is the system group). |
| vstore | vstore | Read-only | Queries vStore information. |